SLIB

The Portable Scheme Library
Version 3b6, February 2020

Aubrey Jaffer

This manual is for SLIB (version 3b6, February 2020), the portable Scheme library.

Copyright © 1993, 1994, 1995, 1996, 1997, 1998, 1999, 2000, 2001, 2002, 2003, 2004, 2005, 2006, 2007, 2008 Free Software Foundation, Inc.

Table of Contents

6 Database Packages 162

1 The Library System

SLIB is a portable library for the programming language *Scheme*. It provides a platform independent framework for using *packages* of Scheme procedures and syntax. As distributed, SLIB contains useful packages for all Scheme implementations. Its catalog can be transparently extended to accomodate packages specific to a site, implementation, user, or directory.

1.1 Feature

SLIB denotes *features* by symbols. SLIB maintains a list of features supported by a Scheme *session*. The set of features provided by a session may change during that session. Some features are properties of the Scheme implementation being used. The following *intrinsic features* detail what sort of numbers are available from an implementation:

- 'inexact
- 'rational
- 'real
- 'complex
- 'bignum

SLIB initialization (in `require.scm`) tests and *provides* any of these numeric features which are appropriate.

Other features correspond to the presence of packages of Scheme procedures or syntax (macros).

`provided?` *feature*   [Function]

> Returns `#t` if *feature* is present in the current Scheme session; otherwise `#f`. More specifically, `provided?` returns `#t` if the symbol *feature* is the `software-type`, the `scheme-implementation-type`[1], or if *feature* has been provided by a module already loaded; and `#f` otherwise.
>
> In some implementations `provided?` tests whether a module has been `require`d by any module or in any thread; other implementations will have `provided?` reflect only the modules `require`d by that particular session or thread.
>
> To work portably in both scenarios, use `provided?` only to test whether intrinsic properties (like those above) are present.
>
> The *feature* argument can also be an expression calling `and`, `or`, and `not` of features. The boolean result of the logical question asked by *feature* is returned.

The generalization of `provided?` for arbitrary features and catalog is `feature-eval`:

`feature-eval` *expression provided?*   [Function]

> Evaluates `and`, `or`, and `not` forms in *expression*, using the values returned by calling *provided?* on the leaf symbols. `feature-eval` returns the boolean result of the logical combinations.

---

[1] scheme-implementation-type is the name symbol of the running Scheme implementation (RScheme, |STk|, Bigloo, chez, Elk, gambit, gauche, guile, JScheme, kawa, MacScheme, MITScheme, PocketScheme, S7, Scheme48, Scheme->C, Scheme48, Scsh, SISC, T, umb-scheme, or Vscm). Dependence on scheme-implementation-type is almost always the wrong way to do things.

`provide` *feature* [Procedure]

Informs SLIB that *feature* is supported in this session.

```
(provided? 'foo)     ⇒ #f
(provide 'foo)
(provided? 'foo)     ⇒ #t
```

1.2 Require

SLIB creates and maintains a *catalog* mapping features to locations of files introducing procedures and syntax denoted by those features.

`*catalog*` [Variable]

Is an association list of features (symbols) and pathnames which will supply those features. The pathname can be either a string or a pair. If pathname is a pair then the first element should be a macro feature symbol, `source`, `compiled`, or one of the other cases described in Section 1.3 [Library Catalogs], page 3. The cdr of the pathname should be either a string or a list.

At the beginning of each section of this manual, there is a line like (`require 'feature`). The Scheme files comprising SLIB are cataloged so that these feature names map to the corresponding files.

SLIB provides a form, `require`, which loads the files providing the requested feature.

`require` *feature* [Procedure]

- If (`provided? feature`) is true, then `require` just returns.
- Otherwise, if *feature* is found in the catalog, then the corresponding files will be loaded and (`provided? feature`) will henceforth return #t. That *feature* is thereafter `provided`.
- Otherwise (*feature* not found in the catalog), an error is signaled.

There is a related form `require-if`, used primarily for enabling compilers to statically include modules which would be dynamically loaded by interpreters.

`require-if` *condition feature* [Procedure]

Requires *feature* if *condition* is true.

The `random` module uses `require-if` to flag `object->string` as a (dynamic) required module.

```
(require 'byte)
(require 'logical)
(require-if 'compiling 'object->string)
```

The `batch` module uses `require-if` to flag `posix-time` as a module to load if the implementation supports large precision exact integers.

```
(require-if '(and bignum compiling) 'posix-time)
```

The catalog can also be queried using `slib:in-catalog?`.

`slib:in-catalog?` *feature*   [Function]
> Returns a `CDR` of the catalog entry if one was found for the symbol *feature* in the alist
> `*catalog*` (and transitively through any symbol aliases encountered). Otherwise,
> returns `#f`. The format of catalog entries is explained in Section 1.3 [Library Catalogs],
> page 3.

1.3 Library Catalogs

Catalog files consist of one or more *association list*s. In the circumstance where a feature
symbol appears in more than one list, the latter list's association is retrieved. Here are the
supported formats for elements of catalog lists:

(*feature* . `<symbol>`)
> Redirects to the feature named *<symbol>*.

(*feature* . `"<path>"`)
> Loads file *<path>*.

(*feature* `source` `"<path>"`)
> `slib:load`s the Scheme source file *<path>*.

(*feature* `compiled` `"<path>"` ...)
> `slib:load-compiled`s the files *<path>* ....

(*feature* `aggregate` `<symbol>` ...)
> `require`s the features *<symbol>* ....

The various macro styles first `require` the named macro package, then just load *<path>* or
load-and-macro-expand *<path>* as appropriate for the implementation.

(*feature* `defmacro` `"<path>"`)
> `defmacro:load`s the Scheme source file *<path>*.

(*feature* `macro-by-example` `"<path>"`)
> `defmacro:load`s the Scheme source file *<path>*.

(*feature* `macro` `"<path>"`)
> `macro:load`s the Scheme source file *<path>*.

(*feature* `macros-that-work` `"<path>"`)
> `macro:load`s the Scheme source file *<path>*.

(*feature* `syntax-case` `"<path>"`)
> `macro:load`s the Scheme source file *<path>*.

(*feature* `syntactic-closures` `"<path>"`)
> `macro:load`s the Scheme source file *<path>*.

1.4 Catalog Creation

At the start of an interactive session no catalog is present, but is created with the first
catalog inquiry (such as `(require 'random)`). Several sources of catalog information are
combined to produce the catalog:

- standard SLIB packages.

- additional packages of interest to this site.

- packages specifically for the variety of Scheme which this session is running. This catalog, if it exists, is the file `implcat` in `implementation-invicinity`, which is created by loading `mkimpcat.scm` in `implementation-invicinity` if it exists.

- packages this user wants to always have available. This catalog is the file `homecat` in the user's *HOME* directory.

- packages germane to working in this (current working) directory. This catalog is the file `usercat` in the directory to which it applies. One would typically `cd` to this directory before starting the Scheme session.

- packages which are part of an application program.

SLIB combines the catalog information which doesn't vary per user into the file `slibcat` in the implementation-vicinity. Therefore `slibcat` needs change only when new software is installed or compiled. Because the actual pathnames of files can differ from installation to installation, SLIB builds a separate catalog for each implementation it is used with.

The definition of `*slib-version*` in SLIB file `require.scm` is checked against the catalog association of `*slib-version*` to ascertain when versions have changed. It is a reasonable practice to change the definition of `*slib-version*` whenever the library is changed. If multiple implementations of Scheme use SLIB, remember that recompiling one `slibcat` will update only that implementation's catalog.

The compilation scripts of Scheme implementations which work with SLIB can automatically trigger catalog compilation by deleting `slibcat` or by invoking `require` of a special feature:

`require` 'new-catalog   [Procedure]
>   This will load `mklibcat`, which compiles and writes a new `slibcat`.

Another special feature of `require` erases SLIB's catalog, forcing it to be reloaded the next time the catalog is queried.

`require` #f   [Procedure]
>   Removes SLIB's catalog information. This should be done before saving an executable image so that, when restored, its catalog will be loaded afresh.

1.5 Catalog Vicinities

Each file in the table below is descibed in terms of its file-system independent *vicinity* (see Section 2.1 [Vicinity], page 11). The entries of a catalog in the table override those of catalogs above it in the table.

`implementation-vicinity slibcat`
>   This file contains the associations for the packages comprising SLIB, the `implcat` and the `sitecat`s. The associations in the other catalogs override those of the standard catalog.

`library-vicinity mklibcat.scm`
>   creates `slibcat`.

`library-vicinity` sitecat

> This file contains the associations specific to an SLIB installation.

`implementation-vicinity` implcat

> This file contains the associations specific to an implementation of Scheme. Different implementations of Scheme should have different `implementation-vicinity`.

`implementation-vicinity` mkimpcat.scm

> if present, creates `implcat`.

`implementation-vicinity` sitecat

> This file contains the associations specific to a Scheme implementation installation.

`home-vicinity` homecat

> This file contains the associations specific to an SLIB user.

`user-vicinity` usercat

> This file contains associations affecting only those sessions whose *working directory* is `user-vicinity`.

Here is an example of a `usercat` catalog. A program in this directory can invoke the 'run' feature with (`require 'run`).

```
;;; "usercat": SLIB catalog additions for SIMSYNCH.    -*-scheme-*-
(
 (simsynch      . "../synch/simsynch.scm")
 (run           . "../synch/run.scm")
 (schlep        . "schlep.scm")
)
```

Copying `usercat` to many directories is inconvenient. Application programs which aren't always run in specially prepared directories can nonetheless register their features during initialization.

`catalog:read` *vicinity catalog*   [Procedure]

> Reads file named by string *catalog* in *vicinity*, resolving all paths relative to *vicinity*, and adds those feature associations to **catalog**.
>
> `catalog:read` would typically be used by an application program having dynamically loadable modules. For instance, to register factoring and other modules in **catalog**, JACAL does:
>
> ```
> (catalog:read (program-vicinity) "jacalcat")
> ```

For an application program there are three appropriate venues for registering its catalog associations:

- in a `usercat` file in the directory where the program runs; or
- in an `implcat` file in the `implementation-vicinity`; or
- in an application program directory; loaded by calling `catalog:read`.

1.6 Compiling Scheme

To use Scheme compilers effectively with SLIB the compiler needs to know which SLIB modules are to be compiled and which symbols are exported from those modules.

The procedures in this section automate the extraction of this information from SLIB modules. They are guaranteed to work on SLIB modules; to use them on other sources, those sources should follow SLIB conventions.

1.6.1 Module Conventions

- All the top-level `require` commands have one quoted argument and are positioned before other Scheme definitions and expressions in the file.
- Any conditionally `require`d SLIB modules[2] also appear at the beginning of their files conditioned on the feature `compiling` using `require-if` (see Section 1.2 [Require], page 2).

  ```
  (require 'logical)
  (require 'multiarg/and-)
  (require-if 'compiling 'sort)
  (require-if 'compiling 'ciexyz)
  ```

- Schmooz-style comments preceding a definition, identify that definition as an exported identifier (see Section 4.15 [Schmooz], page 101). For non-schmooz files, putting ';@' at the beginning of the line immediately preceding the definition (`define`, `define-syntax`, or `defmacro`) suffices.

  ```
  ;@
  (define (identity <obj>) <obj>)
  ```

- Syntax (macro) definitions are grouped at the end of a module file.
- Modules defining macros do not invoke those macros. SLIB macro implementations are exempt from this rule.

  An example of how to expand macro invocations is:

  ```
  (require 'macros-that-work)
  (require 'yasos)
  (require 'pprint-file)
  (pprint-filter-file "collect.scm" macwork:expand)
  ```

1.6.2 Module Manifests

```
(require 'manifest)
```

In some of these examples, *slib:catalog* is the SLIB part of the catalog; it is free of compiled and implementation-specific entries. It would be defined by:

```
(define slib:catalog (cdr (member (assq 'null *catalog*) *catalog*)))
```

`file->requires` *file provided? catalog*   [Function]
> Returns a list of the features `require`d by *file* assuming the predicate *provided?* and association-list *catalog*.

---

[2] There are some functions with internal `require` calls to delay loading modules until they are needed. While this reduces startup latency for interpreters, it can produce headaches for compilers.

```
(define (provided+? . features)
  (lambda (feature)
    (or (memq feature features) (provided? feature))))

(file->requires "obj2str.scm" (provided+? 'compiling) '())
        ⇒ (string-port generic-write)

(file->requires "obj2str.scm" provided? '())
        ⇒ (string-port)
```

feature->requires *feature provided? catalog*   [Function]
Returns a list of the features **require**d by *feature* assuming the predicate *provided?*
and association-list *catalog*.

```
(feature->requires 'batch (provided+? 'compiling) *catalog*)
        ⇒ (tree line-i/o databases parameters string-port
                pretty-print common-list-functions posix-time)

(feature->requires 'batch provided? *catalog*)
        ⇒ (tree line-i/o databases parameters string-port
                pretty-print common-list-functions)

(feature->requires 'batch provided? '((batch . "batch")))
        ⇒ (tree line-i/o databases parameters string-port
                pretty-print common-list-functions)
```

feature->requires* *feature provided? catalog*   [Function]
Returns a list of the features transitively **require**d by *feature* assuming the predicate
provided? and association-list *catalog*.

file->requires* *file provided? catalog*   [Function]
Returns a list of the features transitively **require**d by *file* assuming the predicate
provided? and association-list *catalog*.

file->loads *file*   [Function]
Returns a list of strings naming existing files loaded (load slib:load slib:load-source
macro:load defmacro:load syncase:load synclo:load macwork:load) by *file* or any of
the files it loads.

```
(file->loads (in-vicinity (library-vicinity) "scainit.scm"))
        ⇒ ("/usr/local/lib/slib/scaexpp.scm"
            "/usr/local/lib/slib/scaglob.scm"
            "/usr/local/lib/slib/scaoutp.scm")
```

load->path *exp*   [Function]
Given a (`load '<expr>`), where `<expr>` is a string or vicinity stuff), (`load->path`
`<expr>`) figures a path to the file. `load->path` returns that path if it names an
existing file; otherwise #f.

```
(load->path '(in-vicinity (library-vicinity) "mklibcat"))
        ⇒ "/usr/local/lib/slib/mklibcat.scm"
```

file->definitions *file definer . . .*   [Function]
> Returns a list of the identifier symbols defined by SLIB (or SLIB-style) file *file*. The optional arguments *definers* should be symbols signifying a defining form. If none are supplied, then the symbols `define-operation`, `define`, `define-syntax`, and `defmacro` are captured.

```
(file->definitions "random.scm")
        ⇒ (*random-state* make-random-state
           seed->random-state copy-random-state random
           random:chunk)
```

file->exports *file definer . . .*   [Function]
> Returns a list of the identifier symbols exported (advertised) by SLIB (or SLIB-style) file *file*. The optional arguments *definers* should be symbols signifying a defining form. If none are supplied, then the symbols `define-operation`, `define`, `define-syntax`, and `defmacro` are captured.

```
(file->exports "random.scm")
        ⇒ (make-random-state seed->random-state
            copy-random-state random)

(file->exports "randinex.scm")
        ⇒ (random:solid-sphere! random:hollow-sphere!
            random:normal-vector! random:normal
            random:exp random:uniform)
```

feature->export-alist *feature catalog*   [Function]
> Returns a list of lists; each sublist holding the name of the file implementing *feature*, and the identifier symbols exported (advertised) by SLIB (or SLIB-style) feature *feature*, in *catalog*.

feature->exports *feature catalog*   [Function]
> Returns a list of all exports of *feature*.

In the case of `aggregate` features, more than one file may have export lists to report:

```
(feature->export-alist 'r5rs slib:catalog))
        ⇒ (("/usr/local/lib/slib/values.scm"
             call-with-values values)
           ("/usr/local/lib/slib/mbe.scm"
             define-syntax macro:expand
             macro:load macro:eval)
           ("/usr/local/lib/slib/eval.scm"
             eval scheme-report-environment
             null-environment interaction-environment))

(feature->export-alist 'stdio *catalog*)
```

```
        ⇒ (("/usr/local/lib/slib/scanf.scm"
            fscanf sscanf scanf scanf-read-list)
           ("/usr/local/lib/slib/printf.scm"
            sprintf printf fprintf)
           ("/usr/local/lib/slib/stdio.scm"
            stderr stdout stdin))

   (feature->exports 'stdio slib:catalog)
        ⇒ (fscanf sscanf scanf scanf-read-list
             sprintf printf fprintf stderr stdout stdin)
```

1.6.3 Module Semantics

For the purpose of compiling Scheme code, each top-level `require` makes the identifiers exported by its feature's module `defined` (or defmacroed or defined-syntaxed) within the file (being compiled) headed with those requires.

Top-level occurrences of `require-if` make defined the exports from the module named by the second argument *if* the *feature-expression* first argument is true in the target environment. The target feature `compiling` should be provided during this phase of compilation.

Non-top-level SLIB occurences of `require` and `require-if` of quoted features can be ignored by compilers. The SLIB modules will all have top-level constructs for those features.

Note that aggregate catalog entries import more than one module. Implementations of `require` may or may *not* be transitive; code which uses module exports without requiring the providing module is in error.

In the SLIB modules `modular`, `batch`, `hash`, `common-lisp-time`, `commutative-ring`, `charplot`, `logical`, `common-list-functions`, `coerce` and `break` there is code conditional on features being `provided?`. Most are testing for the presence of features which are intrinsic to implementations (inexacts, bignums, ...).

In all cases these `provided?` tests can be evaluated at compile-time using `feature-eval` (see Section 1.1 [Feature], page 1). The simplest way to compile these constructs may be to treat `provided?` as a macro.

1.6.4 Top-level Variable References

`(require 'top-refs)`

These procedures complement those in Section 1.6.2 [Module Manifests], page 6, by finding the top-level variable references in Scheme source code. They work by traversing expressions and definitions, keeping track of bindings encountered. It is certainly possible to foil these functions, but they return useful information about SLIB source code.

top-refs *obj*   [Function]
    Returns a list of the top-level variables referenced by the Scheme expression *obj*.

top-refs<-file *filename*   [Function]
    filename should be a string naming an existing file containing Scheme source code.
    `top-refs<-file` returns a list of the top-level variable references made by expressions in the file named by *filename*.

Code in modules which *filename* `require`s is not traversed. Code in files loaded from top-level *is* traversed if the expression argument to `load`, `slib:load`, `slib:load-source`, `macro:load`, `defmacro:load`, `synclo:load`, `syncase:load`, or `macwork:load` is a literal string constant or composed of combinations of vicinity functions and string literal constants; and the resulting file exists (possibly with ".scm" appended).

The following function parses an *Info* Index. [3]

`exports<-info-index` *file n ...*   [Function]
> *n ...* must be an increasing series of positive integers. `exports<-info-index` returns a list of all the identifiers appearing in the *n*th ... (info) indexes of *file*. The identifiers have the case that the implementation's `read` uses for symbols. Identifiers containing spaces (eg. `close-base on base-table`) are *not* included. #f is returned if the index is not found.

> Each info index is headed by a '`* Menu:`' line. To list the symbols in the first and third info indexes do:

>   (exports<-info-index "slib.info" 1 3)

1.6.5 Module Analysis

`(require 'vet)`

`vet-slib` *file1 ...*   [Function]
> Using the procedures in the `top-refs` and `manifest` modules, `vet-slib` analyzes each SLIB module and *file1*, ..., reporting about any procedure or macro defined whether it is:

> orphaned   defined, not called, not exported;

> missing   called, not defined, and not exported by its `require`d modules;

> undocumented-export
>   Exported by module, but no index entry in `slib.info`;

> And for the library as a whole:

> documented-unexport
>   Index entry in `slib.info`, but no module exports it.

> This straightforward analysis caught three full days worth of never-executed branches, transitive require assumptions, spelling errors, undocumented procedures, missing procedures, and cyclic dependencies in SLIB.

> The optional arguments *file1*, ... provide a simple way to vet prospective SLIB modules.

---

[3] Although it will work on large info files, feeding it an excerpt is much faster; and has less chance of being confused by unusual text in the info file. This command excerpts the SLIB index into `slib-index.info`:

`info -f slib2d6.info -n "Index" -o slib-index.info`

2 Universal SLIB Procedures

The procedures described in these sections are supported by all implementations as part of the '`*.init`' files or by `require.scm`.

2.1 Vicinity

A vicinity is a descriptor for a place in the file system. Vicinities hide from the programmer the concepts of host, volume, directory, and version. Vicinities express only the concept of a file environment where a file name can be resolved to a file in a system independent manner. Vicinities can even be used on *flat* file systems (which have no directory structure) by having the vicinity express constraints on the file name.

All of these procedures are file-system dependent. Use of these vicinity procedures can make programs file-system *in*dependent.

These procedures are provided by all implementations. On most systems a vicinity is a string.

`make-vicinity` *dirpath* [Function]

    Returns *dirpath* as a vicinity for use as first argument to `in-vicinity`.

`pathname->vicinity` *path* [Function]

    Returns the vicinity containing *path*.

```
(pathname->vicinity "/usr/local/lib/scm/Link.scm")
                ⇒ "/usr/local/lib/scm/"
```

`program-vicinity` [Function]

    Returns the vicinity of the currently loading Scheme code. For an interpreter this would be the directory containing source code. For a compiled system (with multiple files) this would be the directory where the object or executable files are. If no file is currently loading, then the result is undefined. **Warning:** `program-vicinity` can return incorrect values if your program escapes back into a `load` continuation.

`library-vicinity` [Function]

    Returns the vicinity of the shared Scheme library.

`implementation-vicinity` [Function]

    Returns the vicinity of the underlying Scheme implementation. This vicinity will likely contain startup code and messages and a compiler.

`user-vicinity` [Function]

    Returns the vicinity of the current directory of the user. On most systems this is `""` (the empty string).

`home-vicinity` [Function]

    Returns the vicinity of the user's *HOME* directory, the directory which typically contains files which customize a computer environment for a user. If scheme is running without a user (eg. a daemon) or if this concept is meaningless for the platform, then `home-vicinity` returns `#f`.

`vicinity:suffix?` *chr*   [Function]
>   Returns the '`#t`' if *chr* is a vicinity suffix character; and `#f` otherwise. Typical vicinity
>   suffixes are '`/`', '`:`', and '`\`',

`in-vicinity` *vicinity filename*   [Function]
>   Returns a filename suitable for use by `slib:load`, `slib:load-source`,
>   `slib:load-compiled`, `open-input-file`, `open-output-file`, etc. The returned
>   filename is *filename* in *vicinity*. `in-vicinity` should allow *filename* to override
>   *vicinity* when *filename* is an absolute pathname and *vicinity* is equal to the value
>   of (`user-vicinity`). The behavior of `in-vicinity` when *filename* is absolute and
>   *vicinity* is not equal to the value of (`user-vicinity`) is unspecified. For most
>   systems `in-vicinity` can be `string-append`.

`sub-vicinity` *vicinity name*   [Function]
>   Returns the vicinity of *vicinity* restricted to *name*. This is used for large systems
>   where names of files in subsystems could conflict. On systems with directory structure
>   `sub-vicinity` will return a pathname of the subdirectory *name* of *vicinity*.

`with-load-pathname` *path thunk*   [Function]
>   *path* should be a string naming a file being read or loaded. `with-load-pathname`
>   evaluates *thunk* in a dynamic scope where an internal variable is bound to *path*; the
>   internal variable is used for messages and `program-vicinity`. `with-load-pathname`
>   returns the value returned by *thunk*.

2.2 Configuration

These constants and procedures describe characteristics of the Scheme and underlying operating system. They are provided by all implementations.

`char-code-limit`   [Constant]
>   An integer 1 larger that the largest value which can be returned by `char->integer`.

`most-positive-fixnum`   [Constant]
>   In implementations which support integers of practically unlimited size, *most-positive-fixnum* is a large exact integer within the range of exact integers that may result from computing the length of a list, vector, or string.
>
>   In implementations which do not support integers of practically unlimited size, *most-positive-fixnum* is the largest exact integer that may result from computing the length of a list, vector, or string.

`slib:tab`   [Constant]
>   The tab character.

`slib:form-feed`   [Constant]
>   The form-feed character.

`software-type`   [Function]
>   Returns a symbol denoting the generic operating system type. For instance, `unix`,
>   `vms`, `macos`, `amiga`, or `ms-dos`.

slib:report-version   [Function]
>   Displays the versions of SLIB and the underlying Scheme implementation and the
>   name of the operating system. An unspecified value is returned.

>   (slib:report-version) ⇒ slib "3b6" on scm "5b1" on unix

slib:report   [Function]
>   Displays the information of (slib:report-version) followed by almost all the infor-
>   mation neccessary for submitting a problem report. An unspecified value is returned.

slib:report #t   [Function]
>   provides a more verbose listing.

slib:report filename   [Function]
>   Writes the verbose report to file filename.

```
(slib:report)
⇒
slib "3b6" on scm "5b1" on unix
(implementation-vicinity) is "/usr/local/lib/scm/"
(library-vicinity) is "/usr/local/lib/slib/"
(scheme-file-suffix) is ".scm"
loaded slib:features :
        trace alist qp sort
        common-list-functions macro values getopt
        compiled
implementation slib:features :
        bignum complex real rational
        inexact vicinity ed getenv
        tmpnam abort transcript with-file
        ieee-p1178 r4rs rev4-optional-procedures hash
        object-hash delay eval dynamic-wind
        multiarg-apply multiarg/and- logical defmacro
        string-port source current-time record
        rev3-procedures rev2-procedures sun-dl string-case
        array dump char-ready? full-continuation
        system
implementation *catalog* :
        (i/o-extensions compiled "/usr/local/lib/scm/ioext.so")
        ...
```

2.3 Input/Output

These procedures are provided by all implementations.

file-exists? filename   [Function]
>   Returns #t if the specified file exists. Otherwise, returns #f. If the underlying imple-
>   mentation does not support this feature then #f is always returned.

delete-file *filename*   [Function]

Deletes the file specified by *filename*. If *filename* can not be deleted, `#f` is returned. Otherwise, `#t` is returned.

open-file *filename modes*   [Function]

filename should be a string naming a file. `open-file` returns a port depending on the symbol *modes*:

r   an input port capable of delivering characters from the file.

rb   a *binary* input port capable of delivering characters from the file.

w   an output port capable of writing characters to a new file by that name.

wb   a *binary* output port capable of writing characters to a new file by that name.

If an implementation does not distinguish between binary and non-binary files, then it must treat rb as r and wb as w.

If the file cannot be opened, either #f is returned or an error is signalled. For output, if a file with the given name already exists, the effect is unspecified.

port? *obj*   [Function]

Returns `#t` if *obj* is an input or output port, otherwise returns `#f`.

close-port *port*   [Procedure]

Closes the file associated with *port*, rendering the *port* incapable of delivering or accepting characters.

`close-file` has no effect if the file has already been closed. The value returned is unspecified.

call-with-open-ports *proc ports . . .*   [Function]
call-with-open-ports *ports . . . proc*   [Function]

Proc should be a procedure that accepts as many arguments as there are *ports* passed to `call-with-open-ports`. `call-with-open-ports` calls *proc* with *ports* .... If *proc* returns, then the ports are closed automatically and the value yielded by the *proc* is returned. If *proc* does not return, then the ports will not be closed automatically unless it is possible to prove that the ports will never again be used for a read or write operation.

tmpnam   [Function]

Returns a pathname for a file which will likely not be used by any other process. Successive calls to (`tmpnam`) will return different pathnames.

current-error-port   [Function]

Returns the current port to which diagnostic and error output is directed.

force-output   [Procedure]
force-output *port*   [Procedure]

Forces any pending output on *port* to be delivered to the output device and returns an unspecified value. The *port* argument may be omitted, in which case it defaults to the value returned by (`current-output-port`).

`file-position` *port*   [Function]
`file-position` *port* #f   [Function]
>   *port* must be open to a file. `file-position` returns the current position of the
>   character in *port* which will next be read or written. If the implementation does not
>   support file-position, then `#f` is returned.

`file-position` *port k*   [Function]
>   *port* must be open to a file. `file-position` sets the current position in *port* which
>   will next be read or written. If successful, `#t` is returned; otherwise `file-position`
>   returns `#f`.

`output-port-width`   [Function]
`output-port-width` *port*   [Function]
>   Returns the width of *port*, which defaults to (`current-output-port`) if absent. If
>   the width cannot be determined 79 is returned.

`output-port-height`   [Function]
`output-port-height` *port*   [Function]
>   Returns the height of *port*, which defaults to (`current-output-port`) if absent. If
>   the height cannot be determined 24 is returned.

2.4 System

These procedures are provided by all implementations.

`slib:load-source` *name*   [Procedure]
>   Loads a file of Scheme source code from *name* with the default filename extension
>   used in SLIB. For instance if the filename extension used in SLIB is `.scm` then
>   (`slib:load-source "foo"`) will load from file `foo.scm`.

`slib:load-compiled` *name*   [Procedure]
>   On implementations which support separtely loadable compiled modules, loads a
>   file of compiled code from *name* with the implementation's filename extension for
>   compiled code appended.

`slib:load` *name*   [Procedure]
>   Loads a file of Scheme source or compiled code from *name* with the appropriate
>   suffixes appended. If both source and compiled code are present with the appropriate
>   names then the implementation will load just one. It is up to the implementation to
>   choose which one will be loaded.
>
>   If an implementation does not support compiled code then `slib:load` will be identical
>   to `slib:load-source`.

`slib:eval` *obj*   [Procedure]
>   `eval` returns the value of *obj* evaluated in the current top level environment.
>   Section 7.4.11 [Eval], page 251, provides a more general evaluation facility.

`slib:eval-load` *filename eval*   [Procedure]
>   *filename* should be a string. If filename names an existing file, the Scheme source
>   code expressions and definitions are read from the file and *eval* called with them

sequentially. The `slib:eval-load` procedure does not affect the values returned by
`current-input-port`, `current-error-port`, and `current-output-port`.

`slib:warn` *arg1 arg2 . . .*   [Procedure]
> Outputs a warning message containing the arguments.

`slib:error` *arg1 arg2 . . .*   [Procedure]
> Outputs an error message containing the arguments, aborts evaluation of the current
> form and responds in a system dependent way to the error. Typical responses are to
> abort the program or to enter a read-eval-print loop.

`slib:exit` *n*   [Procedure]
`slib:exit`   [Procedure]
> Exits from the Scheme session returning status *n* to the system. If *n* is omitted or `#t`,
> a success status is returned to the system (if possible). If *n* is `#f` a failure is returned
> to the system (if possible). If *n* is an integer, then *n* is returned to the system (if
> possible). If the Scheme session cannot exit, then an unspecified value is returned
> from `slib:exit`.

`browse-url` *url*   [Function]
> Web browsers have become so ubiquitous that programming languagues should sup-
> port a uniform interface to them.

> If a browser is running, `browse-url` causes the browser to display the page specified
> by string *url* and returns `#t`.

> If the browser is not running, `browse-url` starts a browser displaying the argument
> *url*. If the browser starts as a background job, `browse-url` returns `#t` immediately; if
> the browser starts as a foreground job, then `browse-url` returns `#t` when the browser
> exits; otherwise (if no browser) it returns `#f`.

2.5 Miscellany

These procedures are provided by all implementations.

`identity` *x*   [Function]
> *identity* returns its argument.

> Example:

```
(identity 3)
    ⇒ 3
(identity '(foo bar))
    ⇒ (foo bar)
(map identity lst)
    ≡ (copy-list lst)
```

2.5.1 Mutual Exclusion

An *exchanger* is a procedure of one argument regulating mutually exclusive access to a
resource. When a exchanger is called, its current content is returned, while being replaced
by its argument in an atomic operation.

make-exchanger *obj*   [Function]
>   Returns a new exchanger with the argument *obj* as its initial content.

```
(define queue (make-exchanger (list a)))
```

A queue implemented as an exchanger holding a list can be protected from reentrant execution thus:

```
(define (pop queue)
  (let ((lst #f))
    (dynamic-wind
        (lambda () (set! lst (queue #f)))
        (lambda () (and lst (not (null? lst))
                        (let ((ret (car lst)))
                          (set! lst (cdr lst))
                          ret)))
        (lambda () (and lst (queue lst))))))
```

```
(pop queue)            ⇒ a
```

```
(pop queue)            ⇒ #f
```

2.5.2 Legacy

The following procedures were present in Scheme until R4RS (see Section "Language changes " in *Revised(4) Scheme*). They are provided by all SLIB implementations.

t   [Constant]
>   Defined as #t.

nil   [Constant]
>   Defined as #f.

last-pair *l*   [Function]
>   Returns the last pair in the list *l*. Example:

```
(last-pair (cons 1 2))
   ⇒ (1 . 2)
(last-pair '(1 2))
   ⇒ (2)
     ≡ (cons 2 '())
```

3 Scheme Syntax Extension Packages

3.1 Defmacro

Defmacros are supported by all implementations.

gentemp [Function]

> Returns a new (interned) symbol each time it is called. The symbol names are implementation-dependent
>
>> (gentemp) $\Rightarrow$ scm:G0
>> (gentemp) $\Rightarrow$ scm:G1

defmacro:eval *e* [Function]

> Returns the `slib:eval` of expanding all defmacros in scheme expression *e*.

defmacro:load *filename* [Function]

> *filename* should be a string. If filename names an existing file, the `defmacro:load` procedure reads Scheme source code expressions and definitions from the file and evaluates them sequentially. These source code expressions and definitions may contain defmacro definitions. The `defmacro:load` procedure does not affect the values returned by `current-input-port`, `current-error-port`, and `current-output-port`.

defmacro? *sym* [Function]

> Returns `#t` if *sym* has been defined by `defmacro`, `#f` otherwise.

macroexpand-1 *form* [Function]
macroexpand *form* [Function]

> If *form* is a macro call, `macroexpand-1` will expand the macro call once and return it. A *form* is considered to be a macro call only if it is a cons whose `car` is a symbol for which a `defmacro` has been defined.
>
> `macroexpand` is similar to `macroexpand-1`, but repeatedly expands *form* until it is no longer a macro call.

defmacro *name lambda-list form . . .* [Macro]

> When encountered by `defmacro:eval`, `defmacro:macroexpand*`, or `defmacro:load` defines a new macro which will henceforth be expanded when encountered by `defmacro:eval`, `defmacro:macroexpand*`, or `defmacro:load`.

3.1.1 Defmacroexpand

(require 'defmacroexpand)

defmacro:expand* *e* [Function]

> Returns the result of expanding all defmacros in scheme expression *e*.

3.2 R4RS Macros

(`require 'macro`) is the appropriate call if you want R4RS high-level macros but don't care about the low level implementation. If an SLIB R4RS macro implementation is already loaded it will be used. Otherwise, one of the R4RS macros implemetations is loaded.

The SLIB R4RS macro implementations support the following uniform interface:

macro:expand *sexpression*   [Function]
> Takes an R4RS expression, macro-expands it, and returns the result of the macro expansion.

macro:eval *sexpression*   [Function]
> Takes an R4RS expression, macro-expands it, evals the result of the macro expansion, and returns the result of the evaluation.

macro:load *filename*   [Procedure]
> *filename* should be a string. If filename names an existing file, the `macro:load` procedure reads Scheme source code expressions and definitions from the file and evaluates them sequentially. These source code expressions and definitions may contain macro definitions. The `macro:load` procedure does not affect the values returned by `current-input-port`, `current-error-port`, and `current-output-port`.

3.3 Macro by Example

(`require 'macro-by-example`)

A vanilla implementation of *Macro by Example* (Eugene Kohlbecker, R4RS) by Dorai Sitaram, (dorai @ cs.rice.edu) using `defmacro`.

- generating hygienic global `define-syntax` Macro-by-Example macros **cheaply**.
- can define macros which use . . . .
- needn't worry about a lexical variable in a macro definition clashing with a variable from the macro use context
- don't suffer the overhead of redefining the repl if `defmacro` natively supported (most implementations)

3.3.1 Caveat

These macros are not referentially transparent (see Section "Macros" in *Revised(4) Scheme*). Lexically scoped macros (i.e., `let-syntax` and `letrec-syntax`) are not supported. In any case, the problem of referential transparency gains poignancy only when `let-syntax` and `letrec-syntax` are used. So you will not be courting large-scale disaster unless you're using system-function names as local variables with unintuitive bindings that the macro can't use. However, if you must have the full *r4rs* macro functionality, look to the more featureful (but also more expensive) versions of syntax-rules available in slib Section 3.4 [Macros That Work], page 20, Section 3.5 [Syntactic Closures], page 23, and Section 3.6 [Syntax-Case Macros], page 30.

define-syntax *keyword transformer-spec*   [Macro]
> The *keyword* is an identifier, and the *transformer-spec* should be an instance of `syntax-rules`.

The top-level syntactic environment is extended by binding the *keyword* to the specified transformer.

```
(define-syntax let*
  (syntax-rules ()
    ((let* () body1 body2 ...)
     (let () body1 body2 ...))
    ((let* ((name1 val1) (name2 val2) ...)
       body1 body2 ...)
     (let ((name1 val1))
       (let* (( name2 val2) ...)
         body1 body2 ...)))))
```

syntax-rules *literals syntax-rule . . .*   [Macro]

literals is a list of identifiers, and each *syntax-rule* should be of the form

(`pattern` `template`)

where the *pattern* and *template* are as in the grammar above.

An instance of **syntax-rules** produces a new macro transformer by specifying a sequence of hygienic rewrite rules. A use of a macro whose keyword is associated with a transformer specified by **syntax-rules** is matched against the patterns contained in the *syntax-rule*s, beginning with the leftmost *syntax-rule*. When a match is found, the macro use is trancribed hygienically according to the template.

Each pattern begins with the keyword for the macro. This keyword is not involved in the matching and is not considered a pattern variable or literal identifier.

3.4 Macros That Work

(require 'macros-that-work)

Macros That Work differs from the other R4RS macro implementations in that it does not expand derived expression types to primitive expression types.

macro:expand *expression*   [Function]
macwork:expand *expression*   [Function]

Takes an R4RS expression, macro-expands it, and returns the result of the macro expansion.

macro:eval *expression*   [Function]
macwork:eval *expression*   [Function]

macro:eval returns the value of *expression* in the current top level environment. *expression* can contain macro definitions. Side effects of *expression* will affect the top level environment.

macro:load *filename*   [Procedure]
macwork:load *filename*   [Procedure]

filename should be a string. If filename names an existing file, the `macro:load` procedure reads Scheme source code expressions and definitions from the file and evaluates them sequentially. These source code expressions and definitions may contain

macro definitions. The `macro:load` procedure does not affect the values returned by `current-input-port`, `current-error-port`, and `current-output-port`.

References:

The *Revised^4 Report on the Algorithmic Language Scheme* Clinger and Rees [editors]. To appear in LISP Pointers. Also available as a technical report from the University of Oregon, MIT AI Lab, and Cornell.
                    Macros That Work. Clinger and Rees. POPL '91.

The supported syntax differs from the R4RS in that vectors are allowed as patterns and as templates and are not allowed as pattern or template data.

```
transformer spec  ↦  (syntax-rules literals rules)


rules  ↦  ()
         | (rule . rules)


rule  ↦  (pattern template)


pattern  ↦  pattern_var      ; a symbol not in literals
           | symbol          ; a symbol in literals
           | ()
           | (pattern . pattern)
           | (ellipsis_pattern)
           | #(pattern*)                      ; extends R4RS
           | #(pattern* ellipsis_pattern)    ; extends R4RS
           | pattern_datum


template  ↦  pattern_var
           | symbol
           | ()
           | (template2 . template2)
           | #(template*)                     ; extends R4RS
           | pattern_datum


template2  ↦  template
           | ellipsis_template


pattern_datum  ↦  string                      ; no vector
               | character
               | boolean
               | number


ellipsis_pattern  ↦ pattern ...


ellipsis_template  ↦  template ...


pattern_var  ↦  symbol   ; not in literals
```

```
literals  ↦  ()
            |  (symbol . literals)
```

3.4.1 Definitions

Scope of an ellipsis
> Within a pattern or template, the scope of an ellipsis (. . .) is the pattern or template that appears to its left.

Rank of a pattern variable
> The rank of a pattern variable is the number of ellipses within whose scope it appears in the pattern.

Rank of a subtemplate
> The rank of a subtemplate is the number of ellipses within whose scope it appears in the template.

Template rank of an occurrence of a pattern variable
> The template rank of an occurrence of a pattern variable within a template is the rank of that occurrence, viewed as a subtemplate.

Variables bound by a pattern
> The variables bound by a pattern are the pattern variables that appear within it.

Referenced variables of a subtemplate
> The referenced variables of a subtemplate are the pattern variables that appear within it.

Variables opened by an ellipsis template
> The variables opened by an ellipsis template are the referenced pattern variables whose rank is greater than the rank of the ellipsis template.

3.4.2 Restrictions

No pattern variable appears more than once within a pattern.

For every occurrence of a pattern variable within a template, the template rank of the occurrence must be greater than or equal to the pattern variable's rank.

Every ellipsis template must open at least one variable.

For every ellipsis template, the variables opened by an ellipsis template must all be bound to sequences of the same length.

The compiled form of a *rule* is

```
rule  ↦  (pattern template inserted)


pattern  ↦  pattern_var
          |  symbol
          |  ()
          |  (pattern . pattern)
          |  ellipsis_pattern
```

```
                    |  #(pattern)
                    |  pattern_datum

  template  ↦  pattern_var
              |  symbol
              |  ()
              |  (template2 . template2)
              |  #(pattern)
              |  pattern_datum

  template2  ↦  template
             |  ellipsis_template

  pattern_datum  ↦  string
                 |  character
                 |  boolean
                 |  number

  pattern_var  ↦  #(V symbol rank)

  ellipsis_pattern  ↦  #(E pattern pattern_vars)

  ellipsis_template  ↦  #(E template pattern_vars)

  inserted  ↦  ()
            |  (symbol . inserted)

  pattern_vars  ↦  ()
                |  (pattern_var . pattern_vars)

  rank  ↦  exact non-negative integer
```

where V and E are unforgeable values.

The pattern variables associated with an ellipsis pattern are the variables bound by
the pattern, and the pattern variables associated with an ellipsis template are the variables
opened by the ellipsis template.

If the template contains a big chunk that contains no pattern variables or inserted
identifiers, then the big chunk will be copied unnecessarily. That shouldn't matter very
often.

3.5 Syntactic Closures

```
(require 'syntactic-closures)
```

macro:expand *expression*   [Function]
synclo:expand *expression*   [Function]
> Returns scheme code with the macros and derived expression types of *expression* expanded to primitive expression types.

macro:eval *expression*   [Function]
synclo:eval *expression*   [Function]
> `macro:eval` returns the value of *expression* in the current top level environment. *expression* can contain macro definitions. Side effects of *expression* will affect the top level environment.

macro:load *filename*   [Procedure]
synclo:load *filename*   [Procedure]
> *filename* should be a string. If filename names an existing file, the `macro:load` procedure reads Scheme source code expressions and definitions from the file and evaluates them sequentially. These source code expressions and definitions may contain macro definitions. The `macro:load` procedure does not affect the values returned by `current-input-port`, `current-error-port`, and `current-output-port`.

3.5.1 Syntactic Closure Macro Facility

<div align="center">
A Syntactic Closures Macro Facility
by Chris Hanson
9 November 1991
</div>

This document describes *syntactic closures*, a low-level macro facility for the Scheme programming language. The facility is an alternative to the low-level macro facility described in the *Revised^4 Report on Scheme.* This document is an addendum to that report.

The syntactic closures facility extends the BNF rule for *transformer spec* to allow a new keyword that introduces a low-level macro transformer:

> *transformer spec* := (transformer *expression*)

Additionally, the following procedures are added:

```
make-syntactic-closure
capture-syntactic-environment
identifier?
identifier=?
```

The description of the facility is divided into three parts. The first part defines basic terminology. The second part describes how macro transformers are defined. The third part describes the use of *identifiers*, which extend the syntactic closure mechanism to be compatible with `syntax-rules`.

3.5.1.1 Terminology

This section defines the concepts and data types used by the syntactic closures facility.

- *Forms* are the syntactic entities out of which programs are recursively constructed. A form is any expression, any definition, any syntactic keyword, or any syntactic closure. The variable name that appears in a `set!` special form is also a form. Examples of forms:

```
17
#t
car
(+ x 4)
(lambda (x) x)
(define pi 3.14159)
if
define
```

- An *alias* is an alternate name for a given symbol. It can appear anywhere in a form that the symbol could be used, and when quoted it is replaced by the symbol; however, it does not satisfy the predicate `symbol?`. Macro transformers rarely distinguish symbols from aliases, referring to both as identifiers.

- A *syntactic* environment maps identifiers to their meanings. More precisely, it determines whether an identifier is a syntactic keyword or a variable. If it is a keyword, the meaning is an interpretation for the form in which that keyword appears. If it is a variable, the meaning identifies which binding of that variable is referenced. In short, syntactic environments contain all of the contextual information necessary for interpreting the meaning of a particular form.

- A *syntactic closure* consists of a form, a syntactic environment, and a list of identifiers. All identifiers in the form take their meaning from the syntactic environment, except those in the given list. The identifiers in the list are to have their meanings determined later. A syntactic closure may be used in any context in which its form could have been used. Since a syntactic closure is also a form, it may not be used in contexts where a form would be illegal. For example, a form may not appear as a clause in the cond special form. A syntactic closure appearing in a quoted structure is replaced by its form.

3.5.1.2 Transformer Definition

This section describes the `transformer` special form and the procedures `make-syntactic-closure` and `capture-syntactic-environment`.

transformer *expression*   [Syntax]

    Syntax: It is an error if this syntax occurs except as a *transformer spec*.

    Semantics: The *expression* is evaluated in the standard transformer environment to yield a macro transformer as described below. This macro transformer is bound to a macro keyword by the special form in which the `transformer` expression appears (for example, `let-syntax`).

    A *macro transformer* is a procedure that takes two arguments, a form and a syntactic environment, and returns a new form. The first argument, the *input form*, is the form in which the macro keyword occurred. The second argument, the *usage environment*, is the syntactic environment in which the input form occurred. The result of the transformer, the *output form*, is automatically closed in the *transformer environment*, which is the syntactic environment in which the `transformer` expression occurred.

    For example, here is a definition of a push macro using `syntax-rules`:

```
(define-syntax  push
```

```
        (syntax-rules ()
          ((push item list)
           (set! list (cons item list)))))
```

Here is an equivalent definition using `transformer`:

```
      (define-syntax push
        (transformer
         (lambda (exp env)
           (let ((item
                   (make-syntactic-closure env '() (cadr exp)))
                 (list
                   (make-syntactic-closure env '() (caddr exp))))
             '(set! ,list (cons ,item ,list))))))
```

In this example, the identifiers `set!` and `cons` are closed in the transformer environment, and thus will not be affected by the meanings of those identifiers in the usage environment `env`.

Some macros may be non-hygienic by design. For example, the following defines a loop macro that implicitly binds `exit` to an escape procedure. The binding of `exit` is intended to capture free references to `exit` in the body of the loop, so `exit` must be left free when the body is closed:

```
      (define-syntax loop
        (transformer
         (lambda (exp env)
           (let ((body (cdr exp)))
             '(call-with-current-continuation
               (lambda (exit)
                 (let f ()
                   ,@(map (lambda  (exp)
                            (make-syntactic-closure env '(exit)
                                                    exp))
                          body)
                   (f))))))))
```

To assign meanings to the identifiers in a form, use `make-syntactic-closure` to close the form in a syntactic environment.

`make-syntactic-closure` *environment free-names form* [Function]
 environment must be a syntactic environment, *free-names* must be a list of identifiers, and *form* must be a form. `make-syntactic-closure` constructs and returns a syntactic closure of *form* in *environment*, which can be used anywhere that *form* could have been used. All the identifiers used in *form*, except those explicitly excepted by *free-names*, obtain their meanings from *environment*.

Here is an example where *free-names* is something other than the empty list. It is instructive to compare the use of *free-names* in this example with its use in the `loop` example above: the examples are similar except for the source of the identifier being left free.

```
(define-syntax let1
  (transformer
   (lambda (exp env)
     (let ((id (cadr exp))
           (init (caddr exp))
           (exp (cadddr exp)))
       `((lambda (,id)
           ,(make-syntactic-closure env (list id) exp))
         ,(make-syntactic-closure env '() init))))))
```

`let1` is a simplified version of `let` that only binds a single identifier, and whose body consists of a single expression. When the body expression is syntactically closed in its original syntactic environment, the identifier that is to be bound by `let1` must be left free, so that it can be properly captured by the `lambda` in the output form.

To obtain a syntactic environment other than the usage environment, use `capture-syntactic-environment`.

`capture-syntactic-environment` *procedure*   [Function]

> `capture-syntactic-environment` returns a form that will, when transformed, call *procedure* on the current syntactic environment. *procedure* should compute and return a new form to be transformed, in that same syntactic environment, in place of the form.

> An example will make this clear. Suppose we wanted to define a simple `loop-until` keyword equivalent to

```
(define-syntax loop-until
  (syntax-rules ()
    ((loop-until id init test return step)
     (letrec ((loop
                (lambda (id)
                  (if test return (loop step)))))
       (loop init)))))
```

> The following attempt at defining `loop-until` has a subtle bug:

```
(define-syntax loop-until
  (transformer
   (lambda (exp env)
     (let ((id (cadr exp))
           (init (caddr exp))
           (test (cadddr exp))
           (return (cadddr (cdr exp)))
           (step (cadddr (cddr exp)))
           (close
            (lambda (exp free)
              (make-syntactic-closure env free exp))))
       `(letrec ((loop
                   (lambda (,id)
                     (if ,(close test (list id))
```

```
                                 ,(close return (list id))
                                 (loop ,(close step (list id)))))))))
                   (loop ,(close init '()))))))))
```

This definition appears to take all of the proper precautions to prevent unintended captures. It carefully closes the subexpressions in their original syntactic environment and it leaves the `id` identifier free in the `test`, `return`, and `step` expressions, so that it will be captured by the binding introduced by the `lambda` expression. Unfortunately it uses the identifiers `if` and `loop` within that `lambda` expression, so if the user of `loop-until` just happens to use, say, `if` for the identifier, it will be inadvertently captured.

The syntactic environment that `if` and `loop` want to be exposed to is the one just outside the `lambda` expression: before the user's identifier is added to the syntactic environment, but after the identifier loop has been added. `capture-syntactic-environment` captures exactly that environment as follows:

```
      (define-syntax loop-until
        (transformer
         (lambda (exp env)
           (let ((id (cadr exp))
                 (init (caddr exp))
                 (test (cadddr exp))
                 (return (cadddr (cdr exp)))
                 (step (cadddr (cddr exp)))
                 (close
                  (lambda (exp free)
                    (make-syntactic-closure env free exp))))
             `(letrec ((loop
                        ,(capture-syntactic-environment
                          (lambda (env)
                            `(lambda (,id)
                               (,(make-syntactic-closure env '() `if)
                                ,(close test (list id))
                                ,(close return (list id))
                                (,(make-syntactic-closure env '()
                                                          `loop)
                                 ,(close step (list id)))))))))
                   (loop ,(close init '()))))))))
```

In this case, having captured the desired syntactic environment, it is convenient to construct syntactic closures of the identifiers `if` and the `loop` and use them in the body of the `lambda`.

A common use of `capture-syntactic-environment` is to get the transformer environment of a macro transformer:

```
      (transformer
       (lambda (exp env)
         (capture-syntactic-environment
          (lambda (transformer-env)
```

```
               ...)))))
```

3.5.1.3 Identifiers

This section describes the procedures that create and manipulate identifiers. Previous syntactic closure proposals did not have an identifier data type – they just used symbols. The identifier data type extends the syntactic closures facility to be compatible with the high-level `syntax-rules` facility.

As discussed earlier, an identifier is either a symbol or an *alias*. An alias is implemented as a syntactic closure whose *form* is an identifier:

```
(make-syntactic-closure env '() 'a)
    ⇒ an alias
```

Aliases are implemented as syntactic closures because they behave just like syntactic closures most of the time. The difference is that an alias may be bound to a new value (for example by `lambda` or `let-syntax`); other syntactic closures may not be used this way. If an alias is bound, then within the scope of that binding it is looked up in the syntactic environment just like any other identifier.

Aliases are used in the implementation of the high-level facility `syntax-rules`. A macro transformer created by `syntax-rules` uses a template to generate its output form, substituting subforms of the input form into the template. In a syntactic closures implementation, all of the symbols in the template are replaced by aliases closed in the transformer environment, while the output form itself is closed in the usage environment. This guarantees that the macro transformation is hygienic, without requiring the transformer to know the syntactic roles of the substituted input subforms.

`identifier?` *object*   [Function]
    Returns `#t` if *object* is an identifier, otherwise returns `#f`. Examples:

```
        (identifier? 'a)
           ⇒ #t
        (identifier? (make-syntactic-closure env '() 'a))
           ⇒ #t
        (identifier? "a")
           ⇒ #f
        (identifier? #\a)
           ⇒ #f
        (identifier? 97)
           ⇒ #f
        (identifier? #f)
           ⇒ #f
        (identifier? '(a))
           ⇒ #f
        (identifier? '#(a))
           ⇒ #f
```

The predicate `eq?` is used to determine if two identifers are "the same". Thus `eq?` can be used to compare identifiers exactly as it would be used to compare symbols. Often, though, it is useful to know whether two identifiers "mean the same thing".

For example, the `cond` macro uses the symbol `else` to identify the final clause in the conditional. A macro transformer for `cond` cannot just look for the symbol `else`, because the `cond` form might be the output of another macro transformer that replaced the symbol `else` with an alias. Instead the transformer must look for an identifier that "means the same thing" in the usage environment as the symbol `else` means in the transformer environment.

`identifier=?` *environment1 identifier1 environment2 identifier2*   [Function]
   environment1 and *environment2* must be syntactic environments, and *identifier1* and *identifier2* must be identifiers. `identifier=?` returns `#t` if the meaning of *identifier1* in *environment1* is the same as that of *identifier2* in *environment2*, otherwise it returns `#f`. Examples:

```
(let-syntax
    ((foo
      (transformer
       (lambda (form env)
         (capture-syntactic-environment
          (lambda (transformer-env)
            (identifier=? transformer-env 'x env 'x)))))))
  (list (foo)
        (let ((x 3))
          (foo))))
 ⇒ (#t #f)

(let-syntax ((bar foo))
  (let-syntax
      ((foo
        (transformer
         (lambda (form env)
           (capture-syntactic-environment
            (lambda (transformer-env)
              (identifier=? transformer-env 'foo
                            env (cadr form))))))))
    (list (foo foo)
          (foobar))))
 ⇒ (#f #t)
```

3.5.1.4 Acknowledgements

The syntactic closures facility was invented by Alan Bawden and Jonathan Rees. The use of aliases to implement `syntax-rules` was invented by Alan Bawden (who prefers to call them *synthetic names*). Much of this proposal is derived from an earlier proposal by Alan Bawden.

3.6 Syntax-Case Macros

```
(require 'syntax-case)
```

`macro:expand` *expression*   [Function]
`syncase:expand` *expression*   [Function]
> Returns scheme code with the macros and derived expression types of *expression* expanded to primitive expression types.

`macro:eval` *expression*   [Function]
`syncase:eval` *expression*   [Function]
> `macro:eval` returns the value of *expression* in the current top level environment. *expression* can contain macro definitions. Side effects of *expression* will affect the top level environment.

`macro:load` *filename*   [Procedure]
`syncase:load` *filename*   [Procedure]
> *filename* should be a string. If filename names an existing file, the `macro:load` procedure reads Scheme source code expressions and definitions from the file and evaluates them sequentially. These source code expressions and definitions may contain macro definitions. The `macro:load` procedure does not affect the values returned by `current-input-port`, `current-error-port`, and `current-output-port`.

This is version 2.1 of `syntax-case`, the low-level macro facility proposed and implemented by Robert Hieb and R. Kent Dybvig.

This version is further adapted by Harald Hanche-Olsen <hanche @ imf.unit.no> to make it compatible with, and easily usable with, SLIB. Mainly, these adaptations consisted of:

- Removing white space from `expand.pp` to save space in the distribution. This file is not meant for human readers anyway...

- Removed a couple of Chez scheme dependencies.

- Renamed global variables used to minimize the possibility of name conflicts.

- Adding an SLIB-specific initialization file.

- Removing a couple extra files, most notably the documentation (but see below).

If you wish, you can see exactly what changes were done by reading the shell script in the file `syncase.sh`.

The two PostScript files were omitted in order to not burden the SLIB distribution with them. If you do intend to use `syntax-case`, however, you should get these files and print them out on a PostScript printer. They are available with the original `syntax-case` distribution by anonymous FTP in `cs.indiana.edu:/pub/scheme/syntax-case`.

In order to use syntax-case from an interactive top level, execute:

```
(require 'syntax-case)
(require 'repl)
(repl:top-level macro:eval)
```

See the section Repl (see Section 7.5.1 [Repl], page 258) for more information.

To check operation of syntax-case get `cs.indiana.edu:/pub/scheme/syntax-case`, and type

```
(require 'syntax-case)
```

```
(syncase:sanity-check)
```

Beware that `syntax-case` takes a long time to load – about 20s on a SPARCstation SLC (with SCM) and about 90s on a Macintosh SE/30 (with Gambit).

3.6.1 Notes

All R4RS syntactic forms are defined, including `delay`. Along with `delay` are simple definitions for `make-promise` (into which `delay` expressions expand) and `force`.

`syntax-rules` and `with-syntax` (described in *TR356*) are defined.

`syntax-case` is actually defined as a macro that expands into calls to the procedure `syntax-dispatch` and the core form `syntax-lambda`; do not redefine these names.

Several other top-level bindings not documented in TR356 are created:

- the "hooks" in `hooks.ss`
- the `build-` procedures in `output.ss`
- `expand-syntax` (the expander)

The syntax of define has been extended to allow (`define id`), which assigns *id* to some unspecified value.

We have attempted to maintain R4RS compatibility where possible. The incompatibilities should be confined to `hooks.ss`. Please let us know if there is some incompatibility that is not flagged as such.

Send bug reports, comments, suggestions, and questions to Kent Dybvig (dyb @ iu-vax.cs.indiana.edu).

3.7 Define-Structure

```
(require 'structure)
```

Included with the `syntax-case` files was `structure.scm` which defines a macro `define-structure`. Here is its documentation from Gambit-4.0:

define-structure (*name field*. . .)   [special form]
> Record data types similar to Pascal records and C `struct` types can be defined using the `define-structure` special form. The identifier *name* specifies the name of the new data type. The structure name is followed by *k* identifiers naming each field of the record. The `define-structure` expands into a set of definitions of the following procedures:
>
> - 'make-*name*' – A *k* argument procedure which constructs a new record from the value of its *k* fields.
> - '*name*?' – A procedure which tests if its single argument is of the given record type.
> - '*name-field*' – For each field, a procedure taking as its single argument a value of the given record type and returning the content of the corresponding field of the record.

- '*name-field*-`set!`' – For each field, a two argument procedure taking as its first argument a value of the given record type. The second argument gets assigned to the corresponding field of the record and the void object is returned.

Gambit record data types have a printed representation that includes the name of the type and the name and value of each field.

For example:

```
> (require 'syntax-case)
> (require 'repl)
> (repl:top-level macro:eval)
> (require 'structure)
> (define-structure (point x y color))
> (define p (make-point 3 5 'red))
> p
#<point #3 x: 3 y: 5 color: red>
> (point-x p)
3
> (point-color p)
red
> (point-color-set! p 'black)
> p
#<point #3 x: 3 y: 5 color: black>
```

3.8 Define-Record-Type

(require 'define-record-type) or (require 'srfi-9)

    http://srfi.schemers.org/srfi-9/srfi-9.html

define-record-type <*type-name*> (<*constructor-name*>   [Special Form]
        <*field-tag*> ...) <*predicate-name*> <*field-spec*> ...
    Where

            <field-spec> ≡ (<field-tag> <accessor-name>)
                          ≡ (<field-tag> <accessor-name> <modifier-name>)

    define-record-type is a syntax wrapper for the SLIB record module.

3.9 Fluid-Let

(require 'fluid-let)

    Note: fluid-let is not thread-safe. It is better to use Section 3.10 [Parameter Objects], page 34, (srfi-39) or Section 7.1.11 [Dynamic Data Type], page 213, both of which will be made thread-safe in the future.

fluid-let (*bindings* ...) *forms*...   [Syntax]
    (fluid-let ((*variable init*) ...)
        expression expression ...)

The *init*s are evaluated in the current environment (in some unspecified order), the current values of the *variable*s are saved, the results are assigned to the *variable*s, the *expression*s are evaluated sequentially in the current environment, the *variable*s are restored to their original values, and the value of the last *expression* is returned.

The syntax of this special form is similar to that of `let`, but `fluid-let` temporarily rebinds existing *variable*s. Unlike `let`, `fluid-let` creates no new bindings; instead it *assigns* the values of each *init* to the binding (determined by the rules of lexical scoping) of its corresponding *variable*.

3.10 Parameter Objects

```
(require 'srfi-39)
```
> http://srfi.schemers.org/srfi-39/srfi-39.html

3.11 Binding to multiple values

`(require 'receive)` or `(require 'srfi-8)`

receive *formals expression body . . .*   [Special Form]
> http://srfi.schemers.org/srfi-8/srfi-8.html

> `(require 'let-values)` or `(require 'srfi-11)`

let-values ((*formals expression*) ...) *body . . .*   [Special Form]
let-values* ((*formals expression*) ...) *body . . .*   [Special Form]
> http://srfi.schemers.org/srfi-11/srfi-11.html

3.12 Guarded LET* special form

`(require 'and-let*)` or `(require 'srfi-2)`

and-let* *claws body . . .*   [Macro]
> http://srfi.schemers.org/srfi-2/srfi-2.html

3.13 Guarded COND Clause

`(require 'guarded-cond-clause)` or `(require 'srfi-61)`
> http://srfi.schemers.org/srfi-61/srfi-61.html

cond *<clause1> <clause2> . . .*   [library syntax]
> *Syntax:* Each <clause> should be of the form

> `(<test> <expression1> ...)`

> where <test> is any expression. Alternatively, a <clause> may be of the form

> `(<test> => <expression>)`

> The <clause> production in the formal syntax of Scheme as written by R5RS in section 7.1.3 is extended with a new option:

> `<clause> => (<generator> <guard> => <receiver>)`

where <generator>, <guard>, & <receiver> are all <expression>s.

Clauses of this form have the following semantics: <generator> is evaluated. It may return arbitrarily many values. <Guard> is applied to an argument list containing the values in order that <generator> returned. If <guard> returns a true value for that argument list, <receiver> is applied with an equivalent argument list. If <guard> returns a false value, however, the clause is abandoned and the next one is tried.

The last <clause> may be an "else clause," which has the form

(else <expression1> <expression2> ...).

This `port->char-list` procedure accepts an input port and returns a list of all the characters it produces until the end.

```
(define (port->char-list port)
  (cond ((read-char port) char?
         => (lambda (c) (cons c (port->char-list port))))
        (else '())))

(call-with-input-string "foo" port->char-list)  ==>  (#\f #\o #\o)
```

3.14 Yasos

(require 'oop) or (require 'yasos)

'Yet Another Scheme Object System' is a simple object system for Scheme based on the paper by Norman Adams and Jonathan Rees: *Object Oriented Programming in Scheme*, Proceedings of the 1988 ACM Conference on LISP and Functional Programming, July 1988 [ACM #552880].

Another reference is:

Ken Dickey.
Scheming with Objects
AI Expert Volume 7, Number 10 (October 1992), pp. 24-33.
`ftp://ftp.cs.indiana.edu/pub/scheme-repository/doc/pubs/swob.txt`

3.14.1 Terms

Object   Any Scheme data object.

Instance   An instance of the OO system; an *object*.

Operation   A *method*.

Notes: The object system supports multiple inheritance. An instance can inherit from 0 or more ancestors. In the case of multiple inherited operations with the same identity, the operation used is that from the first ancestor which contains it (in

the ancestor `let`). An operation may be applied to any Scheme data object—
not just instances. As code which creates instances is just code, there are no
classes and no meta-*anything*. Method dispatch is by a procedure call a la
CLOS rather than by `send` syntax a la Smalltalk.

Disclaimer:
There are a number of optimizations which can be made. This implementation
is expository (although performance should be quite reasonable). See the L&FP
paper for some suggestions.

3.14.2 Interface

define-operation (*opname self arg* . . .) `default-body`   [Syntax]
Defines a default behavior for data objects which don't handle the operation *opname*.
The default behavior (for an empty *default-body*) is to generate an error.

define-predicate *opname?*   [Syntax]
Defines a predicate *opname?*, usually used for determining the *type* of an object,
such that (`opname? object`) returns `#t` if *object* has an operation *opname?* and `#f`
otherwise.

object ((`name self arg` . . .) `body`) . . .   [Syntax]
Returns an object (an instance of the object system) with operations. Invoking (`name
object arg` . . .) executes the *body* of the *object* with *self* bound to *object* and with
argument(s) *arg*. . . .

object-with-ancestors ((*ancestor1 init1*) . . .) *operation* . . .   [Syntax]
A `let`-like form of `object` for multiple inheritance. It returns an object inheriting the
behaviour of *ancestor1* etc. An operation will be invoked in an ancestor if the object
itself does not provide such a method. In the case of multiple inherited operations
with the same identity, the operation used is the one found in the first ancestor in
the ancestor list.

operate-as *component operation self arg* . . .   [Syntax]
Used in an operation definition (of *self*) to invoke the *operation* in an ancestor *com-
ponent* but maintain the object's identity. Also known as "send-to-super".

print *obj port*   [Procedure]
A default `print` operation is provided which is just (`format port obj`) (see
Section 4.2 [Format], page 46) for non-instances and prints *obj* preceded by
'`#<INSTANCE>`' for instances.

size *obj*   [Function]
The default method returns the number of elements in *obj* if it is a vector, string or
list, `2` for a pair, `1` for a character and by default id an error otherwise. Objects such
as collections (see Section 7.1.10 [Collections], page 210) may override the default in
an obvious way.

3.14.3 Setters

Setters implement *generalized locations* for objects associated with some sort of mutable state. A *getter* operation retrieves a value from a generalized location and the corresponding setter operation stores a value into the location. Only the getter is named – the setter is specified by a procedure call as below. (Dylan uses special syntax.) Typically, but not necessarily, getters are access operations to extract values from Yasos objects (see Section 3.14 [Yasos], page 35). Several setters are predefined, corresponding to getters `car`, `cdr`, `string-ref` and `vector-ref` e.g., `(setter car)` is equivalent to `set-car!`.

This implementation of setters is similar to that in Dylan(TM) (*Dylan: An object-oriented dynamic language*, Apple Computer Eastern Research and Technology). Common LISP provides similar facilities through `setf`.

setter *getter*   [Function]

> Returns the setter for the procedure *getter*. E.g., since `string-ref` is the getter corresponding to a setter which is actually `string-set!`:
>
> ```
> (define foo "foo")
> ((setter string-ref) foo 0 #\F) ; set element 0 of foo
> foo ⇒ "Foo"
> ```

set *place new-value*   [Syntax]

> If *place* is a variable name, `set` is equivalent to `set!`. Otherwise, *place* must have the form of a procedure call, where the procedure name refers to a getter and the call indicates an accessible generalized location, i.e., the call would return a value. The return value of `set` is usually unspecified unless used with a setter whose definition guarantees to return a useful value.
>
> ```
> (set (string-ref foo 2) #\O)  ; generalized location with getter
> foo ⇒ "FoO"
> (set foo "foo")               ; like set!
> foo ⇒ "foo"
> ```

add-setter *getter setter*   [Procedure]

> Add procedures *getter* and *setter* to the (inaccessible) list of valid setter/getter pairs. *setter* implements the store operation corresponding to the *getter* access operation for the relevant state. The return value is unspecified.

remove-setter-for *getter*   [Procedure]

> Removes the setter corresponding to the specified *getter* from the list of valid setters. The return value is unspecified.

define-access-operation *getter-name*   [Syntax]

> Shorthand for a Yasos `define-operation` defining an operation *getter-name* that objects may support to return the value of some mutable state. The default operation is to signal an error. The return value is unspecified.

3.14.4 Examples

```
;;; These definitions for PRINT and SIZE are
;;; already supplied by
```

```scheme
(require 'yasos)

(define-operation (print obj port)
  (format port
          (if (instance? obj) "#<instance>" "~s")
          obj))

(define-operation (size obj)
  (cond
   ((vector? obj) (vector-length obj))
   ((list?   obj) (length obj))
   ((pair?   obj) 2)
   ((string? obj) (string-length obj))
   ((char?   obj) 1)
   (else
    (slib:error "Operation not supported: size" obj))))

(define-predicate cell?)
(define-operation (fetch obj))
(define-operation (store! obj newValue))

(define (make-cell value)
  (object
   ((cell? self) #t)
   ((fetch self) value)
   ((store! self newValue)
    (set! value newValue)
    newValue)
   ((size self) 1)
   ((print self port)
    (format port "#<Cell: ~s>" (fetch self)))))

(define-operation (discard obj value)
  (format #t "Discarding ~s~%" value))

(define (make-filtered-cell value filter)
  (object-with-ancestors
   ((cell (make-cell value)))
   ((store! self newValue)
    (if (filter newValue)
        (store! cell newValue)
        (discard self newValue)))))

(define-predicate array?)
(define-operation (array-ref array index))
(define-operation (array-set! array index value))
```

```
(define (make-array num-slots)
  (let ((anArray (make-vector num-slots)))
    (object
      ((array? self) #t)
      ((size self) num-slots)
      ((array-ref self index)
       (vector-ref  anArray index))
      ((array-set! self index newValue)
       (vector-set! anArray index newValue))
      ((print self port)
       (format port "#<Array ~s>" (size self)))))))

(define-operation (position obj))
(define-operation (discarded-value obj))

(define (make-cell-with-history value filter size)
  (let ((pos 0) (most-recent-discard #f))
    (object-with-ancestors
      ((cell (make-filtered-call value filter))
       (sequence (make-array size)))
      ((array? self) #f)
      ((position self) pos)
      ((store! self newValue)
       (operate-as cell store! self newValue)
       (array-set! self pos newValue)
       (set! pos (+ pos 1)))
      ((discard self value)
       (set! most-recent-discard value))
      ((discarded-value self) most-recent-discard)
      ((print self port)
       (format port "#<Cell-with-history ~s>"
               (fetch self))))))

(define-access-operation fetch)
(add-setter fetch store!)
(define foo (make-cell 1))
(print foo #f)
⇒ "#<Cell: 1>"
(set (fetch foo) 2)
⇒
(print foo #f)
⇒ "#<Cell: 2>"
(fetch foo)
⇒ 2
```

4 Textual Conversion Packages

4.1 Precedence Parsing

(`require 'precedence-parse`) or (`require 'parse`)

This package implements:

- a Pratt style precedence parser;
- a *tokenizer* which congeals tokens according to assigned classes of constituent characters;
- procedures giving direct control of parser rulesets;
- procedures for higher level specification of rulesets.

4.1.1 Precedence Parsing Overview

This package offers improvements over previous parsers.

- Common computer language constructs are concisely specified.
- Grammars can be changed dynamically. Operators can be assigned different meanings within a lexical context.
- Rulesets don't need compilation. Grammars can be changed incrementally.
- Operator precedence is specified by integers.
- All possibilities of bad input are handled[1] and return as much structure as was parsed when the error occured; The symbol ? is substituted for missing input.

The notion of *binding power* may be unfamiliar to those accustomed to BNF grammars.

When two consecutive objects are parsed, the first might be the prefix to the second, or the second might be a suffix of the first. Comparing the left and right binding powers of the two objects decides which way to interpret them.

Objects at each level of syntactic grouping have binding powers.

A syntax tree is not built unless the rules explicitly do so. The call graph of grammar rules effectively instantiate the sytnax tree.

The JACAL symbolic math system (`http://people.csail.mit.edu/jaffer/JACAL`) uses `precedence-parse`. Its grammar definitions in the file `jacal/English.scm` can serve as examples of use.

4.1.2 Rule Types

Here are the higher-level syntax types and an example of each. Precedence considerations are omitted for clarity. See Section 4.1.6 [Grammar Rule Definition], page 44, for full details.

nofix *bye exit*   [Grammar]

      `bye`

    calls the function `exit` with no arguments.

---

[1] How do I know this? I parsed 250kbyte of random input (an e-mail file) with a non-trivial grammar utilizing all constructs.

prefix - *negate* [Grammar]

       `- 42`

    Calls the function `negate` with the argument `42`.

infix - *difference* [Grammar]

       `x - y`

    Calls the function `difference` with arguments `x` and `y`.

nary + *sum* [Grammar]

       `x + y + z`

    Calls the function `sum` with arguments `x`, `y`, and `y`.

postfix *! factorial* [Grammar]

       `5 !`

    Calls the function `factorial` with the argument `5`.

prestfix *set set!* [Grammar]

       `set foo bar`

    Calls the function `set!` with the arguments `foo` and `bar`.

commentfix */* */* [Grammar]

       `/* almost any text here */`

    Ignores the comment delimited by `/*` and `*/`.

matchfix { *list* } [Grammar]

       `{0, 1, 2}`

    Calls the function `list` with the arguments `0`, `1`, and `2`.

inmatchfix (*funcall*) [Grammar]

       `f(x, y)`

    Calls the function `funcall` with the arguments `f`, `x`, and `y`.

delim ; [Grammar]

       `set foo bar;`

    delimits the extent of the restfix operator `set`.

4.1.3 Ruleset Definition and Use

syn-defs [Variable]

    A grammar is built by one or more calls to `prec:define-grammar`. The rules are appended to **syn-defs**. The value of **syn-defs** is the grammar suitable for passing as an argument to `prec:parse`.

syn-ignore-whitespace [Constant]

    Is a nearly empty grammar with whitespace characters set to group 0, which means they will not be made into tokens. Most rulesets will want to start with `*syn-ignore-whitespace*`

In order to start defining a grammar, either

    (set! *syn-defs* '())

or

    (set! *syn-defs* *syn-ignore-whitespace*)

prec:define-grammar *rule1 ...*   [Function]
> Appends *rule1 ...* to **syn-defs**. `prec:define-grammar` is used to define both the character classes and rules for tokens.

Once your grammar is defined, save the value of `*syn-defs*` in a variable (for use when calling `prec:parse`).

    (define my-ruleset *syn-defs*)

prec:parse *ruleset delim column*   [Function]
prec:parse *ruleset delim column port*   [Function]
> The *ruleset* argument must be a list of rules as constructed by `prec:define-grammar` and extracted from **syn-defs**.
>
> The token *delim* may be a character, symbol, or string. A character *delim* argument will match only a character token; i.e. a character for which no token-group is assigned. A symbol or string will match only a token string; i.e. a token resulting from a token group.
>
> `prec:parse` reads a *ruleset* grammar expression delimited by *delim* from the given input *port*. `prec:parse` returns the next object parsable from the given input *port*, updating *port* to point to the first character past the end of the external representation of the object.
>
> For the purpose of reporting problems in error messages, this package keeps track of the *current column*. Its initial value is passed as the third argument to `prec:parse`.
>
> If an end of file is encountered in the input before any characters are found that can begin an object, then an end of file object is returned. If a delimiter (such as *delim*) is found before any characters are found that can begin an object, then `#f` is returned.
>
> The *port* argument may be omitted, in which case it defaults to the value returned by `current-input-port`. It is an error to parse from a closed port.

4.1.4 Token definition

tok:char-group *group chars chars-proc*   [Function]
> The argument *chars* may be a single character, a list of characters, or a string. Each character in *chars* is treated as though `tok:char-group` was called with that character alone.
>
> The argument *chars-proc* must be a procedure of one argument, a list of characters. After `tokenize` has finished accumulating the characters for a token, it calls *chars-proc* with the list of characters. The value returned is the token which `tokenize` returns.

The argument *group* may be an exact integer or a procedure of one character argument. The following discussion concerns the treatment which the tokenizing routine, `tokenize`, will accord to characters on the basis of their groups.

When *group* is a non-zero integer, characters whose group number is equal to or exactly one less than *group* will continue to accumulate. Any other character causes the accumulation to stop (until a new token is to be read).

The *group* of zero is special. These characters are ignored when parsed pending a token, and stop the accumulation of token characters when the accumulation has already begun. Whitespace characters are usually put in group 0.

If *group* is a procedure, then, when triggerd by the occurence of an initial (no accumulation) *chars* character, this procedure will be repeatedly called with each successive character from the input stream until the *group* procedure returns a non-false value.

The following convenient constants are provided for use with `tok:char-group`.

`tok:decimal-digits`   [Constant]
    Is the string `"0123456789"`.

`tok:upper-case`   [Constant]
    Is the string consisting of all upper-case letters (`"ABCDEFGHIJKLMNOPQRSTU-VWXYZ"`).

`tok:lower-case`   [Constant]
    Is the string consisting of all lower-case letters (`"abcdefghijklmnopqrstuvwxyz"`).

`tok:whitespaces`   [Constant]
    Is the string consisting of all characters between 0 and 255 for which `char-whitespace?` returns true.

4.1.5 Nud and Led Definition

This section describes advanced features. You can skip this section on first reading.

The *Null Denotation* (or *nud*) of a token is the procedure and arguments applying for that token when *Left*, an unclaimed parsed expression is not extant.

The *Left Denotation* (or *led*) of a token is the procedure, arguments, and lbp applying for that token when there is a *Left*, an unclaimed parsed expression.

In his paper,

    Pratt, V. R. Top Down Operator Precendence. *SIGACT/SIGPLAN Symposium on Principles of Programming Languages*, Boston, 1973, pages 41-51

the *left binding power* (or *lbp*) was an independent property of tokens. I think this was done in order to allow tokens with NUDs but not LEDs to also be used as delimiters, which was a problem for statically defined syntaxes. It turns out that *dynamically binding* NUDs and LEDs allows them independence.

For the rule-defining procedures that follow, the variable *tk* may be a character, string, or symbol, or a list composed of characters, strings, and symbols. Each element of *tk* is treated as though the procedure were called for each element.

Character *tk* arguments will match only character tokens; i.e. characters for which no token-group is assigned. Symbols and strings will both match token strings; i.e. tokens resulting from token groups.

prec:make-nud *tk sop arg1 ...*   [Function]
> Returns a rule specifying that *sop* be called when *tk* is parsed. If *sop* is a procedure, it is called with *tk* and *arg1* ... as its arguments; the resulting value is incorporated into the expression being built. Otherwise, (`list sop arg1 ...`) is incorporated.

If no NUD has been defined for a token; then if that token is a string, it is converted to a symbol and returned; if not a string, the token is returned.

prec:make-led *tk sop arg1 ...*   [Function]
> Returns a rule specifying that *sop* be called when *tk* is parsed and *left* has an unclaimed parsed expression. If *sop* is a procedure, it is called with *left*, *tk*, and *arg1* ... as its arguments; the resulting value is incorporated into the expression being built. Otherwise, *left* is incorporated.

If no LED has been defined for a token, and *left* is set, the parser issues a warning.

4.1.6 Grammar Rule Definition

Here are procedures for defining rules for the syntax types introduced in Section 4.1.1 [Precedence Parsing Overview], page 40.

For the rule-defining procedures that follow, the variable *tk* may be a character, string, or symbol, or a list composed of characters, strings, and symbols. Each element of *tk* is treated as though the procedure were called for each element.

For procedures prec:delim, ..., prec:prestfix, if the *sop* argument is `#f`, then the token which triggered this rule is converted to a symbol and returned. A false *sop* argument to the procedures prec:commentfix, prec:matchfix, or prec:inmatchfix has a different meaning.

Character *tk* arguments will match only character tokens; i.e. characters for which no token-group is assigned. Symbols and strings will both match token strings; i.e. tokens resulting from token groups.

prec:delim *tk*   [Function]
> Returns a rule specifying that *tk* should not be returned from parsing; i.e. *tk*'s function is purely syntactic. The end-of-file is always treated as a delimiter.

prec:nofix *tk sop*   [Function]
> Returns a rule specifying the following actions take place when *tk* is parsed:
> - If *sop* is a procedure, it is called with no arguments; the resulting value is incorporated into the expression being built. Otherwise, the list of *sop* is incorporated.

prec:prefix *tk sop bp rule1 ...*   [Function]
> Returns a rule specifying the following actions take place when *tk* is parsed:
> - The rules *rule1* ... augment and, in case of conflict, override rules currently in effect.
> - `prec:parse1` is called with binding-power *bp*.

- If *sop* is a procedure, it is called with the expression returned from `prec:parse1`; the resulting value is incorporated into the expression being built. Otherwise, the list of *sop* and the expression returned from `prec:parse1` is incorporated.

- The ruleset in effect before *tk* was parsed is restored; *rule1* ... are forgotten.

`prec:infix` *tk sop lbp bp rule1* ...   [Function]
    Returns a rule declaring the left-binding-precedence of the token *tk* is *lbp* and specifying the following actions take place when *tk* is parsed:

- The rules *rule1* ... augment and, in case of conflict, override rules currently in effect.

- One expression is parsed with binding-power *lbp*. If instead a delimiter is encountered, a warning is issued.

- If *sop* is a procedure, it is applied to the list of *left* and the parsed expression; the resulting value is incorporated into the expression being built. Otherwise, the list of *sop*, the *left* expression, and the parsed expression is incorporated.

- The ruleset in effect before *tk* was parsed is restored; *rule1* ... are forgotten.

`prec:nary` *tk sop bp*   [Function]
    Returns a rule declaring the left-binding-precedence of the token *tk* is *bp* and specifying the following actions take place when *tk* is parsed:

- Expressions are parsed with binding-power *bp* as far as they are interleaved with the token *tk*.

- If *sop* is a procedure, it is applied to the list of *left* and the parsed expressions; the resulting value is incorporated into the expression being built. Otherwise, the list of *sop*, the *left* expression, and the parsed expressions is incorporated.

`prec:postfix` *tk sop lbp*   [Function]
    Returns a rule declaring the left-binding-precedence of the token *tk* is *lbp* and specifying the following actions take place when *tk* is parsed:

- If *sop* is a procedure, it is called with the *left* expression; the resulting value is incorporated into the expression being built. Otherwise, the list of *sop* and the *left* expression is incorporated.

`prec:prestfix` *tk sop bp rule1* ...   [Function]
    Returns a rule specifying the following actions take place when *tk* is parsed:

- The rules *rule1* ... augment and, in case of conflict, override rules currently in effect.

- Expressions are parsed with binding-power *bp* until a delimiter is reached.

- If *sop* is a procedure, it is applied to the list of parsed expressions; the resulting value is incorporated into the expression being built. Otherwise, the list of *sop* and the parsed expressions is incorporated.

- The ruleset in effect before *tk* was parsed is restored; *rule1* ... are forgotten.

`prec:commentfix` *tk stp match rule1* ...   [Function]
    Returns rules specifying the following actions take place when *tk* is parsed:

- The rules *rule1* . . . augment and, in case of conflict, override rules currently in effect.

- Characters are read until and end-of-file or a sequence of characters is read which matches the *string match*.

- If *stp* is a procedure, it is called with the string of all that was read between the *tk* and *match* (exclusive).

- The ruleset in effect before *tk* was parsed is restored; *rule1* . . . are forgotten.

Parsing of commentfix syntax differs from the others in several ways. It reads directly from input without tokenizing; It calls *stp* but does not return its value; nay any value. I added the *stp* argument so that comment text could be echoed.

prec:matchfix *tk sop sep match rule1* . . .   [Function]
Returns a rule specifying the following actions take place when *tk* is parsed:

- The rules *rule1* . . . augment and, in case of conflict, override rules currently in effect.

- A rule declaring the token *match* a delimiter takes effect.

- Expressions are parsed with binding-power 0 until the token *match* is reached. If the token *sep* does not appear between each pair of expressions parsed, a warning is issued.

- If *sop* is a procedure, it is applied to the list of parsed expressions; the resulting value is incorporated into the expression being built. Otherwise, the list of *sop* and the parsed expressions is incorporated.

- The ruleset in effect before *tk* was parsed is restored; *rule1* . . . are forgotten.

prec:inmatchfix *tk sop sep match lbp rule1* . . .   [Function]
Returns a rule declaring the left-binding-precedence of the token *tk* is *lbp* and specifying the following actions take place when *tk* is parsed:

- The rules *rule1* . . . augment and, in case of conflict, override rules currently in effect.

- A rule declaring the token *match* a delimiter takes effect.

- Expressions are parsed with binding-power 0 until the token *match* is reached. If the token *sep* does not appear between each pair of expressions parsed, a warning is issued.

- If *sop* is a procedure, it is applied to the list of *left* and the parsed expressions; the resulting value is incorporated into the expression being built. Otherwise, the list of *sop*, the *left* expression, and the parsed expressions is incorporated.

- The ruleset in effect before *tk* was parsed is restored; *rule1* . . . are forgotten.

4.2 Format (version 3.1)

`(require 'format)` or `(require 'srfi-28)`

4.2.1 Format Interface

`format` *destination format-string . arguments*   [Function]

An almost complete implementation of Common LISP format description according to the CL reference book *Common LISP* from Guy L. Steele, Digital Press. Backward compatible to most of the available Scheme format implementations.

Returns `#t`, `#f` or a string; has side effect of printing according to *format-string*. If *destination* is `#t`, the output is to the current output port and `#t` is returned. If *destination* is `#f`, a formatted string is returned as the result of the call. NEW: If *destination* is a string, *destination* is regarded as the format string; *format-string* is then the first argument and the output is returned as a string. If *destination* is a number, the output is to the current error port if available by the implementation. Otherwise *destination* must be an output port and `#t` is returned.

format-string must be a string. In case of a formatting error format returns `#f` and prints a message on the current output or error port. Characters are output as if the string were output by the `display` function with the exception of those prefixed by a tilde (`~`). For a detailed description of the *format-string* syntax please consult a Common LISP format reference manual. For a test suite to verify this format implementation load `formatst.scm`.

4.2.2 Format Specification (Format version 3.1)

Please consult a Common LISP format reference manual for a detailed description of the format string syntax. For a demonstration of the implemented directives see `formatst.scm`.

This implementation supports directive parameters and modifiers (`:` and `@` characters). Multiple parameters must be separated by a comma (`,`). Parameters can be numerical parameters (positive or negative), character parameters (prefixed by a quote character (`'`), variable parameters (`v`), number of rest arguments parameter (`#`), empty and default parameters. Directive characters are case independent. The general form of a directive is:

directive ::= `~`{*directive-parameter*,}[`:`][`@`]*directive-character*

directive-parameter ::= [[-|+]{0-9}+ | '*character* | v | #]

4.2.2.1 Implemented CL Format Control Directives

Documentation syntax: Uppercase characters represent the corresponding control directive characters. Lowercase characters represent control directive parameter descriptions.

`~A`   Any (print as `display` does).

    `~@A`   left pad.

    `~mincol,colinc,minpad,padcharA`
        full padding.

`~S`   S-expression (print as `write` does).

    `~@S`   left pad.

    `~mincol,colinc,minpad,padcharS`
        full padding.

~D   Decimal.

>   ~@D   print number sign always.

>   ~:D   print comma separated.

>   *~mincol,padchar,commachar*D
>   padding.

~X   Hexadecimal.

>   ~@X   print number sign always.

>   ~:X   print comma separated.

>   *~mincol,padchar,commachar*X
>   padding.

~O   Octal.

>   ~@O   print number sign always.

>   ~:O   print comma separated.

>   *~mincol,padchar,commachar*O
>   padding.

~B   Binary.

>   ~@B   print number sign always.

>   ~:B   print comma separated.

>   *~mincol,padchar,commachar*B
>   padding.

*~n*R   Radix *n*.

>   *~n,mincol,padchar,commachar*R
>   padding.

~@R   print a number as a Roman numeral.

~:@R   print a number as an "old fashioned" Roman numeral.

~:R   print a number as an ordinal English number.

~R   print a number as a cardinal English number.

~P   Plural.

>   ~@P   prints y and ies.

>   ~:P   as ~P but jumps 1 argument backward.

>   ~:@P   as ~@P but jumps 1 argument backward.

~C   Character.

>   ~@C   prints a character as the reader can understand it (i.e. #\ prefixing).

>   ~:C   prints a character as emacs does (eg. ^C for ASCII 03).

~F   Fixed-format floating-point (prints a flonum like *mmm.nnn*).

            *~width,digits,scale,overflowchar,padchar*F
            ~@F   If the number is positive a plus sign is printed.

~E   Exponential floating-point (prints a flonum like *mmm.nnn*E*ee*).

            *~width,digits,exponentdigits,scale,overflowchar,padchar,exponentchar*E
            ~@E   If the number is positive a plus sign is printed.

~G   General floating-point (prints a flonum either fixed or exponential).

            *~width,digits,exponentdigits,scale,overflowchar,padchar,exponentchar*G
            ~@G   If the number is positive a plus sign is printed.

~$ Dollars floating-point (prints a flonum in fixed with signs separated).

            ~digits,scale,width,padchar$
            ~@$ If the number is positive a plus sign is printed.

            ~:@$ A sign is always printed and appears before the padding.

            ~:$ The sign appears before the padding.

~%   Newline.

            ~n%   print *n* newlines.

~&   print newline if not at the beginning of the output line.

            ~n&   prints ~& and then *n-1* newlines.

~| Page Separator.

            ~n| print *n* page separators.

~~   Tilde.

            ~n~   print *n* tildes.

~<newline>
            Continuation Line.

            ~:<newline>
                         newline is ignored, white space left.

            ~@<newline>
                         newline is left, white space ignored.

~T   Tabulation.

            ~@T   relative tabulation.

            *~colnum,colinc*T
                         full tabulation.

~?   Indirection (expects indirect arguments as a list).

            ~@?   extracts indirect arguments from format arguments.

~(*str*~)   Case conversion (converts by `string-downcase`).

            ~:(*str*~)   converts by `string-capitalize`.

~@(*str*~)   converts by `string-capitalize-first`.

~:@(*str*~)
                converts by `string-upcase`.

~*   Argument Jumping (jumps 1 argument forward).

    ~*n**   jumps *n* arguments forward.

    ~:*   jumps 1 argument backward.

    ~*n*:*   jumps *n* arguments backward.

    ~@*   jumps to the 0th argument.

    ~*n*@*   jumps to the *n*th argument (beginning from 0)

~[*str0*~;*str1*~;...~;*strn*~]
                Conditional Expression (numerical clause conditional).

    ~*n*[take argument from *n*.

    ~@[true test conditional.

    ~:[if-else-then conditional.

    ~;   clause separator.

    ~:;   default clause follows.

~{*str*~}   Iteration (args come from the next argument (a list)). Iteration bounding is
                controlled by configuration variables *format:iteration-bounded* and *format:max-iterations*. With both variables default, a maximum of 100 iterations will be
                performed.

    ~*n*{   at most *n* iterations.

    ~:{   args from next arg (a list of lists).

    ~@{   args from the rest of arguments.

    ~:@{   args from the rest args (lists).

~^   Up and out.

    ~*n*^   aborts if $n = 0$

    ~*n*,*m*^   aborts if $n = m$

    ~*n*,*m*,*k*^   aborts if $n <= m <= k$

4.2.2.2 Not Implemented CL Format Control Directives

~:A   print `#f` as an empty list (see below).

~:S   print `#f` as an empty list (see below).

~<~>   Justification.

~:^   (sorry I don't understand its semantics completely)

4.2.2.3 Extended, Replaced and Additional Control Directives

`~`*mincol*`,`*padchar*`,`*commachar*`,`*commawidth*`D`
`~`*mincol*`,`*padchar*`,`*commachar*`,`*commawidth*`X`
`~`*mincol*`,`*padchar*`,`*commachar*`,`*commawidth*`O`
`~`*mincol*`,`*padchar*`,`*commachar*`,`*commawidth*`B`
`~`*n*`,`*mincol*`,`*padchar*`,`*commachar*`,`*commawidth*`R`

> *commawidth* is the number of characters between two comma characters.

`~I`   print a R4RS complex number as `~F~@Fi` with passed parameters for `~F`.

`~Y`   Pretty print formatting of an argument for scheme code lists.

`~K`   Same as `~?`.

`~!`   Flushes the output if format *destination* is a port.

`~_`   Print a `#\space` character

> `~`*n*`_`   print *n* `#\space` characters.

`~/`   Print a `#\tab` character

> `~`*n*`/`   print *n* `#\tab` characters.

`~`*n*`C`   Takes *n* as an integer representation for a character. No arguments are consumed. *n* is converted to a character by `integer->char`. *n* must be a positive decimal number.

`~:S`   Print out readproof. Prints out internal objects represented as `#<...>` as strings `"#<...>"` so that the format output can always be processed by `read`.

`~:A`   Print out readproof. Prints out internal objects represented as `#<...>` as strings `"#<...>"` so that the format output can always be processed by `read`.

`~Q`   Prints information and a copyright notice on the format implementation.

> `~:Q`   prints format version.

`~F, ~E, ~G, ~$`

> may also print number strings, i.e. passing a number as a string and format it accordingly.

4.2.2.4 Configuration Variables

Format has some configuration variables at the beginning of `format.scm` to suit the systems and users needs. There should be no modification necessary for the configuration that comes with SLIB. If modification is desired the variable should be set after the format code is loaded. Format detects automatically if the running scheme system implements floating point numbers and complex numbers.

format:symbol-case-conv

> Symbols are converted by `symbol->string` so the case type of the printed symbols is implementation dependent. `format:symbol-case-conv` is a one arg closure which is either `#f` (no conversion), `string-upcase`, `string-downcase` or `string-capitalize`. (default `#f`)

format:iobj-case-conv

> As *format:symbol-case-conv* but applies for the representation of implementation internal objects. (default `#f`)

format:expch

> The character prefixing the exponent value in `~E` printing. (default `#\E`)

format:iteration-bounded

> When `#t`, a `~{...~}` control will iterate no more than the number of times specified by *format:max-iterations* regardless of the number of iterations implied by modifiers and arguments. When `#f`, a `~{...~}` control will iterate the number of times implied by modifiers and arguments, unless termination is forced by language or system limitations. (default `#t`)

format:max-iterations

> The maximum number of iterations performed by a `~{...~}` control. Has effect only when *format:iteration-bounded* is `#t`. (default 100)

4.2.2.5 Compatibility With Other Format Implementations

SLIB format 2.x:

> See `format.doc`.

SLIB format 1.4:

> Downward compatible except for padding support and `~A`, `~S`, `~P`, `~X` uppercase printing. SLIB format 1.4 uses C-style `printf` padding support which is completely replaced by the CL `format` padding style.

MIT C-Scheme 7.1:

> Downward compatible except for `~`, which is not documented (ignores all characters inside the format string up to a newline character). (7.1 implements `~a`, `~s`, *~newline*, `~~`, `~%`, numerical and variable parameters and `:/@` modifiers in the CL sense).

Elk 1.5/2.0:

> Downward compatible except for `~A` and `~S` which print in uppercase. (Elk implements `~a`, `~s`, `~~`, and `~%` (no directive parameters or modifiers)).

Scheme->C 01nov91:

> Downward compatible except for an optional destination parameter: S2C accepts a format call without a destination which returns a formatted string. This is equivalent to a #f destination in S2C. (S2C implements `~a`, `~s`, `~c`, `~%`, and `~~` (no directive parameters or modifiers)).

This implementation of format is solely useful in the SLIB context because it requires other components provided by SLIB.

4.3 Standard Formatted I/O

4.3.1 stdio

`(require 'stdio)`

requires `printf` and `scanf` and additionally defines the symbols:

stdin   [Variable]
>   Defined to be (`current-input-port`).

stdout   [Variable]
>   Defined to be (`current-output-port`).

stderr   [Variable]
>   Defined to be (`current-error-port`).

4.3.2 Standard Formatted Output

(require 'printf)

printf *format arg1 . . .*   [Procedure]
fprintf *port format arg1 . . .*   [Procedure]
sprintf *str format arg1 . . .*   [Procedure]
sprintf *#f format arg1 . . .*   [Procedure]
sprintf *k format arg1 . . .*   [Procedure]

>   Each function converts, formats, and outputs its *arg1 . . .* arguments according to
>   the control string *format* argument and returns the number of characters output.
>
>   `printf` sends its output to the port (`current-output-port`). `fprintf` sends its
>   output to the port *port*. `sprintf` `string-set!`s locations of the non-constant string
>   argument *str* to the output characters.
>
>   Two extensions of `sprintf` return new strings. If the first argument is `#f`, then the
>   returned string's length is as many characters as specified by the *format* and data; if
>   the first argument is a non-negative integer *k*, then the length of the returned string
>   is also bounded by *k*.
>
>   The string *format* contains plain characters which are copied to the output stream,
>   and conversion specifications, each of which results in fetching zero or more of the
>   arguments *arg1 . . . .* The results are undefined if there are an insufficient number
>   of arguments for the format. If *format* is exhausted while some of the *arg1 . . .*
>   arguments remain unused, the excess *arg1 . . .* arguments are ignored.
>
>   The conversion specifications in a format string have the form:
>
>   % [*flags*] [*width*] [. *precision*] [*type*] *conversion*
>
>   An output conversion specifications consist of an initial '%' character followed in se-
>   quence by:
>
>   - Zero or more *flag characters* that modify the normal behavior of the conversion
>   specification.
>
>   '-'   Left-justify the result in the field. Normally the result is right-
>   justified.
>
>   '+'   For the signed '%d' and '%i' conversions and all inexact conversions,
>   prefix a plus sign if the value is positive.
>
>   ' '   For the signed '%d' and '%i' conversions, if the result doesn't start
>   with a plus or minus sign, prefix it with a space character instead.
>   Since the '+' flag ensures that the result includes a sign, this flag is
>   ignored if both are specified.

‘#’   For inexact conversions, ‘#’ specifies that the result should always include a decimal point, even if no digits follow it. For the ‘%g’ and ‘%G’ conversions, this also forces trailing zeros after the decimal point to be printed where they would otherwise be elided.

For the ‘%o’ conversion, force the leading digit to be ‘0’, as if by increasing the precision. For ‘%x’ or ‘%X’, prefix a leading ‘0x’ or ‘0X’ (respectively) to the result. This doesn't do anything useful for the ‘%d’, ‘%i’, or ‘%u’ conversions. Using this flag produces output which can be parsed by the `scanf` functions with the ‘%i’ conversion (see Section 4.3.3 [Standard Formatted Input], page 56).

‘0’   Pad the field with zeros instead of spaces. The zeros are placed after any indication of sign or base. This flag is ignored if the ‘-’ flag is also specified, or if a precision is specified for an exact converson.

- An optional decimal integer specifying the *minimum field width*. If the normal conversion produces fewer characters than this, the field is padded (with spaces or zeros per the ‘0’ flag) to the specified width. This is a *minimum* width; if the normal conversion produces more characters than this, the field is *not* truncated.

  Alternatively, if the field width is ‘*’, the next argument in the argument list (before the actual value to be printed) is used as the field width. The width value must be an integer. If the value is negative it is as though the ‘-’ flag is set (see above) and the absolute value is used as the field width.

- An optional *precision* to specify the number of digits to be written for numeric conversions and the maximum field width for string conversions. The precision is specified by a period (‘.’) followed optionally by a decimal integer (which defaults to zero if omitted).

  Alternatively, if the precision is ‘.*’, the next argument in the argument list (before the actual value to be printed) is used as the precision. The value must be an integer, and is ignored if negative. If you specify ‘*’ for both the field width and precision, the field width argument precedes the precision argument. The ‘.*’ precision is an enhancement. C library versions may not accept this syntax.

  For the ‘%f’, ‘%e’, and ‘%E’ conversions, the precision specifies how many digits follow the decimal-point character. The default precision is 6. If the precision is explicitly 0, the decimal point character is suppressed.

  For the ‘%g’ and ‘%G’ conversions, the precision specifies how many significant digits to print. Significant digits are the first digit before the decimal point, and all the digits after it. If the precision is 0 or not specified for ‘%g’ or ‘%G’, it is treated like a value of 1. If the value being printed cannot be expressed accurately in the specified number of digits, the value is rounded to the nearest number that fits.

  For exact conversions, if a precision is supplied it specifies the minimum number of digits to appear; leading zeros are produced if necessary. If a precision is not supplied, the number is printed with as many digits as necessary. Converting an exact ‘0’ with an explicit precision of zero produces no characters.

- An optional one of '`l`', '`h`' or '`L`', which is ignored for numeric conversions. It is an error to specify these modifiers for non-numeric conversions.
- A character that specifies the conversion to be applied.

Exact Conversions

'`b`', '`B`'   Print an integer as an unsigned binary number.

  Note: '`%b`' and '`%B`' are SLIB extensions.

'`d`', '`i`'   Print an integer as a signed decimal number. '`%d`' and '`%i`' are synonymous for output, but are different when used with `scanf` for input (see Section 4.3.3 [Standard Formatted Input], page 56).

'`o`'   Print an integer as an unsigned octal number.

'`u`'   Print an integer as an unsigned decimal number.

'`x`', '`X`'   Print an integer as an unsigned hexadecimal number. '`%x`' prints using the digits '`0123456789abcdef`'. '`%X`' prints using the digits '`0123456789ABCDEF`'.

Inexact Conversions

'`f`'   Print a floating-point number in fixed-point notation.

'`e`', '`E`'   Print a floating-point number in exponential notation. '`%e`' prints '`e`' between mantissa and exponont. '`%E`' prints '`E`' between mantissa and exponont.

'`g`', '`G`'   Print a floating-point number in either fixed or exponential notation, whichever is more appropriate for its magnitude. Unless an '`#`' flag has been supplied, trailing zeros after a decimal point will be stripped off. '`%g`' prints '`e`' between mantissa and exponont. '`%G`' prints '`E`' between mantissa and exponent.

'`k`', '`K`'   Print a number like '`%g`', except that an SI prefix is output after the number, which is scaled accordingly. '`%K`' outputs a dot between number and prefix, '`%k`' does not.

Other Conversions

'`c`'   Print a single character. The '`-`' flag is the only one which can be specified. It is an error to specify a precision.

'`s`'   Print a string. The '`-`' flag is the only one which can be specified. A precision specifies the maximum number of characters to output; otherwise all characters in the string are output.

'`a`', '`A`'   Print a scheme expression. The '`-`' flag left-justifies the output. The '`#`' flag specifies that strings and characters should be quoted as by `write` (which can be read using `read`); otherwise, output is as `display` prints. A precision specifies the maximum number of characters to output; otherwise as many characters as needed are output.

  Note: '`%a`' and '`%A`' are SLIB extensions.

'%'   Print a literal '%' character. No argument is consumed. It is an error to
               specify flags, field width, precision, or type modifiers with '%%'.

4.3.3 Standard Formatted Input

`(require 'scanf)`

scanf-read-list *format*   [Function]
scanf-read-list *format port*   [Function]
scanf-read-list *format string*   [Function]

scanf *format arg1 . . .*   [Macro]
fscanf *port format arg1 . . .*   [Macro]
sscanf *str format arg1 . . .*   [Macro]

Each function reads characters, interpreting them according to the control string
format argument.

`scanf-read-list` returns a list of the items specified as far as the input matches
format. `scanf`, `fscanf`, and `sscanf` return the number of items successfully matched
and stored. `scanf`, `fscanf`, and `sscanf` also set the location corresponding to *arg1*
. . . using the methods:

symbol   `set!`

car expression
            `set-car!`

cdr expression
            `set-cdr!`

vector-ref expression
            `vector-set!`

substring expression
            `substring-move-left!`

The argument to a `substring` expression in *arg1* . . . must be a non-constant string.
Characters will be stored starting at the position specified by the second argument
to `substring`. The number of characters stored will be limited by either the position
specified by the third argument to `substring` or the length of the matched string,
whichever is less.

The control string, *format*, contains conversion specifications and other characters
used to direct interpretation of input sequences. The control string contains:

- White-space characters (blanks, tabs, newlines, or formfeeds) that cause input
  to be read (and discarded) up to the next non-white-space character.

- An ordinary character (not '%') that must match the next character of the input
  stream.

- Conversion specifications, consisting of the character '%', an optional assignment
  suppressing character '*', an optional numerical maximum-field width, an op-
  tional 'l', 'h' or 'L' which is ignored, and a conversion code.

Unless the specification contains the 'n' conversion character (described below), a conversion specification directs the conversion of the next input field. The result of a conversion specification is returned in the position of the corresponding argument points, unless '*' indicates assignment suppression. Assignment suppression provides a way to describe an input field to be skipped. An input field is defined as a string of characters; it extends to the next inappropriate character or until the field width, if specified, is exhausted.

> *Note:* This specification of format strings differs from the *ANSI C* and *POSIX* specifications. In SLIB, white space before an input field is not skipped unless white space appears before the conversion specification in the format string. In order to write format strings which work identically with *ANSI C* and SLIB, prepend whitespace to all conversion specifications except '[' and 'c'.

The conversion code indicates the interpretation of the input field; For a suppressed field, no value is returned. The following conversion codes are legal:

'%'   A single % is expected in the input at this point; no value is returned.

'd', 'D'   A decimal integer is expected.

'u', 'U'   An unsigned decimal integer is expected.

'o', 'O'   An octal integer is expected.

'x', 'X'   A hexadecimal integer is expected.

'i'   An integer is expected. Returns the value of the next input item, interpreted according to C conventions; a leading '0' implies octal, a leading '0x' implies hexadecimal; otherwise, decimal is assumed.

'n'   Returns the total number of bytes (including white space) read by `scanf`. No input is consumed by `%n`.

'f', 'F', 'e', 'E', 'g', 'G'
           A floating-point number is expected. The input format for floating-point numbers is an optionally signed string of digits, possibly containing a radix character '.', followed by an optional exponent field consisting of an 'E' or an 'e', followed by an optional '+', '-', or space, followed by an integer.

'c', 'C'   *Width* characters are expected. The normal skip-over-white-space is suppressed in this case; to read the next non-space character, use '%1s'. If a field width is given, a string is returned; up to the indicated number of characters is read.

's', 'S'   A character string is expected The input field is terminated by a white-space character. `scanf` cannot read a null string.

'['   Indicates string data and the normal skip-over-leading-white-space is suppressed. The left bracket is followed by a set of characters, called the scanset, and a right bracket; the input field is the maximal sequence of input characters consisting entirely of characters in the scanset. '^', when

it appears as the first character in the scanset, serves as a complement operator and redefines the scanset as the set of all characters not contained in the remainder of the scanset string. Construction of the scanset follows certain conventions. A range of characters may be represented by the construct first-last, enabling '[0123456789]' to be expressed '[0-9]'. Using this convention, first must be lexically less than or equal to last; otherwise, the dash stands for itself. The dash also stands for itself when it is the first or the last character in the scanset. To include the right square bracket as an element of the scanset, it must appear as the first character (possibly preceded by a '^') of the scanset, in which case it will not be interpreted syntactically as the closing bracket. At least one character must match for this conversion to succeed.

The `scanf` functions terminate their conversions at end-of-file, at the end of the control string, or when an input character conflicts with the control string. In the latter case, the offending character is left unread in the input stream.

4.4 Program and Arguments

4.4.1 Getopt

`(require 'getopt)`

This routine implements Posix command line argument parsing. Notice that returning values through global variables means that `getopt` is *not* reentrant.

Obedience to Posix format for the `getopt` calls sows confusion. Passing *argc* and *argv* as arguments while referencing *optind* as a global variable leads to strange behavior, especially when the calls to `getopt` are buried in other procedures.

Even in C, *argc* can be derived from *argv*; what purpose does it serve beyond providing an opportunity for *argv*/*argc* mismatch? Just such a mismatch existed for years in a SLIB `getopt--` example.

I have removed the *argc* and *argv* arguments to getopt procedures; and replaced them with a global variable:

`*argv*`   [Variable]
Define *argv* with a list of arguments before calling getopt procedures. If you don't want the first (0th) element to be ignored, set *optind* to 0 (after requiring getopt).

`*optind*`   [Variable]
Is the index of the current element of the command line. It is initially one. In order to parse a new command line or reparse an old one, *optind* must be reset.

`*optarg*`   [Variable]
Is set by getopt to the (string) option-argument of the current option.

`getopt` *optstring*   [Function]
Returns the next option letter in *argv* (starting from (`vector-ref argv *optind*`)) that matches a letter in *optstring*. *argv* is a vector or list of strings, the 0th of which getopt usually ignores. *optstring* is a string of recognized option

characters; if a character is followed by a colon, the option takes an argument which may be immediately following it in the string or in the next element of *argv*.

optind is the index of the next element of the *argv* vector to be processed. It is initialized to 1 by `getopt.scm`, and `getopt` updates it when it finishes with each element of *argv*.

`getopt` returns the next option character from *argv* that matches a character in *optstring*, if there is one that matches. If the option takes an argument, `getopt` sets the variable *optarg* to the option-argument as follows:

- If the option was the last character in the string pointed to by an element of *argv*, then *optarg* contains the next element of *argv*, and *optind* is incremented by 2. If the resulting value of *optind* is greater than or equal to (`length *argv*`), this indicates a missing option argument, and `getopt` returns an error indication.

- Otherwise, *optarg* is set to the string following the option character in that element of *argv*, and *optind* is incremented by 1.

If, when `getopt` is called, the string (`vector-ref argv *optind*`) either does not begin with the character `#\-` or is just `"-"`, `getopt` returns `#f` without changing *optind*. If (`vector-ref argv *optind*`) is the string `"--"`, `getopt` returns `#f` after incrementing *optind*.

If `getopt` encounters an option character that is not contained in *optstring*, it returns the question-mark `#\?` character. If it detects a missing option argument, it returns the colon character `#\:` if the first character of *optstring* was a colon, or a question-mark character otherwise. In either case, `getopt` sets the variable *getopt:opt* to the option character that caused the error.

The special option `"--"` can be used to delimit the end of the options; `#f` is returned, and `"--"` is skipped.

RETURN VALUE

`getopt` returns the next option character specified on the command line. A colon `#\:` is returned if `getopt` detects a missing argument and the first character of *optstring* was a colon `#\:`.

A question-mark `#\?` is returned if `getopt` encounters an option character not in *optstring* or detects a missing argument and the first character of *optstring* was not a colon `#\:`.

Otherwise, `getopt` returns `#f` when all command line options have been parsed.

Example:

```
#! /usr/local/bin/scm
(require 'program-arguments)
(require 'getopt)
(define argv (program-arguments))

(define opts ":a:b:cd")
(let loop ((opt (getopt (length argv) argv opts)))
  (case opt
```

```
      ((#\a) (print "option a: " *optarg*))
      ((#\b) (print "option b: " *optarg*))
      ((#\c) (print "option c"))
      ((#\d) (print "option d"))
      ((#\?) (print "error" getopt:opt))
      ((#\:) (print "missing arg" getopt:opt))
      ((#f) (if (< *optind* (length argv))
               (print "argv[" *optind* "]="
                        (list-ref argv *optind*)))
           (set! *optind* (+ *optind* 1))))
  (if (< *optind* (length argv))
      (loop (getopt (length argv) argv opts)))))

  (slib:exit)
```

4.4.2 Getopt—

getopt-- *optstring*   [Function]
> The procedure `getopt--` is an extended version of `getopt` which parses *long op-
> tion names* of the form '`--hold-the-onions`' and '`--verbosity-level=extreme`'.
> `Getopt--` behaves as `getopt` except for non-empty options beginning with '`--`'.
>
> Options beginning with '`--`' are returned as strings rather than characters. If a value
> is assigned (using '`=`') to a long option, `*optarg*` is set to the value. The '`=`' and
> value are not returned as part of the option string.
>
> No information is passed to `getopt--` concerning which long options should be ac-
> cepted or whether such options can take arguments. If a long option did not have
> an argument, `*optarg*` will be set to `#f`. The caller is responsible for detecting and
> reporting errors.

```
      (define opts ":-:b:")
      (define *argv* '("foo" "-b9" "--f1" "--2=" "--g3=35234.342" "--"))
      (define *optind* 1)
      (define *optarg* #f)
      (require 'qp)
      (do ((i 5 (+ -1 i)))
          ((zero? i))
        (let ((opt (getopt-- opts)))
          (print *optind* opt *optarg*)))
      ⊣
      2 #\b "9"
      3 "f1" #f
      4 "2" ""
      5 "g3" "35234.342"
      5 #f "35234.342"
```

4.4.3 Command Line

`(require 'read-command)`

read-command *port*   [Function]

read-command   [Function]

   read-command converts a *command line* into a list of strings suitable for parsing by
   getopt. The syntax of command lines supported resembles that of popular *shell*s.
   read-command updates *port* to point to the first character past the command delim-
   iter.

   If an end of file is encountered in the input before any characters are found that can
   begin an object or comment, then an end of file object is returned.

   The *port* argument may be omitted, in which case it defaults to the value returned
   by **current-input-port**.

   The fields into which the command line is split are delimited by whitespace as de-
   fined by **char-whitespace?**. The end of a command is delimited by end-of-file or
   unescaped semicolon (;) or **newline**. Any character can be literally included in a
   field by escaping it with a backslach (\).

   The initial character and types of fields recognized are:

   '\'   The next character has is taken literally and not interpreted as a field
              delimiter. If \ is the last character before a **newline**, that **newline** is
              just ignored. Processing continues from the characters after the **newline**
              as though the backslash and **newline** were not there.

   '"'   The characters up to the next unescaped " are taken literally, according
              to [R4RS] rules for literal strings (see Section "Strings" in *Revised(4)
              Scheme*).

   '(', '%'' One scheme expression is **read** starting with this character. The **read** ex-
              pression is evaluated, converted to a string (using **display**), and replaces
              the expression in the returned field.

   ';'   Semicolon delimits a command. Using semicolons more than one com-
              mand can appear on a line. Escaped semicolons and semicolons inside
              strings do not delimit commands.

   The comment field differs from the previous fields in that it must be the first character
   of a command or appear after whitespace in order to be recognized. # can be part of
   fields if these conditions are not met. For instance, ab#c is just the field ab#c.

   '#'   Introduces a comment. The comment continues to the end of the line on
              which the semicolon appears. Comments are treated as whitespace by
              read-dommand-line and backslashes before **newline**s in comments are
              also ignored.

read-options-file *filename*   [Function]

   read-options-file converts an *options file* into a list of strings suitable for parsing
   by **getopt**. The syntax of options files is the same as the syntax for command lines,
   except that **newline**s do not terminate reading (only ; or end of file).

   If an end of file is encountered before any characters are found that can begin an
   object or comment, then an end of file object is returned.

4.4.4 Parameter lists

`(require 'parameters)`

Arguments to procedures in scheme are distinguished from each other by their position in the procedure call. This can be confusing when a procedure takes many arguments, many of which are not often used.

A *parameter-list* is a way of passing named information to a procedure. Procedures are also defined to set unused parameters to default values, check parameters, and combine parameter lists.

A *parameter* has the form (parameter-name value1 ...). This format allows for more than one value per parameter-name.

A *parameter-list* is a list of *parameter*s, each with a different *parameter-name*.

`make-parameter-list` *parameter-names*   [Function]
> Returns an empty parameter-list with slots for *parameter-names*.

`parameter-list-ref` *parameter-list parameter-name*   [Function]
> *parameter-name* must name a valid slot of *parameter-list*. `parameter-list-ref` returns the value of parameter *parameter-name* of *parameter-list*.

`remove-parameter` *parameter-name parameter-list*   [Function]
> Removes the parameter *parameter-name* from *parameter-list*. `remove-parameter` does not alter the argument *parameter-list*.
>
> If there are more than one *parameter-name* parameters, an error is signaled.

`adjoin-parameters!` *parameter-list parameter1 ...*   [Procedure]
> Returns *parameter-list* with *parameter1 ...* merged in.

`parameter-list-expand` *expanders parameter-list*   [Procedure]
> *expanders* is a list of procedures whose order matches the order of the *parameter-names* in the call to `make-parameter-list` which created *parameter-list*. For each non-false element of *expanders* that procedure is mapped over the corresponding parameter value and the returned parameter lists are merged into *parameter-list*.
>
> This process is repeated until *parameter-list* stops growing. The value returned from `parameter-list-expand` is unspecified.

`fill-empty-parameters` *defaulters parameter-list*   [Function]
> *defaulters* is a list of procedures whose order matches the order of the *parameter-names* in the call to `make-parameter-list` which created *parameter-list*. `fill-empty-parameters` returns a new parameter-list with each empty parameter replaced with the list returned by calling the corresponding *defaulter* with *parameter-list* as its argument.

`check-parameters` *checks parameter-list*   [Function]
> *checks* is a list of procedures whose order matches the order of the *parameter-names* in the call to `make-parameter-list` which created *parameter-list*.
>
> `check-parameters` returns *parameter-list* if each *check* of the corresponding *parameter-list* returns non-false. If some *check* returns `#f` a warning is signaled.

In the following procedures *arities* is a list of symbols. The elements of `arities` can be:

single   Requires a single parameter.

optional   A single parameter or no parameter is acceptable.

boolean   A single boolean parameter or zero parameters is acceptable.

nary   Any number of parameters are acceptable.

nary1   One or more of parameters are acceptable.

`parameter-list->arglist` *positions arities parameter-list*   [Function]
> Returns *parameter-list* converted to an argument list. Parameters of *arity* type
> `single` and `boolean` are converted to the single value associated with them. The
> other *arity* types are converted to lists of the value(s).
>
> *positions* is a list of positive integers whose order matches the order of the *parameter-
> name*s in the call to `make-parameter-list` which created *parameter-list*. The inte-
> gers specify in which argument position the corresponding parameter should appear.

4.4.5 Getopt Parameter lists

`(require 'getopt-parameters)`

`getopt->parameter-list` *optnames arities types aliases desc . . .*   [Function]
> Returns **argv** converted to a parameter-list. *optnames* are the parameter-names.
> *arities* and *types* are lists of symbols corresponding to *optnames*.
>
> *aliases* is a list of lists of strings or integers paired with elements of *optnames*. Each
> one-character string will be treated as a single '`-`' option by `getopt`. Longer strings
> will be treated as long-named options (see Section 4.4.1 [Getopt], page 58).
>
> If the *aliases* association list has only strings as its `cars`, then all the option-arguments
> after an option (and before the next option) are adjoined to that option.
>
> If the *aliases* association list has integers, then each (string) option will take at most
> one option-argument. Unoptioned arguments are collected in a list. A '`-1`' alias will
> take the last argument in this list; '`+1`' will take the first argument in the list. The
> aliases -2 then +2; -3 then +3; . . . are tried so long as a positive or negative consecutive
> alias is found and arguments remain in the list. Finally a '`0`' alias, if found, absorbs
> any remaining arguments.
>
> In all cases, if unclaimed arguments remain after processing, a warning is signaled
> and #f is returned.

`getopt->arglist` *optnames positions arities types defaulters checks*   [Function]
>   *aliases desc . . .*
> Like `getopt->parameter-list`, but converts **argv** to an argument-list as specified
> by *optnames*, *positions*, *arities*, *types*, *defaulters*, *checks*, and *aliases*. If the options
> supplied violate the *arities* or *checks* constraints, then a warning is signaled and #f
> is returned.

These `getopt` functions can be used with SLIB relational databases. For an example, See
Section 6.1.1 [Using Databases], page 162.

If errors are encountered while processing options, directions for using the options (and argument strings *desc* ...) are printed to `current-error-port`.

```
(begin
  (set! *optind* 1)
  (set! *argv* '("cmd" "-?"))
  (getopt->parameter-list
   '(flag number symbols symbols string flag2 flag3 num2 num3)
   '(boolean optional nary1 nary single boolean boolean nary nary)
   '(boolean integer symbol symbol string boolean boolean integer integer)
   '(("flag" flag)
     ("f" flag)
     ("Flag" flag2)
     ("B" flag3)
     ("optional" number)
     ("o" number)
     ("nary1" symbols)
     ("N" symbols)
     ("nary" symbols)
     ("n" symbols)
     ("single" string)
     ("s" string)
     ("a" num2)
     ("Abs" num3))))
 ⊣
Usage: cmd [OPTION ARGUMENT ...] ...

  -f, --flag
  -o, --optional=<number>
  -n, --nary=<symbols> ...
  -N, --nary1=<symbols> ...
  -s, --single=<string>
      --Flag
  -B
  -a        <num2> ...
      --Abs=<num3> ...

ERROR: getopt->parameter-list "unrecognized option" "-?"
```

4.4.6 Filenames

`(require 'filename)`

filename:match?? *pattern*   [Function]
filename:match-ci?? *pattern*   [Function]
> Returns a predicate which returns a non-false value if its string argument matches (the string) *pattern*, false otherwise. Filename matching is like *glob* expansion described the bash manpage, except that names beginning with '.' are matched and '/' characters are not treated specially.

These functions interpret the following characters specially in *pattern* strings:

'*'   Matches any string, including the null string.

'?'   Matches any single character.

'[...]'   Matches any one of the enclosed characters. A pair of characters sepa-
           rated by a minus sign (-) denotes a range; any character lexically between
           those two characters, inclusive, is matched. If the first character following
           the '[' is a '!' or a '^' then any character not enclosed is matched. A '-'
           or ']' may be matched by including it as the first or last character in the
           set.

`filename:substitute??` *pattern template*   [Function]
`filename:substitute-ci??` *pattern template*   [Function]

> Returns a function transforming a single string argument according to glob patterns
> *pattern* and *template*. *pattern* and *template* must have the same number of wildcard
> specifications, which need not be identical. *pattern* and *template* may have a different
> number of literal sections. If an argument to the function matches *pattern* in the sense
> of `filename:match??` then it returns a copy of *template* in which each wildcard
> specification is replaced by the part of the argument matched by the corresponding
> wildcard specification in *pattern*. A * wildcard matches the longest leftmost string
> possible. If the argument does not match *pattern* then false is returned.
>
> *template* may be a function accepting the same number of string arguments as there
> are wildcard specifications in *pattern*. In the case of a match the result of applying
> *template* to a list of the substrings matched by wildcard specifications will be returned,
> otherwise *template* will not be called and `#f` will be returned.
>
> ```
> ((filename:substitute?? "scm_[0-9]*.html" "scm5c4_??.htm")
>  "scm_10.html")
> ⇒ "scm5c4_10.htm"
> ((filename:substitute?? "??" "beg?mid?end") "AZ")
> ⇒ "begAmidZend"
> ((filename:substitute?? "*na*" "?NA?") "banana")
> ⇒ "banaNA"
> ((filename:substitute?? "?*?" (lambda (s1 s2 s3) (string-append s3 s1)))
>  "ABZ")
> ⇒ "ZA"
> ```

`replace-suffix` *str old new*   [Function]

> *str* can be a string or a list of strings. Returns a new string (or strings) similar to
> `str` but with the suffix string *old* removed and the suffix string *new* appended. If the
> end of *str* does not match *old*, an error is signaled.
>
> ```
> (replace-suffix "/usr/local/lib/slib/batch.scm" ".scm" ".c")
> ⇒ "/usr/local/lib/slib/batch.c"
> ```

`call-with-tmpnam` *proc k*   [Function]
`call-with-tmpnam` *proc*   [Function]

> Calls *proc* with *k* arguments, strings returned by successive calls to `tmpnam`. If *proc*
> returns, then any files named by the arguments to *proc* are deleted automatically and

the value(s) yielded by the *proc* is(are) returned. *k* may be ommited, in which case
it defaults to `1`.

call-with-tmpnam *proc suffix1 . . .*   [Function]
Calls *proc* with strings returned by successive calls to `tmpnam`, each with the cor-
responding *suffix* string appended. If *proc* returns, then any files named by the
arguments to *proc* are deleted automatically and the value(s) yielded by the *proc*
is(are) returned.

4.4.7 Batch

`(require 'batch)`

The batch procedures provide a way to write and execute portable scripts for a variety
of operating systems. Each `batch:` procedure takes as its first argument a parameter-list
(see Section 4.4.4 [Parameter lists], page 62). This parameter-list argument *parms* contains
named associations. Batch currently uses 2 of these:

`batch-port`
> The port on which to write lines of the batch file.

`batch-dialect`
> The syntax of batch file to generate. Currently supported are:
> - unix
> - dos
> - vms
> - amigaos
> - system
> - *unknown*

The 'batch' module uses 2 enhanced relational tables (see Section 6.1.1 [Using
Databases], page 162) to store information linking the names of `operating-system`s to
`batch-dialect`es.

batch:initialize! *database*   [Function]
Defines `operating-system` and `batch-dialect` tables and adds the domain
`operating-system` to the enhanced relational database *database*.

operating-system   [Variable]
Is batch's best guess as to which operating-system it is running under.
`*operating-system*` is set to `(software-type)` (see Section 2.2 [Configuration],
page 12) unless `(software-type)` is `unix`, in which case finer distinctions are made.

batch:call-with-output-script *parms file proc*   [Function]
proc should be a procedure of one argument. If *file* is an output-port,
`batch:call-with-output-script` writes an appropriate header to *file* and then
calls *proc* with *file* as the only argument. If *file* is a string, `batch:call-with-
output-script` opens a output-file of name *file*, writes an appropriate header
to *file*, and then calls *proc* with the newly opened port as the only argument.
Otherwise, `batch:call-with-output-script` acts as if it was called with the result
of `(current-output-port)` as its third argument.

The rest of the `batch:` procedures write (or execute if `batch-dialect` is `system`) commands to the batch port which has been added to *parms* or (`copy-tree parms`) by the code:

    (adjoin-parameters! parms (list 'batch-port port))

`batch:command` *parms string1 string2 . . .*   [Function]
>   Calls `batch:try-command` (below) with arguments, but signals an error if `batch:try-command` returns `#f`.

These functions return a non-false value if the command was successfully translated into the batch dialect and `#f` if not. In the case of the `system` dialect, the value is non-false if the operation suceeded.

`batch:try-command` *parms string1 string2 . . .*   [Function]
>   Writes a command to the `batch-port` in *parms* which executes the program named *string1* with arguments *string2* . . . .

`batch:try-chopped-command` *parms arg1 arg2 . . . list*   [Function]
>   breaks the last argument *list* into chunks small enough so that the command:
>
>   arg1 arg2 ... chunk
>
>   fits withing the platform's maximum command-line length.
>
>   `batch:try-chopped-command` calls `batch:try-command` with the command and returns non-false only if the commands all fit and `batch:try-command` of each command line returned non-false.

`batch:run-script` *parms string1 string2 . . .*   [Function]
>   Writes a command to the `batch-port` in *parms* which executes the batch script named *string1* with arguments *string2* . . . .
>
>   *Note:* `batch:run-script` and `batch:try-command` are not the same for some operating systems (VMS).

`batch:comment` *parms line1 . . .*   [Function]
>   Writes comment lines *line1* . . . to the `batch-port` in *parms*.

`batch:lines->file` *parms file line1 . . .*   [Function]
>   Writes commands to the `batch-port` in *parms* which create a file named *file* with contents *line1* . . . .

`batch:delete-file` *parms file*   [Function]
>   Writes a command to the `batch-port` in *parms* which deletes the file named *file*.

`batch:rename-file` *parms old-name new-name*   [Function]
>   Writes a command to the `batch-port` in *parms* which renames the file *old-name* to *new-name*.

In addition, batch provides some small utilities very useful for writing scripts:

`truncate-up-to` *path char*   [Function]
`truncate-up-to` *path string*   [Function]

truncate-up-to *path charlist*   [Function]

> *path* can be a string or a list of strings. Returns *path* sans any prefixes ending with
> a character of the second argument. This can be used to derive a filename moved
> locally from elsewhere.

>> ```
>> (truncate-up-to "/usr/local/lib/slib/batch.scm" "/")
>> ⇒ "batch.scm"
>> ```

string-join *joiner string1 . . .*   [Function]

> Returns a new string consisting of all the strings *string1* . . . in order appended
> together with the string *joiner* between each adjacent pair.

must-be-first *list1 list2*   [Function]

> Returns a new list consisting of the elements of *list2* ordered so that if some elements
> of *list1* are `equal?` to elements of *list2*, then those elements will appear first and in
> the order of *list1*.

must-be-last *list1 list2*   [Function]

> Returns a new list consisting of the elements of *list1* ordered so that if some elements
> of *list2* are `equal?` to elements of *list1*, then those elements will appear last and in
> the order of *list2*.

os->batch-dialect *osname*   [Function]

> Returns its best guess for the `batch-dialect` to be used for the operating-system
> named *osname*. `os->batch-dialect` uses the tables added to *database* by
> `batch:initialize!`.

Here is an example of the use of most of batch's procedures:

```
(require 'databases)
(require 'parameters)
(require 'batch)
(require 'filename)

(define batch (create-database #f 'alist-table))
(batch:initialize! batch)

(define my-parameters
  (list (list 'batch-dialect (os->batch-dialect *operating-system*))
        (list 'operating-system *operating-system*)
        (list 'batch-port (current-output-port)))) ;gets filled in later

(batch:call-with-output-script
 my-parameters
 "my-batch"
 (lambda (batch-port)
   (adjoin-parameters! my-parameters (list 'batch-port batch-port))
   (and
    (batch:comment my-parameters
```

```
                          "================ Write file with C program.")
        (batch:rename-file my-parameters "hello.c" "hello.c~")
        (batch:lines->file my-parameters "hello.c"
                           "#include <stdio.h>"
                           "int main(int argc, char **argv)"
                           "{"
                           "  printf(\"hello world\\n\");"
                           "  return 0;"
                           "}" )
        (batch:command my-parameters "cc" "-c" "hello.c")
        (batch:command my-parameters "cc" "-o" "hello"
                    (replace-suffix "hello.c" ".c" ".o"))
        (batch:command my-parameters "hello")
        (batch:delete-file my-parameters "hello")
        (batch:delete-file my-parameters "hello.c")
        (batch:delete-file my-parameters "hello.o")
        (batch:delete-file my-parameters "my-batch")
        )))
```

Produces the file `my-batch`:

```
#! /bin/sh
# "my-batch" script created by SLIB/batch Sun Oct 31 18:24:10 1999
# ================ Write file with C program.
mv -f hello.c hello.c~
rm -f hello.c
echo '#include <stdio.h>'>>hello.c
echo 'int main(int argc, char **argv)'>>hello.c
echo '{'>>hello.c
echo '  printf("hello world\n");'>>hello.c
echo '  return 0;'>>hello.c
echo '}'>>hello.c
cc -c hello.c
cc -o hello hello.o
hello
rm -f hello
rm -f hello.c
rm -f hello.o
rm -f my-batch
```

When run, `my-batch` prints:

```
bash$ my-batch
mv: hello.c: No such file or directory
hello world
```

4.5 HTML

```
(require 'html-form)
```

`html:atval` *txt*   [Function]
> Returns a string with character substitutions appropriate to send *txt* as an *attribute-value*.

`html:plain` *txt*   [Function]
> Returns a string with character substitutions appropriate to send *txt* as an *plain-text*.

`html:meta` *name content*   [Function]
> Returns a tag of meta-information suitable for passing as the third argument to `html:head`. The tag produced is '`<META NAME="`*name*`" CONTENT="`*content*`">`'. The string or symbol *name* can be '`author`', '`copyright`', '`keywords`', '`description`', '`date`', '`robots`', ...

`html:http-equiv` *name content*   [Function]
> Returns a tag of HTTP information suitable for passing as the third argument to `html:head`. The tag produced is '`<META HTTP-EQUIV="`*name*`" CONTENT="`*content*`">`'. The string or symbol *name* can be '`Expires`', '`PICS-Label`', '`Content-Type`', '`Refresh`', ...

`html:meta-refresh` *delay uri*   [Function]
`html:meta-refresh` *delay*   [Function]
> Returns a tag suitable for passing as the third argument to `html:head`. If *uri* argument is supplied, then *delay* seconds after displaying the page with this tag, Netscape or IE browsers will fetch and display *uri*. Otherwise, *delay* seconds after displaying the page with this tag, Netscape or IE browsers will fetch and redisplay this page.

`html:head` *title backlink tags* . . .   [Function]
`html:head` *title backlink*   [Function]
`html:head` *title*   [Function]
> Returns header string for an HTML page named *title*. If *backlink* is a string, it is used verbatim between the '`H1`' tags; otherwise *title* is used. If string arguments *tags* ... are supplied, then they are included verbatim within the `<HEAD>` section.

`html:body` *body* . . .   [Function]
> Returns HTML string to end a page.

`html:pre` *line1 line* . . .   [Function]
> Returns the strings *line1*, *lines* as *PRE*formmated plain text (rendered in fixed-width font). Newlines are inserted between *line1*, *lines*. HTML tags ('`<tag>`') within *lines* will be visible verbatim.

`html:comment` *line1 line* . . .   [Function]
> Returns the strings *line1* as HTML comments.

4.6 HTML Forms

`html:form` *method action body* . . .   [Function]
> The symbol *method* is either `get`, `head`, `post`, `put`, or `delete`. The strings *body* form the body of the form. `html:form` returns the HTML *form*.

html:hidden *name value*   [Function]
    Returns HTML string which will cause *name=value* in form.

html:checkbox *pname default*   [Function]
    Returns HTML string for check box.

html:text *pname default size . . .*   [Function]
    Returns HTML string for one-line text box.

html:text-area *pname default-list*   [Function]
    Returns HTML string for multi-line text box.

html:select *pname arity default-list foreign-values*   [Function]
    Returns HTML string for pull-down menu selector.

html:buttons *pname arity default-list foreign-values*   [Function]
    Returns HTML string for any-of selector.

form:submit *submit-label command*   [Function]
form:submit *submit-label*   [Function]
    The string or symbol *submit-label* appears on the button which submits the form.
    If the optional second argument *command* is given, then `*command*=`*command* and
    `*button*=`*submit-label* are set in the query. Otherwise, `*command*=`*submit-label*
    is set in the query.

form:image *submit-label image-src*   [Function]
    The *image-src* appears on the button which submits the form.

form:reset   [Function]
    Returns a string which generates a *reset* button.

form:element *pname arity default-list foreign-values*   [Function]
    Returns a string which generates an INPUT element for the field named *pname*. The
    element appears in the created form with its representation determined by its *arity*
    and domain. For domains which are foreign-keys:

    `single`   select menu

    `optional`   select menu

    `nary`   check boxes

    `nary1`   check boxes

    If the foreign-key table has a field named '`visible-name`', then the contents of that
    field are the names visible to the user for those choices. Otherwise, the foreign-key
    itself is visible.

    For other types of domains:

    `single`   text area

    `optional`   text area

```
boolean     check box

nary        text area

nary1       text area
```

form:delimited *pname doc aliat arity default-list foreign-values*   [Function]
> Returns a HTML string for a form element embedded in a line of a delimited list.
> Apply map `form:delimited` to the list returned by `command->p-specs`.

html:delimited-list *row ...*   [Function]
> Wraps its arguments with delimited-list ('`DL`' command.

get-foreign-choices *tab*   [Function]
> Returns a list of the '`visible-name`' or first fields of table *tab*.

command->p-specs *rdb command-table command*   [Function]
> The symbol *command-table* names a command table in the *rdb* relational database.
> The symbol *command* names a key in *command-table*.
>
> `command->p-specs` returns a list of lists of *pname*, *doc*, *aliat*, *arity*, *default-list*, and
> *foreign-values*. The returned list has one element for each parameter of command
> *command*.
>
> This example demonstrates how to create a HTML-form for the '`build`' command.

```
      (require (in-vicinity (implementation-vicinity) "build.scm"))
      (call-with-output-file "buildscm.html"
        (lambda (port)
          (display
           (string-append
            (html:head 'commands)
            (html:body
             (sprintf #f "<H2>%s:</H2><BLOCKQUOTE>%s</BLOCKQUOTE>\\n"
                      (html:plain 'build)
                      (html:plain ((comtab 'get 'documentation) 'build)))
            (html:form
             'post
             (or "http://localhost:8081/buildscm" "/cgi-bin/build.cgi")
             (apply html:delimited-list
                    (apply map form:delimited
                           (command->p-specs build '*commands* 'build)))
            (form:submit 'build)
            (form:reset))))
          port)))
```

4.7 HTML Tables

`(require 'db->html)`

html:table *options row ...*   [Function]

html:caption *caption align*   [Function]

html:caption *caption*   [Function]
> *align* can be 'top' or 'bottom'.

html:heading *columns*   [Function]
> Outputs a heading row for the currently-started table.

html:href-heading *columns uris*   [Function]
> Outputs a heading row with column-names *columns* linked to URIs *uris*.

html:linked-row-converter *k foreigns*   [Function]
> The positive integer *k* is the primary-key-limit (number of primary-keys) of the table. *foreigns* is a list of the filenames of foreign-key field pages and #f for non foreign-key fields.
>
> **html:linked-row-converter** returns a procedure taking a row for its single argument. This returned procedure returns the html string for that table row.

table-name->filename *table-name*   [Function]
> Returns the symbol *table-name* converted to a filename.

table->linked-html *caption db table-name match-key1 . . .*   [Function]
> Returns HTML string for *db* table *table-name* chopped into 50-row HTML tables. Every foreign-key value is linked to the page (of the table) defining that key.
>
> The optional *match-key1 . . .* arguments restrict actions to a subset of the table. See Section 6.1.2 [Table Operations], page 165.

table->linked-page *db table-name index-filename arg . . .*   [Function]
> Returns a complete HTML page. The string *index-filename* names the page which refers to this one.
>
> The optional *args . . .* arguments restrict actions to a subset of the table. See Section 6.1.2 [Table Operations], page 165.

catalog->html *db caption arg . . .*   [Function]
> Returns HTML string for the catalog table of *db*.

4.7.1 HTML editing tables

A client can modify one row of an editable table at a time. For any change submitted, these routines check if that row has been modified during the time the user has been editing the form. If so, an error page results.

The behavior of edited rows is:

- If no fields are changed, then no change is made to the table.
- If the primary keys equal null-keys (parameter defaults), and no other user has modified that row, then that row is deleted.
- If only primary keys are changed, there are non-key fields, and no row with the new keys is in the table, then the old row is deleted and one with the new keys is inserted.
- If only non-key fields are changed, and that row has not been modified by another user, then the row is changed to reflect the fields.

- If both keys and non-key fields are changed, and no row with the new keys is in the table, then a row is created with the new keys and fields.
- If fields are changed, all fields are primary keys, and no row with the new keys is in the table, then a row is created with the new keys.

After any change to the table, a `sync-database` of the database is performed.

command:modify-table *table-name null-keys update delete retrieve*   [Function]
command:modify-table *table-name null-keys update delete*   [Function]
command:modify-table *table-name null-keys update*   [Function]
command:modify-table *table-name null-keys*   [Function]

> Returns procedure (of *db*) which returns procedure to modify row of *table-name*. *null-keys* is the list of *null* keys indicating the row is to be deleted when any matches its corresponding primary key. Optional arguments *update*, *delete*, and *retrieve* default to the `row:update`, `row:delete`, and `row:retrieve` of *table-name* in *db*.

command:make-editable-table *rdb table-name arg* . . .   [Function]

> Given *table-name* in *rdb*, creates parameter and `*command*` tables for editing one row of *table-name* at a time. `command:make-editable-table` returns a procedure taking a row argument which returns the HTML string for editing that row.
>
> Optional *args* are expressions (lists) added to the call to `command:modify-table`.
>
> The domain name of a column determines the expected arity of the data stored in that column. Domain names ending in:

> '*'   have arity '`nary`';
>
> '+'   have arity '`nary1`'.

html:editable-row-converter *k names edit-point edit-converter*   [Function]

> The positive integer *k* is the primary-key-limit (number of primary-keys) of the table. *names* is a list of the field-names. *edit-point* is the list of primary-keys denoting the row to edit (or #f). *edit-converter* is the procedure called with *k*, *names*, and the row to edit.
>
> `html:editable-row-converter` returns a procedure taking a row for its single argument. This returned procedure returns the html string for that table row.
>
> Each HTML table constructed using `html:editable-row-converter` has first *k* fields (typically the primary key fields) of each row linked to a text encoding of these fields (the result of calling `row->anchor`). The page so referenced typically allows the user to edit fields of that row.

4.7.2 HTML databases

db->html-files *db dir index-filename caption*   [Function]

> *db* must be a relational database. *dir* must be #f or a non-empty string naming an existing sub-directory of the current directory.
>
> `db->html-files` creates an html page for each table in the database *db* in the subdirectory named *dir*, or the current directory if *dir* is #f. The top level page with the catalog of tables (captioned *caption*) is written to a file named *index-filename*.

`db->html-directory` *db dir index-filename*   [Function]
`db->html-directory` *db dir*   [Function]
> *db* must be a relational database. *dir* must be a non-empty string naming an existing
> sub-directory of the current directory or one to be created. The optional string *index-
> filename* names the filename of the top page, which defaults to `index.html`.
>
> `db->html-directory` creates sub-directory *dir* if neccessary, and calls (`db->html-
> files` *db dir index-filename dir*). The '`file:`' URI of *index-filename* is returned.

`db->netscape` *db dir index-filename*   [Function]
`db->netscape` *db dir*   [Function]
> `db->netscape` is just like `db->html-directory`, but calls `browse-url` with the uri
> for the top page after the pages are created.

4.8 HTTP and CGI

(`require 'http`) or (`require 'cgi`)

`http:header` *alist*   [Function]
> Returns a string containing lines for each element of *alist*; the `car` of which is followed
> by ': ', then the `cdr`.

`http:content` *alist body . . .*   [Function]
> Returns the concatenation of strings *body* with the (`http:header` *alist*) and the
> '`Content-Length`' prepended.

`*http:byline*`   [Variable]
> String appearing at the bottom of error pages.

`http:error-page` *status-code reason-phrase html-string . . .*   [Function]
> *status-code* and *reason-phrase* should be an integer and string as specified in *RFC
> 2068*. The returned page (string) will show the *status-code* and *reason-phrase* and
> any additional *html-strings . . .*; with **http:byline** or SLIB's default at the bottom.

`http:forwarding-page` *title dly uri html-string . . .*   [Function]
> The string or symbol *title* is the page title. *dly* is a non-negative integer. The *html-
> strings . . .* are typically used to explain to the user why this page is being forwarded.
>
> `http:forwarding-page` returns an HTML string for a page which automatically for-
> wards to *uri* after *dly* seconds. The returned page (string) contains any *html-strings
> . . .* followed by a manual link to *uri*, in case the browser does not forward automati-
> cally.

`http:serve-query` *serve-proc input-port output-port*   [Function]
> reads the *URI* and *query-string* from *input-port*. If the query is a valid '`"POST"`'
> or '`"GET"`' query, then `http:serve-query` calls *serve-proc* with three arguments,
> the *request-line*, *query-string*, and *header-alist*. Otherwise, `http:serve-query` calls
> *serve-proc* with the *request-line*, #f, and *header-alist*.
>
> If *serve-proc* returns a string, it is sent to *output-port*. If *serve-proc* returns a list
> whose first element is an integer, then an error page with the status integer which

is the first element of the list and strings from the list. If *serve-proc* returns a list whose first element isn't an number, then an error page with the status code 500 and strings from the list. If *serve-proc* returns #f, then a '`Bad Request`' (400) page is sent to *output-port*.

Otherwise, `http:serve-query` replies (to *output-port*) with appropriate HTML describing the problem.

This example services HTTP queries from *port-number*:

```
(define socket (make-stream-socket AF_INET 0))
(and (socket:bind socket port-number) ; AF_INET INADDR_ANY
     (socket:listen socket 10)        ; Queue up to 10 requests.
     (dynamic-wind
        (lambda () #f)
        (lambda ()
          (do ((port (socket:accept socket) (socket:accept socket)))
              ((#f)
            (let ((iport (duplicate-port port "r"))
                  (oport (duplicate-port port "w")))
              (http:serve-query build:serve iport oport)
              (close-port iport)
              (close-port oport))
            (close-port port)))
        (lambda () (close-port socket)))))
```

`cgi:serve-query` *serve-proc*   [Function]
> reads the *URI* and *query-string* from (`current-input-port`). If the query is a valid '`"POST"`' or '`"GET"`' query, then `cgi:serve-query` calls *serve-proc* with three arguments, the *request-line*, *query-string*, and *header-alist*. Otherwise, `cgi:serve-query` calls *serve-proc* with the *request-line*, #f, and *header-alist*.
>
> If *serve-proc* returns a string, it is sent to (`current-ouput-port`). If *serve-proc* returns a list whose first element is an integer, then an error page with the status integer which is the first element of the list and strings from the list. If *serve-proc* returns a list whose first element isn't an number, then an error page with the status code 500 and strings from the list. If *serve-proc* returns #f, then a '`Bad Request`' (400) page is sent to (`current-ouput-port`).
>
> Otherwise, `cgi:serve-query` replies (to (`current-output-port`)) with appropriate HTML describing the problem.

`make-query-alist-command-server` *rdb command-table*   [Function]
`make-query-alist-command-server` *rdb command-table #t*   [Function]
> Returns a procedure of one argument. When that procedure is called with a *query-alist* (as returned by `uri:decode-query`, the value of the '`*command*`' association will be the command invoked in *command-table*. If '`*command*`' is not in the *query-alist* then the value of '`*suggest*`' is tried. If neither name is in the *query-alist*, then the literal value '`*default*`' is tried in *command-table*.

If optional third argument is non-false, then the command is called with just the parameter-list; otherwise, command is called with the arguments described in its table.

4.9 Parsing HTML

`(require 'html-for-each)`

`html-for-each` *file word-proc markup-proc white-proc newline-proc*   [Function]
file is an input port or a string naming an existing file containing HTML text. *word-proc* is a procedure of one argument or #f. *markup-proc* is a procedure of one argument or #f. *white-proc* is a procedure of one argument or #f. *newline-proc* is a procedure of no arguments or #f.

`html-for-each` opens and reads characters from port *file* or the file named by string *file*. Sequential groups of characters are assembled into strings which are either

- enclosed by '`<`' and '`>`' (hypertext markups or comments);
- end-of-line;
- whitespace; or
- none of the above (words).

Procedures are called according to these distinctions in order of the string's occurrence in *file*.

newline-proc is called with no arguments for end-of-line *not within a markup or comment*.

white-proc is called with strings of non-newline whitespace.

markup-proc is called with hypertext markup strings (including '`<`' and '`>`').

word-proc is called with the remaining strings.

`html-for-each` returns an unspecified value.

`html:read-title` *file limit*   [Function]
`html:read-title` *file*   [Function]
file is an input port or a string naming an existing file containing HTML text. If supplied, *limit* must be an integer. *limit* defaults to 1000.

`html:read-title` opens and reads HTML from port *file* or the file named by string *file*, until reaching the (mandatory) 'TITLE' field. `html:read-title` returns the title string with adjacent whitespaces collapsed to one space. `html:read-title` returns #f if the title field is empty, absent, if the first character read from *file* is not '`#\<`', or if the end of title is not found within the first (approximately) *limit* words.

`htm-fields` *htm*   [Function]
htm is a hypertext markup string.

If *htm* is a (hypertext) comment or DTD, then `htm-fields` returns #f. Otherwise `htm-fields` returns the hypertext element string consed onto an association list of the attribute name-symbols and values. If the tag ends with `"/>"`, then `"/"` is appended to the hypertext element string. The name-symbols are created by `string-ci->symbol`. Each value is a string; or #t if the name had no value assigned within the markup.

4.10 URI

`(require 'uri)`

Implements *Uniform Resource Identifiers* (URI) as described in RFC 2396.

`make-uri`   [Function]
`make-uri` *fragment*   [Function]
`make-uri` *query fragment*   [Function]
`make-uri` *path query fragment*   [Function]
`make-uri` *authority path query fragment*   [Function]
`make-uri` *scheme authority path query fragment*   [Function]
> Returns a Uniform Resource Identifier string from component arguments.

`uri:make-path` *path*   [Function]
> Returns a URI string combining the components of list *path*.

`html:anchor` *name*   [Function]
> Returns a string which defines this location in the (HTML) file as *name*. The hyper-
> text '`<A HREF="#name">`' will link to this point.
>
>   (html:anchor "(section 7)")
>   ⇒
>   "<A NAME=\"(section%207)\"></A>"

`html:link` *uri highlighted*   [Function]
> Returns a string which links the *highlighted* text to *uri*.
>
>   (html:link (make-uri "(section 7)") "section 7")
>   ⇒
>   "<A HREF=\"#(section%207)\">section 7</A>"

`html:base` *uri*   [Function]
> Returns a string specifying the *base uri* of a document, for inclusion in the HEAD of
> the document (see Section 4.5 [HTML], page 69).

`html:isindex` *prompt*   [Function]
> Returns a string specifying the search *prompt* of a document, for inclusion in the
> HEAD of the document (see Section 4.5 [HTML], page 69).

`uri->tree` *uri-reference base-tree*   [Function]
`uri->tree` *uri-reference*   [Function]
> Returns a list of 5 elements corresponding to the parts (*scheme authority path query
> fragment*) of string *uri-reference*. Elements corresponding to absent parts are #f.
>
> The *path* is a list of strings. If the first string is empty, then the path is absolute;
> otherwise relative. The optional *base-tree* is a tree as returned by `uri->tree`; and is
> used as the base address for relative URIs.
>
> If the *authority* component is a *Server-based Naming Authority*, then it is a list of
> the *userinfo*, *host*, and *port* strings (or #f). For other types of *authority* components
> the *authority* will be a string.
>
>   (uri->tree "http://www.ics.uci.edu/pub/ietf/uri/#Related")

$\Rightarrow$
```
(http "www.ics.uci.edu" ("" "pub" "ietf" "uri" "") #f "Related")
```

uri:split-fields *txt chr*   [Function]

> Returns a list of *txt* split at each occurrence of *chr*. *chr* does not appear in the returned list of strings.

uri:decode-query *query-string*   [Function]

> Converts a *URI* encoded *query-string* to a query-alist.

`uric:` prefixes indicate procedures dealing with URI-components.

uric:encode *uri-component allows*   [Function]

> Returns a copy of the string *uri-component* in which all *unsafe* octets (as defined in RFC 2396) have been '%' escaped. `uric:decode` decodes strings encoded by `uric:encode`.

uric:decode *uri-component*   [Function]

> Returns a copy of the string *uri-component* in which each '%' escaped characters in *uri-component* is replaced with the character it encodes. This routine is useful for showing URI contents on error pages.

uri:path->keys *path-list ptypes*   [Function]

> *path-list* is a path-list as returned by `uri:split-fields`. `uri:path->keys` returns a list of items returned by `uri:decode-path`, coerced to types *ptypes*.

File-system Locators and Predicates

path->uri *path*   [Function]

> Returns a URI-string for *path* on the local host.

absolute-uri? *str*   [Function]

> Returns #t if *str* is an absolute-URI as indicated by a syntactically valid (per RFC 2396) *scheme*; otherwise returns #f.

absolute-path? *file-name*   [Function]

> Returns #t if *file-name* is a fully specified pathname (does not depend on the current working directory); otherwise returns #f.

null-directory? *str*   [Function]

> Returns #t if changing directory to *str* would leave the current directory unchanged; otherwise returns #f.

glob-pattern? *str*   [Function]

> Returns #t if the string *str* contains characters used for specifying glob patterns, namely '*', '?', or '['.

Before RFC 2396, the *File Transfer Protocol* (FTP) served a similar purpose.

`parse-ftp-address` *uri*   [Function]

> Returns a list of the decoded FTP *uri*; or #f if indecipherable. FTP *Uniform Resource Locator*, *ange-ftp*, and *getit* formats are handled. The returned list has four elements which are strings or #f:
>
>   0. username
>
>   1. password
>
>   2. remote-site
>
>   3. remote-directory

4.11 Parsing XML

(`require 'xml-parse`) or (`require 'ssax`)

The XML standard document referred to in this module is
`http://www.w3.org/TR/1998/REC-xml-19980210.html`.

The present frameworks fully supports the XML Namespaces Recommendation
`http://www.w3.org/TR/REC-xml-names`.

4.11.1 String Glue

`ssax:reverse-collect-str` *list-of-frags*   [Function]

> Given the list of fragments (some of which are text strings), reverse the list and concatenate adjacent text strings. If LIST-OF-FRAGS has zero or one element, the result of the procedure is `equal?` to its argument.

`ssax:reverse-collect-str-drop-ws` *list-of-frags*   [Function]

> Given the list of fragments (some of which are text strings), reverse the list and concatenate adjacent text strings while dropping "unsignificant" whitespace, that is, whitespace in front, behind and between elements. The whitespace that is included in character data is not affected.
>
> Use this procedure to "intelligently" drop "insignificant" whitespace in the parsed SXML. If the strict compliance with the XML Recommendation regarding the whitespace is desired, use the `ssax:reverse-collect-str` procedure instead.

4.11.2 Character and Token Functions

The following functions either skip, or build and return tokens, according to inclusion or delimiting semantics. The list of characters to expect, include, or to break at may vary from one invocation of a function to another. This allows the functions to easily parse even context-sensitive languages.

Exceptions are mentioned specifically. The list of expected characters (characters to skip until, or break-characters) may include an EOF "character", which is coded as symbol *eof*

The input stream to parse is specified as a PORT, which is the last argument.

`ssax:assert-current-char` *char-list string port*   [Function]
> Reads a character from the *port* and looks it up in the *char-list* of expected characters.
> If the read character was found among expected, it is returned. Otherwise, the
> procedure writes a message using *string* as a comment and quits.

`ssax:skip-while` *char-list port*   [Function]
> Reads characters from the *port* and disregards them, as long as they are mentioned
> in the *char-list*. The first character (which may be EOF) peeked from the stream that
> is *not* a member of the *char-list* is returned.

`ssax:init-buffer`   [Function]
> Returns an initial buffer for `ssax:next-token*` procedures. `ssax:init-buffer` may
> allocate a new buffer at each invocation.

`ssax:next-token` *prefix-char-list break-char-list comment-string port*   [Function]
> Skips any number of the prefix characters (members of the *prefix-char-list*), if any,
> and reads the sequence of characters up to (but not including) a break character, one
> of the *break-char-list*.
>
> The string of characters thus read is returned. The break character is left on the
> input stream. *break-char-list* may include the symbol `*eof*`; otherwise, EOF is fatal,
> generating an error message including a specified *comment-string*.

`ssax:next-token-of` is similar to `ssax:next-token` except that it implements an inclusion
rather than delimiting semantics.

`ssax:next-token-of` *inc-charset port*   [Function]
> Reads characters from the *port* that belong to the list of characters *inc-charset*. The
> reading stops at the first character which is not a member of the set. This character
> is left on the stream. All the read characters are returned in a string.

`ssax:next-token-of` *pred port*   [Function]
> Reads characters from the *port* for which *pred* (a procedure of one argument) returns
> non-#f. The reading stops at the first character for which *pred* returns #f. That
> character is left on the stream. All the results of evaluating of *pred* up to #f are
> returned in a string.
>
> *pred* is a procedure that takes one argument (a character or the EOF object) and
> returns a character or #f. The returned character does not have to be the same as
> the input argument to the *pred*. For example,
>
> ```
> (ssax:next-token-of (lambda (c)
>                       (cond ((eof-object? c) #f)
>                             ((char-alphabetic? c) (char-downcase c))
>                             (else #f)))
>                     (current-input-port))
> ```
>
> will try to read an alphabetic token from the current input port, and return it in
> lower case.

`ssax:read-string` *len port*   [Function]

> Reads *len* characters from the *port*, and returns them in a string. If EOF is encountered before *len* characters are read, a shorter string will be returned.

4.11.3 Data Types

`TAG-KIND`

> A symbol 'START', 'END', 'PI', 'DECL', 'COMMENT', 'CDSECT', or 'ENTITY-REF' that identifies a markup token

`UNRES-NAME`

> a name (called GI in the XML Recommendation) as given in an XML document for a markup token: start-tag, PI target, attribute name. If a GI is an NCName, UNRES-NAME is this NCName converted into a Scheme symbol. If a GI is a QName, 'UNRES-NAME' is a pair of symbols: (*PREFIX . LOCALPART*).

`RES-NAME`

> An expanded name, a resolved version of an 'UNRES-NAME'. For an element or an attribute name with a non-empty namespace URI, 'RES-NAME' is a pair of symbols, (*URI-SYMB . LOCALPART*). Otherwise, it's a single symbol.

`ELEM-CONTENT-MODEL`

> A symbol:

> 'ANY'   anything goes, expect an END tag.

> 'EMPTY-TAG'
> > no content, and no END-tag is coming

> 'EMPTY'   no content, expect the END-tag as the next token

> 'PCDATA'   expect character data only, and no children elements

> 'MIXED'

> 'ELEM-CONTENT'

`URI-SYMB`

> A symbol representing a namespace URI – or other symbol chosen by the user to represent URI. In the former case, URI-SYMB is created by %-quoting of bad URI characters and converting the resulting string into a symbol.

`NAMESPACES`

> A list representing namespaces in effect. An element of the list has one of the following forms:

> (*prefix uri-symb . uri-symb*) or
> (*prefix user-prefix . uri-symb*)
> > *user-prefix* is a symbol chosen by the user to represent the URI.

> (#f *user-prefix . uri-symb*)
> > Specification of the user-chosen prefix and a URI-SYMBOL.

> (*DEFAULT* *user-prefix . uri-symb*)
> > Declaration of the default namespace

(*DEFAULT* #f . #f)

> Un-declaration of the default namespace. This notation represents overriding of the previous declaration

A NAMESPACES list may contain several elements for the same *prefix*. The one closest to the beginning of the list takes effect.

ATTLIST

> An ordered collection of (*NAME* . *VALUE*) pairs, where *NAME* is a RES-NAME or an UNRES-NAME. The collection is an ADT.

STR-HANDLER

> A procedure of three arguments: *string1 string2 seed* returning a new *seed*. The procedure is supposed to handle a chunk of character data *string1* followed by a chunk of character data *string2*. *string2* is a short string, often '"\n"' and even '""'.

ENTITIES   An assoc list of pairs:

(`named-entity-name` . `named-entity-body`)

where *named-entity-name* is a symbol under which the entity was declared, *named-entity-body* is either a string, or (for an external entity) a thunk that will return an input port (from which the entity can be read). *named-entity-body* may also be #f. This is an indication that a *named-entity-name* is currently being expanded. A reference to this *named-entity-name* will be an error: violation of the WFC nonrecursion.

XML-TOKEN

> This record represents a markup, which is, according to the XML Recommendation, "takes the form of start-tags, end-tags, empty-element tags, entity references, character references, comments, CDATA section delimiters, document type declarations, and processing instructions."

> kind   a TAG-KIND

> head   an UNRES-NAME. For XML-TOKENs of kinds 'COMMENT and 'CDSECT, the head is #f.

> For example,

```
<P>                 => kind=START,      head=P
</P>                => kind=END,        head=P
<BR/>               => kind=EMPTY-EL,   head=BR
<!DOCTYPE OMF ...>  => kind=DECL,       head=DOCTYPE
<?xml version="1.0"?> => kind=PI,       head=xml
&my-ent;            => kind=ENTITY-REF, head=my-ent
```

> Character references are not represented by xml-tokens as these references are transparently resolved into the corresponding characters.

XML-DECL

> The record represents a datatype of an XML document: the list of declared elements and their attributes, declared notations, list of replacement strings or loading procedures for parsed general entities, etc. Normally an XML-DECL

record is created from a DTD or an XML Schema, although it can be created
and filled in in many other ways (e.g., loaded from a file).

elems   an (assoc) list of decl-elem or #f. The latter instructs the parser
to do no validation of elements and attributes.

decl-elem   declaration of one element:

(`elem-name elem-content decl-attrs`)

elem-name is an UNRES-NAME for the element.

elem-content is an ELEM-CONTENT-MODEL.

decl-attrs is an `ATTLIST`, of (`attr-name . value`) associations.

This element can declare a user procedure to handle parsing of an
element (e.g., to do a custom validation, or to build a hash of IDs
as they're encountered).

decl-attr   an element of an `ATTLIST`, declaration of one attribute:

(`attr-name content-type use-type default-value`)

attr-name is an UNRES-NAME for the declared attribute.

content-type is a symbol: `CDATA`, `NMTOKEN`, `NMTOKENS`, ... or a list
of strings for the enumerated type.

use-type is a symbol: `REQUIRED`, `IMPLIED`, or `FIXED`.

default-value is a string for the default value, or #f if not given.

4.11.4 Low-Level Parsers and Scanners

These procedures deal with primitive lexical units (Names, whitespaces, tags) and with
pieces of more generic productions. Most of these parsers must be called in appropriate
context. For example, `ssax:complete-start-tag` must be called only when the start-tag
has been detected and its GI has been read.

`ssax:skip-s` *port*   [Function]
Skip the S (whitespace) production as defined by

```
[3] S ::= (#x20 | #x09 | #x0D | #x0A)
```

`ssax:skip-s` returns the first not-whitespace character it encounters while scanning
the *port*. This character is left on the input stream.

`ssax:read-ncname` *port*   [Function]
Read a NCName starting from the current position in the *port* and return it as a
symbol.

```
[4] NameChar ::= Letter | Digit | '.' | '-' | '_' | ':'
                    | CombiningChar | Extender
[5] Name ::= (Letter | '_' | ':') (NameChar)*
```

This code supports the XML Namespace Recommendation REC-xml-names, which
modifies the above productions as follows:

```
[4] NCNameChar ::= Letter | Digit | '.' | '-' | '_'
                      | CombiningChar | Extender
```

```
[5] NCName ::= (Letter | '_') (NCNameChar)*
```

As the Rec-xml-names says,

> "An XML document conforms to this specification if all other tokens
> [other than element types and attribute names] in the document which
> are required, for XML conformance, to match the XML production for
> Name, match this specification's production for NCName."

Element types and attribute names must match the production QName, defined below.

ssax:read-qname *port*   [Function]

Read a (namespace-) Qualified Name, QName, from the current position in *port*; and
return an UNRES-NAME.

From REC-xml-names:

```
[6] QName ::= (Prefix ':')? LocalPart
[7] Prefix ::= NCName
[8] LocalPart ::= NCName
```

ssax:read-markup-token *port*   [Function]

This procedure starts parsing of a markup token. The current position in the stream
must be '<'. This procedure scans enough of the input stream to figure out what kind
of a markup token it is seeing. The procedure returns an XML-TOKEN structure
describing the token. Note, generally reading of the current markup is not finished!
In particular, no attributes of the start-tag token are scanned.

Here's a detailed break out of the return values and the position in the PORT when
that particular value is returned:

PI-token

> only PI-target is read. To finish the Processing-Instruction and disregard
> it, call `ssax:skip-pi`. `ssax:read-attributes` may be useful as well (for
> PIs whose content is attribute-value pairs).

END-token

> The end tag is read completely; the current position is right after the
> terminating '>' character.

COMMENT

> is read and skipped completely. The current position is right after '-->'
> that terminates the comment.

CDSECT

> The current position is right after '<!CDATA['. Use `ssax:read-cdata-body` to read the rest.

DECL

> We have read the keyword (the one that follows '<!') identifying this
> declaration markup. The current position is after the keyword (usually a
> whitespace character)

START-token

> We have read the keyword (GI) of this start tag. No attributes are
> scanned yet. We don't know if this tag has an empty content either.
> Use `ssax:complete-start-tag` to finish parsing of the token.

`ssax:skip-pi` *port*   [Function]
The current position is inside a PI. Skip till the rest of the PI

`ssax:read-pi-body-as-string` *port*   [Function]
The current position is right after reading the PITarget. We read the body of PI and
return is as a string. The port will point to the character right after '`?>`' combination
that terminates PI.

```
[16] PI ::= '<?' PITarget (S (Char* - (Char* '?>' Char*)))? '?>'
```

`ssax:skip-internal-dtd` *port*   [Function]
The current pos in the port is inside an internal DTD subset (e.g., after reading '`#\[`'
that begins an internal DTD subset) Skip until the '`]>`' combination that terminates
this DTD.

`ssax:read-cdata-body` *port str-handler seed*   [Function]
This procedure must be called after we have read a string '`<![CDATA[`' that begins a
CDATA section. The current position must be the first position of the CDATA body.
This function reads *lines* of the CDATA body and passes them to a *str-handler*, a
character data consumer.

str-handler is a procedure taking arguments: *string1*, *string2*, and *seed*. The first
string1 argument to *str-handler* never contains a newline; the second *string2* argu-
ment often will. On the first invocation of *str-handler*, *seed* is the one passed to
`ssax:read-cdata-body` as the third argument. The result of this first invocation
will be passed as the *seed* argument to the second invocation of the line consumer,
and so on. The result of the last invocation of the *str-handler* is returned by the
`ssax:read-cdata-body`. Note a similarity to the fundamental *fold* iterator.

Within a CDATA section all characters are taken at their face value, with three
exceptions:

- CR, LF, and CRLF are treated as line delimiters, and passed as a single
  '`#\newline`' to *str-handler*

- '`]]>`' combination is the end of the CDATA section. '`>`' is treated as an
  embedded '`>`' character.

- '`<`' and '`&`' are not specially recognized (and are not expanded)!

`ssax:read-char-ref` *port*   [Function]
```
[66]  CharRef ::=  '&#' [0-9]+ ';'
                 | '&#x' [0-9a-fA-F]+ ';'
```

This procedure must be called after we we have read '`&#`' that introduces a char
reference. The procedure reads this reference and returns the corresponding char.
The current position in PORT will be after the '`;`' that terminates the char reference.

Faults detected:
WFC: XML-Spec.html#wf-Legalchar

According to Section *4.1 Character and Entity References* of the XML Recommendation:

> "[Definition: A character reference refers to a specific character in the ISO/IEC 10646 character set, for example one not directly accessible from available input devices.]"

ssax:handle-parsed-entity *port name entities content-handler*   [Function]
    str-handler seed
Expands and handles a parsed-entity reference.

name is a symbol, the name of the parsed entity to expand. *content-handler* is a procedure of arguments *port*, *entities*, and *seed* that returns a seed. *str-handler* is called if the entity in question is a pre-declared entity.

ssax:handle-parsed-entity returns the result returned by *content-handler* or *str-handler*.

Faults detected:
WFC: XML-Spec.html#wf-entdeclared
WFC: XML-Spec.html#norecursion

attlist-add *attlist name-value*   [Function]
Add a *name-value* pair to the existing *attlist*, preserving its sorted ascending order; and return the new list. Return #f if a pair with the same name already exists in *attlist*

attlist-remove-top *attlist*   [Function]
Given an non-null *attlist*, return a pair of values: the top and the rest.

ssax:read-attributes *port entities*   [Function]
This procedure reads and parses a production *Attribute*.

```
[41] Attribute ::= Name Eq AttValue
[10] AttValue ::=  '"' ([^<&"] | Reference)* '"'
                 | "'" ([^<&'] | Reference)* "'"
[25] Eq ::= S? '=' S?
```

The procedure returns an ATTLIST, of Name (as UNRES-NAME), Value (as string) pairs. The current character on the *port* is a non-whitespace character that is not an NCName-starting character.

Note the following rules to keep in mind when reading an *AttValue*:

> Before the value of an attribute is passed to the application or checked for validity, the XML processor must normalize it as follows:
>
> - A character reference is processed by appending the referenced character to the attribute value.
> - An entity reference is processed by recursively processing the replacement text of the entity. The named entities 'amp', 'lt', 'gt', 'quot', and 'apos' are pre-declared.

- A whitespace character (#x20, #x0D, #x0A, #x09) is processed by appending #x20 to the normalized value, except that only a single #x20 is appended for a "#x0D#x0A" sequence that is part of an external parsed entity or the literal entity value of an internal parsed entity.

- Other characters are processed by appending them to the normalized value.

Faults detected:
WFC: XML-Spec.html#CleanAttrVals
WFC: XML-Spec.html#uniqattspec

`ssax:resolve-name` *port unres-name namespaces apply-default-ns?*   [Function]
Convert an *unres-name* to a RES-NAME, given the appropriate *namespaces* declarations. The last parameter, *apply-default-ns?*, determines if the default namespace applies (for instance, it does not for attribute names).

Per REC-xml-names/#nsc-NSDeclared, the "xml" prefix is considered pre-declared and bound to the namespace name "http://www.w3.org/XML/1998/namespace".

`ssax:resolve-name` tests for the namespace constraints:
`http://www.w3.org/TR/REC-xml-names/#nsc-NSDeclared`

`ssax:complete-start-tag` *tag port elems entities namespaces*   [Function]
Complete parsing of a start-tag markup. `ssax:complete-start-tag` must be called after the start tag token has been read. *tag* is an UNRES-NAME. *elems* is an instance of the ELEMS slot of XML-DECL; it can be #f to tell the function to do *no* validation of elements and their attributes.

`ssax:complete-start-tag` returns several values:

- ELEM-GI: a RES-NAME.

- ATTRIBUTES: element's attributes, an ATTLIST of (RES-NAME . STRING) pairs. The list does NOT include xmlns attributes.

- NAMESPACES: the input list of namespaces amended with namespace (re-)declarations contained within the start-tag under parsing

- ELEM-CONTENT-MODEL

On exit, the current position in *port* will be the first character after '>' that terminates the start-tag markup.

Faults detected:
VC: XML-Spec.html#enum
VC: XML-Spec.html#RequiredAttr
VC: XML-Spec.html#FixedAttr
VC: XML-Spec.html#ValueType
WFC: XML-Spec.html#uniqattspec (after namespaces prefixes are resolved)
VC: XML-Spec.html#elementvalid
WFC: REC-xml-names/#dt-NSName

Note: although XML Recommendation does not explicitly say it, xmlns and xmlns: attributes don't have to be declared (although they can be declared, to specify their default value).

`ssax:read-external-id` *port*   [Function]
>   Parses an ExternalID production:

>   [75] ExternalID ::= 'SYSTEM' S SystemLiteral
> | 'PUBLIC' S PubidLiteral S SystemLiteral
>   [11] SystemLiteral ::= ('"' [^"]* '"') | ("'" [^']* "'")
>   [12] PubidLiteral ::= '"' PubidChar* '"'
> | "'" (PubidChar - "'")* "'"
>   [13] PubidChar ::= #x20 | #x0D | #x0A | [a-zA-Z0-9]
> | [-'()+,./:=?;!*#@$_%]

>   Call `ssax:read-external-id` when an ExternalID is expected; that is, the current character must be either #\S or #\P that starts correspondingly a SYSTEM or PUBLIC token. `ssax:read-external-id` returns the *SystemLiteral* as a string. A *PubidLiteral* is disregarded if present.

4.11.5 Mid-Level Parsers and Scanners

These procedures parse productions corresponding to the whole (document) entity or its higher-level pieces (prolog, root element, etc).

`ssax:scan-misc` *port*   [Function]
>   Scan the Misc production in the context:

>   [1]   document ::= prolog element Misc*
>   [22] prolog ::= XMLDecl? Misc* (doctypedec l Misc*)?
>   [27] Misc ::= Comment | PI | S

>   Call `ssax:scan-misc` in the prolog or epilog contexts. In these contexts, whitespaces are completely ignored. The return value from `ssax:scan-misc` is either a PI-token, a DECL-token, a START token, or *EOF*. Comments are ignored and not reported.

`ssax:read-char-data` *port expect-eof? str-handler iseed*   [Function]
>   Read the character content of an XML document or an XML element.

>   [43] content ::=
>   (element | CharData | Reference | CDSect | PI | Comment)*

>   To be more precise, `ssax:read-char-data` reads CharData, expands CDSect and character entities, and skips comments. `ssax:read-char-data` stops at a named reference, EOF, at the beginning of a PI, or a start/end tag.

>   *expect-eof?* is a boolean indicating if EOF is normal; i.e., the character data may be terminated by the EOF. EOF is normal while processing a parsed entity.

>   *iseed* is an argument passed to the first invocation of *str-handler*.

>   `ssax:read-char-data` returns two results: *seed* and *token*. The *seed* is the result of the last invocation of *str-handler*, or the original *iseed* if *str-handler* was never called.

>   *token* can be either an eof-object (this can happen only if *expect-eof?* was #t), or:

- an xml-token describing a START tag or an END-tag; For a start token, the caller has to finish reading it.
- an xml-token describing the beginning of a PI. It's up to an application to read or skip through the rest of this PI;
- an xml-token describing a named entity reference.

CDATA sections and character references are expanded inline and never returned. Comments are silently disregarded.

As the XML Recommendation requires, all whitespace in character data must be preserved. However, a CR character (#x0D) must be disregarded if it appears before a LF character (#x0A), or replaced by a #x0A character otherwise. See Secs. 2.10 and 2.11 of the XML Recommendation. See also the canonical XML Recommendation.

ssax:assert-token *token kind gi error-cont*   [Function]
> Make sure that *token* is of anticipated *kind* and has anticipated *gi*. Note that the *gi* argument may actually be a pair of two symbols, Namespace-URI or the prefix, and of the localname. If the assertion fails, *error-cont* is evaluated by passing it three arguments: *token kind gi*. The result of *error-cont* is returned.

4.11.6 High-level Parsers

These procedures are to instantiate a SSAX parser. A user can instantiate the parser to do the full validation, or no validation, or any particular validation. The user specifies which PI he wants to be notified about. The user tells what to do with the parsed character and element data. The latter handlers determine if the parsing follows a SAX or a DOM model.

ssax:make-pi-parser *my-pi-handlers*   [Function]
> Create a parser to parse and process one Processing Element (PI).
>
> *my-pi-handlers* is an association list of pairs (`pi-tag` . `pi-handler`) where *pi-tag* is an NCName symbol, the PI target; and *pi-handler* is a procedure taking arguments *port*, *pi-tag*, and *seed*.
>
> *pi-handler* should read the rest of the PI up to and including the combination '`?>`' that terminates the PI. The handler should return a new seed. One of the *pi-tag*s may be the symbol `*DEFAULT*`. The corresponding handler will handle PIs that no other handler will. If the *DEFAULT* *pi-tag* is not specified, `ssax:make-pi-parser` will assume the default handler that skips the body of the PI.
>
> `ssax:make-pi-parser` returns a procedure of arguments *port*, *pi-tag*, and *seed*; that will parse the current PI according to *my-pi-handlers*.

ssax:make-elem-parser *my-new-level-seed my-finish-element*   [Function]
        my-char-data-handler my-pi-handlers
> Create a parser to parse and process one element, including its character content or children elements. The parser is typically applied to the root element of a document.
>
> *my-new-level-seed*
>> is a procedure taking arguments:
>>> *elem-gi attributes namespaces expected-content seed*

where *elem-gi* is a RES-NAME of the element about to be processed.

my-new-level-seed is to generate the seed to be passed to handlers that process the content of the element.

my-finish-element

is a procedure taking arguments:

elem-gi attributes namespaces parent-seed seed

my-finish-element is called when parsing of *elem-gi* is finished. The *seed* is the result from the last content parser (or from *my-new-level-seed* if the element has the empty content). *parent-seed* is the same seed as was passed to *my-new-level-seed*. *my-finish-element* is to generate a seed that will be the result of the element parser.

my-char-data-handler

is a STR-HANDLER as described in Data Types above.

my-pi-handlers

is as described for `ssax:make-pi-handler` above.

The generated parser is a procedure taking arguments:

start-tag-head port elems entities namespaces preserve-ws? seed

The procedure must be called after the start tag token has been read. *start-tag-head* is an UNRES-NAME from the start-element tag. ELEMS is an instance of ELEMS slot of XML-DECL.

Faults detected:
VC: XML-Spec.html#elementvalid
WFC: XML-Spec.html#GIMatch

`ssax:make-parser` *user-handler-tag user-handler . . .*   [Function]
Create an XML parser, an instance of the XML parsing framework. This will be a SAX, a DOM, or a specialized parser depending on the supplied user-handlers.

`ssax:make-parser` takes an even number of arguments; *user-handler-tag* is a symbol that identifies a procedure (or association list for `PROCESSING-INSTRUCTIONS`) (*user-handler*) that follows the tag. Given below are tags and signatures of the corresponding procedures. Not all tags have to be specified. If some are omitted, reasonable defaults will apply.

'`DOCTYPE`'   handler-procedure: *port docname systemid internal-subset? seed*

If *internal-subset?* is #t, the current position in the port is right after we have read '[' that begins the internal DTD subset. We must finish reading of this subset before we return (or must call `skip-internal-dtd` if we aren't interested in reading it). *port* at exit must be at the first symbol after the whole DOCTYPE declaration.

The handler-procedure must generate four values:

elems entities namespaces seed

elems is as defined for the ELEMS slot of XML-DECL. It may be #f to switch off validation. *namespaces* will typically contain *user-prefixes*

for selected *uri-symb*s. The default handler-procedure skips the internal subset, if any, and returns (`values #f '() '() seed`).

'`UNDECL-ROOT`'

>   procedure: *elem-gi seed*

>   where *elem-gi* is an UNRES-NAME of the root element. This procedure is called when an XML document under parsing contains *no* DOCTYPE declaration.

>   The handler-procedure, as a DOCTYPE handler procedure above, must generate four values:

>>   *elems entities namespaces seed*

>   The default handler-procedure returns (values #f '() '() seed)

'`DECL-ROOT`'

>   procedure: *elem-gi seed*

>   where *elem-gi* is an UNRES-NAME of the root element. This procedure is called when an XML document under parsing does contains the DOC-TYPE declaration. The handler-procedure must generate a new *seed* (and verify that the name of the root element matches the doctype, if the handler so wishes). The default handler-procedure is the identity function.

'`NEW-LEVEL-SEED`'

>   procedure: see ssax:make-elem-parser, my-new-level-seed

'`FINISH-ELEMENT`'

>   procedure: see ssax:make-elem-parser, my-finish-element

'`CHAR-DATA-HANDLER`'

>   procedure: see ssax:make-elem-parser, my-char-data-handler

'`PROCESSING-INSTRUCTIONS`'

>   association list as is passed to `ssax:make-pi-parser`. The default value is '()

The generated parser is a procedure of arguments *port* and *seed*.

This procedure parses the document prolog and then exits to an element parser (created by `ssax:make-elem-parser`) to handle the rest.

```
[1]  document ::=  prolog element Misc*
[22] prolog ::= XMLDecl? Misc* (doctypedec | Misc*)?
[27] Misc ::= Comment | PI |  S
[28] doctypedecl ::=  '<!DOCTYPE' S Name (S ExternalID)? S?
              ('[' (markupdecl | PEReference | S)* ']' S?)? '>'
[29] markupdecl ::= elementdecl | AttlistDecl
                        | EntityDecl
                        | NotationDecl | PI
                        | Comment
```

4.11.7 Parsing XML to SXML

`ssax:xml->sxml` *port namespace-prefix-assig* [Function]
> This is an instance of the SSAX parser that returns an SXML representation of the XML document to be read from *port*. *namespace-prefix-assig* is a list of (`user-prefix . uri-string`) that assigns *user-prefix*es to certain namespaces identified by particular *uri-string*s. It may be an empty list. `ssax:xml->sxml` returns an SXML tree. The port points out to the first character after the root element.

4.12 Printing Scheme

4.12.1 Generic-Write

`(require 'generic-write)`

    `generic-write` is a procedure that transforms a Scheme data value (or Scheme program expression) into its textual representation and prints it. The interface to the procedure is sufficiently general to easily implement other useful formatting procedures such as pretty printing, output to a string and truncated output.

`generic-write` *obj display? width output* [Procedure]

> *obj*   Scheme data value to transform.
>
> *display?*   Boolean, controls whether characters and strings are quoted.
>
> *width*   Extended boolean, selects format:
>
> > #f   single line format
> >
> > integer > 0
> > > pretty-print (value = max nb of chars per line)
>
> *output*   Procedure of 1 argument of string type, called repeatedly with successive substrings of the textual representation. This procedure can return `#f` to stop the transformation.

> The value returned by `generic-write` is undefined.

> Examples:

> > `(write obj)` ≡ `(generic-write obj #f #f display-string)`
> > `(display obj)` ≡ `(generic-write obj #t #f display-string)`

> where

> > `display-string` ≡
> > `(lambda (s) (for-each write-char (string->list s)) #t)`

4.12.2 Object-To-String

`(require 'object->string)`

`object->string` *obj* [Function]
> Returns the textual representation of *obj* as a string.

`object->limited-string` *obj limit* [Function]
> Returns the textual representation of *obj* as a string of length at most *limit*.

4.12.3 Pretty-Print

`(require 'pretty-print)`

pretty-print *obj*   [Procedure]

pretty-print *obj port*   [Procedure]

    `pretty-prints` *obj* on *port*. If *port* is not specified, `current-output-port` is used.

    Example:

```
(pretty-print '((1 2 3 4 5) (6 7 8 9 10) (11 12 13 14 15)
                (16 17 18 19 20) (21 22 23 24 25)))
   ⊣ ((1 2 3 4 5)
   ⊣  (6 7 8 9 10)
   ⊣  (11 12 13 14 15)
   ⊣  (16 17 18 19 20)
   ⊣  (21 22 23 24 25))
```

pretty-print->string *obj*   [Procedure]

pretty-print->string *obj width*   [Procedure]

    Returns the string of *obj* `pretty-print`ed in *width* columns. If *width* is not specified,
    `(output-port-width)` is used.

    Example:

```
(pretty-print->string '((1 2 3 4 5) (6 7 8 9 10) (11 12 13 14 15)
                        (16 17 18 19 20) (21 22 23 24 25)))
⇒
"((1 2 3 4 5)
 (6 7 8 9 10)
 (11 12 13 14 15)
 (16 17 18 19 20)
 (21 22 23 24 25))
"
```

```
(pretty-print->string '((1 2 3 4 5) (6 7 8 9 10) (11 12 13 14 15)
                        (16 17 18 19 20) (21 22 23 24 25))
                      16)
⇒
"((1 2 3 4 5)
 (6 7 8 9 10)
 (11
  12
  13
  14
  15)
 (16
  17
  18
  19
  20)
 (21
  22
  23
  24
  25))
"
```

```
(require 'pprint-file)
```

pprint-file *infile*   [Procedure]
pprint-file *infile outfile*   [Procedure]
> Pretty-prints all the code in *infile*. If *outfile* is specified, the output goes to *outfile*,
> otherwise it goes to (`current-output-port`).

pprint-filter-file *infile proc outfile*   [Function]
pprint-filter-file *infile proc*   [Function]
> *infile* is a port or a string naming an existing file. Scheme source code expressions and
> definitions are read from the port (or file) and *proc* is applied to them sequentially.
>
> *outfile* is a port or a string. If no *outfile* is specified then `current-output-port` is
> assumed. These expanded expressions are then `pretty-print`ed to this port.
>
> Whitepsace and comments (introduced by `;`) which are not part of scheme expressions
> are reproduced in the output. This procedure does not affect the values returned by
> `current-input-port`, `current-error-port`, and `current-output-port`.

   `pprint-filter-file` can be used to pre-compile macro-expansion and thus can re-
duce loading time. The following will write into `exp-code.scm` the result of expanding all
defmacros in `code.scm`.

```
(require 'pprint-file)
(require 'defmacroexpand)
(defmacro:load "my-macros.scm")
(pprint-filter-file "code.scm" defmacro:expand* "exp-code.scm")
```

4.13 Time and Date

If (provided? 'current-time):

The procedures `current-time`, `difftime`, and `offset-time` deal with a *calendar time* datatype which may or may not be disjoint from other Scheme datatypes.

current-time   [Function]
>   Returns the time since 00:00:00 GMT, January 1, 1970, measured in seconds. Note that the reference time is different from the reference time for `get-universal-time` in Section 4.13.3 [Common-Lisp Time], page 99.

difftime *caltime1 caltime0*   [Function]
>   Returns the difference (number of seconds) between twe calendar times: *caltime1 - caltime0*. *caltime0* may also be a number.

offset-time *caltime offset*   [Function]
>   Returns the calendar time of *caltime* offset by *offset* number of seconds (`+ caltime offset`).

4.13.1 Time Zone

(require 'time-zone)

TZ-string   [Data Format]
>   POSIX standards specify several formats for encoding time-zone rules.

> **:<*pathname*>**
>>   If the first character of <*pathname*> is '/', then <*pathname*> specifies the absolute pathname of a tzfile(5) format time-zone file. Otherwise, <*pathname*> is interpreted as a pathname within *tzfile:vicinity* (/usr/lib/zoneinfo/) naming a tzfile(5) format time-zone file.

> **<*std*><*offset*>**
>>   The string <*std*> consists of 3 or more alphabetic characters. <*offset*> specifies the time difference from GMT. The <*offset*> is positive if the local time zone is west of the Prime Meridian and negative if it is east. <*offset*> can be the number of hours or hours and minutes (and optionally seconds) separated by ':'. For example, `-4:30`.

> **<*std*><*offset*><*dst*>**
>>   <*dst*> is the at least 3 alphabetic characters naming the local daylight-savings-time.

> **<*std*><*offset*><*dst*><*doffset*>**
>>   <*doffset*> specifies the offset from the Prime Meridian when daylight-savings-time is in effect.

> The non-tzfile formats can optionally be followed by transition times specifying the day and time when a zone changes from standard to daylight-savings and back again.

> **,<*date*>/<*time*>,<*date*>/<*time*>**
>>   The <*time*>s are specified like the <*offset*>s above, except that leading '+' and '-' are not allowed.

Each **<***date***>** has one of the formats:

J<*day***>**   specifies the Julian day with **<***day***>** between 1 and 365. February 29 is never counted and cannot be referenced.

<*day***>**   This specifies the Julian day with n between 0 and 365. February 29 is counted in leap years and can be specified.

M<*month***>.<***week***>.<***day***>**

This specifies day **<***day***>** (0 <= **<***day***>** <= 6) of week **<***week***>** (1 <= **<***week***>** <= 5) of month **<***month***>** (1 <= **<***month***>** <= 12). Week 1 is the first week in which day d occurs and week 5 is the last week in which day **<***day***>** occurs. Day 0 is a Sunday.

time-zone   [Data Type]

is a datatype encoding how many hours from Greenwich Mean Time the local time is, and the *Daylight Savings Time* rules for changing it.

time-zone *TZ-string*   [Function]

Creates and returns a time-zone object specified by the string *TZ-string*. If **time-zone** cannot interpret *TZ-string*, **#f** is returned.

tz:params *caltime tz*   [Function]

tz is a time-zone object. **tz:params** returns a list of three items:

0. An integer. 0 if standard time is in effect for timezone *tz* at *caltime*; 1 if daylight savings time is in effect for timezone *tz* at *caltime*.

1. The number of seconds west of the Prime Meridian timezone *tz* is at *caltime*.

2. The name for timezone *tz* at *caltime*.

tz:params is unaffected by the default timezone; inquiries can be made of any timezone at any calendar time.

tz:std-offset *tz*   [Function]

tz is a time-zone object. **tz:std-offset** returns the number of seconds west of the Prime Meridian timezone *tz* is.

The rest of these procedures and variables are provided for POSIX compatability. Because of shared state they are not thread-safe.

tzset   [Function]

Returns the default time-zone.

tzset *tz*   [Function]

Sets (and returns) the default time-zone to *tz*.

tzset *TZ-string*   [Function]

Sets (and returns) the default time-zone to that specified by *TZ-string*.

tzset also sets the variables **timezone**, *daylight?*, and *tzname*. This function is automatically called by the time conversion procedures which depend on the time zone (see Section 4.13 [Time and Date], page 96).

timezone   [Variable]
>   Contains the difference, in seconds, between Greenwich Mean Time and local stan-
>   dard time (for example, in the U.S. Eastern time zone (EST), timezone is 5*60*60).
>   *timezone* is initialized by `tzset`.

daylight?   [Variable]
>   is `#t` if the default timezone has rules for *Daylight Savings Time. Note: daylight?*
>   does not tell you when Daylight Savings Time is in effect, just that the default zone
>   sometimes has Daylight Savings Time.

tzname   [Variable]
>   is a vector of strings. Index 0 has the abbreviation for the standard timezone; If
>   *daylight?*, then index 1 has the abbreviation for the Daylight Savings timezone.

4.13.2 Posix Time

>   `(require 'posix-time)`

Calendar-Time   [Data Type]
>   is a datatype encapsulating time.

Coordinated Universal Time   [Data Type]
>   (abbreviated *UTC*) is a vector of integers representing time:

>   0. seconds (0 - 61)
>   1. minutes (0 - 59)
>   2. hours since midnight (0 - 23)
>   3. day of month (1 - 31)
>   4. month (0 - 11). Note difference from `decode-universal-time`.
>   5. the number of years since 1900. Note difference from `decode-universal-time`.
>   6. day of week (0 - 6)
>   7. day of year (0 - 365)
>   8. 1 for daylight savings, 0 for regular time

gmtime *caltime*   [Function]
>   Converts the calendar time *caltime* to UTC and returns it.

localtime *caltime tz*   [Function]
>   Returns *caltime* converted to UTC relative to timezone *tz*.

localtime *caltime*   [Function]
>   converts the calendar time *caltime* to a vector of integers expressed relative to the
>   user's time zone. `localtime` sets the variable **timezone** with the difference be-
>   tween Coordinated Universal Time (UTC) and local standard time in seconds (see
>   Section 4.13.1 [Time Zone], page 96).

gmktime *univtime*   [Function]
>   Converts a vector of integers in GMT Coordinated Universal Time (UTC) format to
>   a calendar time.

mktime *univtime*   [Function]
>   Converts a vector of integers in local Coordinated Universal Time (UTC) format to
>   a calendar time.

mktime *univtime tz*   [Function]
>   Converts a vector of integers in Coordinated Universal Time (UTC) format (relative
>   to time-zone *tz*) to calendar time.

asctime *univtime*   [Function]
>   Converts the vector of integers *caltime* in Coordinated Universal Time (UTC) format
>   into a string of the form `"Wed Jun 30 21:49:08 1993"`.

gtime *caltime*   [Function]
ctime *caltime*   [Function]
ctime *caltime tz*   [Function]
>   Equivalent to `(asctime (gmtime caltime))`, `(asctime (localtime caltime))`, and
>   `(asctime (localtime caltime tz))`, respectively.

4.13.3 Common-Lisp Time

>   `(require 'common-lisp-time)`

get-decoded-time   [Function]
>   Equivalent to `(decode-universal-time (get-universal-time))`.

get-universal-time   [Function]
>   Returns the current time as *Universal Time*, number of seconds since 00:00:00 Jan 1,
>   1900 GMT. Note that the reference time is different from `current-time`.

decode-universal-time *univtime*   [Function]
>   Converts *univtime* to *Decoded Time* format. Nine values are returned:
>
>   0. seconds (0 - 61)
>
>   1. minutes (0 - 59)
>
>   2. hours since midnight
>
>   3. day of month
>
>   4. month (1 - 12). Note difference from `gmtime` and `localtime`.
>
>   5. year (A.D.). Note difference from `gmtime` and `localtime`.
>
>   6. day of week (0 - 6)
>
>   7. #t for daylight savings, #f otherwise
>
>   8. hours west of GMT (-24 - +24)
>
>   Notice that the values returned by `decode-universal-time` do not match the argu-
>   ments to `encode-universal-time`.

`encode-universal-time` *second minute hour date month year* [Function]
`encode-universal-time` *second minute hour date month year* [Function]
      time-zone

    Converts the arguments in Decoded Time format to Universal Time format. If *time-zone* is not specified, the returned time is adjusted for daylight saving time. Otherwise, no adjustment is performed.

    Notice that the values returned by `decode-universal-time` do not match the arguments to `encode-universal-time`.

4.13.4 ISO 8601

`(require 'iso-8601)`

`time->iso-8601` *time* [Function]
    time is the time in seconds since 00:00:00 GMT, January 1, 1970. `time->iso-8601` returns an expanded ISO 8601 format string for the date and time.

`time->iso8601` *time* [Function]
    time is a time in seconds since 00:00:00 GMT, January 1, 1970. `time->iso8601` returns a compact ISO 8601 format string for the date and time.

`iso-8601->time` *str* [Function]
    str is a string in ISO 8601 format, either compact or expanded. `iso-8601->time` returns that time in seconds since 00:00:00 GMT, January 1, 1970.

4.13.5 Time Infrastructure

`(require 'time-core)`

`time:gmtime` *tm* [Function]
`time:invert` *decoder target* [Function]
`time:split` *t tm_isdst tm_gmtoff tm_zone* [Function]
   `(require 'tzfile)`

`tzfile:read` *path* [Function]

4.14 NCBI-DNA

`(require 'ncbi-dma)`

`ncbi:read-dna-sequence` *port* [Function]
    Reads the NCBI-format DNA sequence following the word '`ORIGIN`' from *port*.

`ncbi:read-file` *file* [Function]
    Reads the NCBI-format DNA sequence following the word '`ORIGIN`' from *file*.

`mrna<-cdna` *str* [Function]
    Replaces 'T' with 'U' in *str*

`codons<-cdna` *cdna* [Function]
    Returns a list of three-letter symbol codons comprising the protein sequence encoded by *cdna* starting with its first occurence of '`atg`'.

`protein<-cdna` *cdna*   [Function]
> Returns a list of three-letter symbols for the protein sequence encoded by *cdna* starting with its first occurence of '`atg`'.

`p<-cdna` *cdna*   [Function]
> Returns a string of one-letter amino acid codes for the protein sequence encoded by *cdna* starting with its first occurence of '`atg`'.

These cDNA count routines provide a means to check the nucleotide sequence with the '`BASE COUNT`' line preceding the sequence from NCBI.

`cdna:base-count` *cdna*   [Function]
> Returns a list of counts of '`a`', '`c`', '`g`', and '`t`' occurrencing in *cdna*.

`cdna:report-base-count` *cdna*   [Function]
> Prints the counts of '`a`', '`c`', '`g`', and '`t`' occurrencing in *cdna*.

4.15 Schmooz

Schmooz is a simple, lightweight markup language for interspersing Texinfo documentation with Scheme source code. Schmooz does not create the top level Texinfo file; it creates '`txi`' files which can be imported into the documentation using the Texinfo command '`@include`'.

`(require 'schmooz)` defines the function `schmooz`, which is used to process files. Files containing schmooz documentation should not contain `(require 'schmooz)`.

`schmooz` *filename*.scm . . .   [Procedure]
> *Filename*.scm should be a string ending with '`.scm`' naming an existing file containing Scheme source code. `schmooz` extracts top-level comments containing schmooz commands from *filename*.scm and writes the converted Texinfo source to a file named *filename*.txi.

`schmooz` *filename*.texi . . .   [Procedure]
`schmooz` *filename*.tex . . .   [Procedure]
`schmooz` *filename*.txi . . .   [Procedure]
> *Filename* should be a string naming an existing file containing Texinfo source code. For every occurrence of the string '`@include filename.txi`' within that file, `schmooz` calls itself with the argument '`filename.scm`'.

Schmooz comments are distinguished (from non-schmooz comments) by their first line, which must start with an at-sign (@) preceded by one or more semicolons (;). A schmooz comment ends at the first subsequent line which does *not* start with a semicolon. Currently schmooz comments are recognized only at top level.

Schmooz comments are copied to the Texinfo output file with the leading contiguous semicolons removed. Certain character sequences starting with at-sign are treated specially. Others are copied unchanged.

A schmooz comment starting with '`@body`' must be followed by a Scheme definition. All comments between the '`@body`' line and the definition will be included in a Texinfo

definition, either a '`@defun`' or a '`@defvar`', depending on whether a procedure or a variable is being defined.

Within the text of that schmooz comment, at-sign followed by '`0`' will be replaced by `@code{procedure-name}` if the following definition is of a procedure; or `@var{variable}` if defining a variable.

An at-sign followed by a non-zero digit will expand to the variable citation of that numbered argument: '`@var{argument-name}`'.

If more than one definition follows a '`@body`' comment line without an intervening blank or comment line, then those definitions will be included in the same Texinfo definition using '`@defvarx`' or '`@defunx`', depending on whether the first definition is of a variable or of a procedure.

Schmooz can figure out whether a definition is of a procedure if it is of the form:

'`(define (<identifier> <arg> ...) <expression>)`'

or if the left hand side of the definition is some form ending in a lambda expression. Obviously, it can be fooled. In order to force recognition of a procedure definition, start the documentation with '`@args`' instead of '`@body`'. '`@args`' should be followed by the argument list of the function being defined, which may be enclosed in parentheses and delimited by whitespace, (as in Scheme), enclosed in braces and separated by commas, (as in Texinfo), or consist of the remainder of the line, separated by whitespace.

For example:

```
;;@args arg1 args ...
;;@0 takes argument @1 and any number of @2
(define myfun (some-function-returning-magic))
```

Will result in:

```
@defun myfun arg1 args @dots{}

@code{myfun} takes argument @var{arg1} and any number of @var{args}
@end defun
```

'`@args`' may also be useful for indicating optional arguments by name. If '`@args`' occurs inside a schmooz comment section, rather than at the beginning, then it will generate a '`@defunx`' line with the arguments supplied.

If the first at-sign in a schmooz comment is immediately followed by whitespace, then the comment will be expanded to whatever follows that whitespace. If the at-sign is followed by a non-whitespace character then the at-sign will be included as the first character of the expansion. This feature is intended to make it easy to include Texinfo directives in schmooz comments.

5 Mathematical Packages

5.1 Bit-Twiddling

`(require 'logical)` or `(require 'srfi-60)`

The bit-twiddling functions are made available through the use of the `logical` package. `logical` is loaded by inserting `(require 'logical)` before the code that uses these functions. These functions behave as though operating on integers in two's-complement representation.

5.1.1 Bitwise Operations

`logand` *n1 ...*   [Function]
`bitwise-and` *n1 ...*   [Function]
> Returns the integer which is the bit-wise AND of the integer arguments.
>
> Example:
>
> ```
> (number->string (logand #b1100 #b1010) 2)
>     ⇒ "1000"
> ```

`logior` *n1 ...*   [Function]
`bitwise-ior` *n1 ...*   [Function]
> Returns the integer which is the bit-wise OR of the integer arguments.
>
> Example:
>
> ```
> (number->string (logior #b1100 #b1010) 2)
>     ⇒ "1110"
> ```

`logxor` *n1 ...*   [Function]
`bitwise-xor` *n1 ...*   [Function]
> Returns the integer which is the bit-wise XOR of the integer arguments.
>
> Example:
>
> ```
> (number->string (logxor #b1100 #b1010) 2)
>     ⇒ "110"
> ```

`lognot` *n*   [Function]
`bitwise-not` *n*   [Function]
> Returns the integer which is the one's-complement of the integer argument.
>
> Example:
>
> ```
> (number->string (lognot #b10000000) 2)
>     ⇒ "-10000001"
> (number->string (lognot #b0) 2)
>     ⇒ "-1"
> ```

bitwise-if *mask n0 n1*   [Function]
bitwise-merge *mask n0 n1*   [Function]
> Returns an integer composed of some bits from integer *n0* and some from integer *n1*.
> A bit of the result is taken from *n0* if the corresponding bit of integer *mask* is 1 and
> from *n1* if that bit of *mask* is 0.

logtest *j k*   [Function]
any-bits-set? *j k*   [Function]
> ```
> (logtest j k) ≡ (not (zero? (logand j k)))
> ```
>
> ```
> (logtest #b0100 #b1011) ⇒ #f
> (logtest #b0100 #b0111) ⇒ #t
> ```

5.1.2 Integer Properties

logcount *n*   [Function]
> Returns the number of bits in integer *n*. If integer is positive, the 1-bits in its binary
> representation are counted. If negative, the 0-bits in its two's-complement binary
> representation are counted. If 0, 0 is returned.
>
> Example:
>
> ```
> (logcount #b10101010)
>     ⇒ 4
> (logcount 0)
>     ⇒ 0
> (logcount -2)
>     ⇒ 1
> ```

On `discuss@r6rs.org` Ben Harris credits Simon Tatham with the idea to have
`bitwise-bit-count` return a negative count for negative inputs. Alan Bawden came up
with the succinct invariant.

bitwise-bit-count *n*   [Function]
> If *n* is non-negative, this procedure returns the number of 1 bits in the two's-
> complement representation of *n*. Otherwise it returns the result of the following
> computation:
>
> ```
> (bitwise-not (bitwise-bit-count (bitwise-not n)))
> ```

integer-length *n*   [Function]
> Returns the number of bits neccessary to represent *n*.
>
> Example:
>
> ```
> (integer-length #b10101010)
>     ⇒ 8
> (integer-length 0)
>     ⇒ 0
> (integer-length #b1111)
>     ⇒ 4
> ```

log2-binary-factors *n*   [Function]
first-set-bit *n*   [Function]

Returns the number of factors of two of integer *n*. This value is also the bit-index of the least-significant '1' bit in *n*.

```
(require 'printf)
(do ((idx 0 (+ 1 idx)))
        ((> idx 16))
      (printf "%s(%3d) ==> %-5d %s(%2d) ==> %-5d\n"
                'log2-binary-factors
                (- idx) (log2-binary-factors (- idx))
                'log2-binary-factors
                idx (log2-binary-factors idx)))
⊣
log2-binary-factors(  0) ==> -1    log2-binary-factors( 0) ==> -1
log2-binary-factors( -1) ==> 0     log2-binary-factors( 1) ==> 0
log2-binary-factors( -2) ==> 1     log2-binary-factors( 2) ==> 1
log2-binary-factors( -3) ==> 0     log2-binary-factors( 3) ==> 0
log2-binary-factors( -4) ==> 2     log2-binary-factors( 4) ==> 2
log2-binary-factors( -5) ==> 0     log2-binary-factors( 5) ==> 0
log2-binary-factors( -6) ==> 1     log2-binary-factors( 6) ==> 1
log2-binary-factors( -7) ==> 0     log2-binary-factors( 7) ==> 0
log2-binary-factors( -8) ==> 3     log2-binary-factors( 8) ==> 3
log2-binary-factors( -9) ==> 0     log2-binary-factors( 9) ==> 0
log2-binary-factors(-10) ==> 1     log2-binary-factors(10) ==> 1
log2-binary-factors(-11) ==> 0     log2-binary-factors(11) ==> 0
log2-binary-factors(-12) ==> 2     log2-binary-factors(12) ==> 2
log2-binary-factors(-13) ==> 0     log2-binary-factors(13) ==> 0
log2-binary-factors(-14) ==> 1     log2-binary-factors(14) ==> 1
log2-binary-factors(-15) ==> 0     log2-binary-factors(15) ==> 0
log2-binary-factors(-16) ==> 4     log2-binary-factors(16) ==> 4
```

5.1.3 Bit Within Word

logbit? *index n*   [Function]
bit-set? *index n*   [Function]

```
(logbit? index n) ≡ (logtest (expt 2 index) n)

(logbit? 0 #b1101) ⇒ #t
(logbit? 1 #b1101) ⇒ #f
(logbit? 2 #b1101) ⇒ #t
(logbit? 3 #b1101) ⇒ #t
(logbit? 4 #b1101) ⇒ #f
```

copy-bit *index from bit*   [Function]

Returns an integer the same as *from* except in the *index*th bit, which is 1 if *bit* is #t and 0 if *bit* is #f.

Example:

```
(number->string (copy-bit 0 0 #t) 2)      ⇒ "1"
(number->string (copy-bit 2 0 #t) 2)      ⇒ "100"
(number->string (copy-bit 2 #b1111 #f) 2)  ⇒ "1011"
```

5.1.4 Field of Bits

bit-field *n start end*   [Function]

Returns the integer composed of the *start* (inclusive) through *end* (exclusive) bits of *n*. The *start*th bit becomes the 0-th bit in the result.

Example:

```
(number->string (bit-field #b1101101010 0 4) 2)
    ⇒ "1010"
(number->string (bit-field #b1101101010 4 9) 2)
    ⇒ "10110"
```

copy-bit-field *to from start end*   [Function]

Returns an integer the same as *to* except possibly in the *start* (inclusive) through *end* (exclusive) bits, which are the same as those of *from*. The 0-th bit of *from* becomes the *start*th bit of the result.

Example:

```
(number->string (copy-bit-field #b1101101010 0 0 4) 2)
        ⇒ "1101100000"
(number->string (copy-bit-field #b1101101010 -1 0 4) 2)
        ⇒ "1101101111"
(number->string (copy-bit-field #b110100100010000 -1 5 9) 2)
        ⇒ "110100111110000"
```

ash *n count*   [Function]
arithmetic-shift *n count*   [Function]

Returns an integer equivalent to (`inexact->exact (floor (* n (expt 2 count)))`).

Example:

```
(number->string (ash #b1 3) 2)
    ⇒ "1000"
(number->string (ash #b1010 -1) 2)
    ⇒ "101"
```

rotate-bit-field *n count start end*   [Function]

Returns *n* with the bit-field from *start* to *end* cyclically permuted by *count* bits towards high-order.

Example:

```
(number->string (rotate-bit-field #b0100 3 0 4) 2)
    ⇒ "10"
(number->string (rotate-bit-field #b0100 -1 0 4) 2)
    ⇒ "10"
(number->string (rotate-bit-field #b110100100010000 -1 5 9) 2)
    ⇒ "110100010010000"
```

```
(number->string (rotate-bit-field #b110100100010000 1 5 9) 2)
   ⇒ "110100000110000"
```

reverse-bit-field *n start end*   [Function]
> Returns *n* with the order of bits *start* to *end* reversed.

```
(number->string (reverse-bit-field #xa7 0 8) 16)
   ⇒ "e5"
```

5.1.5 Bits as Booleans

integer->list *k len*   [Function]
integer->list *k*   [Function]
> `integer->list` returns a list of *len* booleans corresponding to each bit of the non-negative integer *k*. #t is coded for each 1; #f for 0. The *len* argument defaults to (`integer-length k`).

list->integer *list*   [Function]
> `list->integer` returns an integer formed from the booleans in the list *list*, which must be a list of booleans. A 1 bit is coded for each #t; a 0 bit for #f.

```
(list->integer (integer->list k))  ⇒  k
```

booleans->integer *bool1 ...*   [Function]
> Returns the integer coded by the *bool1 ...* arguments.

5.2 Modular Arithmetic

(require 'modular)

extended-euclid *n1 n2*   [Function]
> Returns a list of 3 integers (`d x y`) such that d = gcd(*n1*, *n2*) = *n1* * x + *n2* * y.

symmetric:modulus *m*   [Function]
> For odd positive integer *m*, returns an object suitable for passing as the first argument to `modular:` procedures, directing them to return a symmetric modular number, ie. an *n* such that

```
(<= (quotient m -2) n (quotient m 2)
```

modular:characteristic *modulus*   [Function]
> Returns the non-negative integer characteristic of the ring formed when *modulus* is used with `modular:` procedures.

modular:normalize *modulus n*   [Function]
> Returns the integer (`modulo n (modular:characteristic modulus)`) in the representation specified by *modulus*.

The rest of these functions assume normalized arguments; That is, the arguments are constrained by the following table:

For all of these functions, if the first argument (*modulus*) is:

positive?
>   Integers mod *modulus*. The result is between 0 and *modulus*.

zero?   The arguments are treated as integers. An integer is returned.

Otherwise, if *modulus* is a value returned by (`symmetric:modulus radix`), then the arguments and result are treated as members of the integers modulo *radix*, but with *symmetric* representation; i.e.

>   (<= (quotient `radix` 2) `n` (quotient (- -1 `radix`) 2)

If all the arguments are fixnums the computation will use only fixnums.

modular:invertable? *modulus k*   [Function]
>   Returns #t if there exists an integer n such that $k * n \equiv 1$ mod *modulus*, and #f otherwise.

modular:invert *modulus n2*   [Function]
>   Returns an integer n such that $1 = (n * n2)$ mod *modulus*. If *n2* has no inverse mod *modulus* an error is signaled.

modular:negate *modulus n2*   [Function]
>   Returns $(-n2)$ mod *modulus*.

modular:+ *modulus n2 n3*   [Function]
>   Returns $(n2 + n3)$ mod *modulus*.

modular:- *modulus n2 n3*   [Function]
>   Returns $(n2 - n3)$ mod *modulus*.

modular:* *modulus n2 n3*   [Function]
>   Returns $(n2 * n3)$ mod *modulus*.

>   The Scheme code for modular:* with negative *modulus* is not completed for fixnum-only implementations.

modular:expt *modulus n2 n3*   [Function]
>   Returns $(n2 \ \hat{} \ n3)$ mod *modulus*.

5.3 Irrational Integer Functions

(require 'math-integer)

integer-expt *n1 n2*   [Function]
>   Returns *n1* raised to the power *n2* if that result is an exact integer; otherwise signals an error.

>   (integer-expt 0 *n2*)

>   returns 1 for *n2* equal to 0; returns 0 for positive integer *n2*; signals an error otherwise.

`integer-log` *base k* [Function]

> Returns the largest exact integer whose power of *base* is less than or equal to *k*. If *base* or *k* is not a positive exact integer, then `integer-log` signals an error.

`integer-sqrt` *k* [Function]

> For non-negative integer *k* returns the largest integer whose square is less than or equal to *k*; otherwise signals an error.

`quotient` *n1 n2* [Function]
`remainder` *n1 n2* [Function]
`modulo` *n1 n2* [Function]

> are redefined so that they accept only exact-integer arguments.

`round-quotient` *n1 n2* [Function]

> Returns the quotient of *n1* and *n2* rounded toward even.

>   (quotient 3 2)   ⇒ 1
>   (round-quotient 3 2) ⇒ 2

5.4 Irrational Real Functions

`(require 'math-real)`

   Although this package defines real and complex functions, it is safe to load into an integer-only implementation; those functions will be defined to #f.

`real-exp` *x* [Function]
`real-ln` *x* [Function]
`real-log` *y x* [Function]
`real-sin` *x* [Function]
`real-cos` *x* [Function]
`real-tan` *x* [Function]
`real-asin` *x* [Function]
`real-acos` *x* [Function]
`real-atan` *x* [Function]
`atan` *y x* [Function]

> These procedures are part of every implementation that supports general real numbers; they compute the usual transcendental functions. 'real-ln' computes the natural logarithm of *x*; 'real-log' computes the logarithm of *x* base *y*, which is (/ (real-ln x) (real-ln y)). If arguments *x* and *y* are not both real; or if the correct result would not be real, then these procedures signal an error.

`real-sqrt` *x* [Function]

> For non-negative real *x* the result will be its positive square root; otherwise an error will be signaled.

`real-expt` *x1 x2* [Function]

> Returns *x1* raised to the power *x2* if that result is a real number; otherwise signals an error.

```
(real-expt 0.0 x2)
```

- returns 1.0 for *x2* equal to 0.0;
- returns 0.0 for positive real *x2*;
- signals an error otherwise.

quo *x1 x2*   [Function]
rem *x1 x2*   [Function]
mod *x1 x2*   [Function]

> *x2* should be non-zero.

```
(quo x1 x2)                    ==> n_q
(rem x1 x2)                    ==> x_r
(mod x1 x2)                    ==> x_m
```

> where n_q is $x1/x2$ rounded towards zero, $0 < |x\_r| < |x2|$, $0 < |x\_m| < |x2|$, x_r and x_m differ from *x1* by a multiple of *x2*, x_r has the same sign as *x1*, and x_m has the same sign as *x2*.

> From this we can conclude that for *x2* not equal to 0,

```
(= x1 (+ (* x2 (quo x1 x2))
         (rem x1 x2)))
                               ==>  #t
```

> provided all numbers involved in that computation are exact.

```
(quo 2/3 1/5)                      ==>  3
(mod 2/3 1/5)                      ==>  1/15

(quo .666 1/5)                     ==>  3.0
(mod .666 1/5)                     ==>  65.99999999999995e-3
```

ln *z*   [Function]

> These procedures are part of every implementation that supports general real numbers. 'Ln' computes the natural logarithm of *z*

> In general, the mathematical function ln is multiply defined. The value of ln *z* is defined to be the one whose imaginary part lies in the range from -pi (exclusive) to pi (inclusive).

abs *x*   [Function]

> For real argument *x*, 'Abs' returns the absolute value of *x*' otherwise it signals an error.

```
(abs -7)                           ==>  7
```

make-rectangular *x1 x2*   [Function]
make-polar *x3 x4*   [Function]

> These procedures are part of every implementation that supports general complex numbers. Suppose *x1*, *x2*, *x3*, and *x4* are real numbers and *z* is a complex number such that

$$z = x1 + x2\text{i} = x3 \;.\; \text{e\textasciicircum i } x4$$

Then

```
(make-rectangular x1 x2)              ==> z
(make-polar x3 x4)                    ==> z
```

where -pi < x_angle <= pi with x_angle = $x4$ + 2pi n for some integer n.

If an argument is not real, then these procedures signal an error.

5.5 Prime Numbers

`(require 'factor)`

`prime:prngs`   [Variable]

    prime:prngs is the random-state (see Section 5.6 [Random Numbers], page 112) used by these procedures. If you call these procedures from more than one thread (or from interrupt), `random` may complain about reentrant calls.

    Note: The prime test and generation procedures implement (or use) the Solovay-Strassen primality test. See

- Robert Solovay and Volker Strassen, *A Fast Monte-Carlo Test for Primality*, SIAM Journal on Computing, 1977, pp 84-85.

`jacobi-symbol` *p q*   [Function]

    Returns the value (+1, −1, or 0) of the Jacobi-Symbol of exact non-negative integer *p* and exact positive odd integer *q*.

`prime:trials`   [Variable]

    prime:trials the maximum number of iterations of Solovay-Strassen that will be done to test a number for primality.

`prime?` *n*   [Function]

    Returns `#f` if *n* is composite; `#t` if *n* is prime. There is a slight chance (`expt 2 (- prime:trials)`) that a composite will return `#t`.

`primes<` *start count*   [Function]

    Returns a list of the first *count* prime numbers less than *start*. If there are fewer than *count* prime numbers less than *start*, then the returned list will have fewer than *start* elements.

`primes>` *start count*   [Function]

    Returns a list of the first *count* prime numbers greater than *start*.

`factor` *k*   [Function]

    Returns a list of the prime factors of *k*. The order of the factors is unspecified. In order to obtain a sorted list do (`sort! (factor k) <`).

5.6 Random Numbers

A pseudo-random number generator is only as good as the tests it passes. George Marsaglia of Florida State University developed a battery of tests named *DIEHARD* (`http://stat.fsu.edu/~geo/diehard.html`). `diehard.c` has a bug which the patch `http://groups.csail.mit.edu/mac/ftpdir/users/jaffer/diehard.c.pat` corrects.

SLIB's PRNG generates 8 bits at a time. With the degenerate seed '0', the numbers generated pass DIEHARD; but when bits are combined from sequential bytes, tests fail. With the seed '`http://swissnet.ai.mit.edu/~jaffer/SLIB.html`', all of those tests pass.

5.6.1 Exact Random Numbers

(`require 'random`)

random *n state*   [Function]
random *n*   [Function]

> *n* must be an exact positive integer. `random` returns an exact integer between zero (inclusive) and *n* (exclusive). The values returned by `random` are uniformly distributed from 0 to *n*.

> The optional argument *state* must be of the type returned by (`seed->random-state`) or (`make-random-state`). It defaults to the value of the variable `*random-state*`. This object is used to maintain the state of the pseudo-random-number generator and is altered as a side effect of calls to `random`.

random-state   [Variable]

> Holds a data structure that encodes the internal state of the random-number generator that `random` uses by default. The nature of this data structure is implementation-dependent. It may be printed out and successfully read back in, but may or may not function correctly as a random-number state object in another implementation.

copy-random-state *state*   [Function]

> Returns a new copy of argument *state*.

copy-random-state   [Function]

> Returns a new copy of `*random-state*`.

seed->random-state *seed*   [Function]

> Returns a new object of type suitable for use as the value of the variable `*random-state*` or as a second argument to `random`. The number or string *seed* is used to initialize the state. If `seed->random-state` is called twice with arguments which are `equal?`, then the returned data structures will be `equal?`. Calling `seed->random-state` with unequal arguments will nearly always return unequal states.

make-random-state   [Function]
make-random-state *obj*   [Function]

> Returns a new object of type suitable for use as the value of the variable `*random-state*` or as a second argument to `random`. If the optional argument *obj* is given, it should be a printable Scheme object; the first 50 characters of its printed

representation will be used as the seed. Otherwise the value of `*random-state*` is used as the seed.

5.6.2 Inexact Random Numbers

`(require 'random-inexact)`

`random:uniform`   [Function]
`random:uniform` *state*   [Function]
> Returns an uniformly distributed inexact real random number in the range between 0 and 1.

`random:exp`   [Function]
`random:exp` *state*   [Function]
> Returns an inexact real in an exponential distribution with mean 1. For an exponential distribution with mean *u* use `(* u (random:exp))`.

`random:normal`   [Function]
`random:normal` *state*   [Function]
> Returns an inexact real in a normal distribution with mean 0 and standard deviation 1. For a normal distribution with mean *m* and standard deviation *d* use `(+ m (* d (random:normal)))`.

`random:normal-vector!` *vect*   [Procedure]
`random:normal-vector!` *vect state*   [Procedure]
> Fills *vect* with inexact real random numbers which are independent and standard normally distributed (i.e., with mean 0 and variance 1).

`random:hollow-sphere!` *vect*   [Procedure]
`random:hollow-sphere!` *vect state*   [Procedure]
> Fills *vect* with inexact real random numbers the sum of whose squares is equal to 1.0. Thinking of *vect* as coordinates in space of dimension n = `(vector-length vect)`, the coordinates are uniformly distributed over the surface of the unit n-shere.

`random:solid-sphere!` *vect*   [Procedure]
`random:solid-sphere!` *vect state*   [Procedure]
> Fills *vect* with inexact real random numbers the sum of whose squares is less than 1.0. Thinking of *vect* as coordinates in space of dimension n = `(vector-length vect)`, the coordinates are uniformly distributed within the unit n-shere. The sum of the squares of the numbers is returned.

5.7 Discrete Fourier Transform

`(require 'dft)` or `(require 'Fourier-transform)`

`fft` and `fft-1` compute the Fast-Fourier-Transforms (O(n*log(n))) of arrays whose dimensions are all powers of 2.

`sft` and `sft-1` compute the Discrete-Fourier-Transforms for all combinations of dimensions (O(n^2)).

sft *array prot*   [Function]
sft *array*   [Function]
> *array* is an array of positive rank. **sft** returns an array of type *prot* (defaulting to *array*) of complex numbers comprising the *Discrete Fourier Transform* of *array*.

sft-1 *array prot*   [Function]
sft-1 *array*   [Function]
> *array* is an array of positive rank. **sft-1** returns an array of type *prot* (defaulting to *array*) of complex numbers comprising the inverse Discrete Fourier Transform of *array*.

fft *array prot*   [Function]
fft *array*   [Function]
> *array* is an array of positive rank whose dimensions are all powers of 2. **fft** returns an array of type *prot* (defaulting to *array*) of complex numbers comprising the Discrete Fourier Transform of *array*.

fft-1 *array prot*   [Function]
fft-1 *array*   [Function]
> *array* is an array of positive rank whose dimensions are all powers of 2. **fft-1** returns an array of type *prot* (defaulting to *array*) of complex numbers comprising the inverse Discrete Fourier Transform of *array*.

dft and dft-1 compute the discrete Fourier transforms using the best method for decimating each dimension.

dft *array prot*   [Function]
dft *array*   [Function]
> **dft** returns an array of type *prot* (defaulting to *array*) of complex numbers comprising the Discrete Fourier Transform of *array*.

dft-1 *array prot*   [Function]
dft-1 *array*   [Function]
> **dft-1** returns an array of type *prot* (defaulting to *array*) of complex numbers comprising the inverse Discrete Fourier Transform of *array*.

(fft-1 (fft array)) will return an array of values close to *array*.

```
(fft '#(1 0+i -1 0-i 1 0+i -1 0-i)) ⇒

#(0.0 0.0 0.0+628.0783185208527e-18i 0.0
   0.0 0.0 8.0-628.0783185208527e-18i 0.0)

(fft-1 '#(0 0 0 0 0 0 8 0)) ⇒

#(1.0 -61.23031769111886e-18+1.0i -1.0 61.23031769111886e-18-1.0i
   1.0 -61.23031769111886e-18+1.0i -1.0 61.23031769111886e-18-1.0i)
```

5.8 Cyclic Checksum

(`require 'crc`) Cyclic Redundancy Checks using Galois field GF(2) polynomial arithmetic
are used for error detection in many data transmission and storage applications.

The generator polynomials for various CRC protocols are availble from many sources. But
the polynomial is just one of many parameters which must match in order for a CRC
implementation to interoperate with existing systems:

- the byte-order and bit-order of the data stream;

- whether the CRC or its inverse is being calculated;

- the initial CRC value; and

- whether and where the CRC value is appended (inverted or non-inverted) to the data
  stream.

The performance of a particular CRC polynomial over packets of given sizes varies widely.
In terms of the probability of undetected errors, some uses of extant CRC polynomials are
suboptimal by several orders of magnitude.

If you are considering CRC for a new application, consult the following article to find the
optimum CRC polynomial for your range of data lengths:

- Philip Koopman and Tridib Chakravarty,
  "Cyclic Redundancy Code (CRC) Polynomial Selection For Embedded Networks",
  The International Conference on Dependable Systems and Networks, DSN-2004.

`http://www.ece.cmu.edu/~koopman/roses/dsn04/koopman04_crc_poly_embedded.pdf`

There is even some controversy over the polynomials themselves.

crc-32-polynomial   [Constant]
> For   CRC-32,   http://www2.sis.pitt.edu/~jkabara/tele-2100/lect08.html   gives
> x^32+x^26+x^23+x^16+x^12+x^11+x^10+x^8+x^7+x^5+x^4+x^2+x^1+1.

> But   http://www.cs.ncl.ac.uk/people/harry.whitfield/home.formal/CRCs.html,
> http://duchon.umuc.edu/Web_Pages/duchon/99_f_cm435/ShiftRegister.htm,
> http://spinroot.com/spin/Doc/Book91_PDF/ch3.pdf, http://www.erg.abdn.ac.uk/users/gorry/course/
> pages/crc.html,   http://www.rad.com/networks/1994/err_con/crc_most.htm,   and
> http://www.gpfn.sk.ca/~rhg/csc8550s02/crc.html, http://www.nobugconsulting.ro/crc.php
> give x^32+x^26+x^23+x^22+x^16+x^12+x^11+x^10+x^8+x^7+x^5+x^4+x^2+x+1.

> SLIB `crc-32-polynomial` uses the latter definition.

crc-ccitt-polynomial   [Constant]
> http://www.math.grin.edu/~rebelsky/Courses/CS364/2000S/Outlines/outline.12.html,
> http://duchon.umuc.edu/Web_Pages/duchon/99_f_cm435/ShiftRegister.htm,
> http://www.cs.ncl.ac.uk/people/harry.whitfield/home.formal/CRCs.html,
> http://www2.sis.pitt.edu/~jkabara/tele-2100/lect08.html, and http://www.gpfn.sk.ca/~rhg/csc8550s02/
> give CRC-CCITT: x^16+x^12+x^5+1.

`crc-16-polynomial`   [Constant]

>   http://www.math.grin.edu/~rebelsky/Courses/CS364/2000S/Outlines/outline.12.html,
>   http://duchon.umuc.edu/Web_Pages/duchon/99_f_cm435/ShiftRegister.htm,
>   http://www.cs.ncl.ac.uk/people/harry.whitfield/home.formal/CRCs.html,
>   http://www.gpfn.sk.ca/~rhg/csc8550s02/crc.html, and http://www.usb.org/developers/data/crcdes.pdf
>   give CRC-16: x^16+x^15+x^2+1.

`crc-12-polynomial`   [Constant]

>   http://www.math.grin.edu/~rebelsky/Courses/CS364/2000S/Outlines/outline.12.html,
>   http://www.cs.ncl.ac.uk/people/harry.whitfield/home.formal/CRCs.html,
>   http://www.it.iitb.ac.in/it605/lectures/link/node4.html, and http://spinroot.com/spin/Doc/Book91_PI
>   give CRC-12: x^12+x^11+x^3+x^2+1.

>   But http://www.ffldusoe.edu/Faculty/Denenberg/Topics/Networks/Error_Detection_Correction/crc.htm
>   http://duchon.umuc.edu/Web_Pages/duchon/99_f_cm435/ShiftRegister.htm,
>   http://www.eng.uwi.tt/depts/elec/staff/kimal/errorcc.html, http://www.ee.uwa.edu.au/~roberto/teach
>   http://www.gpfn.sk.ca/~rhg/csc8550s02/crc.html, and http://www.efg2.com/Lab/Mathematics/CRC.h
>   give CRC-12: x^12+x^11+x^3+x^2+x+1.

>   These differ in bit 1 and calculations using them return different values. With citations
>   near evenly split, it is hard to know which is correct. Thanks to Philip Koopman for
>   breaking the tie in favor of the latter (#xC07).

`crc-10-polynomial`   [Constant]

>   http://www.math.grin.edu/~rebelsky/Courses/CS364/2000S/Outlines/outline.12.html
>   gives   CRC-10:   x^10+x^9+x^5+x^4+1;   but   http://cell-relay.indiana.edu/cell-
>   relay/publications/software/CRC/crc10.html, http://www.it.iitb.ac.in/it605/lectures/link/node4.html,
>   http://www.gpfn.sk.ca/~rhg/csc8550s02/crc.html, http://www.techfest.com/networking/atm/atm.htm,
>   http://www.protocols.com/pbook/atmcell2.htm, and http://www.nobugconsulting.ro/crc.php
>   give CRC-10: x^10+x^9+x^5+x^4+x+1.

`crc-08-polynomial`   [Constant]

>   http://www.math.grin.edu/~rebelsky/Courses/CS364/2000S/Outlines/outline.12.html,
>   http://www.cs.ncl.ac.uk/people/harry.whitfield/home.formal/CRCs.html,
>   http://www.it.iitb.ac.in/it605/lectures/link/node4.html, and http://www.nobugconsulting.ro/crc.php
>   give CRC-8: x^8+x^2+x^1+1

`atm-hec-polynomial`   [Constant]

>   http://cell-relay.indiana.edu/cell-relay/publications/software/CRC/32bitCRC.tutorial.html
>   and http://www.gpfn.sk.ca/~rhg/csc8550s02/crc.html give ATM HEC: x^8+x^2+x+1.

`dowcrc-polynomial`   [Constant]

>   http://www.cs.ncl.ac.uk/people/harry.whitfield/home.formal/CRCs.html   gives
>   DOWCRC: x^8+x^5+x^4+1.

`usb-token-polynomial`   [Constant]

>   http://www.usb.org/developers/data/crcdes.pdf and http://www.nobugconsulting.ro/crc.php
>   give USB-token: x^5+x^2+1.

Each of these polynomial constants is a string of '1's and '0's, the exponent of each power
of x in descending order.

`crc:make-table` *poly*   [Function]

> *poly* must be string of '1's and '0's beginning with '1' and having length greater than 8. `crc:make-table` returns a vector of 256 integers, such that:
>
> ```
> (set! crc
>       (logxor (ash (logand (+ -1 (ash 1 (- deg 8)))) crc) 8)
>               (vector-ref crc-table
>                           (logxor (ash crc (- 8 deg)) byte))))
> ```

will compute the *crc* with the 8 additional bits in *byte*; where *crc* is the previous accumulated CRC value, *deg* is the degree of *poly*, and *crc-table* is the vector returned by `crc:make-table`.

If the implementation does not support *deg*-bit integers, then `crc:make-table` returns #f.

`cksum` *file*   [Function]

> Computes the P1003.2/D11.2 (POSIX.2) 32-bit checksum of *file*.
>
> ```
> (require 'crc)
> (cksum (in-vicinity (library-vicinity) "ratize.scm"))
> ⇒ 157103930
> ```

`cksum` *port*   [Function]

> Computes the checksum of the bytes read from *port* until the end-of-file.

`cksum-string`, which returns the P1003.2/D11.2 (POSIX.2) 32-bit checksum of the bytes in *str*, can be defined as follows:

```
(require 'string-port)
(define (cksum-string str) (call-with-input-string str cksum))
```

`crc16` *file*   [Function]

> Computes the USB data-packet (16-bit) CRC of *file*.

`crc16` *port*   [Function]

> Computes the USB data-packet (16-bit) CRC of the bytes read from *port* until the end-of-file.
>
> `crc16` calculates the same values as the crc16.pl program given in http://www.usb.org/developers/data/crcdes.pdf.

`crc5` *file*   [Function]

> Computes the USB token (5-bit) CRC of *file*.

`crc5` *port*   [Function]

> Computes the USB token (5-bit) CRC of the bytes read from *port* until the end-of-file.
>
> `crc5` calculates the same values as the crc5.pl program given in http://www.usb.org/developers/data/crcdes.pdf.

5.9 Graphing

5.9.1 Character Plotting

`(require 'charplot)`

`charplot:dimensions`   [Variable]

A list of the maximum height (number of lines) and maximum width (number of columns) for the graph, its scales, and labels.

The default value for *charplot:dimensions* is the `output-port-height` and `output-port-width` of `current-output-port`.

`plot` *coords x-label y-label*   [Procedure]

coords is a list or vector of coordinates, lists of x and y coordinates. *x-label* and *y-label* are strings with which to label the x and y axes.

Example:

```
(require 'charplot)
(set! charplot:dimensions '(20 55))

(define (make-points n)
  (if (zero? n)
      '()
      (cons (list (/ n 6) (sin (/ n 6))) (make-points (1- n)))))

(plot (make-points 40) "x" "Sin(x)")
⊣
   Sin(x)     _____
        1|-         ****                             |
         |        **      **                         |
     0.75|-     *          *                         |
         |     *            *                        |
      0.5|-   *              *                       |
         |   *                                      *|
     0.25|-                      *                *  |
         | *                      *                  |
        0|------------------*----------------*--|
         |                                   *      |
    -0.25|-                     *              *     |
         |                     *               *    |
     -0.5|-                   *                      |
         |                  *              *         |
    -0.75|-                 *            *           |
         |                   **      **              |
       -1|-                     ****                 |
         |:_____._____:_____._____:_____._____:____|
            x              2         4         6
```
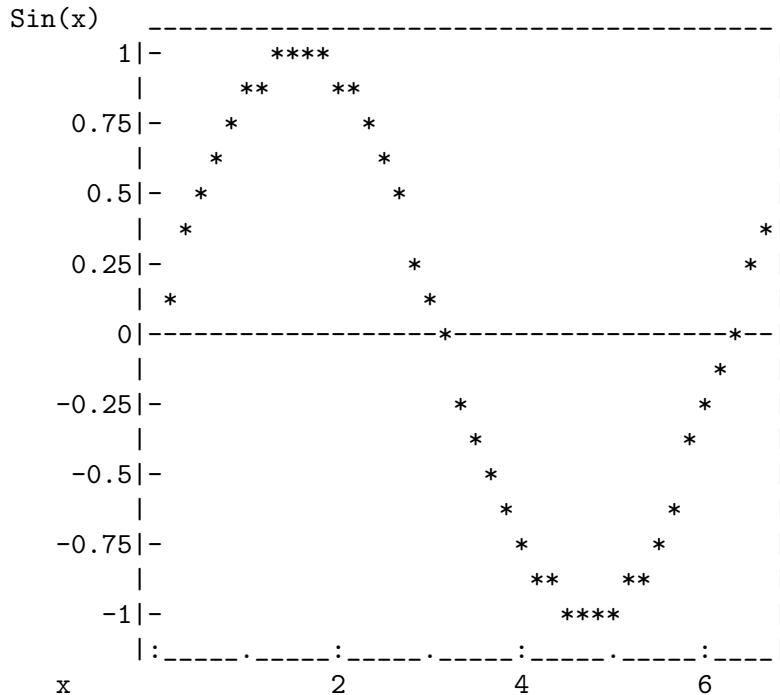
`plot` *func x1 x2*   [Procedure]
`plot` *func x1 x2 npts*   [Procedure]

> Plots the function of one argument *func* over the range *x1* to *x2*. If the optional
> integer argument *npts* is supplied, it specifies the number of points to evaluate *func*
> at.

```
(plot sin 0 (* 2 pi))
⊣
        _____
    1|-:           ****                          |
     | :         **      **                       |
 0.75|-:       *           *                      |
     | :      *             *                     |
  0.5|-:    **               **                   |
     | : *                    *                   |
 0.25|-:**                      **                |
     | :*                        *                |
    0|-*-----------------------*------------------|
     | :                        *               * |
-0.25|-:                         **            **  |
     | :                          *          *    |
 -0.5|-:                           *        **     |
     | :                            *       *      |
-0.75|-:                             *     **       |
     | :                              **  **        |
   -1|-:                                ****        |
     |_:_____._____:_____._____:_____._____:___|
        0           2           4           6
```
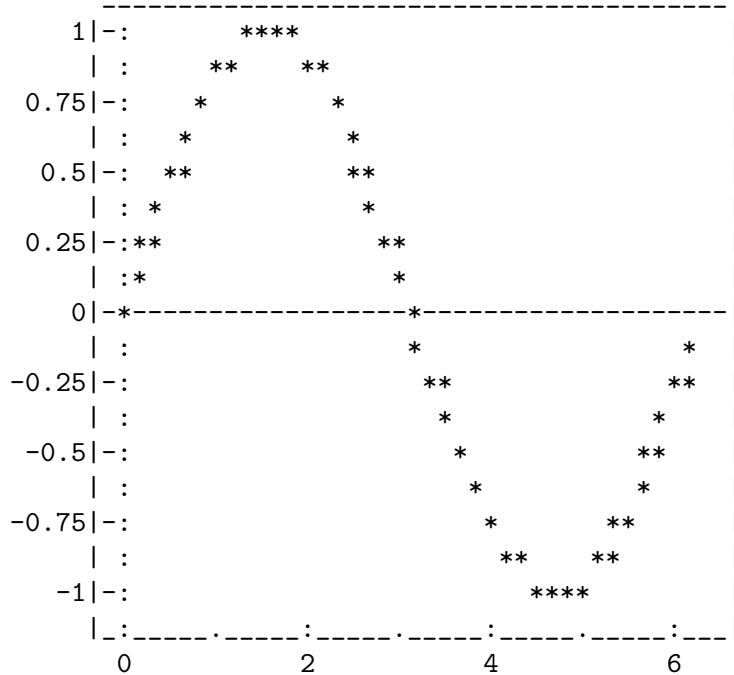
`histograph` *data label*   [Procedure]

> Creates and displays a histogram of the numerical values contained in vector or list
> *data*

```
(require 'random-inexact)
(histograph (do ((idx 99 (+ -1 idx))
                 (lst '() (cons (* .02 (random:normal)) lst)))
                ((negative? idx) lst))
            "normal")
⊣
```

```
      ----------------------------------------
   8|-                   :    I                |
    |                    :    I                |
   7|-            I  I :    I                |
    |             I  I :    I                |
   6|-          III I :I   I                |
    |           III I :I   I                |
   5|-          IIIIIIIIII I                |
    |           IIIIIIIIII I                |
   4|-          IIIIIIIIIII                 |
    |           IIIIIIIIIII                 |
   3|-I    I I  IIIIIIIIIII  II     I        |
    | I    I I  IIIIIIIIIII  II     I        |
   2|-I     I I IIIIIIIIIIIIIII     I        |
    | I     I I IIIIIIIIIIIIIII     I        |
   1|-II I I IIIIIIIIIIIIIIIIIIII   I I I    |
    | II I I IIIIIIIIIIIIIIIIIIII   I I I    |
   0|-IIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIII----|
    |__.____:____.____:____.____:____.___|
 normal      -0.025      0      0.025     0.05
```
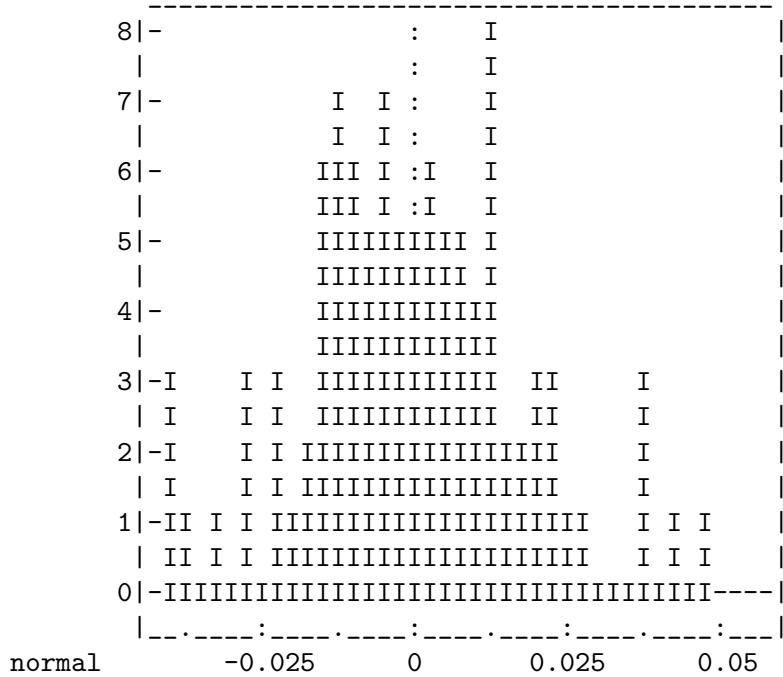
5.9.2 PostScript Graphing

`(require 'eps-graph)`

This is a graphing package creating encapsulated-PostScript files. Its motivations and design choice are described in `http://people.csail.mit.edu/jaffer/Docupage/grapheps`

A dataset to be plotted is taken from a 2-dimensional array. Corresponding coordinates are in rows. Coordinates from any pair of columns can be plotted.

---

create-postscript-graph *filename.eps size elt1 . . .*   [Function]
> *filename.eps* should be a string naming an output file to be created. *size* should be an exact integer, a list of two exact integers, or #f. *elt1*, ... are values returned by graphing primitives described here.
>
> **create-postscript-graph** creates an *Encapsulated-PostScript* file named *filename.eps* containing graphs as directed by the *elt1*, ... arguments.
>
> The size of the graph is determined by the *size* argument. If a list of two integers, they specify the width and height. If one integer, then that integer is the width and the height is 3/4 of the width. If #f, the graph will be 800 by 600.

These graphing procedures should be called as arguments to **create-postscript-graph**. The order of these arguments is significant; PostScript graphics state is affected serially from the first *elt* argument to the last.

---

whole-page   [Function]
> Pushes a rectangle for the whole encapsulated page onto the PostScript stack. This pushed rectangle is an implicit argument to **partition-page** or **setup-plot**.

5.9.2.1 Column Ranges

A *range* is a list of two numbers, the minimum and the maximum. Ranges can be given explicity or computed in PostScript by `column-range`.

`column-range` *array k*   [Function]
  Returns the range of values in 2-dimensional *array* column *k*.

`pad-range` *range p*   [Function]
  Expands *range* by *p*/100 on each end.

`snap-range` *range*   [Function]
  Expands *range* to round number of ticks.

`combine-ranges` *range1 range2 . . .*   [Function]
  Returns the minimal range covering all *range1*, *range2*, ...

`setup-plot` *x-range y-range pagerect*   [Function]
`setup-plot` *x-range y-range*   [Function]
  x-range and *y-range* should each be a list of two numbers or the value returned by `pad-range`, `snap-range`, or `combine-range`. *pagerect* is the rectangle bounding the graph to be drawn; if missing, the rectangle from the top of the PostScript stack is popped and used.

  Based on the given ranges, `setup-plot` sets up scaling and margins for making a graph. The margins are sized proportional to the *fontheight* value at the time of the call to setup-plot. `setup-plot` sets two variables:

  plotrect   The region where data points will be plotted.

  graphrect   The *pagerect* argument to `setup-plot`. Includes plotrect, legends, etc.

5.9.2.2 Drawing the Graph

`plot-column` *array x-column y-column proc3s*   [Function]
  Plots points with x coordinate in *x-column* of *array* and y coordinate *y-column* of *array*. The symbol *proc3s* specifies the type of glyph or drawing style for presenting these coordinates.

The glyphs and drawing styles available are:

`line`   Draws line connecting points in order.

`mountain`   Fill area below line connecting points.

`cloud`   Fill area above line connecting points.

`impulse`   Draw line from x-axis to each point.

`bargraph`   Draw rectangle from x-axis to each point.

`disc`   Solid round dot.

`point`   Minimal point – invisible if linewidth is 0.

`square`   Square box.

`diamond`	Square box at 45.o
`plus`	Plus sign.
`cross`	X sign.
`triup`	Triangle pointing upward
`tridown`	Triangle pointing downward
`pentagon`	Five sided polygon
`circle`	Hollow circle

`plot-text-column` *array x-column y-column t-column proc3s*   [Function]
> Plots text in *t-column* of *array* at x coordinate in *x-column* of *array* and y coordinate *y-column* of *array*. The symbol *proc3s* specifies the offset of the text from the specified coordinates.

The offsets available are:

`above`	Draws the text centered above at the point.
`center`	Draws the text centered at the point.
`below`	Draws the text centered below the point.
`left`	Draws the text to the left of the point.
`right`	Draws the text to the right of the point.

All the offsets other than `center` are calculated to keep the text clear of a glyph drawn at the same coordinates. If you need more or less clearance, use `set-glyphsize`.

5.9.2.3 Graphics Context

`in-graphic-context` *arg ...*   [Function]
> Saves the current graphics state, executes *args*, then restores to saved graphics state.

`set-color` *color*   [Function]
> *color* should be a string naming a Resene color, a saturate color, or a number between 0 and 100.
>
> `set-color` sets the PostScript color to the color of the given string, or a grey value between black (0) and white (100).

`set-font` *font height*   [Function]
`set-font` *font encoding height*   [Function]
> *font* should be a (case-sensitive) string naming a PostScript font. *height* should be a positive real number. *encoding* should name a PostScript encoding such as '`ISOLatin1Encoding`'.
>
> `set-font` Changes the current PostScript font to *font* with the *encoding* encoding, and height equal to *height*. The default font is '`Helvetica`' (12pt). The default encoding is '`StandardEncoding`'.

The base set of PostScript fonts is:

Times   Times-Italic   Times-Bold   Times-BoldItalic
Helvetica   Helvetica-Oblique   Helvetica-Bold   Helvetica-BoldOblique
Courier   Courier-Oblique   Courier-Bold   Courier-BoldOblique
Symbol

The base set of PostScript encodings is:

StandardEncoding   ISOLatin1Encoding   ExpertEncoding
ExpertSubsetEncoding   SymbolEncoding

Line parameters do no affect fonts; they do effect glyphs.

set-linewidth *w*   [Function]
> The default linewidth is 1. Setting it to 0 makes the lines drawn as skinny as possible. Linewidth must be much smaller than glyphsize for readable glyphs.

set-linedash *j k*   [Function]
> Lines are drawn *j*-on *k*-off.

set-linedash *j*   [Function]
> Lines are drawn *j*-on *j*-off.

set-linedash   [Function]
> Turns off dashing.

set-glyphsize *w*   [Function]
> Sets the (PostScript) variable glyphsize to *w*. The default glyphsize is 6.

The effects of `clip-to-rect` are also part of the graphic context.

5.9.2.4 Rectangles

A *rectangle* is a list of 4 numbers; the first two elements are the x and y coordinates of lower left corner of the rectangle. The other two elements are the width and height of the rectangle.

whole-page   [Function]
> Pushes a rectangle for the whole encapsulated page onto the PostScript stack. This pushed rectangle is an implicit argument to `partition-page` or `setup-plot`.

partition-page *xparts yparts*   [Function]
> Pops the rectangle currently on top of the stack and pushes *xparts * yparts* sub-rectangles onto the stack in decreasing y and increasing x order. If you are drawing just one graph, then you don't need `partition-page`.

plotrect   [Variable]
> The rectangle where data points should be plotted. *plotrect* is set by `setup-plot`.

graphrect   [Variable]
> The *pagerect* argument of the most recent call to `setup-plot`. Includes plotrect, legends, etc.

`fill-rect` *rect*   [Function]
>   fills *rect* with the current color.

`outline-rect` *rect*   [Function]
>   Draws the perimiter of *rect* in the current color.

`clip-to-rect` *rect*   [Function]
>   Modifies the current graphics-state so that nothing will be drawn outside of the rect-
>   angle *rect*. Use `in-graphic-context` to limit the extent of `clip-to-rect`.

5.9.2.5 Legending

`title-top` *title subtitle*   [Function]
`title-top` *title*   [Function]
>   Puts a *title* line and an optional *subtitle* line above the `graphrect`.

`title-bottom` *title subtitle*   [Function]
`title-bottom` *title*   [Function]
>   Puts a *title* line and an optional *subtitle* line below the `graphrect`.

`topedge`   [Variable]
`bottomedge`   [Variable]
>   These edge coordinates of `graphrect` are suitable for passing as the first argument
>   to `rule-horizontal`.

`leftedge`   [Variable]
`rightedge`   [Variable]
>   These edge coordinates of `graphrect` are suitable for passing as the first argument
>   to `rule-vertical`.

`set-margin-templates` *left right*   [Function]
>   The margin-templates are strings whose displayed width is used to reserve space for
>   the left and right side numerical legends. The default values are `"-.0123456789"`.

`rule-vertical` *x-coord text tick-width*   [Function]
>   Draws a vertical ruler with X coordinate *x-coord* and labeled with string *text*. If
>   *tick-width* is positive, then the ticks are *tick-width* long on the right side of *x-coord*;
>   and *text* and numeric legends are on the left. If *tick-width* is negative, then the ticks
>   are *-tick-width* long on the left side of *x-coord*; and *text* and numeric legends are on
>   the right.

`rule-horizontal` *y-coord text tick-height*   [Function]
>   Draws a horizontal ruler with Y coordinate *y-coord* and labeled with string *text*. If
>   *tick-height* is positive, then the ticks are *tick-height* long on the top side of *y-coord*;
>   and *text* and numeric legends are on the bottom. If *tick-height* is negative, then
>   the ticks are *-tick-height* long on the bottom side of *y-coord*; and *text* and numeric
>   legends are on the top.

`y-axis`   [Function]
>   Draws the y-axis.

`x-axis`   [Function]
>   Draws the x-axis.

`grid-verticals`   [Function]
>   Draws vertical lines through `graphrect` at each tick on the vertical ruler.

`grid-horizontals`   [Function]
>   Draws horizontal lines through `graphrect` at each tick on the horizontal ruler.

5.9.2.6 Legacy Plotting

`graph:dimensions`   [Variable]
>   A list of the width and height of the graph to be plotted using `plot`.

`plot` *func x1 x2 npts*   [Function]
`plot` *func x1 x2*   [Function]
>   Creates and displays using (`system "gv tmp.eps"`) an encapsulated PostScript graph
>   of the function of one argument *func* over the range *x1* to *x2*. If the optional integer
>   argument *npts* is supplied, it specifies the number of points to evaluate *func* at.

`plot` *x1 x2 npts func1 func2 ...*   [Function]
>   Creates and displays an encapsulated PostScript graph of the one-argument functions
>   *func1*, *func2*, ... over the range *x1* to *x2* at *npts* points.

`plot` *coords x-label y-label*   [Function]
>   *coords* is a list or vector of coordinates, lists of x and y coordinates. *x-label* and
>   *y-label* are strings with which to label the x and y axes.

5.9.2.7 Example Graph

The file `am1.5.html`, a table of solar irradiance, is fetched with '`wget`' if it isn't already in
the working directory. The file is read and stored into an array, *irradiance*.

   `create-postscript-graph` is then called to create an encapsulated-PostScript file,
`solarad.eps`. The size of the page is set to 600 by 300. `whole-page` is called and leaves
the rectangle on the PostScript stack. `setup-plot` is called with a literal range for x and
computes the range for column 1.

   Two calls to `top-title` are made so a different font can be used for the lower half.
`in-graphic-context` is used to limit the scope of the font change. The graphing area is
outlined and a rule drawn on the left side.

   Because the X range was intentionally reduced, `in-graphic-context` is called and
`clip-to-rect` limits drawing to the plotting area. A black line is drawn from data column
1. That line is then overlayed with a mountain plot of the same column colored "Bright
Sun".

   After returning from the `in-graphic-context`, the bottom ruler is drawn. Had it been
drawn earlier, all its ticks would have been painted over by the mountain plot.

   The color is then changed to '`seagreen`' and the same graphrect is setup again, this
time with a different Y scale, 0 to 1000. The graphic context is again clipped to *plotrect*,

linedash is set, and column 2 is plotted as a dashed line. Finally the rightedge is ruled. Having the line and its scale both in green helps disambiguate the scales.

```
(require 'eps-graph)
(require 'line-i/o)
(require 'string-port)

(define irradiance
  (let ((url "http://www.pv.unsw.edu.au/am1.5.html")
        (file "am1.5.html"))
    (define (read->list line)
      (define elts '())
      (call-with-input-string line
        (lambda (iprt) (do ((elt (read iprt) (read iprt)))
                           ((eof-object? elt) elts)
                         (set! elts (cons elt elts))))))
    (if (not (file-exists? file))
        (system (string-append "wget -c -O" file " " url)))
    (call-with-input-file file
      (lambda (iprt)
        (define lines '())
        (do ((line (read-line iprt) (read-line iprt)))
            ((eof-object? line)
             (let ((nra (make-array (A:floR64b)
                                    (length lines)
                                    (length (car lines)))))
               (do ((lns lines (cdr lns))
                    (idx (+ -1 (length lines)) (+ -1 idx)))
                   ((null? lns) nra)
                 (do ((kdx (+ -1 (length (car lines))) (+ -1 kdx))
                      (lst (car lns) (cdr lst)))
                     ((null? lst))
                   (array-set! nra (car lst) idx kdx)))))
          (if (and (positive? (string-length line))
                   (char-numeric? (string-ref line 0)))
              (set! lines (cons (read->list line) lines))))))))

(let ((xrange '(.25 2.5)))
  (create-postscript-graph
   "solarad.eps" '(600 300)
   (whole-page)
   (setup-plot xrange (column-range irradiance 1))
   (title-top
    "Solar Irradiance   http://www.pv.unsw.edu.au/am1.5.html")
   (in-graphic-context
    (set-font "Helvetica-Oblique" 12)
    (title-top
```

```
              ""
                "Key Centre for Photovoltaic Engineering UNSW - Air Mass 1.5 Global Spectrum"))
            (outline-rect plotrect)
            (rule-vertical leftedge "W/(m^2.um)" 10)
            (in-graphic-context (clip-to-rect plotrect)
                                (plot-column irradiance 0 1 'line)
                                (set-color "Bright Sun")
                                (plot-column irradiance 0 1 'mountain)
                                )
            (rule-horizontal bottomedge "Wavelength in .um" 5)
            (set-color 'seagreen)

            (setup-plot xrange '(0 1000) graphrect)
            (in-graphic-context (clip-to-rect plotrect)
                                (set-linedash 5 2)
                                (plot-column irradiance 0 2 'line))
            (rule-vertical rightedge "Integrated .W/(m^2)" -10)
            ))

      (system "gv solarad.eps")
```

5.10 Solid Modeling

```
(require 'solid)
```

`http://people.csail.mit.edu/jaffer/Solid/#Example` gives an example use of this package.

vrml *node . . .*   [Function]
>   Returns the VRML97 string (including header) of the concatenation of strings *nodes*, ...

vrml-append *node1 node2 . . .*   [Function]
>   Returns the concatenation with interdigitated newlines of strings *node1*, *node2*, ...

vrml-to-file *file node . . .*   [Function]
>   Writes to file named *file* the VRML97 string (including header) of the concatenation of strings *nodes*, ...

world:info *title info . . .*   [Function]
>   Returns a VRML97 string setting the title of the file in which it appears to *title*. Additional strings *info*, . . . are comments.

VRML97 strings passed to `vrml` and `vrml-to-file` as arguments will appear in the resulting VRML code. This string turns off the headlight at the viewpoint:

```
  " NavigationInfo {headlight FALSE}"
```

scene:panorama *front right back left top bottom*   [Function]
>   Specifies the distant images on the inside faces of the cube enclosing the virtual world.

scene:sphere *colors angles*   [Function]

> *colors* is a list of color objects. Each may be of type Section 5.11.1 [Color Data-Type], page 134, a 24-bit sRGB integer, or a list of 3 numbers between 0.0 and 1.0.
>
> *angles* is a list of non-increasing angles the same length as *colors*. Each angle is between 90 and -90 degrees. If 90 or -90 are not elements of *angles*, then the color at the zenith and nadir are taken from the colors paired with the angles nearest them.
>
> **scene:sphere** fills horizontal bands with interpolated colors on the background sphere encasing the world.

scene:sky-and-dirt   [Function]

> Returns a blue and brown background sphere encasing the world.

scene:sky-and-grass   [Function]

> Returns a blue and green background sphere encasing the world.

scene:sun *latitude julian-day hour turbidity strength*   [Function]
scene:sun *latitude julian-day hour turbidity*   [Function]

> *latitude* is the virtual place's latitude in degrees. *julian-day* is an integer from 0 to 366, the day of the year. *hour* is a real number from 0 to 24 for the time of day; 12 is noon. *turbidity* is the degree of fogginess described in See Section 5.11.7 [Daylight], page 149.
>
> **scene:sun** returns a bright yellow, distant sphere where the sun would be at *hour* on *julian-day* at *latitude*. If *strength* is positive, included is a light source of *strength* (default 1).

scene:overcast *latitude julian-day hour turbidity strength*   [Function]
scene:overcast *latitude julian-day hour turbidity*   [Function]

> *latitude* is the virtual place's latitude in degrees. *julian-day* is an integer from 0 to 366, the day of the year. *hour* is a real number from 0 to 24 for the time of day; 12 is noon. *turbidity* is the degree of cloudiness described in See Section 5.11.7 [Daylight], page 149.
>
> **scene:overcast** returns an overcast sky as it might look at *hour* on *julian-day* at *latitude*. If *strength* is positive, included is an ambient light source of *strength* (default 1).

Viewpoints are objects in the virtual world, and can be transformed individually or with solid objects.

scene:viewpoint *name distance compass pitch*   [Function]
scene:viewpoint *name distance compass*   [Function]

> Returns a viewpoint named *name* facing the origin and placed *distance* from it. *compass* is a number from 0 to 360 giving the compass heading. *pitch* is a number from -90 to 90, defaulting to 0, specifying the angle from the horizontal.

scene:viewpoints *proximity*   [Function]

> Returns 6 viewpoints, one at the center of each face of a cube with sides 2 * *proximity*, centered on the origin.

Light Sources

In VRML97, lights shine only on objects within the same children node and descendants of that node. Although it would have been convenient to let light direction be rotated by `solid:rotation`, this restricts a rotated light's visibility to objects rotated with it.

To workaround this limitation, these directional light source procedures accept either Cartesian or spherical coordinates for direction. A spherical coordinate is a list (`theta azimuth`); where *theta* is the angle in degrees from the zenith, and *azimuth* is the angle in degrees due west of south.

It is sometimes useful for light sources to be brighter than '1'. When *intensity* arguments are greater than 1, these functions gang multiple sources to reach the desired strength.

`light:ambient` *color intensity*   [Function]
`light:ambient` *color*   [Function]

> Ambient light shines on all surfaces with which it is grouped.
>
> *color* is a an object of type Section 5.11.1 [Color Data-Type], page 134, a 24-bit sRGB integer, or a list of 3 numbers between 0.0 and 1.0. If *color* is #f, then the default color will be used. *intensity* is a real non-negative number defaulting to '1'.
>
> `light:ambient` returns a light source or sources of *color* with total strength of *intensity* (or 1 if omitted).

`light:directional` *color direction intensity*   [Function]
`light:directional` *color direction*   [Function]
`light:directional` *color*   [Function]

> Directional light shines parallel rays with uniform intensity on all objects with which it is grouped.
>
> *color* is a an object of type Section 5.11.1 [Color Data-Type], page 134, a 24-bit sRGB integer, or a list of 3 numbers between 0.0 and 1.0. If *color* is #f, then the default color will be used.
>
> *direction* must be a list or vector of 2 or 3 numbers specifying the direction to this light. If *direction* has 2 numbers, then these numbers are the angle from zenith and the azimuth in degrees; if *direction* has 3 numbers, then these are taken as a Cartesian vector specifying the direction to the light source. The default direction is upwards; thus its light will shine down.
>
> *intensity* is a real non-negative number defaulting to '1'.
>
> `light:directional` returns a light source or sources of *color* with total strength of *intensity*, shining from *direction*.

`light:beam` *attenuation radius aperture peak*   [Function]
`light:beam` *attenuation radius aperture*   [Function]
`light:beam` *attenuation radius*   [Function]
`light:beam` *attenuation*   [Function]

> *attenuation* is a list or vector of three nonnegative real numbers specifying the reduction of intensity, the reduction of intensity with distance, and the reduction of intensity as the square of distance. *radius* is the distance beyond which the light does not shine. *radius* defaults to '100'.

aperture is a real number between 0 and 180, the angle centered on the light's axis through which it sheds some light. *peak* is a real number between 0 and 90, the angle of greatest illumination.

`light:point` *location color intensity beam*   [Function]
`light:point` *location color intensity*   [Function]
`light:point` *location color*   [Function]
`light:point` *location*   [Function]

> Point light radiates from *location*, intensity decreasing with distance, towards all objects with which it is grouped.

> *color* is a an object of type Section 5.11.1 [Color Data-Type], page 134, a 24-bit sRGB integer, or a list of 3 numbers between 0.0 and 1.0. If *color* is #f, then the default color will be used. *intensity* is a real non-negative number defaulting to '1'. *beam* is a structure returned by `light:beam` or #f.

> `light:point` returns a light source or sources at *location* of *color* with total strength *intensity* and *beam* properties. Note that the pointlight itself is not visible. To make it so, place an object with emissive appearance at *location*.

`light:spot` *location direction color intensity beam*   [Function]
`light:spot` *location direction color intensity*   [Function]
`light:spot` *location direction color*   [Function]
`light:spot` *location direction*   [Function]
`light:spot` *location*   [Function]

> Spot light radiates from *location* towards *direction*, intensity decreasing with distance, illuminating objects with which it is grouped.

> *direction* must be a list or vector of 2 or 3 numbers specifying the direction to this light. If *direction* has 2 numbers, then these numbers are the angle from zenith and the azimuth in degrees; if *direction* has 3 numbers, then these are taken as a Cartesian vector specifying the direction to the light source. The default direction is upwards; thus its light will shine down.

> *color* is a an object of type Section 5.11.1 [Color Data-Type], page 134, a 24-bit sRGB integer, or a list of 3 numbers between 0.0 and 1.0. If *color* is #f, then the default color will be used.

> *intensity* is a real non-negative number defaulting to '1'.

> `light:spot` returns a light source or sources at *location* of *direction* with total strength *color*. Note that the spotlight itself is not visible. To make it so, place an object with emissive appearance at *location*.

Object Primitives

`solid:box` *geometry appearance*   [Function]
`solid:box` *geometry*   [Function]

> *geometry* must be a number or a list or vector of three numbers. If *geometry* is a number, the `solid:box` returns a cube with sides of length *geometry* centered on the origin. Otherwise, `solid:box` returns a rectangular box with dimensions *geometry*

centered on the origin. *appearance* determines the surface properties of the returned
object.

solid:lumber *geometry appearance*   [Function]
> Returns a box of the specified *geometry*, but with the y-axis of a texture specified in
> *appearance* being applied along the longest dimension in *geometry*.

solid:cylinder *radius height appearance*   [Function]
solid:cylinder *radius height*   [Function]
> Returns a right cylinder with dimensions (`abs` `radius`) and (`abs` `height`) centered
> on the origin. If *height* is positive, then the cylinder ends will be capped. If *radius*
> is negative, then only the ends will appear. *appearance* determines the surface
> properties of the returned object.

solid:disk *radius thickness appearance*   [Function]
solid:disk *radius thickness*   [Function]
> *thickness* must be a positive real number. **solid:disk** returns a circular disk with
> dimensions *radius* and *thickness* centered on the origin. *appearance* determines the
> surface properties of the returned object.

solid:cone *radius height appearance*   [Function]
solid:cone *radius height*   [Function]
> Returns an isosceles cone with dimensions *radius* and *height* centered on the origin.
> *appearance* determines the surface properties of the returned object.

solid:pyramid *side height appearance*   [Function]
solid:pyramid *side height*   [Function]
> Returns an isosceles pyramid with dimensions *side* and *height* centered on the origin.
> *appearance* determines the surface properties of the returned object.

solid:sphere *radius appearance*   [Function]
solid:sphere *radius*   [Function]
> Returns a sphere of radius *radius* centered on the origin. *appearance* determines the
> surface properties of the returned object.

solid:ellipsoid *geometry appearance*   [Function]
solid:ellipsoid *geometry*   [Function]
> *geometry* must be a number or a list or vector of three numbers. If *geometry* is a
> number, the **solid:ellipsoid** returns a sphere of diameter *geometry* centered on the
> origin. Otherwise, **solid:ellipsoid** returns an ellipsoid with diameters *geometry*
> centered on the origin. *appearance* determines the surface properties of the returned
> object.

solid:polyline *coordinates appearance*   [Function]
solid:polyline *coordinates*   [Function]
> *coordinates* must be a list or vector of coordinate lists or vectors specifying the x, y,
> and z coordinates of points. **solid:polyline** returns lines connecting successive pairs
> of points. If called with one argument, then the polyline will be white. If *appearance*

is given, then the polyline will have its emissive color only; being black if *appearance* does not have an emissive color.

The following code will return a red line between points at (1 2 3) and (4 5 6):

```
(solid:polyline '((1 2 3) (4 5 6)) (solid:color #f 0 #f 0 '(1 0 0)))
```

`solid:prism` *xz-array y appearance*   [Function]
`solid:prism` *xz-array y*   [Function]

> *xz-array* must be an *n*-by-2 array holding a sequence of coordinates tracing a non-intersecting clockwise loop in the x-z plane. `solid:prism` will close the sequence if the first and last coordinates are not the same.

> `solid:prism` returns a capped prism *y* long.

`solid:basrelief` *width height depth colorray appearance*   [Function]
`solid:basrelief` *width height depth appearance*   [Function]
`solid:basrelief` *width height depth*   [Function]

> One of *width*, *height*, or *depth* must be a 2-dimensional array; the others must be real numbers giving the length of the basrelief in those dimensions. The rest of this description assumes that *height* is an array of heights.

> `solid:basrelief` returns a *width* by *depth* basrelief solid with heights per array *height* with the buttom surface centered on the origin.

> If present, *appearance* determines the surface properties of the returned object. If present, *colorray* must be an array of objects of type Section 5.11.1 [Color Data-Type], page 134, 24-bit sRGB integers or lists of 3 numbers between 0.0 and 1.0.

> If *colorray*'s dimensions match *height*, then each element of *colorray* paints its corresponding vertex of *height*. If *colorray* has all dimensions one smaller than *height*, then each element of *colorray* paints the corresponding face of *height*. Other dimensions for *colorray* are in error.

`solid:text` *fontstyle str len appearance*   [Function]
`solid:text` *fontstyle str len*   [Function]

> *fontstyle* must be a value returned by `solid:font`.

> *str* must be a string or list of strings.

> *len* must be #f, a nonnegative integer, or list of nonnegative integers.

> *appearance*, if given, determines the surface properties of the returned object.

> `solid:text` returns a two-sided, flat text object positioned in the Z=0 plane of the local coordinate system

Surface Attributes

`solid:color` *diffuseColor ambientIntensity specularColor shininess*   [Function]
>   *emissiveColor transparency*
`solid:color` *diffuseColor ambientIntensity specularColor shininess*   [Function]
>   *emissiveColor*
`solid:color` *diffuseColor ambientIntensity specularColor shininess*   [Function]
`solid:color` *diffuseColor ambientIntensity specularColor*   [Function]

`solid:color` *diffuseColor ambientIntensity*   [Function]
`solid:color` *diffuseColor*   [Function]
> Returns an *appearance*, the optical properties of the objects with which it is associated. *ambientIntensity*, *shininess*, and *transparency* must be numbers between 0 and 1. *diffuseColor*, *specularColor*, and *emissiveColor* are objects of type Section 5.11.1 [Color Data-Type], page 134, 24-bit sRGB integers or lists of 3 numbers between 0.0 and 1.0. If a color argument is omitted or #f, then the default color will be used.

`solid:texture` *image color scale rotation center translation*   [Function]
`solid:texture` *image color scale rotation center*   [Function]
`solid:texture` *image color scale rotation*   [Function]
`solid:texture` *image color scale*   [Function]
`solid:texture` *image color*   [Function]
`solid:texture` *image*   [Function]
> Returns an *appearance*, the optical properties of the objects with which it is associated. *image* is a string naming a JPEG or PNG image resource. *color* is #f, a color, or the string returned by `solid:color`. The rest of the optional arguments specify 2-dimensional transforms applying to the *image*.
>
> *scale* must be #f, a number, or list or vector of 2 numbers specifying the scale to apply to *image*. *rotation* must be #f or the number of degrees to rotate *image*. *center* must be #f or a list or vector of 2 numbers specifying the center of *image* relative to the *image* dimensions. *translation* must be #f or a list or vector of 2 numbers specifying the translation to apply to *image*.

`solid:font` *family style justify size spacing language direction*   [Function]
> Returns a fontstyle object suitable for passing as an argument to `solid:text`. Any of the arguments may be #f, in which case its default value, which is first in each list of allowed values, is used.
>
> *family* is a case-sensitive string naming a font; '`SERIF`', '`SANS`', and '`TYPEWRITER`' are supported at the minimum.
>
> *style* is a case-sensitive string '`PLAIN`', '`BOLD`', '`ITALIC`', or '`BOLDITALIC`'.
>
> *justify* is a case-sensitive string '`FIRST`', '`BEGIN`', '`MIDDLE`', or '`END`'; or a list of one or two case-sensitive strings (same choices). The mechanics of *justify* get complicated; it is explained by tables 6.2 to 6.7 of `http://www.web3d.org/x3d/specifications/vrml/ISO-IEC-14772-IS-VRML97WithAmendment1/part1/nodesRef.html#Table6.2`
>
> *size* is the extent, in the non-advancing direction, of the text. *size* defaults to 1.
>
> *spacing* is the ratio of the line (or column) offset to *size*. *spacing* defaults to 1.
>
> *language* is the RFC-1766 language name.
>
> *direction* is a list of two numbers: (*x y*). If (`>` (`abs` *x*) (`abs` *y*)), then the text will be arrayed horizontally; otherwise vertically. The direction in which characters are arrayed is determined by the sign of the major axis: positive *x* being left-to-right; positive *y* being top-to-bottom.

Aggregating Objects

solid:center-row-of *number solid spacing*   [Function]
> Returns a row of *number solid* objects spaced evenly *spacing* apart.

solid:center-array-of *number-a number-b solid spacing-a*   [Function]
> *spacing-b*
> Returns *number-b* rows, *spacing-b* apart, of *number-a solid* objects *spacing-a* apart.

solid:center-pile-of *number-a number-b number-c solid*   [Function]
> *spacing-a spacing-b spacing-c*
> Returns *number-c* planes, *spacing-c* apart, of *number-b* rows, *spacing-b* apart, of *number-a solid* objects *spacing-a* apart.

solid:arrow *center*   [Function]
> *center* must be a list or vector of three numbers. Returns an upward pointing metallic arrow centered at *center*.

solid:arrow   [Function]
> Returns an upward pointing metallic arrow centered at the origin.

Spatial Transformations

solid:translation *center solid* . . .   [Function]
> *center* must be a list or vector of three numbers. **solid:translation** Returns an aggregate of *solids*, . . . with their origin moved to *center*.

solid:scale *scale solid* . . .   [Function]
> *scale* must be a number or a list or vector of three numbers. **solid:scale** Returns an aggregate of *solids*, . . . scaled per *scale*.

solid:rotation *axis angle solid* . . .   [Function]
> *axis* must be a list or vector of three numbers. **solid:rotation** Returns an aggregate of *solids*, . . . rotated *angle* degrees around the axis *axis*.

5.11 Color

`http://people.csail.mit.edu/jaffer/Color`

The goals of this package are to provide methods to specify, compute, and transform colors in a core set of additive color spaces. The color spaces supported should be sufficient for working with the color data encountered in practice and the literature.

5.11.1 Color Data-Type

`(require 'color)`

color? *obj*   [Function]
> Returns #t if *obj* is a color.

color? *obj typ*   [Function]

    Returns #t if *obj* is a color of color-space *typ*. The symbol *typ* must be one of:

- CIEXYZ
- RGB709
- L*a*b*
- L*u*v*
- sRGB
- e-sRGB
- L*C*h

make-color *space arg* . . .   [Function]

    Returns a color of type *space*.

- For *space* arguments `CIEXYZ`, `RGB709`, and `sRGB`, the sole *arg* is a list of three numbers.
- For *space* arguments `L*a*b*`, `L*u*v*`, and `L*C*h`, *arg* is a list of three numbers optionally followed by a whitepoint.
- For `xRGB`, *arg* is an integer.
- For `e-sRGB`, the arguments are as for `e-sRGB->color`.

color-space *color*   [Function]

    Returns the symbol for the color-space in which *color* is embedded.

color-precision *color*   [Function]

    For colors in digital color-spaces, `color-precision` returns the number of bits used for each of the R, G, and B channels of the encoding. Otherwise, `color-precision` returns #f

color-white-point *color*   [Function]

    Returns the white-point of *color* in all color-spaces except CIEXYZ.

convert-color *color space white-point*   [Function]
convert-color *color space*   [Function]
convert-color *color e-sRGB precision*   [Function]

    Converts *color* into *space* at optional *white-point*.

5.11.1.1 External Representation

Each color encoding has an external, case-insensitive representation. To ensure portability, the white-point for all color strings is D65.[1]

Color Space	External Representation
CIEXYZ	CIEXYZ:<*X*>/<*Y*>/<*Z*>

---

[1] Readers may recognize these color string formats from Xlib. X11's color management system was doomed by its fiction that CRT monitors' (and X11 default) color-spaces were linear RGBi. Unable to shed this legacy, the only practical way to view pictures on X is to ignore its color management system and use an sRGB monitor. In this implementation the device-independent RGB709 and sRGB spaces replace the device-dependent RGBi and RGB spaces of Xlib.

RGB709	RGBi:<*R*>/<*G*>/<*B*>
L*a*b*	CIELAB:<*L*>/<*a*>/<*b*>
L*u*v*	CIELuv:<*L*>/<*u*>/<*v*>
L*C*h	CIELCh:<*L*>/<*C*>/<*h*>

The X, Y, Z, L, a, b, u, v, C, h, R, G, and B fields are (Scheme) real numbers within the appropriate ranges.

Color Space	External Representation
sRGB	sRGB:<*R*>/<*G*>/<*B*>
e-sRGB10	e-sRGB10:<*R*>/<*G*>/<*B*>
e-sRGB12	e-sRGB12:<*R*>/<*G*>/<*B*>
e-sRGB16	e-sRGB16:<*R*>/<*G*>/<*B*>

The R, G, and B, fields are non-negative exact decimal integers within the appropriate ranges.

Several additional syntaxes are supported by `string->color`:

Color Space	External Representation
sRGB	sRGB:<*RRGGBB*>
sRGB	#<*RRGGBB*>
sRGB	0x<*RRGGBB*>
sRGB	#x<*RRGGBB*>

Where *RRGGBB* is a non-negative six-digit hexadecimal number.

color->string *color*   [Function]
    Returns a string representation of *color*.

string->color *string*   [Function]
    Returns the color represented by *string*. If *string* is not a syntactically valid notation
    for a color, then `string->color` returns #f.

5.11.1.2 White

We experience color relative to the illumination around us. CIEXYZ coordinates, although subject to uniform scaling, are objective. Thus other color spaces are specified relative to a *white point* in CIEXYZ coordinates.

The white point for digital color spaces is set to D65. For the other spaces a *white-point* argument can be specified. The default if none is specified is the white-point with which the color was created or last converted; and D65 if none has been specified.

D65   [Constant]
    Is the color of 6500.K (blackbody) illumination. D65 is close to the average color of
    daylight.

D50   [Constant]
    Is the color of 5000.K (blackbody) illumination. D50 is the color of indoor lighting
    by incandescent bulbs, whose filaments have temperatures around 5000.K.

5.11.2 Color Spaces

Measurement-based Color Spaces

The *tristimulus* color spaces are those whose component values are proportional measurements of light intensity. The CIEXYZ(1931) system provides 3 sets of spectra to dot-product with a spectrum of interest. The result of those dot-products is coordinates in CIEXYZ space. All tristimuls color spaces are related to CIEXYZ by linear transforms, namely matrix multiplication. Of the color spaces listed here, CIEXYZ and RGB709 are tristimulus spaces.

`CIEXYZ`   [Color Space]

> The CIEXYZ color space covers the full *gamut*. It is the basis for color-space conversions.
>
> CIEXYZ is a list of three inexact numbers between 0.0 and 1.1. '(0. 0. 0.) is black; '(1. 1. 1.) is white.

`ciexyz->color` *xyz*   [Function]

> *xyz* must be a list of 3 numbers. If *xyz* is valid CIEXYZ coordinates, then `ciexyz->color` returns the color specified by *xyz*; otherwise returns #f.

`color:ciexyz` *x y z*   [Function]

> Returns the CIEXYZ color composed of *x*, *y*, *z*. If the coordinates do not encode a valid CIEXYZ color, then an error is signaled.

`color->ciexyz` *color*   [Function]

> Returns the list of 3 numbers encoding *color* in CIEXYZ.

`RGB709`   [Color Space]

> BT.709-4 (03/00) *Parameter values for the HDTV standards for production and international programme exchange* specifies parameter values for chromaticity, sampling, signal format, frame rates, etc., of high definition television signals.
>
> An RGB709 color is represented by a list of three inexact numbers between 0.0 and 1.0. '(0. 0. 0.) is black '(1. 1. 1.) is white.

`rgb709->color` *rgb*   [Function]

> *rgb* must be a list of 3 numbers. If *rgb* is valid RGB709 coordinates, then `rgb709->color` returns the color specified by *rgb*; otherwise returns #f.

`color:rgb709` *r g b*   [Function]

> Returns the RGB709 color composed of *r*, *g*, *b*. If the coordinates do not encode a valid RGB709 color, then an error is signaled.

`color->rgb709` *color*   [Function]

> Returns the list of 3 numbers encoding *color* in RGB709.

Perceptual Uniformity

Although properly encoding the chromaticity, tristimulus spaces do not match the logarithmic response of human visual systems to intensity. Minimum detectable differences between colors correspond to a smaller range of distances (6:1) in the L*a*b* and L*u*v* spaces than in tristimulus spaces (80:1). For this reason, color distances are computed in L*a*b* (or L*C*h).

`L*a*b*`   [Color Space]
>   Is a CIE color space which better matches the human visual system's perception of
>   color. It is a list of three numbers:
>
>   - 0 <= L* <= 100 (CIE *Lightness*)
>   - -500 <= a* <= 500
>   - -200 <= b* <= 200

`l*a*b*->color` *L*a*b* white-point*   [Function]
>   *L*a*b** must be a list of 3 numbers. If *L*a*b** is valid L*a*b* coordinates, then
>   `l*a*b*->color` returns the color specified by *L*a*b**; otherwise returns #f.

`color:l*a*b*` *L* a* b* white-point*   [Function]
>   Returns the L*a*b* color composed of *L**, *a**, *b** with *white-point*.

`color:l*a*b*` *L* a* b**   [Function]
>   Returns the L*a*b* color composed of *L**, *a**, *b**. If the coordinates do not encode
>   a valid L*a*b* color, then an error is signaled.

`color->l*a*b*` *color white-point*   [Function]
>   Returns the list of 3 numbers encoding *color* in L*a*b* with *white-point*.

`color->l*a*b*` *color*   [Function]
>   Returns the list of 3 numbers encoding *color* in L*a*b*.

`L*u*v*`   [Color Space]
>   Is another CIE encoding designed to better match the human visual system's percep-
>   tion of color.

`l*u*v*->color` *L*u*v* white-point*   [Function]
>   *L*u*v** must be a list of 3 numbers. If *L*u*v** is valid L*u*v* coordinates, then
>   `l*u*v*->color` returns the color specified by *L*u*v**; otherwise returns #f.

`color:l*u*v*` *L* u* v* white-point*   [Function]
>   Returns the L*u*v* color composed of *L**, *u**, *v** with *white-point*.

`color:l*u*v*` *L* u* v**   [Function]
>   Returns the L*u*v* color composed of *L**, *u**, *v**. If the coordinates do not encode
>   a valid L*u*v* color, then an error is signaled.

`color->l*u*v*` *color white-point*   [Function]
>   Returns the list of 3 numbers encoding *color* in L*u*v* with *white-point*.

`color->l*u*v*` *color*   [Function]
>   Returns the list of 3 numbers encoding *color* in L*u*v*.

Cylindrical Coordinates

HSL (Hue Saturation Lightness), HSV (Hue Saturation Value), HSI (Hue Saturation Inten-
sity) and HCI (Hue Chroma Intensity) are cylindrical color spaces (with angle hue). But
these spaces are all defined in terms device-dependent RGB spaces.

One might wonder if there is some fundamental reason why intuitive specification of color must be device-dependent. But take heart! A cylindrical system can be based on L*a*b* and is used for predicting how close colors seem to observers.

L*C*h   [Color Space]

Expresses the *a and b* of L*a*b* in polar coordinates. It is a list of three numbers:

- 0 <= L* <= 100 (CIE *Lightness*)
- C* (CIE *Chroma*) is the distance from the neutral (gray) axis.
- 0 <= h <= 360 (CIE *Hue*) is the angle.

The colors by quadrant of h are:

0	red, orange, yellow	90
90	yellow, yellow-green, green	180
180	green, cyan (blue-green), blue	270
270	blue, purple, magenta	360

l*c*h->color *L*C*h white-point*   [Function]

L^*C^*h must be a list of 3 numbers. If L^*C^*h is valid L*C*h coordinates, then **l*c*h->color** returns the color specified by L^*C^*h; otherwise returns #f.

color:l*c*h *L* C* h white-point*   [Function]

Returns the L*C*h color composed of *L*, *C**, *h* with *white-point*.

color:l*c*h *L* C* h*   [Function]

Returns the L*C*h color composed of *L*, *C**, *h*. If the coordinates do not encode a valid L*C*h color, then an error is signaled.

color->l*c*h *color white-point*   [Function]

Returns the list of 3 numbers encoding *color* in L*C*h with *white-point*.

color->l*c*h *color*   [Function]

Returns the list of 3 numbers encoding *color* in L*C*h.

Digital Color Spaces

The color spaces discussed so far are impractical for image data because of numerical precision and computational requirements. In 1998 the IEC adopted *A Standard Default Color Space for the Internet - sRGB* (`http://www.w3.org/Graphics/Color/sRGB`). sRGB was cleverly designed to employ the 24-bit (256x256x256) color encoding already in widespread use; and the 2.2 gamma intrinsic to CRT monitors.

Conversion from CIEXYZ to digital (sRGB) color spaces is accomplished by conversion first to a RGB709 tristimulus space with D65 white-point; then each coordinate is individually subjected to the same non-linear mapping. Inverse operations in the reverse order create the inverse transform.

sRGB   [Color Space]

    Is "A Standard Default Color Space for the Internet". Most display monitors will work fairly well with sRGB directly. Systems using ICC profiles [2] should work very well with sRGB.

srgb->color *rgb*   [Function]

    rgb must be a list of 3 numbers. If *rgb* is valid sRGB coordinates, then **srgb->color** returns the color specified by *rgb*; otherwise returns #f.

color:srgb *r g b*   [Function]

    Returns the sRGB color composed of *r*, *g*, *b*. If the coordinates do not encode a valid sRGB color, then an error is signaled.

xRGB   [Color Space]

    Represents the equivalent sRGB color with a single 24-bit integer. The most significant 8 bits encode red, the middle 8 bits blue, and the least significant 8 bits green.

color->srgb *color*   [Function]

    Returns the list of 3 integers encoding *color* in sRGB.

color->xrgb *color*   [Function]

    Returns the 24-bit integer encoding *color* in sRGB.

xrgb->color *k*   [Function]

    Returns the sRGB color composed of the 24-bit integer *k*.

e-sRGB   [Color Space]

    Is "Photography - Electronic still picture imaging - Extended sRGB color encoding" (PIMA 7667:2001). It extends the gamut of sRGB; and its higher precision numbers provide a larger dynamic range.

    A triplet of integers represent e-sRGB colors. Three precisions are supported:

    e-sRGB10   0 to 1023

    e-sRGB12   0 to 4095

    e-sRGB16   0 to 65535

e-srgb->color *precision rgb*   [Function]

    precision must be the integer 10, 12, or 16. *rgb* must be a list of 3 numbers. If *rgb* is valid e-sRGB coordinates, then **e-srgb->color** returns the color specified by *rgb*; otherwise returns #f.

---

[2] A comprehensive encoding of transforms between CIEXYZ and device color spaces is the International Color Consortium profile format, ICC.1:1998-09:

    The intent of this format is to provide a cross-platform device profile format. Such device profiles can be used to translate color data created on one device into another device's native color space.

`color:e-srgb` *10 r g b*   [Function]
>   Returns the e-sRGB10 color composed of integers *r*, *g*, *b*.

`color:e-srgb` *12 r g b*   [Function]
>   Returns the e-sRGB12 color composed of integers *r*, *g*, *b*.

`color:e-srgb` *16 r g b*   [Function]
>   Returns the e-sRGB16 color composed of integers *r*, *g*, *b*. If the coordinates do not
>   encode a valid e-sRGB color, then an error is signaled.

`color->e-srgb` *precision color*   [Function]
>   *precision* must be the integer 10, 12, or 16. `color->e-srgb` returns the list of 3
>   integers encoding *color* in sRGB10, sRGB12, or sRGB16.

5.11.3 Spectra

The following functions compute colors from spectra, scale color luminance, and extract
chromaticity. XYZ is used in the names of procedures for unnormalized colors; the coor-
dinates of CIEXYZ colors are constrained as described in Section 5.11.2 [Color Spaces],
page 136.

```
(require 'color-space)
```

A spectrum may be represented as:

- A procedure of one argument accepting real numbers from 380e-9 to 780e-9, the wave-
  length in meters; or

- A vector of real numbers representing intensity samples evenly spaced over some range
  of wavelengths overlapping the range 380e-9 to 780e-9.

CIEXYZ values are calculated as dot-product with the X, Y (Luminance), and Z *Spectral
Tristimulus Values*. The files `cie1931.xyz` and `cie1964.xyz` in the distribution contain
these CIE-defined values.

`cie1964`   [Feature]
>   Loads the Spectral Tristimulus Values *CIE 1964 Supplementary Standard Colorimet-
>   ric Observer*, defining *cie:x-bar*, *cie:y-bar*, and *cie:z-bar*.

`cie1931`   [Feature]
>   Loads the Spectral Tristimulus Values *CIE 1931 Supplementary Standard Colorimet-
>   ric Observer*, defining *cie:x-bar*, *cie:y-bar*, and *cie:z-bar*.

`ciexyz`   [Feature]
>   Requires Spectral Tristimulus Values, defaulting to cie1931, defining *cie:x-bar*, *cie:y-
>   bar*, and *cie:z-bar*.

(`require 'cie1964`) or (`require 'cie1931`) will `load-ciexyz` specific values used by the
following spectrum conversion procedures. The spectrum conversion procedures (`require
'ciexyz`) to assure that a set is loaded.

read-cie-illuminant *path*   [Function]

    path must be a string naming a file consisting of 107 numbers for 5.nm intervals from 300.nm to 830.nm. `read-cie-illuminant` reads (using Scheme `read`) these numbers and returns a length 107 vector filled with them.

```
(define CIE:SI-D65
  (read-CIE-illuminant (in-vicinity (library-vicinity) "ciesid65.dat")))
(spectrum->XYZ CIE:SI-D65 300e-9 830e-9)
⇒ (25.108569422374994 26.418013465625001 28.764075683374993)
```

read-normalized-illuminant *path*   [Function]

    path must be a string naming a file consisting of 107 numbers for 5.nm intervals from 300.nm to 830.nm. `read-normalized-illuminant` reads (using Scheme `read`) these numbers and returns a length 107 vector filled with them, normalized so that `spectrum->XYZ` of the illuminant returns its whitepoint.

    CIE Standard Illuminants A and D65 are included with SLIB:

```
(define CIE:SI-A
  (read-normalized-illuminant (in-vicinity (library-vicinity) "ciesia.dat")))
(define CIE:SI-D65
  (read-normalized-illuminant (in-vicinity (library-vicinity) "ciesid65.dat")))
(spectrum->XYZ CIE:SI-A 300e-9 830e-9)
⇒ (1.098499460820401 999.9999999999998e-3 355.8173930654951e-3)
(CIEXYZ->sRGB (spectrum->XYZ CIE:SI-A 300e-9 830e-9))
⇒ (255 234 133)
(spectrum->XYZ CIE:SI-D65 300e-9 830e-9)
⇒ (950.4336673552745e-3 1.0000000000000002 1.0888053986649182)
(CIEXYZ->sRGB (spectrum->XYZ CIE:SI-D65 300e-9 830e-9))
⇒ (255 255 255)
```

illuminant-map *proc siv*   [Function]

    siv must be a one-dimensional array or vector of 107 numbers. `illuminant-map` returns a vector of length 107 containing the result of applying *proc* to each element of *siv*.

illuminant-map->XYZ *proc siv*   [Function]

```
(spectrum->XYZ (illuminant-map proc siv) 300e-9 830e-9)
```

spectrum->XYZ *proc*   [Function]

    proc must be a function of one argument. `spectrum->XYZ` computes the CIEXYZ(1931) values for the spectrum returned by *proc* when called with arguments from 380e-9 to 780e-9, the wavelength in meters.

spectrum->XYZ *spectrum x1 x2*   [Function]

    x1 and *x2* must be positive real numbers specifying the wavelengths (in meters) corresponding to the zeroth and last elements of vector or list *spectrum*. `spectrum->XYZ` returns the CIEXYZ(1931) values for a light source with spectral values proportional to the elements of *spectrum* at evenly spaced wavelengths between *x1* and *x2*.

Compute the colors of 6500.K and 5000.K blackbody radiation:

```
(require 'color-space)
(define xyz (spectrum->XYZ (blackbody-spectrum 6500)))
(define y_n (cadr xyz))
(map (lambda (x) (/ x y_n)) xyz)
    ⇒ (0.9687111145512467 1.0 1.1210875945303613)


(define xyz (spectrum->XYZ (blackbody-spectrum 5000)))
(map (lambda (x) (/ x y_n)) xyz)
    ⇒ (0.2933441826889158 0.2988931825387761 0.25783646831201573)
```

`spectrum->chromaticity` *proc*   [Function]
`spectrum->chromaticity` *spectrum x1 x2*   [Function]
> Computes the chromaticity for the given spectrum.

`wavelength->XYZ` *w*   [Function]
> *w* must be a number between 380e-9 to 780e-9. `wavelength->XYZ` returns (unnormalized) XYZ values for a monochromatic light source with wavelength *w*.

`wavelength->chromaticity` *w*   [Function]
> *w* must be a number between 380e-9 to 780e-9. `wavelength->chromaticity` returns the chromaticity for a monochromatic light source with wavelength *w*.

`blackbody-spectrum` *temp*   [Function]
`blackbody-spectrum` *temp span*   [Function]
> Returns a procedure of one argument (wavelength in meters), which returns the radiance of a black body at *temp*.

> The optional argument *span* is the wavelength analog of bandwidth. With the default *span* of 1.nm (1e-9.m), the values returned by the procedure correspond to the power of the photons with wavelengths *w* to *w*+1e-9.

`temperature->XYZ` *x*   [Function]
> The positive number *x* is a temperature in degrees kelvin. `temperature->XYZ` computes the unnormalized CIEXYZ(1931) values for the spectrum of a black body at temperature *x*.

> Compute the chromaticities of 6500.K and 5000.K blackbody radiation:

```
(require 'color-space)
(XYZ->chromaticity (temperature->XYZ 6500))
    ⇒ (0.3135191660557008 0.3236456786200268)

(XYZ->chromaticity (temperature->XYZ 5000))
    ⇒ (0.34508082841161052 0.3516084965163377)
```

`temperature->chromaticity` *x*   [Function]
> The positive number *x* is a temperature in degrees kelvin. `temperature->cromaticity` computes the chromaticity for the spectrum of a black body at temperature *x*.

Compute the chromaticities of 6500.K and 5000.K blackbody radiation:

```
(require 'color-space)
(temperature->chromaticity 6500)
    ⇒ (0.3135191660557008 0.3236456786200268)

(temperature->chromaticity 5000)
    ⇒ (0.34508082841161052 0.3516084965163377)
```

XYZ->chromaticity *xyz*   [Function]
> Returns a two element list: the x and y components of *xyz* normalized to 1 (= x + y + z).

chromaticity->CIEXYZ *x y*   [Function]
> Returns the list of x, and y, 1 - y - x.

chromaticity->whitepoint *x y*   [Function]
> Returns the CIEXYZ(1931) values having luminosity 1 and chromaticity x and y.

Many color datasets are expressed in *xyY* format; chromaticity with CIE luminance (Y). But xyY is not a CIE standard like CIEXYZ, CIELAB, and CIELUV. Although chrominance is well defined, the luminance component is sometimes scaled to 1, sometimes to 100, but usually has no obvious range. With no given whitepoint, the only reasonable course is to ascertain the luminance range of a dataset and normalize the values to lie from 0 to 1.

XYZ->xyY *xyz*   [Function]
> Returns a three element list: the x and y components of *XYZ* normalized to 1, and CIE luminance *Y*.

xyY->XYZ *xyY*   [Function]

xyY:normalize-colors *colors*   [Function]
> *colors* is a list of xyY triples. `xyY:normalize-colors` scales each chromaticity so it sums to 1 or less; and divides the *Y* values by the maximum *Y* in the dataset, so all lie between 0 and 1.

xyY:normalize-colors *colors n*   [Function]
> If *n* is positive real, then `xyY:normalize-colors` divides the *Y* values by *n* times the maximum *Y* in the dataset.
>
> If *n* is an exact non-positive integer, then `xyY:normalize-colors` divides the *Y* values by the maximum of the *Y*s in the dataset excepting the -*n* largest *Y* values.
>
> In all cases, returned *Y* values are limited to lie from 0 to 1.

Why would one want to normalize to other than 1? If the sun or its reflection is the brightest object in a scene, then normalizing to its luminance will tend to make the rest of the scene very dark. As with photographs, limiting the specular highlights looks better than darkening everything else.

The results of measurements being what they are, `xyY:normalize-colors` is extremely tolerant. Negative numbers are replaced with zero, and chromaticities with sums greater than one are scaled to sum to one.

5.11.4 Color Difference Metrics

`(require 'color-space)`

The low-level metric functions operate on lists of 3 numbers, lab1, lab2, lch1, or lch2.

`(require 'color)`

The wrapped functions operate on objects of type color, color1 and color2 in the function entries.

`L*a*b*:DE*` *lab1 lab2*   [Function]
    Returns the Euclidean distance between *lab1* and *lab2*.

`CIE:DE*` *color1 color2 white-point*   [Function]
`CIE:DE*` *color1 color2*   [Function]
    Returns the Euclidean distance in L*a*b* space between *color1* and *color2*.

`L*a*b*:DE*94` *lab1 lab2 parametric-factors*   [Function]
`L*a*b*:DE*94` *lab1 lab2*   [Function]
`CIE:DE*94` *color1 color2 parametric-factors*   [Function]
`CIE:DE*94` *color1 color2*   [Function]
    Measures distance in the L*a*b* color-space. The three axes are individually scaled in their contributions to the total distance.

    `DE*94` is not symmetrical in its arguments. *lab1* is the "reference" color and *lab2* is the "sample" color.

    The CIE has defined reference conditions under which the metric with default parameters can be expected to perform well. These are:

- The specimens are homogeneous in colour.
- The colour difference (CIELAB) is `<=` 5 units.
- They are placed in direct edge contact.
- Each specimen subtends an angle of `>`4 degrees to the assessor, whose colour vision is normal.
- They are illuminated at 1000 lux, and viewed against a background of uniform grey, with L* of 50, under illumination simulating D65.

    The *parametric-factors* argument is a list of 3 quantities kL, kC and kH. *parametric-factors* independently adjust each colour-difference term to account for any deviations from the reference viewing conditions. Under the reference conditions explained above, the default is kL = kC = kH = 1.

The Color Measurement Committee of The Society of Dyers and Colorists in Great Britain created a more sophisticated color-distance function for use in judging the consistency of dye lots. With CMC:DE* it is possible to use a single value pass/fail tolerance for all shades.

`CMC-DE` *lch1 lch2 parametric-factors*   [Function]
`CMC-DE` *lch1 lch2 l c*   [Function]
`CMC-DE` *lch1 lch2 l*   [Function]
`CMC-DE` *lch1 lch2*   [Function]
`CMC:DE*` *color1 color2 l c*   [Function]

CMC:DE* *color1 color2*   [Function]

> **CMC:DE** is a L*C*h metric. The *parametric-factors* argument is a list of 2 numbers *l*
> and *c*. *l* and *c* parameterize this metric. 1 and 1 are recommended for perceptibility;
> the default, 2 and 1, for acceptability.

5.11.5 Color Conversions

This package contains the low-level color conversion and color metric routines operating on
lists of 3 numbers. There is no type or range checking.

> ```
> (require 'color-space)
> ```

CIEXYZ:D65   [Constant]

> Is the color of 6500.K (blackbody) illumination. D65 is close to the average color of
> daylight.

CIEXYZ:D50   [Constant]

> Is the color of 5000.K (blackbody) illumination. D50 is the color of indoor lighting
> by incandescent bulbs.

CIEXYZ:A   [Constant]
CIEXYZ:B   [Constant]
CIEXYZ:C   [Constant]
CIEXYZ:E   [Constant]

> CIE 1931 illuminants normalized to 1 = y.

color:linear-transform *matrix row*   [Function]

CIEXYZ->RGB709 *xyz*   [Function]
RGB709->CIEXYZ *srgb*   [Function]

CIEXYZ->L*u*v* *xyz white-point*   [Function]
CIEXYZ->L*u*v* *xyz*   [Function]
L*u*v*->CIEXYZ $L*u*v*$ *white-point*   [Function]
L*u*v*->CIEXYZ $L*u*v*$ [Function]

> The *white-point* defaults to CIEXYZ:D65.

CIEXYZ->L*a*b* *xyz white-point*   [Function]
CIEXYZ->L*a*b* *xyz*   [Function]
L*a*b*->CIEXYZ $L*a*b*$ *white-point*   [Function]
L*a*b*->CIEXYZ $L*a*b*$ [Function]

> The XYZ *white-point* defaults to CIEXYZ:D65.

L*a*b*->L*C*h $L*a*b*$ [Function]
L*C*h->L*a*b* $L*C*h$ [Function]

CIEXYZ->sRGB *xyz*   [Function]
sRGB->CIEXYZ *srgb*   [Function]

CIEXYZ->xRGB *xyz*   [Function]
xRGB->CIEXYZ *srgb*   [Function]

sRGB->xRGB *xyz*   [Function]

xRGB->sRGB *srgb*   [Function]

CIEXYZ->e-sRGB *n xyz*   [Function]
e-sRGB->CIEXYZ *n srgb*   [Function]

sRGB->e-sRGB *n srgb*   [Function]
e-sRGB->sRGB *n srgb*   [Function]
>   The integer *n* must be 10, 12, or 16. Because sRGB and e-sRGB use the same RGB709
>   chromaticities, conversion between them is simpler than conversion through CIEXYZ.

Do not convert e-sRGB precision through `e-sRGB->sRGB` then `sRGB->e-sRGB` – values would
be truncated to 8-bits!

e-sRGB->e-sRGB *n1 srgb n2*   [Function]
>   The integers *n1* and *n2* must be 10, 12, or 16. `e-sRGB->e-sRGB` converts *srgb* to
>   e-sRGB of precision *n2*.

5.11.6 Color Names

`(require 'color-names)`

Rather than ballast the color dictionaries with numbered grays, `file->color-dictionary`
discards them. They are provided through the `grey` procedure:

grey *k*   [Function]
>   Returns `(inexact->exact (round (* k 2.55)))`, the X11 color grey<*k*>.

A color dictionary is a database table relating *canonical* color-names to color-strings (see
Section 5.11.1 [Color Data-Type], page 134).

The column names in a color dictionary are unimportant; the first field is the key, and the
second is the color-string.

color-name:canonicalize *name*   [Function]
>   Returns a downcased copy of the string or symbol *name* with '`_`', '`-`', and whitespace
>   removed.

color-name->color *name table1 table2 . . .*   [Function]
>   *table1*, *table2*, . . . must be color-dictionary tables. `color-name->color` searches for
>   the canonical form of *name* in *table1*, *table2*, . . . in order; returning the color-string
>   of the first matching record; #f otherwise.

color-dictionaries->lookup *table1 table2 . . .*   [Function]
>   *table1*, *table2*, . . . must be color-dictionary tables. `color-dictionaries->lookup`
>   returns a procedure which searches for the canonical form of its string argument in
>   *table1*, *table2*, . . .; returning the color-string of the first matching record; and #f
>   otherwise.

color-dictionary *name rdb base-table-type*   [Function]
>   *rdb* must be a string naming a relational database file; and the symbol *name* a table
>   therein. The database will be opened as *base-table-type*. `color-dictionary` returns
>   the read-only table *name* in database *name* if it exists; #f otherwise.

`color-dictionary` *name rdb*   [Function]
>   *rdb* must be an open relational database or a string naming a relational database file; and the symbol *name* a table therein. `color-dictionary` returns the read-only table *name* in database *name* if it exists; #f otherwise.

`load-color-dictionary` *name rdb base-table-type*   [Function]
`load-color-dictionary` *name rdb*   [Function]
>   *rdb* must be a string naming a relational database file; and the symbol *name* a table therein. If the symbol *base-table-type* is provided, the database will be opened as *base-table-type*. `load-color-dictionary` creates a top-level definition of the symbol *name* to a lookup procedure for the color dictionary *name* in *rdb*.

>   The value returned by `load-color-dictionary` is unspecified.

Dictionary Creation

`(require 'color-database)`

`file->color-dictionary` *file table-name rdb base-table-type*   [Function]
`file->color-dictionary` *file table-name rdb*   [Function]
>   *rdb* must be an open relational database or a string naming a relational database file, *table-name* a symbol, and the string *file* must name an existing file with color-names and their corresponding xRGB (6-digit hex) values. `file->color-dictionary` creates a table *table-name* in *rdb* and enters the associations found in *file* into it.

`url->color-dictionary` *url table-name rdb base-table-type*   [Function]
`url->color-dictionary` *url table-name rdb*   [Function]
>   *rdb* must be an open relational database or a string naming a relational database file and *table-name* a symbol. `url->color-dictionary` retrieves the resource named by the string *url* using the *wget* program; then calls `file->color-dictionary` to enter its associations in *table-name* in *url*.

This section has detailed the procedures for creating and loading color dictionaries. So where are the dictionaries to load?

>   `http://people.csail.mit.edu/jaffer/Color/Dictionaries.html`

Describes and evaluates several color-name dictionaries on the web. The following procedure creates a database containing two of these dictionaries.

`make-slib-color-name-db`   [Function]
>   Creates an alist-table relational database in library-vicinity containing the *Resene* and *saturate* color-name dictionaries.

>   If the files `resenecolours.txt`, `nbs-iscc.txt`, and `saturate.txt` exist in the library-vicinity, then they used as the source of color-name data. Otherwise, `make-slib-color-name-db` calls url->color-dictionary with the URLs of appropriate source files.

The Short List

`(require 'saturate)`

saturate *name*   [Function]

> Looks for *name* among the 19 saturated colors from *Approximate Colors on CIE Chromaticity Diagram*:

reddish orange	orange	yellowish orange	yellow
greenish yellow	yellow green	yellowish green	green
bluish green	blue green	greenish blue	blue
purplish blue	bluish purple	purple	reddish purple
red purple	purplish red	red	

> (`http://people.csail.mit.edu/jaffer/Color/saturate.pdf`). If *name* is found, the corresponding color is returned. Otherwise #f is returned. Use saturate only for light source colors.

Resene Paints Limited, New Zealand's largest privately-owned and operated paint manufacturing company, has generously made their *Resene RGB Values List* available.

```
(require 'resene)
```

resene *name*   [Function]

> Looks for *name* among the 1300 entries in the Resene color-name dictionary (`http://people.csail.mit.edu/jaffer/Color/resene.pdf`). If *name* is found, the corresponding color is returned. Otherwise #f is returned. The *Resene RGB Values List* is an excellent source for surface colors.

If you include the *Resene RGB Values List* in binary form in a program, then you must include its license with your program:

> Resene RGB Values List
> For further information refer to http://www.resene.co.nz
> Copyright Resene Paints Ltd 2001

> Permission to copy this dictionary, to modify it, to redistribute it, to distribute modified versions, and to use it for any purpose is granted, subject to the following restrictions and understandings.

> 1. Any text copy made of this dictionary must include this copyright notice in full.
>
> 2. Any redistribution in binary form must reproduce this copyright notice in the documentation or other materials provided with the distribution.
>
> 3. Resene Paints Ltd makes no warranty or representation that this dictionary is error-free, and is under no obligation to provide any services, by way of maintenance, update, or otherwise.
>
> 4. There shall be no use of the name of Resene or Resene Paints Ltd in any advertising, promotional, or sales literature without prior written consent in each case.
>
> 5. These RGB colour formulations may not be used to the detriment of Resene Paints Ltd.

5.11.7 Daylight

```
(require 'daylight)
```

This package calculates the colors of sky as detailed in:
`http://www.cs.utah.edu/vissim/papers/sunsky/sunsky.pdf`
A Practical Analytic Model for Daylight
A. J. Preetham, Peter Shirley, Brian Smits

solar-hour *julian-day hour*   [Function]
> Returns the solar-time in hours given the integer *julian-day* in the range 1 to 366, and the local time in hours.

> To be meticulous, subtract 4 minutes for each degree of longitude west of the standard meridian of your time zone.

solar-declination *julian-day*   [Function]

solar-polar *declination latitude solar-hour*   [Function]
> Returns a list of *theta_s*, the solar angle from the zenith, and *phi_s*, the solar azimuth. $0 <= theta\_s$ measured in degrees. *phi_s* is measured in degrees from due south; west of south being positive.

In the following procedures, the number $0 <= theta\_s <= 90$ is the solar angle from the zenith in degrees.

Turbidity is a measure of the fraction of scattering due to haze as opposed to molecules. This is a convenient quantity because it can be estimated based on visibility of distant objects. This model fails for turbidity values less than 1.3.

```
      -----------------------------------------------------------
  512|-:                                                         |
     | * pure-air                                                |
  256|-:**                                                       |
     | : ** exceptionally-clear                                  |
  128|-:   *                                                     |
     | :    **                                                   |
   64|-:      *                                                  |
     | :        ** very-clear                                    |
   32|-:        **                                               |
     | :          **                                             |
   16|-:           *** clear                                     |
     | :           ****                                          |
    8|-:              ****                                       |
     | :                **** light-haze                          |
    4|-:                 ****                                    |
     | :                   ******                                |
    2|-:                       ******* haze          thin-|
     | :                         **********    fog |
    1|-:-----------------------------------------------*******--|
     |_:____.____:____.____:____.____:____.____:____.____:____.____:_|
        1        2        4        8       16       32       64
           Meterorological range (km) versus Turbidity
```
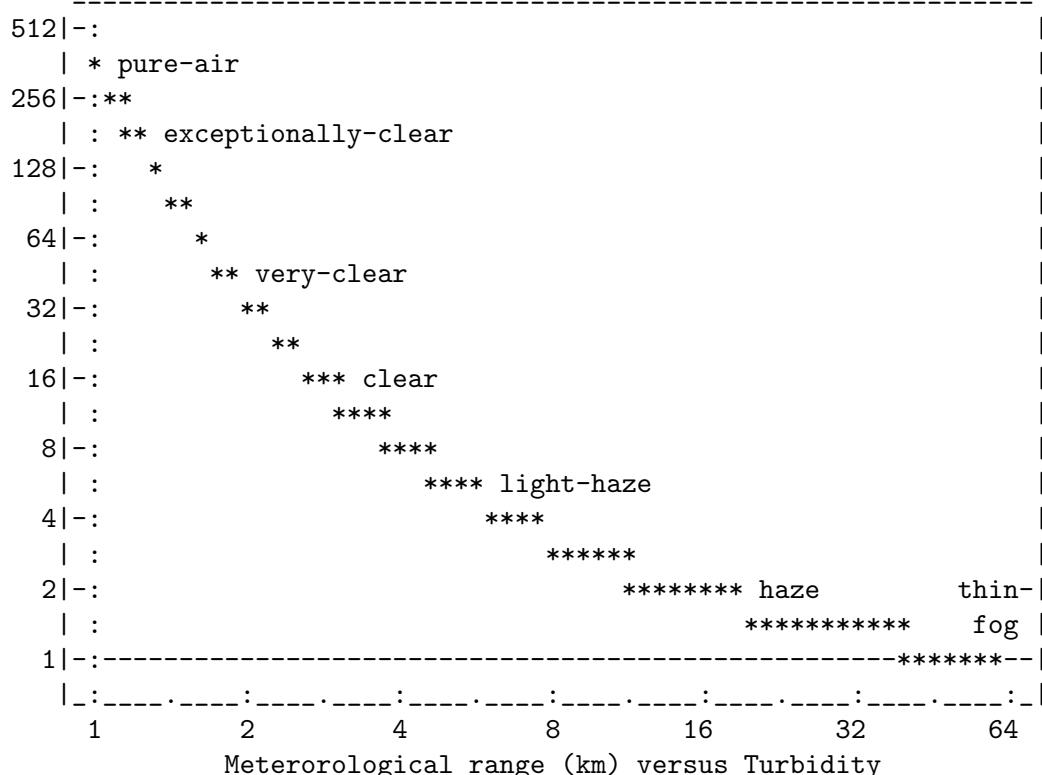
sunlight-spectrum *turbidity theta_s*   [Function]
>   Returns a vector of 41 values, the spectrum of sunlight from 380.nm to 790.nm for a
>   given *turbidity* and *theta_s*.

sunlight-chromaticity *turbidity theta_s*   [Function]
>   Given *turbidity* and *theta_s*, **sunlight-chromaticity** returns the CIEXYZ triple for
>   color of sunlight scaled to be just inside the RGB709 gamut.

zenith-xyy *turbidity theta_s*   [Function]
>   Returns the xyY (chromaticity and luminance) at the zenith. The Luminance has
>   units kcd/m^2.

overcast-sky-color-xyy *turbidity theta_s*   [Function]
>   *turbidity* is a positive real number expressing the amount of light scattering. The
>   real number *theta_s* is the solar angle from the zenith in degrees.
>
>   **overcast-sky-color-xyy** returns a function of one angle *theta*, the angle from the
>   zenith of the viewing direction (in degrees); and returning the xyY value for light
>   coming from that elevation of the sky.

clear-sky-color-xyy *turbidity theta_s phi_s*   [Function]
sky-color-xyy *turbidity theta_s phi_s*   [Function]
>   *turbidity* is a positive real number expressing the amount of light scattering. The
>   real number *theta_s* is the solar angle from the zenith in degrees. The real number
>   *phi_s* is the solar angle from south.
>
>   **clear-sky-color-xyy** returns a function of two angles, *theta* and *phi* which specify
>   the angles from the zenith and south meridian of the viewing direction (in degrees);
>   returning the xyY value for light coming from that direction of the sky.
>
>   **sky-color-xyY** calls **overcast-sky-color-xyY** for *turbidity* <= 20; otherwise the
>   **clear-sky-color-xyy** function.

5.12 Root Finding

`(require 'root)`

In the Newton method, divide the *df/dx* argument by the multiplicity of the desired
root in order to preserve quadratic convergence.

newton:find-integer-root *f df/dx x0*   [Function]
>   Given integer valued procedure *f*, its derivative (with respect to its argument) *df/dx*,
>   and initial integer value *x0* for which *df/dx(x0)* is non-zero, returns an integer *x* for
>   which $f(x)$ is closer to zero than either of the integers adjacent to *x*; or returns **#f** if
>   such an integer can't be found.
>
>   To find the closest integer to a given integer's square root:

```
(define (integer-sqrt y)
  (newton:find-integer-root
   (lambda (x) (- (* x x) y))
   (lambda (x) (* 2 x))
   (ash 1 (quotient (integer-length y) 2))))
```

```
(integer-sqrt 15) ⇒ 4
```

newton:find-root *f df/dx x0 prec*   [Function]

> Given real valued procedures *f*, *df/dx* of one (real) argument, initial real value *x0* for which *df/dx(x0)* is non-zero, and positive real number *prec*, returns a real *x* for which $\text{abs}(f(x))$ is less than *prec*; or returns **#f** if such a real can't be found.
>
> If *prec* is instead a negative integer, **newton:find-root** returns the result of -*prec* iterations.

H. J. Orchard, *The Laguerre Method for Finding the Zeros of Polynomials*, IEEE Transactions on Circuits and Systems, Vol. 36, No. 11, November 1989, pp 1377-1381.

> There are 2 errors in Orchard's Table II. Line k=2 for starting value of 1000+j0 should have Z_k of 1.0475 + j4.1036 and line k=2 for starting value of 0+j1000 should have Z_k of 1.0988 + j4.0833.

laguerre:find-root *f df/dz ddf/dz^2 z0 prec*   [Function]

> Given complex valued procedure *f* of one (complex) argument, its derivative (with respect to its argument) *df/dx*, its second derivative *ddf/dz^2*, initial complex value *z0*, and positive real number *prec*, returns a complex number *z* for which $\text{magnitude}(f(z))$ is less than *prec*; or returns **#f** if such a number can't be found.
>
> If *prec* is instead a negative integer, **laguerre:find-root** returns the result of -*prec* iterations.

laguerre:find-polynomial-root *deg f df/dz ddf/dz^2 z0 prec*   [Function]

> Given polynomial procedure *f* of integer degree *deg* of one argument, its derivative (with respect to its argument) *df/dx*, its second derivative *ddf/dz^2*, initial complex value *z0*, and positive real number *prec*, returns a complex number *z* for which $\text{magnitude}(f(z))$ is less than *prec*; or returns **#f** if such a number can't be found.
>
> If *prec* is instead a negative integer, **laguerre:find-polynomial-root** returns the result of -*prec* iterations.

secant:find-root *f x0 x1 prec*   [Function]
secant:find-bracketed-root *f x0 x1 prec*   [Function]

> Given a real valued procedure *f* and two real valued starting points *x0* and *x1*, returns a real *x* for which **(abs (f x))** is less than *prec*; or returns **#f** if such a real can't be found.
>
> If *x0* and *x1* are chosen such that they bracket a root, that is
>
> ```
> (or (< (f x0) 0 (f x1))
>     (< (f x1) 0 (f x0)))
> ```
>
> then the root returned will be between *x0* and *x1*, and *f* will not be passed an argument outside of that interval.
>
> **secant:find-bracketed-root** will return **#f** unless *x0* and *x1* bracket a root.
>
> The secant method is used until a bracketing interval is found, at which point a modified *regula falsi* method is used.

If *prec* is instead a negative integer, `secant:find-root` returns the result of -*prec* iterations.

If *prec* is a procedure it should accept 5 arguments: *x0 f0 x1 f1* and *count*, where *f0* will be (`f x0`), *f1* (`f x1`), and *count* the number of iterations performed so far. *prec* should return non-false if the iteration should be stopped.

5.13 Minimizing

`(require 'minimize)`

The Golden Section Search[3] algorithm finds minima of functions which are expensive to compute or for which derivatives are not available. Although optimum for the general case, convergence is slow, requiring nearly 100 iterations for the example (x^3-2x-5).

If the derivative is available, Newton-Raphson is probably a better choice. If the function is inexpensive to compute, consider approximating the derivative.

`golden-section-search` *f x0 x1 prec*   [Function]
> x_0 are x_1 real numbers. The (single argument) procedure *f* is unimodal over the open interval (x_0, x_1). That is, there is exactly one point in the interval for which the derivative of *f* is zero.
>
> `golden-section-search` returns a pair (x . $f(x)$) where $f(x)$ is the minimum. The *prec* parameter is the stop criterion. If *prec* is a positive number, then the iteration continues until x is within *prec* from the true value. If *prec* is a negative integer, then the procedure will iterate -*prec* times or until convergence. If *prec* is a procedure of seven arguments, *x0*, *x1*, *a*, *b*, *fa*, *fb*, and *count*, then the iterations will stop when the procedure returns `#t`.
>
> Analytically, the minimum of x^3-2x-5 is 0.816497.
>
> ```
> (define func (lambda (x) (+ (* x (+ (* x x) -2)) -5)))
> (golden-section-search func 0 1 (/ 10000))
>       ==> (816.4883855245578e-3 . -6.0886621077391165)
> (golden-section-search func 0 1 -5)
>       ==> (819.6601125010515e-3 . -6.088637561916407)
> (golden-section-search func 0 1
>                         (lambda (a b c d e f g ) (= g 500)))
>       ==> (816.4965933140557e-3 . -6.088662107903635)
> ```

5.14 The Limit

`limit` *proc x1 x2 k*   [library procedure]
`limit` *proc x1 x2*   [library procedure]
> *Proc* must be a procedure taking a single inexact real argument. K is the number of points on which *proc* will be called; it defaults to 8.
>
> If *x1* is finite, then *Proc* must be continuous on the half-open interval:

---

[3] David Kahaner, Cleve Moler, and Stephen Nash *Numerical Methods and Software* Prentice-Hall, 1989, ISBN 0-13-627258-4

(*x1* .. *x1+x2*]

And *x2* should be chosen small enough so that *proc* is expected to be monotonic or constant on arguments between *x1* and *x1* + *x2*.

`Limit` computes the limit of *proc* as its argument approaches *x1* from *x1* + *x2*. `Limit` returns a real number or real infinity or '`#f`'.

If *x1* is not finite, then *x2* must be a finite nonzero real with the same sign as *x1*; in which case `limit` returns:

`(limit (lambda (x) (proc (/ x))) 0.0 (/ x2) k)`

`Limit` examines the magnitudes of the differences between successive values returned by *proc* called with a succession of numbers from *x1+x2/k* to *x1*.

If the magnitudes of differences are monotonically decreasing, then then the limit is extrapolated from the degree n polynomial passing through the samples returned by *proc*.

If the magnitudes of differences are increasing as fast or faster than a hyperbola matching at *x1+x2*, then a real infinity with sign the same as the differences is returned.

If the magnitudes of differences are increasing more slowly than the hyperbola matching at *x1+x2*, then the limit is extrapolated from the quadratic passing through the three samples closest to *x1*.

If the magnitudes of differences are not monotonic or are not completely within one of the above categories, then #f is returned.

```
;; constant
(limit (lambda (x) (/ x x)) 0 1.0e-9)          ==> 1.0
(limit (lambda (x) (expt 0 x)) 0 1.0e-9)       ==> 0.0
(limit (lambda (x) (expt 0 x)) 0 -1.0e-9)      ==> +inf.0
;; linear
(limit + 0 976.5625e-6)                        ==> 0.0
(limit - 0 976.5625e-6)                        ==> 0.0
;; vertical point of inflection
(limit sqrt 0 1.0e-18)                         ==> 0.0
(limit (lambda (x) (* x (log x))) 0 1.0e-9)    ==> -102.70578127633066e-12
(limit (lambda (x) (/ x (log x))) 0 1.0e-9)    ==> 96.12123142321669e-15
;; limits tending to infinity
(limit + +inf.0 1.0e9)                         ==> +inf.0
(limit + -inf.0 -1.0e9)                        ==> -inf.0
(limit / 0 1.0e-9)                             ==> +inf.0
(limit / 0 -1.0e-9)                            ==> -inf.0
(limit (lambda (x) (/ (log x) x)) 0 1.0e-9)    ==> -inf.0
(limit (lambda (x) (/ (magnitude (log x)) x)) 0 -1.0e-9)
                                               ==> -inf.0
;; limit doesn't exist
(limit sin +inf.0 1.0e9)                       ==> #f
(limit (lambda (x) (sin (/ x))) 0 1.0e-9)      ==> #f
```

```
(limit (lambda (x) (sin (/ x))) 0 -1.0e-9)        ==> #f
(limit (lambda (x) (/ (log x) x)) 0 -1.0e-9)      ==> #f
;; conditionally convergent - return #f
(limit (lambda (x) (/ (sin x) x)) +inf.0 1.0e222)
                                                  ==> #f
;; asymptotes
(limit / -inf.0 -1.0e222)                         ==> 0.0
(limit / +inf.0 1.0e222)                          ==> 0.0
(limit (lambda (x) (expt x x)) 0 1.0e-18)         ==> 1.0
(limit (lambda (x) (sin (/ x))) +inf.0 1.0e222)   ==> 0.0
(limit (lambda (x) (/ (+ (exp (/ x)) 1))) 0 1.0e-9)
                                                  ==> 0.0
(limit (lambda (x) (/ (+ (exp (/ x)) 1))) 0 -1.0e-9)
                                                  ==> 1.0
(limit (lambda (x) (real-part (expt (tan x) (cos x)))) (/ pi 2) 1.0e-9)
                                                  ==> 1.0
;; This example from the 1979 Macsyma manual grows so rapidly
;;  that x2 must be less than 41.  It correctly returns e^2.
(limit (lambda (x) (expt (+ x (exp x) (exp (* 2 x))) (/ x))) +inf.0 40)
                                                  ==> 7.3890560989306504
;; LIMIT can calculate the proper answer when evaluation
;; of the function at the limit point does not:
(tan (atan +inf.0))                               ==> 16.331778728383844e15
(limit tan (atan +inf.0) -1.0e-15)                ==> +inf.0
(tan (atan +inf.0))                               ==> 16.331778728383844e15
(limit tan (atan +inf.0) 1.0e-15)                 ==> -inf.0
((lambda (x) (expt (exp (/ -1 x)) x)) 0)          ==> 1.0
(limit (lambda (x) (expt (exp (/ -1 x)) x)) 0 1.0e-9)
                                                  ==> 0.0
```

5.15 Commutative Rings

Scheme provides a consistent and capable set of numeric functions. Inexacts implement a field; integers a commutative ring (and Euclidean domain). This package allows one to use basic Scheme numeric functions with symbols and non-numeric elements of commutative rings.

```
(require 'commutative-ring)
```

The *commutative-ring* package makes the procedures +, -, *, /, and ^ *careful* in the sense that any non-numeric arguments they do not reduce appear in the expression output. In order to see what working with this package is like, self-set all the single letter identifiers (to their corresponding symbols).

```
(define a 'a)
...
(define z 'z)
```

Or just (require 'self-set). Now try some sample expressions:

```
(+ (+ a b) (- a b)) ⇒ (* a 2)
(* (+ a b) (+ a b)) ⇒ (^ (+ a b) 2)
(* (+ a b) (- a b)) ⇒ (* (+ a b) (- a b))
(* (- a b) (- a b)) ⇒ (^ (- a b) 2)
(* (- a b) (+ a b)) ⇒ (* (+ a b) (- a b))
(/ (+ a b) (+ c d)) ⇒ (/ (+ a b) (+ c d))
(^ (+ a b) 3) ⇒ (^ (+ a b) 3)
(^ (+ a 2) 3) ⇒ (^ (+ 2 a) 3)
```

Associative rules have been applied and repeated addition and multiplication converted to multiplication and exponentiation.

We can enable distributive rules, thus expanding to sum of products form:

```
(set! *ruleset* (combined-rulesets distribute* distribute/))
```

```
(* (+ a b) (+ a b)) ⇒ (+ (* 2 a b) (^ a 2) (^ b 2))
(* (+ a b) (- a b)) ⇒ (- (^ a 2) (^ b 2))
(* (- a b) (- a b)) ⇒ (- (+ (^ a 2) (^ b 2)) (* 2 a b))
(* (- a b) (+ a b)) ⇒ (- (^ a 2) (^ b 2))
(/ (+ a b) (+ c d)) ⇒ (+ (/ a (+ c d)) (/ b (+ c d)))
(/ (+ a b) (- c d)) ⇒ (+ (/ a (- c d)) (/ b (- c d)))
(/ (- a b) (- c d)) ⇒ (- (/ a (- c d)) (/ b (- c d)))
(/ (- a b) (+ c d)) ⇒ (- (/ a (+ c d)) (/ b (+ c d)))
(^ (+ a b) 3) ⇒ (+ (* 3 a (^ b 2)) (* 3 b (^ a 2)) (^ a 3) (^ b 3))
(^ (+ a 2) 3) ⇒ (+ 8 (* a 12) (* (^ a 2) 6) (^ a 3))
```

Use of this package is not restricted to simple arithmetic expressions:

```
(require 'determinant)
```

```
(determinant '((a b c) (d e f) (g h i))) ⇒
(- (+ (* a e i) (* b f g) (* c d h)) (* a f h) (* b d i) (* c e g))
```

Currently, only +, -, *, /, and ^ support non-numeric elements. Expressions with - are converted to equivalent expressions without -, so behavior for - is not defined separately. / expressions are handled similarly.

This list might be extended to include `quotient`, `modulo`, `remainder`, `lcm`, and `gcd`; but these work only for the more restrictive Euclidean (Unique Factorization) Domain.

5.16 Rules and Rulesets

The *commutative-ring* package allows control of ring properties through the use of *rulesets*.

`*ruleset*`   [Variable]

> Contains the set of rules currently in effect. Rules defined by `cring:define-rule` are stored within the value of *ruleset* at the time `cring:define-rule` is called. If *ruleset* is #f, then no rules apply.

`make-ruleset` *rule1 ...*   [Function]
`make-ruleset` *name rule1 ...*   [Function]
>   Returns a new ruleset containing the rules formed by applying `cring:define-rule`
>   to each 4-element list argument *rule*. If the first argument to `make-ruleset` is a
>   symbol, then the database table created for the new ruleset will be named *name*.
>   Calling `make-ruleset` with no rule arguments creates an empty ruleset.

`combined-rulesets` *ruleset1 ...*   [Function]
`combined-rulesets` *name ruleset1 ...*   [Function]
>   Returns a new ruleset containing the rules contained in each ruleset argument *ruleset*.
>   If the first argument to `combined-ruleset` is a symbol, then the database table
>   created for the new ruleset will be named *name*. Calling `combined-ruleset` with no
>   ruleset arguments creates an empty ruleset.

Two rulesets are defined by this package.

`distribute*`   [Constant]
>   Contains the ruleset to distribute multiplication over addition and subtraction.

`distribute/`   [Constant]
>   Contains the ruleset to distribute division over addition and subtraction.
>
>   Take care when using both *distribute** and *distribute/* simultaneously. It is possible
>   to put `/` into an infinite loop.

You can specify how sum and product expressions containing non-numeric elements
simplify by specifying the rules for `+` or `*` for cases where expressions involving objects
reduce to numbers or to expressions involving different non-numeric elements.

`cring:define-rule` *op sub-op1 sub-op2 reduction*   [Function]
>   Defines a rule for the case when the operation represented by symbol *op* is applied
>   to lists whose `cars` are *sub-op1* and *sub-op2*, respectively. The argument *reduction*
>   is a procedure accepting 2 arguments which will be lists whose `cars` are *sub-op1* and
>   *sub-op2*.

`cring:define-rule` *op sub-op1 'identity reduction*   [Function]
>   Defines a rule for the case when the operation represented by symbol *op* is applied to
>   a list whose `car` is *sub-op1*, and some other argument. *Reduction* will be called with
>   the list whose `car` is *sub-op1* and some other argument.
>
>   If *reduction* returns `#f`, the reduction has failed and other reductions will be tried.
>   If *reduction* returns a non-false value, that value will replace the two arguments in
>   arithmetic (`+`, `-`, and `*`) calculations involving non-numeric elements.
>
>   The operations `+` and `*` are assumed commutative; hence both orders of arguments
>   to *reduction* will be tried if necessary.
>
>   The following rule is the definition for distributing `*` over `+`.
>
>    ```
>    (cring:define-rule
>     '* '+ 'identity
>     (lambda (exp1 exp2)
>       (apply + (map (lambda (trm) (* trm exp2)) (cdr exp1)))))))
>    ```

5.17 How to Create a Commutative Ring

The first step in creating your commutative ring is to write procedures to create elements of
the ring. A non-numeric element of the ring must be represented as a list whose first element
is a symbol or string. This first element identifies the type of the object. A convenient and
clear convention is to make the type-identifying element be the same symbol whose top-level
value is the procedure to create it.

```
(define (n . list1)
  (cond ((and (= 2 (length list1))
              (eq? (car list1) (cadr list1)))
          0)
        ((not (term< (first list1) (last1 list1)))
         (apply n (reverse list1)))
        (else (cons 'n list1))))

(define (s x y) (n x y))

(define (m . list1)
  (cond ((neq? (first list1) (term_min list1))
         (apply m (cyclicrotate list1)))
        ((term< (last1 list1) (cadr list1))
         (apply m (reverse (cyclicrotate list1))))
        (else (cons 'm list1))))
```

Define a procedure to multiply 2 non-numeric elements of the ring. Other multipli-
catons are handled automatically. Objects for which rules have *not* been defined are not
changed.

```
(define (n*n ni nj)
  (let ((list1 (cdr ni)) (list2 (cdr nj)))
    (cond ((null? (intersection list1 list2)) #f)
          ((and (eq? (last1 list1) (first list2))
                (neq? (first list1) (last1 list2)))
           (apply n (splice list1 list2)))
          ((and (eq? (first list1) (first list2))
                (neq? (last1 list1) (last1 list2)))
           (apply n (splice (reverse list1) list2)))
          ((and (eq? (last1 list1) (last1 list2))
                (neq? (first list1) (first list2)))
           (apply n (splice list1 (reverse list2))))
          ((and (eq? (last1 list1) (first list2))
                (eq? (first list1) (last1 list2)))
           (apply m (cyclicsplice list1 list2)))
          ((and (eq? (first list1) (first list2))
                (eq? (last1 list1) (last1 list2)))
           (apply m (cyclicsplice (reverse list1) list2)))
          (else #f))))
```

Test the procedures to see if they work.

```
;;; where cyclicrotate(list) is cyclic rotation of the list one step
;;; by putting the first element at the end
(define (cyclicrotate list1)
  (append (rest list1) (list (first list1))))
;;; and where term_min(list) is the element of the list which is
;;; first in the term ordering.
(define (term_min list1)
  (car (sort list1 term<)))
(define (term< sym1 sym2)
  (string<? (symbol->string sym1) (symbol->string sym2)))
(define first car)
(define rest cdr)
(define (last1 list1) (car (last-pair list1)))
(define (neq? obj1 obj2) (not (eq? obj1 obj2)))
;;; where splice is the concatenation of list1 and list2 except that their
;;; common element is not repeated.
(define (splice list1 list2)
  (cond ((eq? (last1 list1) (first list2))
         (append list1 (cdr list2)))
        (else (slib:error 'splice list1 list2))))
;;; where cyclicsplice is the result of leaving off the last element of
;;; splice(list1,list2).
(define (cyclicsplice list1 list2)
  (cond ((and (eq? (last1 list1) (first list2))
              (eq? (first list1) (last1 list2)))
         (butlast (splice list1 list2) 1))
        (else (slib:error 'cyclicsplice list1 list2))))
```

```
(N*N (S a b) (S a b)) ⇒ (m a b)
```

Then register the rule for multiplying type N objects by type N objects.

```
(cring:define-rule '* 'N 'N N*N))
```

Now we are ready to compute!

```
(define (t)
  (define detM
    (+ (* (S g b)
          (+ (* (S f d)
                (- (* (S a f) (S d g)) (* (S a g) (S d f))))
             (* (S f f)
                (- (* (S a g) (S d d)) (* (S a d) (S d g))))
             (* (S f g)
                (- (* (S a d) (S d f)) (* (S a f) (S d d))))))
       (* (S g d)
          (+ (* (S f b)
                (- (* (S a g) (S d f)) (* (S a f) (S d g))))
             (* (S f f)
```

```
                        (- (* (S a b) (S d g)) (* (S a g) (S d b))))
                (* (S f g)
                   (- (* (S a f) (S d b)) (* (S a b) (S d f)))))))
           (* (S g f)
              (+ (* (S f b)
                    (- (* (S a d) (S d g)) (* (S a g) (S d d))))
                 (* (S f d)
                    (- (* (S a g) (S d b)) (* (S a b) (S d g))))
                 (* (S f g)
                    (- (* (S a b) (S d d)) (* (S a d) (S d b))))))
           (* (S g g)
              (+ (* (S f b)
                    (- (* (S a f) (S d d)) (* (S a d) (S d f))))
                 (* (S f d)
                    (- (* (S a b) (S d f)) (* (S a f) (S d b))))
                 (* (S f f)
                    (- (* (S a d) (S d b)) (* (S a b) (S d d)))))))))
     (* (S b e) (S c a) (S e c)
        detM
        ))
  (pretty-print (t))
  ⊣
  (- (+ (m a c e b d f g)
        (m a c e b d g f)
        (m a c e b f d g)
        (m a c e b f g d)
        (m a c e b g d f)
        (m a c e b g f d))
     (* 2 (m a b e c) (m d f g))
     (* (m a c e b d) (m f g))
     (* (m a c e b f) (m d g))
     (* (m a c e b g) (m d f)))
```

5.18 Matrix Algebra

`(require 'determinant)`

A Matrix can be either a list of lists (rows) or an array. Unlike linear-algebra texts, this package uses 0-based coordinates.

matrix->lists *matrix*   [Function]
    Returns the list-of-lists form of *matrix*.

matrix->array *matrix*   [Function]
    Returns the array form of *matrix*.

determinant *matrix*   [Function]
    matrix must be a square matrix. `determinant` returns the determinant of *matrix*.

```
        (require 'determinant)
```

```
(determinant '((1 2) (3 4))) ⇒ -2
(determinant '((1 2 3) (4 5 6) (7 8 9))) ⇒ 0
```

transpose *matrix*   [Function]
>   Returns a copy of *matrix* flipped over the diagonal containing the 1,1 element.

matrix:sum *m1 m2*   [Function]
>   Returns the element-wise sum of matricies *m1* and *m2*.

matrix:difference *m1 m2*   [Function]
>   Returns the element-wise difference of matricies *m1* and *m2*.

matrix:product *m1 m2*   [Function]
>   Returns the product of matrices *m1* and *m2*.

matrix:product *m1 z*   [Function]
>   Returns matrix *m1* times scalar *z*.

matrix:product *z m1*   [Function]
>   Returns matrix *m1* times scalar *z*.

matrix:inverse *matrix*   [Function]
>   *matrix* must be a square matrix. If *matrix* is singular, then `matrix:inverse` returns
>   #f; otherwise `matrix:inverse` returns the `matrix:product` inverse of *matrix*.

6 Database Packages

6.1 Relational Database

`(require 'relational-database)`

    This package implements a database system inspired by the Relational Model (*E. F. Codd, A Relational Model of Data for Large Shared Data Banks*). An SLIB relational database implementation can be created from any Section 6.2.1 [Base Table], page 179, implementation.

    Why relational database? For motivations and design issues see
`http://people.csail.mit.edu/jaffer/DBManifesto.html`.

6.1.1 Using Databases

`(require 'databases)`

This enhancement wraps a utility layer on `relational-database` which provides:

- Identification of open databases by filename.
- Automatic sharing of open (immutable) databases.
- Automatic loading of base-table package when creating a database.
- Detection and automatic loading of the appropriate base-table package when opening a database.
- Table and data definition from Scheme lists.

Database Sharing

Auto-sharing refers to a call to the procedure `open-database` returning an already open database (procedure), rather than opening the database file a second time.

> *Note:* Databases returned by `open-database` do not include wrappers applied by packages like Section 6.1.4 [Embedded Commands], page 169. But wrapped databases do work as arguments to these functions.

When a database is created, it is mutable by the creator and not auto-sharable. A database opened mutably is also not auto-sharable. But any number of readers can (open) share a non-mutable database file.

This next set of procedures mirror the whole-database methods in Section 6.2.4 [Database Operations], page 187. Except for `create-database`, each procedure will accept either a filename or database procedure for its first argument.

create-database *filename base-table-type*   [Function]
    filename should be a string naming a file; or `#f`. *base-table-type* must be a symbol naming a feature which can be passed to `require`. `create-database` returns a new, open relational database (with base-table type *base-table-type*) associated with *filename*, or a new ephemeral database if *filename* is `#f`.

    `create-database` is the only run-time use of require in SLIB which crosses module boundaries. When *base-table-type* is `require`d by `create-database`; it adds an association of *base-table-type* with its *relational-system* procedure to *mdbm:*databases**.

alist-table is the default base-table type:

```
(require 'databases)
(define my-rdb (create-database "my.db" 'alist-table))
```

Only `alist-table` and base-table modules which have been `required` will dispatch correctly from the `open-database` procedures. Therefore, either pass two arguments to `open-database`, or require the base-table of your database file uses before calling `open-database` with one argument.

open-database! *rdb base-table-type*   [Procedure]

> Returns *mutable* open relational database or #f.

open-database *rdb base-table-type*   [Function]

> Returns an open relational database associated with *rdb*. The database will be opened with base-table type *base-table-type*).

open-database *rdb*   [Function]

> Returns an open relational database associated with *rdb*. `open-database` will attempt to deduce the correct base-table-type.

write-database *rdb filename*   [Function]

> Writes the mutable relational-database *rdb* to *filename*.

sync-database *rdb*   [Function]

> Writes the mutable relational-database *rdb* to the filename it was opened with.

solidify-database *rdb*   [Function]

> Syncs *rdb* and makes it immutable.

close-database *rdb*   [Function]

> *rdb* will only be closed when the count of `open-database` - `close-database` calls for *rdb* (and its filename) is 0. `close-database` returns #t if successful; and #f otherwise.

mdbm:report   [Function]

> Prints a table of open database files. The columns are the base-table type, number of opens, '!' for mutable, the filename, and the lock certificate (if locked).

```
(mdbm:report)
⊣
  alist-table 003   /usr/local/lib/slib/clrnamdb.scm
  alist-table 001 ! sdram.db jaffer@aubrey.jaffer.3166:1038628199
```

Opening Tables

open-table *rdb table-name*   [Function]

> *rdb* must be a relational database and *table-name* a symbol.

> `open-table` returns a "methods" procedure for an existing relational table in *rdb* if it exists and can be opened for reading, otherwise returns #f.

`open-table!` *rdb table-name*   [Procedure]

>   *rdb* must be a relational database and *table-name* a symbol.
>
>   `open-table!` returns a "methods" procedure for an existing relational table in *rdb* if
>   it exists and can be opened in mutable mode, otherwise returns `#f`.

Defining Tables

`define-domains` *rdb row5 . . .*   [Function]

>   Adds the domain rows *row5 . . .* to the '`*domains-data*`' table in *rdb*. The format
>   of the row is given in Section 6.2.2 [Catalog Representation], page 184.
>
>    ```
>    (define-domains rdb '(permittivity #f complex? c64 #f))
>    ```

`add-domain` *rdb row5*   [Function]

>   Use `define-domains` instead.

`define-tables` *rdb spec-0 . . .*   [Function]

>   Adds tables as specified in *spec-0 . . .* to the open relational-database *rdb*. Each *spec*
>   has the form:
>
>   > (<name> <descriptor-name> <descriptor-name> <rows>)
>
>   or
>
>   > (<name> <primary-key-fields> <other-fields> <rows>)
>
>   where <name> is the table name, <descriptor-name> is the symbol name of a descriptor
>   table, <primary-key-fields> and <other-fields> describe the primary keys and other
>   fields respectively, and <rows> is a list of data rows to be added to the table.
>
>   <primary-key-fields> and <other-fields> are lists of field descriptors of the form:
>
>   > (<column-name> <domain>)
>
>   or
>
>   > (<column-name> <domain> <column-integrity-rule>)
>
>   where <column-name> is the column name, <domain> is the domain of the column,
>   and <column-integrity-rule> is an expression whose value is a procedure of one argu-
>   ment (which returns `#f` to signal an error).
>
>   If <domain> is not a defined domain name and it matches the name of this table or
>   an already defined (in one of *spec-0 . . .*) single key field table, a foreign-key domain
>   will be created for it.

Listing Tables

`list-table-definition` *rdb table-name*   [Function]

>   If symbol *table-name* exists in the open relational-database *rdb*, then returns a list of
>   the table-name, its primary key names and domains, its other key names and domains,
>   and the table's records (as lists). Otherwise, returns #f.
>
>   The list returned by `list-table-definition`, when passed as an argument to
>   `define-tables`, will recreate the table.

6.1.2 Table Operations

These are the descriptions of the methods available from an open relational table. A method is retrieved from a table by calling the table with the symbol name of the operation. For example:

```
((plat 'get 'processor) 'djgpp) ⇒ i386
```

Some operations described below require primary key arguments. Primary keys arguments are denoted *key1 key2 ....* It is an error to call an operation for a table which takes primary key arguments with the wrong number of primary keys for that table.

get *column-name*   [Operation on `relational-table`]

> Returns a procedure of arguments *key1 key2 ...* which returns the value for the *column-name* column of the row associated with primary keys *key1*, *key2 ...* if that row exists in the table, or `#f` otherwise.

```
((plat 'get 'processor) 'djgpp) ⇒ i386
((plat 'get 'processor) 'be-os) ⇒ #f
```

6.1.2.1 Single Row Operations

The term *row* used below refers to a Scheme list of values (one for each column) in the order specified in the descriptor (table) for this table. Missing values appear as `#f`. Primary keys must not be missing.

row:insert   [Operation on `relational-table`]

> Adds the row *row* to this table. If a row for the primary key(s) specified by *row* already exists in this table an error is signaled. The value returned is unspecified.

```
(define telephone-table-desc
        ((my-database 'create-table) 'telephone-table-desc))
(define ndrp (telephone-table-desc 'row:insert))
(ndrp '(1 #t name #f string))
(ndrp '(2 #f telephone
           (lambda (d)
             (and (string? d) (> (string-length d) 2)
                  (every
                   (lambda (c)
                     (memv c '(#\0 #\1 #\2 #\3 #\4 #\5 #\6 #\7 #\8 #\9
                                   #\+ #\( #\space #\) #\-)))
                   (string->list d))))
           string))
```

row:update   [Operation on `relational-table`]

> Returns a procedure of one argument, *row*, which adds the row, *row*, to this table. If a row for the primary key(s) specified by *row* already exists in this table, it will be overwritten. The value returned is unspecified.

row:retrieve   [Operation on `relational-table`]

> Returns a procedure of arguments *key1 key2 ...* which returns the row associated with primary keys *key1*, *key2 ...* if it exists, or `#f` otherwise.

```
((plat 'row:retrieve) 'linux) ⇒ (linux i386 linux gcc)
((plat 'row:retrieve) 'multics) ⇒ #f
```

row:remove   [Operation on `relational-table`]
> Returns a procedure of arguments *key1 key2* . . . which removes and returns the row associated with primary keys *key1*, *key2* . . . if it exists, or `#f` otherwise.

row:delete   [Operation on `relational-table`]
> Returns a procedure of arguments *key1 key2* . . . which deletes the row associated with primary keys *key1*, *key2* . . . if it exists. The value returned is unspecified.

6.1.2.2 Match-Keys

The (optional) *match-key1* . . . arguments are used to restrict actions of a whole-table operation to a subset of that table. Those procedures (returned by methods) which accept match-key arguments will accept any number of match-key arguments between zero and the number of primary keys in the table. Any unspecified *match-key* arguments default to `#f`.

The *match-key1* . . . restrict the actions of the table command to those records whose primary keys each satisfy the corresponding *match-key* argument. The arguments and their actions are:

> `#f`   The false value matches any key in the corresponding position.

> an object of type procedure
> > This procedure must take a single argument, the key in the corresponding position. Any key for which the procedure returns a non-false value is a match; Any key for which the procedure returns a `#f` is not.

> other values
> > Any other value matches only those keys `equal?` to it.

get* *column-name*   [Operation on `relational-table`]
> Returns a procedure of optional arguments *match-key1* . . . which returns a list of the values for the specified column for all rows in this table. The optional *match-key1* . . . arguments restrict actions to a subset of the table.

```
((plat 'get* 'processor)) ⇒
(i386 i8086 i386 i8086 i386 i386 i8086 m68000
 m68000 m68000 m68000 m68000 powerpc)

((plat 'get* 'processor) #f) ⇒
(i386 i8086 i386 i8086 i386 i386 i8086 m68000
 m68000 m68000 m68000 m68000 powerpc)

(define (a-key? key)
   (char=? #\a (string-ref (symbol->string key) 0)))

((plat 'get* 'processor) a-key?) ⇒
```

```
(m68000 m68000 m68000 m68000 m68000 powerpc)

((plat 'get* 'name) a-key?) ⇒
(atari-st-turbo-c atari-st-gcc amiga-sas/c-5.10
 amiga-aztec amiga-dice-c aix)
```

6.1.2.3 Multi-Row Operations

`row:retrieve*`   [Operation on `relational-table`]

Returns a procedure of optional arguments *match-key1* ... which returns a list of
all rows in this table. The optional *match-key1* ... arguments restrict actions to a
subset of the table. For details see See Section 6.1.2.2 [Match-Keys], page 166.

```
((plat 'row:retrieve*) a-key?) ⇒
((atari-st-turbo-c m68000 atari turbo-c)
 (atari-st-gcc m68000 atari gcc)
 (amiga-sas/c-5.10 m68000 amiga sas/c)
 (amiga-aztec m68000 amiga aztec)
 (amiga-dice-c m68000 amiga dice-c)
 (aix powerpc aix -))
```

`row:remove*`   [Operation on `relational-table`]

Returns a procedure of optional arguments *match-key1* ... which removes and returns
a list of all rows in this table. The optional *match-key1* ... arguments restrict actions
to a subset of the table.

`row:delete*`   [Operation on `relational-table`]

Returns a procedure of optional arguments *match-key1* ... which Deletes all rows
from this table. The optional *match-key1* ... arguments restrict deletions to a subset
of the table. The value returned is unspecified. The descriptor table and catalog entry
for this table are not affected.

`for-each-row`   [Operation on `relational-table`]

Returns a procedure of arguments *proc match-key1* ... which calls *proc* with each
row in this table. The optional *match-key1* ... arguments restrict actions to a subset
of the table. For details see See Section 6.1.2.2 [Match-Keys], page 166.

Note that `row:insert*` and `row:update*` do *not* use match-keys.

`row:insert*`   [Operation on `relational-table`]

Returns a procedure of one argument, *rows*, which adds each row in the list of rows,
rows, to this table. If a row for the primary key specified by an element of *rows*
already exists in this table, an error is signaled. The value returned is unspecified.

`row:update*`   [Operation on `relational-table`]

Returns a procedure of one argument, *rows*, which adds each row in the list of rows,
rows, to this table. If a row for the primary key specified by an element of *rows*
already exists in this table, it will be overwritten. The value returned is unspecified.

6.1.2.4 Indexed Sequential Access Methods

Indexed Sequential Access Methods are a way of arranging database information so that records can be accessed both by key and by key sequence (ordering). *ISAM* is not part of Codd's relational model.

Associative memory in B-Trees is an example of a database implementation which can support a native key ordering. SLIB's `alist-table` implementation uses `sort` to implement `for-each-row-in-order`, but does not support `isam-next` and `isam-prev`.

The multi-primary-key ordering employed by these operations is the lexicographic collation of those primary-key fields in their given order. For example:

```
(12 a 34) < (12 a 36) < (12 b 1) < (13 a 0)
```

6.1.2.5 Sequential Index Operations

The following procedures are individually optional depending on the base-table implememtation. If an operation is *not* supported, then calling the table with that operation symbol will return false.

`for-each-row-in-order`   [Operation on `relational-table`]
> Returns a procedure of arguments *proc match-key1* . . . which calls *proc* with each *row* in this table in the (implementation-dependent) natural, repeatable ordering for rows. The optional *match-key1* . . . arguments restrict actions to a subset of the table. For details see See Section 6.1.2.2 [Match-Keys], page 166.

`isam-next`   [Operation on `relational-table`]
> Returns a procedure of arguments *key1 key2* . . . which returns the key-list identifying the lowest record higher than *key1 key2* . . . which is stored in the relational-table; or false if no higher record is present.

`isam-next` *column-name*   [Operation on `relational-table`]
> The symbol *column-name* names a key field. In the list returned by `isam-next`, that field, or a field to its left, will be changed. This allows one to skip over less significant key fields.

`isam-prev`   [Operation on `relational-table`]
> Returns a procedure of arguments *key1 key2* . . . which returns the key-list identifying the highest record less than *key1 key2* . . . which is stored in the relational-table; or false if no lower record is present.

`isam-prev` *column-name*   [Operation on `relational-table`]
> The symbol *column-name* names a key field. In the list returned by `isam-next`, that field, or a field to its left, will be changed. This allows one to skip over less significant key fields.

> For example, if a table has key fields:

```
(col1 col2)
(9 5)
(9 6)
(9 7)
```

```
(9 8)
(12 5)
(12 6)
(12 7)
```

Then:

```
((table 'isam-next)        '(9 5))        ⇒ (9 6)
((table 'isam-next 'col2) '(9 5))        ⇒ (9 6)
((table 'isam-next 'col1) '(9 5))        ⇒ (12 5)
((table 'isam-prev)        '(12 7))       ⇒ (12 6)
((table 'isam-prev 'col2) '(12 7))       ⇒ (12 6)
((table 'isam-prev 'col1) '(12 7))       ⇒ (9 8)
```

6.1.2.6 Table Administration

`column-names`   [Operation on `relational-table`]
`column-foreigns`   [Operation on `relational-table`]
`column-domains`   [Operation on `relational-table`]
`column-types`   [Operation on `relational-table`]

Return a list of the column names, foreign-key table names, domain names, or type names respectively for this table. These 4 methods are different from the others in that the list is returned, rather than a procedure to obtain the list.

`primary-limit`   [Operation on `relational-table`]

Returns the number of primary keys fields in the relations in this table.

`close-table`   [Operation on `relational-table`]

Subsequent operations to this table will signal an error.

6.1.3 Database Interpolation

`(require 'database-interpolate)`

Indexed sequential access methods allow finding the keys (having associations) closest to a given value. This facilitates the interpolation of associations between those in the table.

`interpolate-from-table` *table column*   [Function]

Table should be a relational table with one numeric primary key field which supports the `isam-prev` and `isam-next` operations. *column* should be a symbol or exact positive integer designating a numerically valued column of *table*.

`interpolate-from-table` calculates and returns a value proportionally intermediate between its values in the next and previous key records contained in *table*. For keys larger than all the stored keys the value associated with the largest stored key is used. For keys smaller than all the stored keys the value associated with the smallest stored key is used.

6.1.4 Embedded Commands

`(require 'database-commands)`

This enhancement wraps a utility layer on `relational-database` which provides:

• Automatic execution of initialization commands stored in database.

- Transparent execution of database commands stored in *commands* table in database.

When an enhanced relational-database is called with a symbol which matches a *name* in the *commands* table, the associated procedure expression is evaluated and applied to the enhanced relational-database. A procedure should then be returned which the user can invoke on (optional) arguments.

The command *initialize* is special.   If present in the *commands* table, open-database or open-database! will return the value of the *initialize* command. Notice that arbitrary code can be run when the *initialize* procedure is automatically applied to the enhanced relational-database.

Note also that if you wish to shadow or hide from the user relational-database methods described in Section 6.2.4 [Database Operations], page 187, this can be done by a dispatch in the closure returned by the *initialize* expression rather than by entries in the *commands* table if it is desired that the underlying methods remain accessible to code in the *commands* table.

6.1.4.1 Database Extension

wrap-command-interface *rdb*   [Function]
> Returns relational database *rdb* wrapped with additional commands defined in its *commands* table.

add-command-tables *rdb*   [Function]
> The relational database *rdb* must be mutable. *add-command-tables* adds a *command* table to *rdb*; then returns (wrap-command-interface *rdb*).

define-*commands* *rdb spec-0* ...   [Function]
> Adds commands to the *commands* table as specified in *spec-0* ... to the open relational-database *rdb*. Each *spec* has the form:

>> ((<name> <rdb>) "comment" <expression1> <expression2> ...)

> or

>> ((<name> <rdb>) <expression1> <expression2> ...)

> where <name> is the command name, <rdb> is a formal passed the calling relational database, "comment" describes the command, and <expression1>, <expression1>, ... are the body of the procedure.

> define-*commands* adds to the *commands* table a command <name>:

>> (lambda (<name> <rdb>) <expression1> <expression2> ...)

open-command-database *filename*   [Function]
open-command-database *filename base-table-type*   [Function]
> Returns an open enhanced relational database associated with *filename*. The database will be opened with base-table type *base-table-type*) if supplied. If *base-table-type* is not supplied, open-command-database will attempt to deduce the correct base-table-type. If the database can not be opened or if it lacks the *commands* table, #f is returned.

open-command-database! *filename*   [Function]
open-command-database! *filename base-table-type*   [Function]
> Returns *mutable* open enhanced relational database . . .

open-command-database *database*   [Function]
> Returns *database* if it is an immutable relational database; #f otherwise.

open-command-database! *database*   [Function]
> Returns *database* if it is a mutable relational database; #f otherwise.

6.1.4.2 Command Intrinsics

Some commands are defined in all extended relational-databases. The are called just like Section 6.2.4 [Database Operations], page 187.

add-domain *domain-row*   [Operation on `relational-database`]
> Adds *domain-row* to the *domains* table if there is no row in the domains table associated with key (`car domain-row`) and returns #t. Otherwise returns #f.

> For the fields and layout of the domain table, See Section 6.2.2 [Catalog Representation], page 184. Currently, these fields are

> - domain-name
> - foreign-table
> - domain-integrity-rule
> - type-id
> - type-param

> The following example adds 3 domains to the 'build' database. 'Optstring' is either a string or #f. filename is a string and build-whats is a symbol.

```
(for-each (build 'add-domain)
          '((optstring #f
                       (lambda (x) (or (not x) (string? x)))
                       string
                       #f)
            (filename #f #f string #f)
            (build-whats #f #f symbol #f)))
```

delete-domain *domain-name*   [Operation on `relational-database`]
> Removes and returns the *domain-name* row from the *domains* table.

domain-checker *domain*   [Operation on `relational-database`]
> Returns a procedure to check an argument for conformance to domain *domain*.

6.1.4.3 Define-tables Example

The following example shows a new database with the name of `foo.db` being created with tables describing processor families and processor/os/compiler combinations. The database is then solidified; saved and changed to immutable.

```
(require 'databases)
```

```
(define my-rdb (create-database "foo.db" 'alist-table))
(define-tables my-rdb
  '(processor-family
    ((family    atom))
    ((also-ran  processor-family))
    ((m68000            #f)
     (m68030            m68000)
     (i386              i8086)
     (i8086             #f)
     (powerpc           #f)))

  '(platform
    ((name      symbol))
    ((processor processor-family)
     (os        symbol)
     (compiler  symbol))
    ((aix                 powerpc aix     -)
     (amiga-dice-c        m68000  amiga   dice-c)
     (amiga-aztec         m68000  amiga   aztec)
     (amiga-sas/c-5.10 m68000  amiga   sas/c)
     (atari-st-gcc        m68000  atari   gcc)
     (atari-st-turbo-c m68000  atari   turbo-c)
     (borland-c-3.1       i8086   ms-dos  borland-c)
     (djgpp               i386    ms-dos  gcc)
     (linux               i386    linux   gcc)
     (microsoft-c         i8086   ms-dos  microsoft-c)
     (os/2-emx            i386    os/2    gcc)
     (turbo-c-2           i8086   ms-dos  turbo-c)
     (watcom-9.0          i386    ms-dos  watcom))))

(solidify-database my-rdb)
```

6.1.4.4 The *commands* Table

The table *commands* in an *enhanced* relational-database has the fields (with domains):

```
PRI name          symbol
    parameters    parameter-list
    procedure     expression
    documentation string
```

The `parameters` field is a foreign key (domain `parameter-list`) of the *catalog-data* table and should have the value of a table described by *parameter-columns*. This `parameter-list` table describes the arguments suitable for passing to the associated command. The intent of this table is to be of a form such that different user-interfaces (for instance, pull-down menus or plain-text queries) can operate from the same table. A `parameter-list` table has the following fields:

```
PRI index       ordinal
    name        symbol
    arity       parameter-arity
    domain      domain
    defaulter   expression
    expander    expression
    documentation string
```

The `arity` field can take the values:

single   Requires a single parameter of the specified domain.

optional   A single parameter of the specified domain or zero parameters is acceptable.

boolean   A single boolean parameter or zero parameters (in which case `#f` is substituted)
           is acceptable.

nary   Any number of parameters of the specified domain are acceptable. The argu-
           ment passed to the command function is always a list of the parameters.

nary1   One or more of parameters of the specified domain are acceptable. The argu-
           ment passed to the command function is always a list of the parameters.

The `domain` field specifies the domain which a parameter or parameters in the `index`th
field must satisfy.

The `defaulter` field is an expression whose value is either `#f` or a procedure of one
argument (the parameter-list) which returns a *list* of the default value or values as appro-
priate. Note that since the `defaulter` procedure is called every time a default parameter
is needed for this column, *sticky* defaults can be implemented using shared state with the
domain-integrity-rule.

6.1.4.5 Command Service

make-command-server *rdb table-name*   [Function]
      Returns a procedure of 2 arguments, a (symbol) command and a call-back procedure.
      When this returned procedure is called, it looks up *command* in table *table-name*
      and calls the call-back procedure with arguments:

      command   The *command*

      command-value
                  The result of evaluating the expression in the *procedure* field of *table-
                  name* and calling it with *rdb*.

      parameter-name
                  A list of the *official* name of each parameter. Corresponds to the `name`
                  field of the *command*'s parameter-table.

      positions A list of the positive integer index of each parameter. Corresponds to the
                   `index` field of the *command*'s parameter-table.

      arities   A list of the arities of each parameter. Corresponds to the `arity` field
                   of the *command*'s parameter-table. For a description of `arity` see table
                   above.

types   A list of the type name of each parameter. Correspnds to the `type-id` field of the contents of the `domain` of the *command*'s parameter-table.

defaulters A list of the defaulters for each parameter.   Corresponds to the `defaulters` field of the *command*'s parameter-table.

domain-integrity-rules

   A list of procedures (one for each parameter) which tests whether a value for a parameter is acceptable for that parameter. The procedure should be called with each datum in the list for `nary` arity parameters.

aliases   A list of lists of (alias parameter-name). There can be more than one alias per *parameter-name*.

For information about parameters, See Section 4.4.4 [Parameter lists], page 62.

6.1.4.6 Command Example

Here is an example of setting up a command with arguments and parsing those arguments from a `getopt` style argument list (see Section 4.4.1 [Getopt], page 58).

```
(require 'database-commands)
(require 'databases)
(require 'getopt-parameters)
(require 'parameters)
(require 'getopt)
(require 'fluid-let)
(require 'printf)

(define my-rdb (add-command-tables (create-database #f 'alist-table)))

(define-tables my-rdb
  '(foo-params
    *parameter-columns*
    *parameter-columns*
    ((1 single-string single string
        (lambda (pl) '("str")) #f "single string")
     (2 nary-symbols nary symbol
        (lambda (pl) '()) #f "zero or more symbols")
     (3 nary1-symbols nary1 symbol
        (lambda (pl) '(symb)) #f "one or more symbols")
     (4 optional-number optional ordinal
        (lambda (pl) '()) #f "zero or one number")
     (5 flag boolean boolean
        (lambda (pl) '(#f)) #f "a boolean flag")))
  '(foo-pnames
    ((name string))
    ((parameter-index ordinal))
    (("s" 1)
     ("single-string" 1)
```

```
            ("n" 2)
            ("nary-symbols" 2)
            ("N" 3)
            ("nary1-symbols" 3)
            ("o" 4)
            ("optional-number" 4)
            ("f" 5)
            ("flag" 5)))
     '(my-commands
        ((name symbol))
        ((parameters parameter-list)
         (parameter-names parameter-name-translation)
         (procedure expression)
         (documentation string))
        ((foo
          foo-params
          foo-pnames
          (lambda (rdb) (lambda args (print args)))
          "test command arguments")))))

    (define (dbutil:serve-command-line rdb command-table command argv)
      (set! *argv* (if (vector? argv) (vector->list argv) argv))
      ((make-command-server rdb command-table)
       command
       (lambda (comname comval options positions
                        arities types defaulters dirs aliases)
         (apply comval (getopt->arglist options positions
                        arities types defaulters dirs aliases)))))

    (define (cmd . opts)
      (fluid-let ((*optind* 1))
        (printf "%-34s ⇒ "
                (call-with-output-string
                  (lambda (pt) (write (cons 'cmd opts) pt))))
        (set! opts (cons "cmd" opts))
        (force-output)
        (dbutil:serve-command-line
         my-rdb 'my-commands 'foo (length opts) opts)))

    (cmd)                              ⇒ ("str" () (symb) () #f)
    (cmd "-f")                         ⇒ ("str" () (symb) () #t)
    (cmd "--flag")                     ⇒ ("str" () (symb) () #t)
    (cmd "-o177")                      ⇒ ("str" () (symb) (177) #f)
    (cmd "-o" "177")                   ⇒ ("str" () (symb) (177) #f)
    (cmd "--optional" "621")           ⇒ ("str" () (symb) (621) #f)
    (cmd "--optional=621")             ⇒ ("str" () (symb) (621) #f)
    (cmd "-s" "speciality")            ⇒ ("speciality" () (symb) () #f)
```

```
(cmd "-sspeciality")                    ⇒ ("speciality" () (symb) () #f)
(cmd "--single" "serendipity")          ⇒ ("serendipity" () (symb) () #f)
(cmd "--single=serendipity")            ⇒ ("serendipity" () (symb) () #f)
(cmd "-n" "gravity" "piety")            ⇒ ("str" () (piety gravity) () #f)
(cmd "-ngravity" "piety")               ⇒ ("str" () (piety gravity) () #f)
(cmd "--nary" "chastity")               ⇒ ("str" () (chastity) () #f)
(cmd "--nary=chastity" "")              ⇒ ("str" () ( chastity) () #f)
(cmd "-N" "calamity")                   ⇒ ("str" () (calamity) () #f)
(cmd "-Ncalamity")                      ⇒ ("str" () (calamity) () #f)
(cmd "--nary1" "surety")                ⇒ ("str" () (surety) () #f)
(cmd "--nary1=surety")                  ⇒ ("str" () (surety) () #f)
(cmd "-N" "levity" "fealty")            ⇒ ("str" () (fealty levity) () #f)
(cmd "-Nlevity" "fealty")               ⇒ ("str" () (fealty levity) () #f)
(cmd "--nary1" "surety" "brevity") ⇒ ("str" () (brevity surety) () #f)
(cmd "--nary1=surety" "brevity")        ⇒ ("str" () (brevity surety) () #f)
(cmd "-?")
⊣
Usage: cmd [OPTION ARGUMENT ...] ...

  -f, --flag
  -o, --optional[=]<number>
  -n, --nary[=]<symbols> ...
  -N, --nary1[=]<symbols> ...
  -s, --single[=]<string>

ERROR: getopt->parameter-list "unrecognized option" "-?"
```

6.1.5 Database Macros

`(require 'within-database)`

The object-oriented programming interface to SLIB relational databases has failed to support clear, understandable, and modular code-writing for database applications.

This seems to be a failure of the object-oriented paradigm where the type of an object is not manifest (or even traceable) in source code.

`within-database`, along with the 'databases' package, reorganizes high-level database functions toward a more declarative style. Using this package, one can tag database table and command declarations for emacs:

```
etags -lscheme -r'/ *(define-\(command\|table\) (\([^; \t]+\))/\2/' \
      source1.scm ...
```

6.1.5.1 Within-database

within-database *database statement-1 ...*   [Function]

within-database creates a lexical scope in which the commands `define-table` and `define-command` create tables and `*commands*`-table entries respectively in open relational database *database*. The expressions in 'within-database' form are executed in order.

within-database Returns *database*.

define-command (<name> <rdb>) "comment" <expression1> <expression2>   [Syntax]
        ...
define-command (<name> <rdb>) <expression1> <expression2> ...   [Syntax]
      Adds to the *commands* table a command <name>:

        (lambda (<name> <rdb>) <expression1> <expression2> ...)

define-table <name> <descriptor-name> <descriptor-name> <rows>   [Syntax]
define-table <name> <primary-key-fields> <other-fields> <rows>   [Syntax]
      where <name> is the table name, <descriptor-name> is the symbol name of a descriptor
      table, <primary-key-fields> and <other-fields> describe the primary keys and other
      fields respectively, and <rows> is a list of data rows to be added to the table.

      <primary-key-fields> and <other-fields> are lists of field descriptors of the form:

        (<column-name> <domain>)

      or

        (<column-name> <domain> <column-integrity-rule>)

      where <column-name> is the column name, <domain> is the domain of the column,
      and <column-integrity-rule> is an expression whose value is a procedure of one argu-
      ment (which returns #f to signal an error).

      If <domain> is not a defined domain name and it matches the name of this table or
      an already defined (in one of *spec-0* ...) single key field table, a foreign-key domain
      will be created for it.

add-macro-support *database*   [Function]
      The relational database *database* must be mutable. **add-macro-support** adds a
      macros table and **define-macro** macro to *database*; then *database* is returned.

define-macro (<name> arg1 ...) "comment" <expression1> <expression2>   [Syntax]
        ...
define-macro (<name> arg1 ...) <expression1> <expression2> ...   [Syntax]
      Adds a macro <name> to the *macros*.

      Note: within-database creates lexical scope where not only define-command and
      define-table, but every command and macro are defined, ie.:

```
(within-database my-rdb
  (define-command (message rdb)
    (lambda (msg)
      (display "message: ")
      (display msg)
      (newline)))
  (message "Defining FOO...")
  ;; ... defining FOO ...
  (message "Defining BAR...")
  ;; ... defining BAR ...
  )
```

6.1.5.2 Within-database Example

Here is an example of `within-database` macros:

```
(require 'within-database)

(define my-rdb
  (add-command-tables
   (create-database "foo.db" 'alist-table)))

(within-database my-rdb
  (define-command (*initialize* rdb)
    "Print Welcome"
    (display "Welcome")
    (newline)
    rdb)
  (define-command (without-documentation rdb)
    (display "without-documentation called")
    (newline))
  (define-table (processor-family
                 ((family   atom))
                 ((also-ran processor-family)))
    (m68000  #f)
    (m68030  m68000)
    (i386    i8086)
    (i8086   #f)
    (powerpc #f))
  (define-table (platform
                 ((name symbol))
                 ((processor processor-family)
                  (os         symbol)
                  (compiler   symbol)))
    (aix              powerpc aix     -)
    ;; ...
    (amiga-aztec      m68000  amiga   aztec)
    (amiga-sas/c-5.10 m68000  amiga   sas/c)
    (atari-st-gcc     m68000  atari   gcc)
    ;; ...
    (watcom-9.0       i386    ms-dos  watcom))
  (define-command (get-processor rdb)
    "Get processor for given platform."
    (((rdb 'open-table) 'platform #f) 'get 'processor)))

(close-database my-rdb)

(set! my-rdb (open-command-database! "foo.db"))
  ⊣
Welcome
```

```
(my-rdb 'without-documentation)
⊣
without-documentation called

((my-rdb 'get-processor) 'amiga-sas/c-5.10)
⇒ m68000

(close-database my-rdb)
```

6.1.6 Database Browser

(require 'database-browse)

browse *database*   [Procedure]
>   Prints the names of all the tables in *database* and sets browse's default to *database*.

browse   [Procedure]
>   Prints the names of all the tables in the default database.

browse *table-name*   [Procedure]
>   For each record of the table named by the symbol *table-name*, prints a line composed of all the field values.

browse *pathname*   [Procedure]
>   Opens the database named by the string *pathname*, prints the names of all its tables, and sets browse's default to the database.

browse *database table-name*   [Procedure]
>   Sets browse's default to *database* and prints the records of the table named by the symbol *table-name*.

browse *pathname table-name*   [Procedure]
>   Opens the database named by the string *pathname* and sets browse's default to it; **browse** prints the records of the table named by the symbol *table-name*.

6.2 Relational Infrastructure

6.2.1 Base Table

A *base-table* is the primitive database layer upon which SLIB relational databases are built. At the minimum, it must support the types integer, symbol, string, and boolean. The base-table may restrict the size of integers, symbols, and strings it supports.

A base table implementation is available as the value of the identifier naming it (eg. *alist-table*) after requiring the symbol of that name.

alist-table   [Feature]
>   (require 'alist-table)

Association-list base tables support all Scheme types and are suitable for small data-bases. In order to be retrieved after being written to a file, the data stored should include only objects which are readable and writeable in the Scheme implementation.

The *alist-table* base-table implementation is included in the SLIB distribution.

WB is a B-tree database package with SCM interfaces. Being disk-based, WB databases readily store and access hundreds of megabytes of data. WB comes with two base-table embeddings.

`wb-table`   [Feature]

> `(require 'wb-table)`
>
> `wb-table` supports scheme expressions for keys and values whose text representations are less than 255 characters in length. See Section "wb-table" in *WB*.

`rwb-isam`   [Feature]

> `(require 'rwb-isam)`
>
> *rwb-isam* is a sophisticated base-table implementation built on WB and SCM which uses binary numerical formats for key and non-key fields. It supports IEEE floating-point and fixed-precision integer keys with the correct numerical collation order.

This rest of this section documents the interface for a base table implementation from which the Section 6.1 [Relational Database], page 162, package constructs a Relational system. It will be of interest primarily to those wishing to port or write new base-table implementations.

`*base-table-implementations*`   [Variable]

> To support automatic dispatch for `open-database`, each base-table module adds an association to **base-table-implementations** when loaded. This association is the list of the base-table symbol and the value returned by (`make-relational-system` `base-table`).

6.2.1.1 The Base

All of these functions are accessed through a single procedure by calling that procedure with the symbol name of the operation. A procedure will be returned if that operation is supported and `#f` otherwise. For example:

```
(require 'alist-table)
(define my-base (alist-table 'make-base))
my-base          ⇒ *a procedure*
(define foo (alist-table 'foo))
foo              ⇒ #f
```

`make-base` *filename key-dimension column-types*   [Operation on `base-table`]

> Returns a new, open, low-level database (collection of tables) associated with *filename*. This returned database has an empty table associated with *catalog-id*. The positive integer *key-dimension* is the number of keys composed to make a *primary-key* for the catalog table. The list of symbols *column-types* describes the types of each column for that table. If the database cannot be created as specified, `#f` is returned.

Calling the `close-base` method on this database and possibly other operations will cause *filename* to be written to. If *filename* is `#f` a temporary, non-disk based database will be created if such can be supported by the base table implelentation.

open-base *filename mutable*   [Operation on `base-table`]

Returns an open low-level database associated with *filename*. If *mutable* is `#t`, this database will have methods capable of effecting change to the database. If *mutable* is `#f`, only methods for inquiring the database will be available. If the database cannot be opened as specified `#f` is returned.

Calling the `close-base` (and possibly other) method on a *mutable* database will cause *filename* to be written to.

write-base *lldb filename*   [Operation on `base-table`]

Causes the low-level database *lldb* to be written to *filename*. If the write is successful, also causes *lldb* to henceforth be associated with *filename*. Calling the `close-database` (and possibly other) method on *lldb* may cause *filename* to be written to. If *filename* is `#f` this database will be changed to a temporary, non-disk based database if such can be supported by the underlying base table implelentation. If the operations completed successfully, `#t` is returned. Otherwise, `#f` is returned.

sync-base *lldb*   [Operation on `base-table`]

Causes the file associated with the low-level database *lldb* to be updated to reflect its current state. If the associated filename is `#f`, no action is taken and `#f` is returned. If this operation completes successfully, `#t` is returned. Otherwise, `#f` is returned.

close-base *lldb*   [Operation on `base-table`]

Causes the low-level database *lldb* to be written to its associated file (if any). If the write is successful, subsequent operations to *lldb* will signal an error. If the operations complete successfully, `#t` is returned. Otherwise, `#f` is returned.

6.2.1.2 Base Tables

make-table *lldb key-dimension column-types*   [Operation on `base-table`]

Returns the ordinal *base-id* for a new base table, otherwise returns `#f`. The base table can then be opened using (`open-table` `lldb` `base-id`). The positive integer *key-dimension* is the number of keys composed to make a *primary-key* for this table. The list of symbols *column-types* describes the types of each column.

open-table *lldb base-id key-dimension column-types*   [Operation on `base-table`]

Returns a *handle* for an existing base table in the low-level database *lldb* if that table exists and can be opened in the mode indicated by *mutable*, otherwise returns `#f`.

As with `make-table`, the positive integer *key-dimension* is the number of keys composed to make a *primary-key* for this table. The list of symbols *column-types* describes the types of each column.

kill-table *lldb base-id key-dimension column-types*   [Operation on `base-table`]

Returns `#t` if the base table associated with *base-id* was removed from the low level database *lldb*, and `#f` otherwise.

catalog-id   [Operation on `base-table`]
>   A constant *base-id* ordinal suitable for passing as a parameter to `open-table`. *catalog-id* will be used as the base table for the system catalog.

6.2.1.3 Base Field Types

supported-type? *symbol*   [Operation on `base-table`]
>   Returns `#t` if *symbol* names a type allowed as a column value by the implementation, and `#f` otherwise. At a minimum, an implementation must support the types `integer`, `ordinal`, `symbol`, `string`, and `boolean`.

supported-key-type? *symbol*   [Operation on `base-table`]
>   Returns `#t` if *symbol* names a type allowed as a key value by the implementation, and `#f` otherwise. At a minimum, an implementation must support the types `ordinal`, and `symbol`.

An *ordinal* is an exact positive integer. The other types are standard Scheme.

6.2.1.4 Composite Keys

make-keyifier-1 *type*   [Operation on `base-table`]
>   Returns a procedure which accepts a single argument which must be of type *type*. This returned procedure returns an object suitable for being a *key* argument in the functions whose descriptions follow.
>
>   Any 2 arguments of the supported type passed to the returned function which are not `equal?` must result in returned values which are not `equal?`.

make-list-keyifier *key-dimension types*   [Operation on `base-table`]
>   The list of symbols *types* must have at least *key-dimension* elements. Returns a procedure which accepts a list of length *key-dimension* and whose types must corresopond to the types named by *types*. This returned procedure combines the elements of its list argument into an object suitable for being a *key* argument in the functions whose descriptions follow.
>
>   Any 2 lists of supported types (which must at least include symbols and non-negative integers) passed to the returned function which are not `equal?` must result in returned values which are not `equal?`.

make-key-extractor *key-dimension types*   [Operation on `base-table`]
>   *column-number*
>   Returns a procedure which accepts objects produced by application of the result of (`make-list-keyifier` *key-dimension types*). This procedure returns a *key* which is `equal?` to the *column-number*th element of the list which was passed to create *composite-key*. The list *types* must have at least *key-dimension* elements.

make-key->list *key-dimension types*   [Operation on `base-table`]
>   Returns a procedure which accepts objects produced by application of the result of (`make-list-keyifier` *key-dimension types*). This procedure returns a list of *key*s which are elementwise `equal?` to the list which was passed to create *composite-key*.

6.2.1.5 Base Record Operations

In the following functions, the *key* argument can always be assumed to be the value returned by a call to a *keyify* routine.

present? *handle key*   [Operation on `base-table`]

> Returns a non-`#f` value if there is a row associated with *key* in the table opened in *handle* and `#f` otherwise.

make-getter *key-dimension types*   [Operation on `base-table`]

> Returns a procedure which takes arguments *handle* and *key*. This procedure returns a list of the non-primary values of the relation (in the base table opened in *handle*) whose primary key is *key* if it exists, and `#f` otherwise.

`make-getter-1` is a new operation. The relational-database module works with older base-table implementations by using `make-getter`.

make-getter-1 *key-dimension types index*   [Operation on `base-table`]

> Returns a procedure which takes arguments *handle* and *key*. This procedure returns the value of the *index*th field (in the base table opened in *handle*) whose primary key is *key* if it exists, and `#f` otherwise.
>
> *index* must be larger than *key-dimension*.

make-putter *key-dimension types*   [Operation on `base-table`]

> Returns a procedure which takes arguments *handle* and *key* and *value-list*. This procedure associates the primary key *key* with the values in *value-list* (in the base table opened in *handle*) and returns an unspecified value.

delete *handle key*   [Operation on `base-table`]

> Removes the row associated with *key* from the table opened in *handle*. An unspecified value is returned.

6.2.1.6 Match Keys

A *match-keys* argument is a list of length equal to the number of primary keys. The *match-keys* restrict the actions of the table command to those records whose primary keys all satisfy the corresponding element of the *match-keys* list. The elements and their actions are:

> `#f`   The false value matches any key in the corresponding position.
>
> an object of type procedure
> > This procedure must take a single argument, the key in the corresponding position. Any key for which the procedure returns a non-false value is a match; Any key for which the procedure returns a `#f` is not.
>
> other values
> > Any other value matches only those keys `equal?` to it.

6.2.1.7 Aggregate Base Operations

The *key-dimension* and *column-types* arguments are needed to decode the composite-keys for matching with *match-keys*.

delete* *handle key-dimension column-types*   [Operation on `base-table`]
        match-keys
>   Removes all rows which satisfy *match-keys* from the table opened in *handle*. An unspecified value is returned.

for-each-key *handle procedure key-dimension*   [Operation on `base-table`]
        column-types match-keys
>   Calls *procedure* once with each *key* in the table opened in *handle* which satisfy *match-keys* in an unspecified order. An unspecified value is returned.

map-key *handle procedure key-dimension*   [Operation on `base-table`]
        column-types match-keys
>   Returns a list of the values returned by calling *procedure* once with each *key* in the table opened in *handle* which satisfy *match-keys* in an unspecified order.

6.2.1.8 Base ISAM Operations

These operations are optional for a Base-Table implementation.

ordered-for-each-key *handle procedure*   [Operation on `base-table`]
        key-dimension column-types match-keys
>   Calls *procedure* once with each *key* in the table opened in *handle* which satisfy *match-keys* in the natural order for the types of the primary key fields of that table. An unspecified value is returned.

make-nexter *handle key-dimension column-types*   [Operation on `base-table`]
        index
>   Returns a procedure of arguments *key1 key2* . . . which returns the key-list identifying the lowest record higher than *key1 key2* . . . which is stored in the base-table and which differs in column *index* or a lower indexed key; or false if no higher record is present.

make-prever *handle key-dimension column-types*   [Operation on `base-table`]
        index
>   Returns a procedure of arguments *key1 key2* . . . which returns the key-list identifying the highest record less than *key1 key2* . . . which is stored in the base-table and which differs in column *index* or a lower indexed key; or false if no higher record is present.

6.2.2 Catalog Representation

Each database (in an implementation) has a *system catalog* which describes all the user accessible tables in that database (including itself).

The system catalog base table has the following fields. `PRI` indicates a primary key for that table.

```
PRI table-name
    column-limit              the highest column number
    coltab-name               descriptor table name
    bastab-id                 data base table identifier
    user-integrity-rule
    view-procedure            A scheme thunk which, when called,
                              produces a handle for the view.  coltab
                              and bastab are specified if and only if
                              view-procedure is not.
```

Descriptors for base tables (not views) are tables (pointed to by system catalog). Descriptor (base) tables have the fields:

```
PRI column-number             sequential integers from 1
    primary-key?              boolean TRUE for primary key components
    column-name
    column-integrity-rule
    domain-name
```

A *primary key* is any column marked as `primary-key?` in the corresponding descriptor table. All the `primary-key?` columns must have lower column numbers than any non-`primary-key?` columns. Every table must have at least one primary key. Primary keys must be sufficient to distinguish all rows from each other in the table. All of the system defined tables have a single primary key.

A *domain* is a category describing the allowable values to occur in a column. It is described by a (base) table with the fields:

```
PRI domain-name
    foreign-table
    domain-integrity-rule
    type-id
    type-param
```

The *type-id* field value is a symbol. This symbol may be used by the underlying base table implementation in storing that field.

If the `foreign-table` field is non-`#f` then that field names a table from the catalog. The values for that domain must match a primary key of the table referenced by the *type-param* (or `#f`, if allowed). This package currently does not support composite foreign-keys.

The types for which support is planned are:

```
atom
symbol
string                  [<length>]
number                  [<base>]
money                   <currency>
date-time
boolean

foreign-key             <table-name>
expression
virtual                 <expression>
```

6.2.3 Relational Database Objects

This object-oriented interface is deprecated for typical database applications; Section 6.1.1
[Using Databases], page 162, provides an application programmer interface which is easier
to understand and use.

---

make-relational-system *base-table-implementation*   [Function]

> Returns a procedure implementing a relational database using the *base-table-implementation*.

> All of the operations of a base table implementation are accessed through a procedure
> defined by `require`ing that implementation. Similarly, all of the operations of the
> relational database implementation are accessed through the procedure returned by
> `make-relational-system`. For instance, a new relational database could be created
> from the procedure returned by `make-relational-system` by:

```
(require 'alist-table)
(define relational-alist-system
        (make-relational-system alist-table))
(define create-alist-database
        (relational-alist-system 'create-database))
(define my-database
        (create-alist-database "mydata.db"))
```

What follows are the descriptions of the methods available from relational system returned
by a call to `make-relational-system`.

---

create-database *filename*   [Operation on `relational-system`]

> Returns an open, nearly empty relational database associated with *filename*. The only
> tables defined are the system catalog and domain table. Calling the `close-database`
> method on this database and possibly other operations will cause *filename* to be
> written to. If *filename* is `#f` a temporary, non-disk based database will be created if
> such can be supported by the underlying base table implelentation. If the database
> cannot be created as specified `#f` is returned. For the fields and layout of descriptor
> tables, Section 6.2.2 [Catalog Representation], page 184,

open-database *filename mutable?*   [Operation on `relational-system`]

> Returns an open relational database associated with *filename*. If *mutable?* is `#t`, this database will have methods capable of effecting change to the database. If *mutable?* is `#f`, only methods for inquiring the database will be available. Calling the `close-database` (and possibly other) method on a *mutable?* database will cause *filename* to be written to. If the database cannot be opened as specified `#f` is returned.

6.2.4 Database Operations

This object-oriented interface is deprecated for typical database applications; Section 6.1.1 [Using Databases], page 162, provides an application programmer interface which is easier to understand and use.

These are the descriptions of the methods available from an open relational database. A method is retrieved from a database by calling the database with the symbol name of the operation. For example:

```
(define my-database
        (create-alist-database "mydata.db"))
(define telephone-table-desc
        ((my-database 'create-table) 'telephone-table-desc))
```

close-database   [Operation on `relational-database`]

> Causes the relational database to be written to its associated file (if any). If the write is successful, subsequent operations to this database will signal an error. If the operations completed successfully, `#t` is returned. Otherwise, `#f` is returned.

write-database *filename*   [Operation on `relational-database`]

> Causes the relational database to be written to *filename*. If the write is successful, also causes the database to henceforth be associated with *filename*. Calling the `close-database` (and possibly other) method on this database will cause *filename* to be written to. If *filename* is `#f` this database will be changed to a temporary, non-disk based database if such can be supported by the underlying base table implelentation. If the operations completed successfully, `#t` is returned. Otherwise, `#f` is returned.

sync-database   [Operation on `relational-database`]

> Causes any pending updates to the database file to be written out. If the operations completed successfully, `#t` is returned. Otherwise, `#f` is returned.

solidify-database   [Operation on `relational-database`]

> Causes any pending updates to the database file to be written out. If the writes completed successfully, then the database is changed to be immutable and `#t` is returned. Otherwise, `#f` is returned.

table-exists? *table-name*   [Operation on `relational-database`]

> Returns `#t` if *table-name* exists in the system catalog, otherwise returns `#f`.

open-table *table-name mutable?*   [Operation on `relational-database`]

> Returns a *methods* procedure for an existing relational table in this database if it exists and can be opened in the mode indicated by *mutable?*, otherwise returns `#f`.

These methods will be present only in mutable databases.

delete-table *table-name*   [Operation on `relational-database`]
> Removes and returns the *table-name* row from the system catalog if the table or view associated with *table-name* gets removed from the database, and `#f` otherwise.

create-table *table-desc-name*   [Operation on `relational-database`]
> Returns a methods procedure for a new (open) relational table for describing the columns of a new base table in this database, otherwise returns `#f`. For the fields and layout of descriptor tables, See Section 6.2.2 [Catalog Representation], page 184.

create-table *table-name table-desc-name*   [Operation on `relational-database`]
> Returns a methods procedure for a new (open) relational table with columns as described by *table-desc-name*, otherwise returns `#f`.

create-view *??*   [Operation on `relational-database`]
project-table *??*   [Operation on `relational-database`]
restrict-table *??*   [Operation on `relational-database`]
cart-prod-tables *??*   [Operation on `relational-database`]
> Not yet implemented.

6.3 Weight-Balanced Trees

`(require 'wt-tree)`

Balanced binary trees are a useful data structure for maintaining large sets of ordered objects or sets of associations whose keys are ordered. MIT Scheme has an comprehensive implementation of weight-balanced binary trees which has several advantages over the other data structures for large aggregates:

- In addition to the usual element-level operations like insertion, deletion and lookup, there is a full complement of collection-level operations, like set intersection, set union and subset test, all of which are implemented with good orders of growth in time and space. This makes weight balanced trees ideal for rapid prototyping of functionally derived specifications.

- An element in a tree may be indexed by its position under the ordering of the keys, and the ordinal position of an element may be determined, both with reasonable efficiency.

- Operations to find and remove minimum element make weight balanced trees simple to use for priority queues.

- The implementation is *functional* rather than *imperative*. This means that operations like 'inserting' an association in a tree do not destroy the old tree, in much the same way that `(+ 1 x)` modifies neither the constant 1 nor the value bound to `x`. The trees are referentially transparent thus the programmer need not worry about copying the trees. Referential transparency allows space efficiency to be achieved by sharing subtrees.

These features make weight-balanced trees suitable for a wide range of applications, especially those that require large numbers of sets or discrete maps. Applications that have a few global databases and/or concentrate on element-level operations like insertion and lookup are probably better off using hash-tables or red-black trees.

The *size* of a tree is the number of associations that it contains. Weight balanced binary trees are balanced to keep the sizes of the subtrees of each node within a constant factor of each other. This ensures logarithmic times for single-path operations (like lookup and insertion). A weight balanced tree takes space that is proportional to the number of associations in the tree. For the current implementation, the constant of proportionality is six words per association.

Weight balanced trees can be used as an implementation for either discrete sets or discrete maps (associations). Sets are implemented by ignoring the datum that is associated with the key. Under this scheme if an associations exists in the tree this indicates that the key of the association is a member of the set. Typically a value such as (), `#t` or `#f` is associated with the key.

Many operations can be viewed as computing a result that, depending on whether the tree arguments are thought of as sets or maps, is known by two different names. An example is `wt-tree/member?`, which, when regarding the tree argument as a set, computes the set membership operation, but, when regarding the tree as a discrete map, `wt-tree/member?` is the predicate testing if the map is defined at an element in its domain. Most names in this package have been chosen based on interpreting the trees as sets, hence the name `wt-tree/member?` rather than `wt-tree/defined-at?`.

The weight balanced tree implementation is a run-time-loadable option. To use weight balanced trees, execute

```
(load-option 'wt-tree)
```

once before calling any of the procedures defined here.

6.3.1 Construction of Weight-Balanced Trees

Binary trees require there to be a total order on the keys used to arrange the elements in the tree. Weight balanced trees are organized by *types*, where the type is an object encapsulating the ordering relation. Creating a tree is a two-stage process. First a tree type must be created from the predicate which gives the ordering. The tree type is then used for making trees, either empty or singleton trees or trees from other aggregate structures like association lists. Once created, a tree 'knows' its type and the type is used to test compatibility between trees in operations taking two trees. Usually a small number of tree types are created at the beginning of a program and used many times throughout the program's execution.

`make-wt-tree-type` *key<?*   [procedure+]

This procedure creates and returns a new tree type based on the ordering predicate *key<?*. *Key<?* must be a total ordering, having the property that for all key values `a`, `b` and `c`:

```
(key<? a a)                        ⇒ #f
(and (key<? a b) (key<? b a))      ⇒ #f
(if (and (key<? a b) (key<? b c))
    (key<? a c)
    #t)                            ⇒ #t
```

Two key values are assumed to be equal if neither is less than the other by *key<?*.

Each call to `make-wt-tree-type` returns a distinct value, and trees are only compatible if their tree types are `eq?`. A consequence is that trees that are intended to be used in binary tree operations must all be created with a tree type originating from the same call to `make-wt-tree-type`.

`number-wt-type`   [variable+]

A standard tree type for trees with numeric keys. `Number-wt-type` could have been defined by

        (define number-wt-type (make-wt-tree-type <))

`string-wt-type`   [variable+]

A standard tree type for trees with string keys. `String-wt-type` could have been defined by

        (define string-wt-type (make-wt-tree-type string<?))

`make-wt-tree` *wt-tree-type*   [procedure+]

This procedure creates and returns a newly allocated weight balanced tree. The tree is empty, i.e. it contains no associations. *Wt-tree-type* is a weight balanced tree type obtained by calling `make-wt-tree-type`; the returned tree has this type.

`singleton-wt-tree` *wt-tree-type key datum*   [procedure+]

This procedure creates and returns a newly allocated weight balanced tree. The tree contains a single association, that of *datum* with *key*. *Wt-tree-type* is a weight balanced tree type obtained by calling `make-wt-tree-type`; the returned tree has this type.

`alist->wt-tree` *tree-type alist*   [procedure+]

Returns a newly allocated weight-balanced tree that contains the same associations as *alist*. This procedure is equivalent to:

        (lambda (type alist)
          (let ((tree (make-wt-tree type)))
            (for-each (lambda (association)
                        (wt-tree/add! tree
                                      (car association)
                                      (cdr association)))
                      alist)
            tree))

6.3.2 Basic Operations on Weight-Balanced Trees

This section describes the basic tree operations on weight balanced trees. These operations are the usual tree operations for insertion, deletion and lookup, some predicates and a procedure for determining the number of associations in a tree.

`wt-tree/empty?` *wt-tree*   [procedure+]

Returns `#t` if *wt-tree* contains no associations, otherwise returns `#f`.

`wt-tree/size` *wt-tree*   [procedure+]
>   Returns the number of associations in *wt-tree*, an exact non-negative integer. This
>   operation takes constant time.

`wt-tree/add` *wt-tree key datum*   [procedure+]
>   Returns a new tree containing all the associations in *wt-tree* and the association of
>   *datum* with *key*. If *wt-tree* already had an association for *key*, the new association
>   overrides the old. The average and worst-case times required by this operation are
>   proportional to the logarithm of the number of associations in *wt-tree*.

`wt-tree/add!` *wt-tree key datum*   [procedure+]
>   Associates *datum* with *key* in *wt-tree* and returns an unspecified value. If *wt-tree*
>   already has an association for *key*, that association is replaced. The average and
>   worst-case times required by this operation are proportional to the logarithm of the
>   number of associations in *wt-tree*.

`wt-tree/member?` *key wt-tree*   [procedure+]
>   Returns `#t` if *wt-tree* contains an association for *key*, otherwise returns `#f`. The aver-
>   age and worst-case times required by this operation are proportional to the logarithm
>   of the number of associations in *wt-tree*.

`wt-tree/lookup` *wt-tree key default*   [procedure+]
>   Returns the datum associated with *key* in *wt-tree*. If *wt-tree* doesn't contain an
>   association for *key*, *default* is returned. The average and worst-case times required
>   by this operation are proportional to the logarithm of the number of associations in
>   *wt-tree*.

`wt-tree/delete` *wt-tree key*   [procedure+]
>   Returns a new tree containing all the associations in *wt-tree*, except that if *wt-tree*
>   contains an association for *key*, it is removed from the result. The average and worst-
>   case times required by this operation are proportional to the logarithm of the number
>   of associations in *wt-tree*.

`wt-tree/delete!` *wt-tree key*   [procedure+]
>   If *wt-tree* contains an association for *key* the association is removed. Returns an
>   unspecified value. The average and worst-case times required by this operation are
>   proportional to the logarithm of the number of associations in *wt-tree*.

6.3.3 Advanced Operations on Weight-Balanced Trees

In the following the *size* of a tree is the number of associations that the tree contains, and
a *smaller* tree contains fewer associations.

`wt-tree/split<` *wt-tree bound*   [procedure+]
>   Returns a new tree containing all and only the associations in *wt-tree* which have
>   a key that is less than *bound* in the ordering relation of the tree type of *wt-tree*.
>   The average and worst-case times required by this operation are proportional to the
>   logarithm of the size of *wt-tree*.

`wt-tree/split>` *wt-tree bound*   [procedure+]

Returns a new tree containing all and only the associations in *wt-tree* which have a key that is greater than *bound* in the ordering relation of the tree type of *wt-tree*. The average and worst-case times required by this operation are proportional to the logarithm of size of *wt-tree*.

`wt-tree/union` *wt-tree-1 wt-tree-2*   [procedure+]

Returns a new tree containing all the associations from both trees. This operation is asymmetric: when both trees have an association for the same key, the returned tree associates the datum from *wt-tree-2* with the key. Thus if the trees are viewed as discrete maps then `wt-tree/union` computes the map override of *wt-tree-1* by *wt-tree-2*. If the trees are viewed as sets the result is the set union of the arguments. The worst-case time required by this operation is proportional to the sum of the sizes of both trees. If the minimum key of one tree is greater than the maximum key of the other tree then the time required is at worst proportional to the logarithm of the size of the larger tree.

`wt-tree/intersection` *wt-tree-1 wt-tree-2*   [procedure+]

Returns a new tree containing all and only those associations from *wt-tree-1* which have keys appearing as the key of an association in *wt-tree-2*. Thus the associated data in the result are those from *wt-tree-1*. If the trees are being used as sets the result is the set intersection of the arguments. As a discrete map operation, `wt-tree/intersection` computes the domain restriction of *wt-tree-1* to (the domain of) *wt-tree-2*. The time required by this operation is never worse that proportional to the sum of the sizes of the trees.

`wt-tree/difference` *wt-tree-1 wt-tree-2*   [procedure+]

Returns a new tree containing all and only those associations from *wt-tree-1* which have keys that *do not* appear as the key of an association in *wt-tree-2*. If the trees are viewed as sets the result is the asymmetric set difference of the arguments. As a discrete map operation, it computes the domain restriction of *wt-tree-1* to the complement of (the domain of) *wt-tree-2*. The time required by this operation is never worse that proportional to the sum of the sizes of the trees.

`wt-tree/subset?` *wt-tree-1 wt-tree-2*   [procedure+]

Returns `#t` iff the key of each association in *wt-tree-1* is the key of some association in *wt-tree-2*, otherwise returns `#f`. Viewed as a set operation, `wt-tree/subset?` is the improper subset predicate. A proper subset predicate can be constructed:

```
(define (proper-subset? s1 s2)
  (and (wt-tree/subset? s1 s2)
       (< (wt-tree/size s1) (wt-tree/size s2))))
```

As a discrete map operation, `wt-tree/subset?` is the subset test on the domain(s) of the map(s). In the worst-case the time required by this operation is proportional to the size of *wt-tree-1*.

`wt-tree/set-equal?` *wt-tree-1 wt-tree-2*   [procedure+]

Returns `#t` iff for every association in *wt-tree-1* there is an association in *wt-tree-2* that has the same key, and *vice versa*.

Viewing the arguments as sets `wt-tree/set-equal?` is the set equality predicate. As a map operation it determines if two maps are defined on the same domain.

This procedure is equivalent to

```
(lambda (wt-tree-1 wt-tree-2)
  (and (wt-tree/subset? wt-tree-1 wt-tree-2
        (wt-tree/subset? wt-tree-2 wt-tree-1)))
```

In the worst-case the time required by this operation is proportional to the size of the smaller tree.

`wt-tree/fold` *combiner initial wt-tree*   [procedure+]

This procedure reduces *wt-tree* by combining all the associations, using an reverse in-order traversal, so the associations are visited in reverse order. *Combiner* is a procedure of three arguments: a key, a datum and the accumulated result so far. Provided *combiner* takes time bounded by a constant, `wt-tree/fold` takes time proportional to the size of *wt-tree*.

A sorted association list can be derived simply:

```
(wt-tree/fold  (lambda (key datum list)
                 (cons (cons key datum) list))
               '()
               wt-tree))
```

The data in the associations can be summed like this:

```
(wt-tree/fold  (lambda (key datum sum) (+ sum datum))
               0
               wt-tree)
```

`wt-tree/for-each` *action wt-tree*   [procedure+]

This procedure traverses the tree in-order, applying *action* to each association. The associations are processed in increasing order of their keys. *Action* is a procedure of two arguments which take the key and datum respectively of the association. Provided *action* takes time bounded by a constant, `wt-tree/for-each` takes time proportional to in the size of *wt-tree*. The example prints the tree:

```
(wt-tree/for-each (lambda (key value)
                    (display (list key value)))
                  wt-tree))
```

`wt-tree/union-merge` *wt-tree-1 wt-tree-2 merge*   [procedure+]

Returns a new tree containing all the associations from both trees. If both trees have an association for the same key, the datum associated with that key in the result tree is computed by applying the procedure *merge* to the key, the value from *wt-tree-1* and the value from *wt-tree-2*. *Merge* is of the form

```
(lambda (key datum-1 datum-2) ...)
```

If some key occurs only in one tree, that association will appear in the result tree without being processed by *merge*, so for this operation to make sense, either *merge* must have both a right and left identity that correspond to the association being

absent in one of the trees, or some guarantee must be made, for example, all the keys in one tree are known to occur in the other.

These are all reasonable procedures for *merge*

```
(lambda (key val1 val2) (+ val1 val2))
(lambda (key val1 val2) (append val1 val2))
(lambda (key val1 val2) (wt-tree/union val1 val2))
```

However, a procedure like

```
(lambda (key val1 val2) (- val1 val2))
```

would result in a subtraction of the data for all associations with keys occuring in both trees but associations with keys occuring in only the second tree would be copied, not negated, as is presumably be intent. The programmer might ensure that this never happens.

This procedure has the same time behavior as `wt-tree/union` but with a slightly worse constant factor. Indeed, `wt-tree/union` might have been defined like this:

```
(define (wt-tree/union tree1 tree2)
  (wt-tree/union-merge tree1 tree2
                       (lambda (key val1 val2) val2)))
```

The *merge* procedure takes the *key* as a parameter in case the data are not independent of the key.

6.3.4 Indexing Operations on Weight-Balanced Trees

Weight balanced trees support operations that view the tree as sorted sequence of associations. Elements of the sequence can be accessed by position, and the position of an element in the sequence can be determined, both in logarithmic time.

wt-tree/index *wt-tree index*   [procedure+]
wt-tree/index-datum *wt-tree index*   [procedure+]
wt-tree/index-pair *wt-tree index*   [procedure+]
    Returns the 0-based *index*th association of *wt-tree* in the sorted sequence under the tree's ordering relation on the keys. `wt-tree/index` returns the *index*th key, `wt-tree/index-datum` returns the datum associated with the *index*th key and `wt-tree/index-pair` returns a new pair (`key . datum`) which is the `cons` of the *index*th key and its datum. The average and worst-case times required by this operation are proportional to the logarithm of the number of associations in the tree.

    These operations signal an error if the tree is empty, if *index*`<0`, or if *index* is greater than or equal to the number of associations in the tree.

    Indexing can be used to find the median and maximum keys in the tree as follows:

```
median:  (wt-tree/index wt-tree (quotient (wt-tree/size wt-tree) 2))

maximum: (wt-tree/index wt-tree (-1+ (wt-tree/size wt-tree)))
```

wt-tree/rank *wt-tree key*   [procedure+]
    Determines the 0-based position of *key* in the sorted sequence of the keys under the tree's ordering relation, or `#f` if the tree has no association with for *key*. This

procedure returns either an exact non-negative integer or `#f`. The average and worst-case times required by this operation are proportional to the logarithm of the number of associations in the tree.

`wt-tree/min` *wt-tree*   [procedure+]

`wt-tree/min-datum` *wt-tree*   [procedure+]

`wt-tree/min-pair` *wt-tree*   [procedure+]

Returns the association of *wt-tree* that has the least key under the tree's ordering relation. `wt-tree/min` returns the least key, `wt-tree/min-datum` returns the datum associated with the least key and `wt-tree/min-pair` returns a new pair (`key . datum`) which is the `cons` of the minimum key and its datum. The average and worst-case times required by this operation are proportional to the logarithm of the number of associations in the tree.

These operations signal an error if the tree is empty. They could be written

```
(define (wt-tree/min tree)        (wt-tree/index tree 0))
(define (wt-tree/min-datum tree)  (wt-tree/index-datum tree 0))
(define (wt-tree/min-pair tree)   (wt-tree/index-pair tree 0))
```

`wt-tree/delete-min` *wt-tree*   [procedure+]

Returns a new tree containing all of the associations in *wt-tree* except the association with the least key under the *wt-tree*'s ordering relation. An error is signalled if the tree is empty. The average and worst-case times required by this operation are proportional to the logarithm of the number of associations in the tree. This operation is equivalent to

```
(wt-tree/delete wt-tree (wt-tree/min wt-tree))
```

`wt-tree/delete-min!` *wt-tree*   [procedure+]

Removes the association with the least key under the *wt-tree*'s ordering relation. An error is signalled if the tree is empty. The average and worst-case times required by this operation are proportional to the logarithm of the number of associations in the tree. This operation is equivalent to

```
(wt-tree/delete! wt-tree (wt-tree/min wt-tree))
```

7 Other Packages

7.1 Data Structures

7.1.1 Arrays

`(require 'array)` or `(require 'srfi-63)`

array? *obj* [Function]

Returns `#t` if the *obj* is an array, and `#f` if not.

Note: Arrays are not disjoint from other Scheme types. Vectors and possibly strings also satisfy `array?`. A disjoint array predicate can be written:

```
(define (strict-array? obj)
  (and (array? obj) (not (string? obj)) (not (vector? obj))))
```

equal? *obj1 obj2* [Function]

Returns `#t` if *obj1* and *obj2* have the same rank and dimensions and the corresponding elements of *obj1* and *obj2* are `equal?`.

`equal?` recursively compares the contents of pairs, vectors, strings, and *arrays*, applying `eqv?` on other objects such as numbers and symbols. A rule of thumb is that objects are generally `equal?` if they print the same. `equal?` may fail to terminate if its arguments are circular data structures.

```
(equal? 'a 'a)                          ⇒   #t
(equal? '(a) '(a))                      ⇒   #t
(equal? '(a (b) c)
        '(a (b) c))                     ⇒   #t
(equal? "abc" "abc")                    ⇒   #t
(equal? 2 2)                            ⇒   #t
(equal? (make-vector 5 'a)
        (make-vector 5 'a))             ⇒   #t
(equal? (make-array (A:fixN32b 4) 5 3)
        (make-array (A:fixN32b 4) 5 3)) ⇒   #t
(equal? (make-array '#(foo) 3 3)
        (make-array '#(foo) 3 3))       ⇒   #t
(equal? (lambda (x) x)
        (lambda (y) y))                 ⇒   unspecified
```

array-rank *obj* [Function]

Returns the number of dimensions of *obj*. If *obj* is not an array, 0 is returned.

array-dimensions *array* [Function]

Returns a list of dimensions.

```
(array-dimensions (make-array '#() 3 5))
    ⇒ (3 5)
```

make-array *prototype k1 . . .*   [Function]

> Creates and returns an array of type *prototype* with dimensions *k1*, . . . and filled
> with elements from *prototype*. *prototype* must be an array, vector, or string. The
> implementation-dependent type of the returned array will be the same as the type of
> *prototype*; except if that would be a vector or string with rank not equal to one, in
> which case some variety of array will be returned.
>
> If the *prototype* has no elements, then the initial contents of the returned array are
> unspecified. Otherwise, the returned array will be filled with the element at the origin
> of *prototype*.

create-array *prototype k1 . . .*   [Function]

> `create-array` is an alias for `make-array`.

make-shared-array *array mapper k1 . . .*   [Function]

> `make-shared-array` can be used to create shared subarrays of other arrays. The
> *mapper* is a function that translates coordinates in the new array into coordinates in
> the old array. A *mapper* must be linear, and its range must stay within the bounds
> of the old array, but it can be otherwise arbitrary. A simple example:

```
(define fred (make-array '#(#f) 8 8))
(define freds-diagonal
  (make-shared-array fred (lambda (i) (list i i)) 8))
(array-set! freds-diagonal 'foo 3)
(array-ref fred 3 3)
   ⇒ FOO
(define freds-center
  (make-shared-array fred (lambda (i j) (list (+ 3 i) (+ 3 j)))
                     2 2))
(array-ref freds-center 0 0)
   ⇒ FOO
```

list->array *rank proto list*   [Function]

> *list* must be a rank-nested list consisting of all the elements, in row-major order, of
> the array to be created.
>
> `list->array` returns an array of rank *rank* and type *proto* consisting of all the
> elements, in row-major order, of *list*. When *rank* is 0, *list* is the lone array element;
> not necessarily a list.

```
(list->array 2 '#() '((1 2) (3 4)))
              ⇒ #2A((1 2) (3 4))
(list->array 0 '#() 3)
              ⇒ #0A 3
```

array->list *array*   [Function]

> Returns a rank-nested list consisting of all the elements, in row-major order, of *array*.
> In the case of a rank-0 array, `array->list` returns the single element.

```
(array->list #2A((ho ho ho) (ho oh oh)))
              ⇒ ((ho ho ho) (ho oh oh))
```

```
(array->list #0A ho)
            ⇒ ho
```

vector->array *vect proto dim1* . . .   [Function]

    vect must be a vector of length equal to the product of exact nonnegative integers *dim1*, ...

    vector->array returns an array of type *proto* consisting of all the elements, in row-major order, of *vect*. In the case of a rank-0 array, *vect* has a single element.

```
(vector->array #(1 2 3 4) #() 2 2)
            ⇒ #2A((1 2) (3 4))
(vector->array '#(3) '#())
            ⇒ #0A 3
```

array->vector *array*   [Function]

    Returns a new vector consisting of all the elements of *array* in row-major order.

```
(array->vector #2A ((1 2)( 3 4)))
            ⇒ #(1 2 3 4)
(array->vector #0A ho)
            ⇒ #(ho)
```

array-in-bounds? *array index1* . . .   [Function]

    Returns **#t** if its arguments would be acceptable to **array-ref**.

array-ref *array k1* . . .   [Function]

    Returns the (*k1*, . . .) element of *array*.

array-set! *array obj k1* . . .   [Procedure]

    Stores *obj* in the (*k1*, . . .) element of *array*. The value returned by **array-set!** is unspecified.

These functions return a prototypical uniform-array enclosing the optional argument (which must be of the correct type). If the uniform-array type is supported by the implementation, then it is returned; defaulting to the next larger precision type; resorting finally to vector.

A:floC128b *z*   [Function]
A:floC128b   [Function]

    Returns an inexact 128.bit flonum complex uniform-array prototype.

A:floC64b *z*   [Function]
A:floC64b   [Function]

    Returns an inexact 64.bit flonum complex uniform-array prototype.

A:floC32b *z*   [Function]
A:floC32b   [Function]

    Returns an inexact 32.bit flonum complex uniform-array prototype.

A:floC16b *z*   [Function]
A:floC16b   [Function]

    Returns an inexact 16.bit flonum complex uniform-array prototype.

`A:floR128b` *x*   [Function]
`A:floR128b`   [Function]
    Returns an inexact 128.bit flonum real uniform-array prototype.

`A:floR64b` *x*   [Function]
`A:floR64b`   [Function]
    Returns an inexact 64.bit flonum real uniform-array prototype.

`A:floR32b` *x*   [Function]
`A:floR32b`   [Function]
    Returns an inexact 32.bit flonum real uniform-array prototype.

`A:floR16b` *x*   [Function]
`A:floR16b`   [Function]
    Returns an inexact 16.bit flonum real uniform-array prototype.

`A:floR128d` *q*   [Function]
`A:floR128d`   [Function]
    Returns an exact 128.bit decimal flonum rational uniform-array prototype.

`A:floR64d` *q*   [Function]
`A:floR64d`   [Function]
    Returns an exact 64.bit decimal flonum rational uniform-array prototype.

`A:floR32d` *q*   [Function]
`A:floR32d`   [Function]
    Returns an exact 32.bit decimal flonum rational uniform-array prototype.

`A:fixZ64b` *n*   [Function]
`A:fixZ64b`   [Function]
    Returns an exact binary fixnum uniform-array prototype with at least 64 bits of
    precision.

`A:fixZ32b` *n*   [Function]
`A:fixZ32b`   [Function]
    Returns an exact binary fixnum uniform-array prototype with at least 32 bits of
    precision.

`A:fixZ16b` *n*   [Function]
`A:fixZ16b`   [Function]
    Returns an exact binary fixnum uniform-array prototype with at least 16 bits of
    precision.

`A:fixZ8b` *n*   [Function]
`A:fixZ8b`   [Function]
    Returns an exact binary fixnum uniform-array prototype with at least 8 bits of pre-
    cision.

`A:fixN64b` *k*   [Function]
`A:fixN64b`   [Function]
>   Returns an exact non-negative binary fixnum uniform-array prototype with at least
>   64 bits of precision.

`A:fixN32b` *k*   [Function]
`A:fixN32b`   [Function]
>   Returns an exact non-negative binary fixnum uniform-array prototype with at least
>   32 bits of precision.

`A:fixN16b` *k*   [Function]
`A:fixN16b`   [Function]
>   Returns an exact non-negative binary fixnum uniform-array prototype with at least
>   16 bits of precision.

`A:fixN8b` *k*   [Function]
`A:fixN8b`   [Function]
>   Returns an exact non-negative binary fixnum uniform-array prototype with at least
>   8 bits of precision.

`A:bool` *bool*   [Function]
`A:bool`   [Function]
>   Returns a boolean uniform-array prototype.

7.1.2 Subarrays

`(require 'subarray)`

subarray *array select* . . .   [Function]
>   selects a subset of an array. For $0 <= j <$ n, *selectj* is either an integer, a list of two
>   integers within the range for the *j*th index, or #f.
>
>   When *selectj* is a list of two integers, then the *j*th index is restricted to that subrange
>   in the returned array.
>
>   When *selectj* is #f, then the full range of the *j*th index is accessible in the returned
>   array. An elided argument is equivalent to #f.
>
>   When *selectj* is an integer, then the rank of the returned array is less than *array*, and
>   only elements whose *j*th index equals *selectj* are shared.
>
> ```
> > (define ra '#2A((a b c) (d e f)))
> #<unspecified>
> > (subarray ra 0 #f)
> #1A(a b c)
> > (subarray ra 1 #f)
> #1A(d e f)
> > (subarray ra #f 1)
> #1A(b e)
> > (subarray ra '(0 1) #f)
> #2A((a b c) (d e f))
> > (subarray ra #f '(0 1))
> ```

```
        #2A((a b) (d e))
        > (subarray ra #f '(1 2))
        #2A((b c) (e f))
        > (subarray ra #f '(2 1))
        #2A((c b) (f e))
```

Arrays can be reflected (reversed) using `subarray`:

```
        > (subarray '#1A(a b c d e) '(4 0))
        #1A(e d c b a)
```

array-trim *array trim . . .*   [Function]

Returns a subarray sharing contents with *array* except for slices removed from either side of each dimension. Each of the *trims* is an exact integer indicating how much to trim. A positive *s* trims the data from the lower end and reduces the upper bound of the result; a negative *s* trims from the upper end and increases the lower bound.

For example:

```
        (array-trim '#(0 1 2 3 4) 1)  ⇒ #1A(1 2 3 4)
        (array-trim '#(0 1 2 3 4) -1) ⇒ #1A(0 1 2 3)

        (require 'array-for-each)
        (define (centered-difference ra)
          (array-map ra - (array-trim ra 1) (array-trim ra -1)))

        (centered-difference '#(0 1 3 5 9 22))
          ⇒ #(1 2 2 4 13)
```

7.1.3 Array Mapping

`(require 'array-for-each)`

array-map! *array0 proc array1 . . .*   [Procedure]

array1, . . . must have the same number of dimensions as *array0* and have a range for each index which includes the range for the corresponding index in *array0*. *proc* is applied to each tuple of elements of *array1* . . . and the result is stored as the corresponding element in *array0*. The value returned is unspecified. The order of application is unspecified.

array-map *prototype proc array1 array2 . . .*   [Function]

array2, . . . must have the same number of dimensions as *array1* and have a range for each index which includes the range for the corresponding index in *array1*. *proc* is applied to each tuple of elements of *array1*, *array2*, . . . and the result is stored as the corresponding element in a new array of type *prototype*. The new array is returned. The order of application is unspecified.

array-for-each *proc array0 . . .*   [Function]

proc is applied to each tuple of elements of *array0* . . . in row-major order. The value returned is unspecified.

array-indexes *array*   [Function]

>   Returns an array of lists of indexes for *array* such that, if *li* is a list of indexes for which *array* is defined, (equal? *li* (apply array-ref (array-indexes *array*) *li*)).

array-index-for-each *array proc*   [Function]

>   applies *proc* to the indices of each element of *array* in turn. The value returned and the order of application are unspecified.

>   One can implement *array-index-map!* as

```
(define (array-index-map! ra fun)
  (array-index-for-each
   ra
   (lambda is (apply array-set! ra (apply fun is) is))))
```

array-index-map! *array proc*   [Procedure]

>   applies *proc* to the indices of each element of *array* in turn, storing the result in the corresponding element. The value returned and the order of application are unspecified.

>   One can implement *array-indexes* as

```
(define (array-indexes array)
    (let ((ra (apply make-array '#() (array-dimensions array))))
      (array-index-map! ra (lambda x x))
      ra))
```

>   Another example:

```
(define (apl:index-generator n)
    (let ((v (make-vector n 1)))
      (array-index-map! v (lambda (i) i))
      v))
```

array:copy! *destination source*   [Procedure]

>   Copies every element from vector or array *source* to the corresponding element of *destination*. *destination* must have the same rank as *source*, and be at least as large in each dimension. The order of copying is unspecified.

7.1.4 Array Interpolation

(require 'array-interpolate)

interpolate-array-ref *ra x1 ... xj*   [Function]

>   *ra* must be an array of rank j containing numbers. `interpolate-array-ref` returns a value interpolated from the nearest j-dimensional cube of elements of *ra*.

```
(interpolate-array-ref '#2A:fixZ32b((1 2 3) (4 5 6)) 1 0.1)
                            ==> 4.1
(interpolate-array-ref '#2A:fixZ32b((1 2 3) (4 5 6)) 0.5 0.25)
                            ==> 2.75
```

`resample-array!` *ra1 ra2*   [Procedure]

> *ra1* and *ra2* must be numeric arrays of equal rank. `resample-array!` sets *ra1* to
> values interpolated from *ra2* such that the values of elements at the corners of *ra1*
> and *ra2* are equal.

```
(define ra (make-array (A:fixZ32b) 2 2))
(resample-array! ra '#2A:fixZ32b((1 2 3) (4 5 6)))
ra              ==>  #2A:fixZ32b((1 3) (4 6))
(define ra (make-array (A:floR64b) 3 2))
(resample-array! ra '#2A:fixZ32b((1 2 3) (4 5 6)))
ra              ==>  #2A:floR64b((1.0 3.0) (2.5 4.5) (4.0 6.0))
```

7.1.5 Association Lists

`(require 'alist)`

Alist functions provide utilities for treating a list of key-value pairs as an associative
database. These functions take an equality predicate, *pred*, as an argument. This predicate
should be repeatable, symmetric, and transitive.

Alist functions can be used with a secondary index method such as hash tables for
improved performance.

`predicate->asso` *pred*   [Function]

> Returns an *association function* (like `assq`, `assv`, or `assoc`) corresponding to *pred*.
> The returned function returns a key-value pair whose key is `pred`-equal to its first
> argument or `#f` if no key in the alist is *pred*-equal to the first argument.

`alist-inquirer` *pred*   [Function]

> Returns a procedure of 2 arguments, *alist* and *key*, which returns the value associated
> with *key* in *alist* or `#f` if *key* does not appear in *alist*.

`alist-associator` *pred*   [Function]

> Returns a procedure of 3 arguments, *alist*, *key*, and *value*, which returns an alist
> with *key* and *value* associated. Any previous value associated with *key* will be lost.
> This returned procedure may or may not have side effects on its *alist* argument. An
> example of correct usage is:

```
(define put (alist-associator string-ci=?))
(define alist '())
(set! alist (put alist "Foo" 9))
```

`alist-remover` *pred*   [Function]

> Returns a procedure of 2 arguments, *alist* and *key*, which returns an alist with an
> association whose *key* is key removed. This returned procedure may or may not have
> side effects on its *alist* argument. An example of correct usage is:

```
(define rem (alist-remover string-ci=?))
(set! alist (rem alist "foo"))
```

`alist-map` *proc alist*   [Function]

> Returns a new association list formed by mapping *proc* over the keys and values of
> *alist*. *proc* must be a function of 2 arguments which returns the new value part.

alist-for-each *proc alist*   [Function]
>   Applies *proc* to each pair of keys and values of *alist*. *proc* must be a function of 2
>   arguments. The returned value is unspecified.

7.1.6 Byte

`(require 'byte)`

Some algorithms are expressed in terms of arrays of small integers. Using Scheme strings to implement these arrays is not portable vis-a-vis the correspondence between integers and characters and non-ascii character sets. These functions abstract the notion of a *byte*.

byte-ref *bytes k*   [Function]
>   *k* must be a valid index of *bytes*. `byte-ref` returns byte *k* of *bytes* using zero-origin
>   indexing.

byte-set! *bytes k byte*   [Procedure]
>   *k* must be a valid index of *bytes*, and *byte* must be a small nonnegative integer.
>   `byte-set!` stores *byte* in element *k* of *bytes* and returns an unspecified value.

make-bytes *k byte*   [Function]
make-bytes *k*   [Function]
>   `make-bytes` returns a newly allocated byte-array of length *k*. If *byte* is given, then
>   all elements of the byte-array are initialized to *byte*, otherwise the contents of the
>   byte-array are unspecified.

bytes-length *bytes*   [Function]
>   `bytes-length` returns length of byte-array *bytes*.

bytes *byte* ...   [Function]
>   Returns a newly allocated byte-array composed of the small nonnegative arguments.

list->bytes *bytes*   [Function]
>   `list->bytes` returns a newly allocated byte-array formed from the small nonnegative
>   integers in the list *bytes*.

bytes->list *bytes*   [Function]
>   `bytes->list` returns a newly allocated list of the bytes that make up the given byte-
>   array.

`Bytes->list` and `list->bytes` are inverses so far as `equal?` is concerned.

bytes->string *bytes*   [Function]
>   Returns a new string formed from applying `integer->char` to each byte in
>   `bytes->string`. Note that this may signal an error for bytes having values between
>   128 and 255.

string->bytes *string*   [Function]
>   Returns a new byte-array formed from applying `char->integer` to each character in
>   `string->bytes`. Note that this may signal an error if an integer is larger than 255.

bytes-copy *bytes*   [Function]
>   Returns a newly allocated copy of the given *bytes*.

subbytes *bytes start end*   [Function]
>   *bytes* must be a bytes, and *start* and *end* must be exact integers satisfying
>
>   0 <= *start* <= *end* <= (bytes-length *bytes*).

>   subbytes returns a newly allocated bytes formed from the bytes of *bytes* beginning
>   with index *start* (inclusive) and ending with index *end* (exclusive).

bytes-reverse! *bytes*   [Procedure]
>   Reverses the order of byte-array *bytes*.

bytes-reverse *bytes*   [Function]
>   Returns a newly allocated bytes-array consisting of the elements of *bytes* in reverse
>   order.

Input and output of bytes should be with ports opened in *binary* mode (see Section 2.3
[Input/Output], page 13). Calling open-file with 'rb or 'wb modes argument will return
a binary port if the Scheme implementation supports it.

write-byte *byte port*   [Function]
write-byte *byte*   [Function]
>   Writes the byte *byte* (not an external representation of the byte) to the given *port*
>   and returns an unspecified value. The *port* argument may be omitted, in which case
>   it defaults to the value returned by current-output-port.

read-byte *port*   [Function]
read-byte   [Function]
>   Returns the next byte available from the input *port*, updating the *port* to point to the
>   following byte. If no more bytes are available, an end-of-file object is returned. *port*
>   may be omitted, in which case it defaults to the value returned by current-input-
>   port.

When reading and writing binary numbers with read-bytes and write-bytes, the sign
of the length argument determines the endianness (order) of bytes. Positive treats them
as big-endian, the first byte input or output is highest order. Negative treats them as
little-endian, the first byte input or output is the lowest order.

Once read in, SLIB treats byte sequences as big-endian. The multi-byte sequences produced
and used by number conversion routines see Section 7.1.7 [Byte/Number Conversions],
page 206, are always big-endian.

read-bytes *n port*   [Function]
read-bytes *n*   [Function]
>   read-bytes returns a newly allocated bytes-array filled with (abs *n*) bytes read from
>   *port*. If *n* is positive, then the first byte read is stored at index 0; otherwise the last
>   byte read is stored at index 0. Note that the length of the returned byte-array will
>   be less than (abs *n*) if *port* reaches end-of-file.

>   *port* may be omitted, in which case it defaults to the value returned by
>   current-input-port.

write-bytes *bytes n port*   [Function]
write-bytes *bytes n*   [Function]
>   write-bytes writes (`abs n`) bytes to output-port *port*. If *n* is positive, then the first
>   byte written is index 0 of *bytes*; otherwise the last byte written is index 0 of *bytes*.
>   write-bytes returns an unspecified value.
>
>   *port* may be omitted, in which case it defaults to the value returned by
>   current-output-port.

subbytes-read! and subbytes-write provide lower-level procedures for reading and writing blocks of bytes. The relative size of *start* and *end* determines the order of writing.

subbytes-read! *bts start end port*   [Procedure]
subbytes-read! *bts start end*   [Procedure]
>   Fills *bts* with up to (`abs (- start end)`) bytes read from *port*. The first byte read
>   is stored at index *bts*. subbytes-read! returns the number of bytes read.
>
>   *port* may be omitted, in which case it defaults to the value returned by
>   current-input-port.

subbytes-write *bts start end port*   [Function]
subbytes-write *bts start end*   [Function]
>   subbytes-write writes (`abs (- start end)`) bytes to output-port *port*. The first
>   byte written is index *start* of *bts*. subbytes-write returns the number of bytes
>   written.
>
>   *port* may be omitted, in which case it defaults to the value returned by
>   current-output-port.

7.1.7 Byte/Number Conversions

(`require 'byte-number`)

The multi-byte sequences produced and used by numeric conversion routines are always big-endian. Endianness can be changed during reading and writing bytes using read-bytes and write-bytes See Section 7.1.6 [Byte], page 204.

The sign of the length argument to bytes/integer conversion procedures determines the signedness of the number.

bytes->integer *bytes n*   [Function]
>   Converts the first (`abs n`) bytes of big-endian *bytes* array to an integer. If *n* is
>   negative then the integer coded by the bytes are treated as two's-complement (can
>   be negative).

```
(bytes->integer (bytes   0   0   0  15) -4)   ⇒            15
(bytes->integer (bytes   0   0   0  15)  4)   ⇒            15
(bytes->integer (bytes 255 255 255 255) -4)   ⇒            -1
(bytes->integer (bytes 255 255 255 255)  4)   ⇒    4294967295
(bytes->integer (bytes 128   0   0   0) -4)   ⇒ -2147483648
(bytes->integer (bytes 128   0   0   0)  4)   ⇒  2147483648
```

`integer->bytes` *n len*   [Function]

>   Converts the integer *n* to a byte-array of (`abs n`) bytes. If *n* and *len* are both negative,
>   then the bytes in the returned array are coded two's-complement.

```
(bytes->list (integer->bytes          15 -4))  ⇒ (0 0 0 15)
(bytes->list (integer->bytes          15  4))  ⇒ (0 0 0 15)
(bytes->list (integer->bytes          -1 -4))  ⇒ (255 255 255 255)
(bytes->list (integer->bytes  4294967295  4))  ⇒ (255 255 255 255)
(bytes->list (integer->bytes -2147483648 -4))  ⇒ (128 0 0 0)
(bytes->list (integer->bytes  2147483648  4))  ⇒ (128 0 0 0)
```

`bytes->ieee-float` *bytes*   [Function]

>   *bytes* must be a 4-element byte-array. `bytes->ieee-float` calculates and returns
>   the value of *bytes* interpreted as a big-endian IEEE 4-byte (32-bit) number.

```
(bytes->ieee-float (bytes    0    0 0 0))  ⇒  0.0
(bytes->ieee-float (bytes #x80    0 0 0))  ⇒ -0.0
(bytes->ieee-float (bytes #x40    0 0 0))  ⇒  2.0
(bytes->ieee-float (bytes #x40 #xd0 0 0))  ⇒  6.5
(bytes->ieee-float (bytes #xc0 #xd0 0 0))  ⇒ -6.5

(bytes->ieee-float (bytes    0 #x80 0 0))  ⇒ 11.754943508222875e-39
(bytes->ieee-float (bytes    0 #x40 0 0))  ⇒  5.877471754111437e-39
(bytes->ieee-float (bytes    0    0 0 1))  ⇒  1.401298464324817e-45

(bytes->ieee-float (bytes #xff #x80 0 0))  ⇒ -inf.0
(bytes->ieee-float (bytes #x7f #x80 0 0))  ⇒ +inf.0
(bytes->ieee-float (bytes #x7f #x80 0 1))  ⇒  0/0
(bytes->ieee-float (bytes #x7f #xc0 0 0))  ⇒  0/0
```

`bytes->ieee-double` *bytes*   [Function]

>   *bytes* must be a 8-element byte-array. `bytes->ieee-double` calculates and returns
>   the value of *bytes* interpreted as a big-endian IEEE 8-byte (64-bit) number.

```
(bytes->ieee-double (bytes    0    0 0 0 0 0 0 0))  ⇒  0.0
(bytes->ieee-double (bytes #x80    0 0 0 0 0 0 0))  ⇒ -0.0
(bytes->ieee-double (bytes #x40    0 0 0 0 0 0 0))  ⇒  2.0
(bytes->ieee-double (bytes #x40 #x1A 0 0 0 0 0 0))  ⇒  6.5
(bytes->ieee-double (bytes #xC0 #x1A 0 0 0 0 0 0))  ⇒ -6.5

(bytes->ieee-double (bytes 0 8 0 0 0 0 0 0)) ⇒ 11.125369292536006e-309
(bytes->ieee-double (bytes 0 4 0 0 0 0 0 0)) ⇒  5.562684646268003e-309
(bytes->ieee-double (bytes 0 0 0 0 0 0 0 1)) ⇒  4.0e-324

(bytes->ieee-double (list->bytes '(127 239 255 255 255 255 255 255))) 179.769313486231
(bytes->ieee-double (bytes #xFF #xF0 0 0 0 0 0 0))  ⇒ -inf.0
(bytes->ieee-double (bytes #x7F #xF0 0 0 0 0 0 0))  ⇒ +inf.0
(bytes->ieee-double (bytes #x7F #xF8 0 0 0 0 0 0))  ⇒  0/0
```

`ieee-float->bytes` *x*   [Function]

Returns a 4-element byte-array encoding the IEEE single-precision floating-point of
x.

```
(bytes->list (ieee-float->bytes  0.0))                       ⇒ (0      0 0 0)
(bytes->list (ieee-float->bytes -0.0))                       ⇒ (128    0 0 0)
(bytes->list (ieee-float->bytes  2.0))                       ⇒ (64     0 0 0)
(bytes->list (ieee-float->bytes  6.5))                       ⇒ (64  208 0 0)
(bytes->list (ieee-float->bytes -6.5))                       ⇒ (192 208 0 0)

(bytes->list (ieee-float->bytes 11.754943508222875e-39))     ⇒ (  0 128 0 0)
(bytes->list (ieee-float->bytes  5.877471754111438e-39))     ⇒ (  0  64 0 0)
(bytes->list (ieee-float->bytes  1.401298464324817e-45))     ⇒ (  0   0 0 1)

(bytes->list (ieee-float->bytes -inf.0))                     ⇒ (255 128 0 0)
(bytes->list (ieee-float->bytes +inf.0))                     ⇒ (127 128 0 0)
(bytes->list (ieee-float->bytes  0/0))                       ⇒ (127 192 0 0)
```

`ieee-double->bytes` *x*   [Function]

Returns a 8-element byte-array encoding the IEEE double-precision floating-point of
x.

```
(bytes->list (ieee-double->bytes  0.0)) ⇒ (0     0 0 0 0 0 0 0)
(bytes->list (ieee-double->bytes -0.0)) ⇒ (128   0 0 0 0 0 0 0)
(bytes->list (ieee-double->bytes  2.0)) ⇒ (64    0 0 0 0 0 0 0)
(bytes->list (ieee-double->bytes  6.5)) ⇒ (64   26 0 0 0 0 0 0)
(bytes->list (ieee-double->bytes -6.5)) ⇒ (192  26 0 0 0 0 0 0)

(bytes->list (ieee-double->bytes 11.125369292536006e-309))
                               ⇒ (  0   8 0 0 0 0 0 0)
(bytes->list (ieee-double->bytes  5.562684646268003e-309))
                               ⇒ (  0   4 0 0 0 0 0 0)
(bytes->list (ieee-double->bytes  4.0e-324))
                               ⇒ (  0   0 0 0 0 0 0 1)

(bytes->list (ieee-double->bytes -inf.0)) ⇒ (255 240 0 0 0 0 0 0)
(bytes->list (ieee-double->bytes +inf.0)) ⇒ (127 240 0 0 0 0 0 0)
(bytes->list (ieee-double->bytes  0/0)) ⇒ (127 248 0 0 0 0 0 0)
```

Byte Collation Order

The `string<?` ordering of big-endian byte-array representations of fixed and IEEE floating-
point numbers agrees with the numerical ordering only when those numbers are non-
negative.

Straighforward modification of these formats can extend the byte-collating order to work
for their entire ranges. This agreement enables the full range of numbers as keys in *indexed-
sequential-access-method* databases.

`integer-byte-collate!` *byte-vector*   [Procedure]

> Modifies sign bit of *byte-vector* so that `string<?` ordering of two's-complement byte-vectors matches numerical order. `integer-byte-collate!` returns *byte-vector* and is its own functional inverse.

`integer-byte-collate` *byte-vector*   [Function]

> Returns copy of *byte-vector* with sign bit modified so that `string<?` ordering of two's-complement byte-vectors matches numerical order. `integer-byte-collate` is its own functional inverse.

`ieee-byte-collate!` *byte-vector*   [Procedure]

> Modifies *byte-vector* so that `string<?` ordering of IEEE floating-point byte-vectors matches numerical order. `ieee-byte-collate!` returns *byte-vector*.

`ieee-byte-decollate!` *byte-vector*   [Procedure]

> Given *byte-vector* modified by `ieee-byte-collate!`, reverses the *byte-vector* modifications.

`ieee-byte-collate` *byte-vector*   [Function]

> Returns copy of *byte-vector* encoded so that `string<?` ordering of IEEE floating-point byte-vectors matches numerical order.

`ieee-byte-decollate` *byte-vector*   [Function]

> Given *byte-vector* returned by `ieee-byte-collate`, reverses the *byte-vector* modifications.

7.1.8 MAT-File Format

`(require 'matfile)`

> `http: / / www . mathworks . com / access / helpdesk / help / pdf_doc / matlab / matfile_format.pdf`

This package reads MAT-File Format version 4 (MATLAB) binary data files. MAT-files written from big-endian or little-endian computers having IEEE format numbers are currently supported. Support for files written from VAX or Cray machines could also be added.

The numeric and text matrix types handled; support for *sparse* matrices awaits a sample file.

`matfile:read` *filename*   [Function]

> *filename* should be a string naming an existing file containing a MATLAB Version 4 MAT-File. The `matfile:read` procedure reads matrices from the file and returns a list of the results; a list of the name string and array for each matrix.

`matfile:load` *filename*   [Function]

> *filename* should be a string naming an existing file containing a MATLAB Version 4 MAT-File. The `matfile:load` procedure reads matrices from the file and defines the `string-ci->symbol` for each matrix to its corresponding array. `matfile:load` returns a list of the symbols defined.

7.1.9 Portable Image Files

`(require 'pnm)`

`pnm:type-dimensions` *path*   [Function]

> The string *path* must name a *portable bitmap graphics* file. `pnm:type-dimensions` returns a list of 4 items:
>
> 1. A symbol describing the type of the file named by *path*.
> 2. The image width in pixels.
> 3. The image height in pixels.
> 4. The maximum value of pixels assume in the file.
>
> The current set of file-type symbols is:
>
> pbm
> pbm-raw   Black-and-White image; pixel values are 0 or 1.
>
> pgm
> pgm-raw   Gray (monochrome) image; pixel values are from 0 to *maxval* specified in file header.
>
> ppm
> ppm-raw   RGB (full color) image; red, green, and blue interleaved pixel values are from 0 to *maxval*

`pnm:image-file->array` *path array*   [Function]

> Reads the *portable bitmap graphics* file named by *path* into *array*. *array* must be the correct size and type for *path*. *array* is returned.

`pnm:image-file->array` *path*   [Function]

> `pnm:image-file->array` creates and returns an array with the *portable bitmap graphics* file named by *path* read into it.

`pnm:array-write` *type array maxval path comment . . .*   [Function]

> Writes the contents of *array* to a *type* image file named *path*. The file will have pixel values between 0 and *maxval*, which must be compatible with *type*. For 'pbm' files, *maxval* must be '1'. *comment*s are included in the file header.

7.1.10 Collections

`(require 'collect)`

Routines for managing collections. Collections are aggregate data structures supporting iteration over their elements, similar to the Dylan(TM) language, but with a different interface. They have *elements* indexed by corresponding *keys*, although the keys may be implicit (as with lists).

New types of collections may be defined as YASOS objects (see Section 3.14 [Yasos], page 35). They must support the following operations:

- `(collection? self)` (always returns `#t`);
- `(size self)` returns the number of elements in the collection;

- (print *self port*) is a specialized print operation for the collection which prints a suitable representation on the given *port* or returns it as a string if *port* is #t;
- (gen-elts *self*) returns a thunk which on successive invocations yields elements of *self* in order or gives an error if it is invoked more than (size *self*) times;
- (gen-keys *self*) is like gen-elts, but yields the collection's keys in order.

They might support specialized for-each-key and for-each-elt operations.

collection? *obj*   [Function]
>   A predicate, true initially of lists, vectors and strings. New sorts of collections must answer #t to collection?.

map-elts *proc collection1 . . .*   [Procedure]
do-elts *proc collection1 . . .*   [Procedure]
>   *proc* is a procedure taking as many arguments as there are *collections* (at least one). The *collections* are iterated over in their natural order and *proc* is applied to the elements yielded by each iteration in turn. The order in which the arguments are supplied corresponds to te order in which the *collections* appear. do-elts is used when only side-effects of *proc* are of interest and its return value is unspecified. map-elts returns a collection (actually a vector) of the results of the applications of *proc*.
>
>   Example:
>
>   (map-elts + (list 1 2 3) (vector 1 2 3))
>   ⇒ #(2 4 6)

map-keys *proc collection1 . . .*   [Procedure]
do-keys *proc collection1 . . .*   [Procedure]
>   These are analogous to map-elts and do-elts, but each iteration is over the *collections' keys* rather than their elements.
>
>   Example:
>
>   (map-keys + (list 1 2 3) (vector 1 2 3))
>   ⇒ #(0 2 4)

for-each-key *collection proc*   [Procedure]
for-each-elt *collection proc*   [Procedure]
>   These are like do-keys and do-elts but only for a single collection; they are potentially more efficient.

reduce *proc seed collection1 . . .*   [Function]
>   A generalization of the list-based reduce-init (see Section 7.2.1.3 [Lists as sequences], page 226) to collections.
>
>   Examples:
>
>   (reduce + 0 (vector 1 2 3))
>   ⇒ 6
>   (reduce union '() '((a b c) (b c d) (d a)))
>   ⇒ (c b d a).
>
>   Reduce called with two arguments will work as does the procedure of the same name from See Section 7.2.1 [Common List Functions], page 221.

any? *pred collection1 . . .*   [Function]

>   A generalization of the list-based `some` (see Section 7.2.1.3 [Lists as sequences], page 226) to collections.
>
>   Example:
>
>   (any? odd? (list 2 3 4 5))
>   ⇒ #t

every? *pred collection1 . . .*   [Function]

>   A generalization of the list-based `every` (see Section 7.2.1.3 [Lists as sequences], page 226) to collections.
>
>   Example:
>
>   (every? collection? '((1 2) #(1 2)))
>   ⇒ #t

empty? *collection*   [Function]

>   Returns `#t` iff there are no elements in *collection*.
>
>   `(empty? collection)` ≡ `(zero? (size collection))`

size *collection*   [Function]

>   Returns the number of elements in *collection*.

Setter *list-ref*   [Function]

>   See Section 3.14.3 [Setters], page 37, for a definition of *setter*. N.B. `(setter list-ref)` doesn't work properly for element 0 of a list.

Here is a sample collection: `simple-table` which is also a `table`.

```
(define-predicate TABLE?)
(define-operation (LOOKUP table key failure-object))
(define-operation (ASSOCIATE! table key value)) ;; returns key
(define-operation (REMOVE! table key))           ;; returns value

(define (MAKE-SIMPLE-TABLE)
  (let ( (table (list)) )
    (object
     ;; table behaviors
     ((TABLE? self) #t)
     ((SIZE self) (size table))
     ((PRINT self port) (format port "#<SIMPLE-TABLE>"))
     ((LOOKUP self key failure-object)
      (cond
       ((assq key table) => cdr)
       (else failure-object)
       ))
     ((ASSOCIATE! self key value)
      (cond
       ((assq key table)
```

```
                => (lambda (bucket) (set-cdr! bucket value) key))
                (else
                 (set! table (cons (cons key value) table))
                 key)
                ))
              ((REMOVE! self key);; returns old value
               (cond
                ((null? table) (slib:error "TABLE:REMOVE! Key not found: " key))
                ((eq? key (caar table))
                 (let ( (value (cdar table)) )
                   (set! table (cdr table))
                   value)
                 )
                (else
                 (let loop ( (last table) (this (cdr table)) )
                   (cond
                    ((null? this)
                     (slib:error "TABLE:REMOVE! Key not found: " key))
                    ((eq? key (caar this))
                     (let ( (value (cdar this)) )
                       (set-cdr! last (cdr this))
                       value)
                     )
                    (else
                     (loop (cdr last) (cdr this)))
                    ) ) )
                ))
              ;; collection behaviors
              ((COLLECTION? self) #t)
              ((GEN-KEYS self) (collect:list-gen-elts (map car table)))
              ((GEN-ELTS self) (collect:list-gen-elts (map cdr table)))
              ((FOR-EACH-KEY self proc)
               (for-each (lambda (bucket) (proc (car bucket))) table)
               )
              ((FOR-EACH-ELT self proc)
               (for-each (lambda (bucket) (proc (cdr bucket))) table)
               ) ) ) )
```

7.1.11 Dynamic Data Type

```
(require 'dynamic)
```

make-dynamic *obj*   [Function]

> Create and returns a new *dynamic* whose global value is *obj*.

dynamic? *obj*   [Function]

> Returns true if and only if *obj* is a dynamic. No object satisfying `dynamic?` satisfies
> any of the other standard type predicates.

dynamic-ref *dyn*   [Function]

    Return the value of the given dynamic in the current dynamic environment.

dynamic-set! *dyn obj*   [Procedure]

    Change the value of the given dynamic to *obj* in the current dynamic environment.
The returned value is unspecified.

call-with-dynamic-binding *dyn obj thunk*   [Function]

    Invoke and return the value of the given thunk in a new, nested dynamic environment
in which the given dynamic has been bound to a new location whose initial contents
are the value *obj*. This dynamic environment has precisely the same extent as the
invocation of the thunk and is thus captured by continuations created within that
invocation and re-established by those continuations when they are invoked.

    The `dynamic-bind` macro is not implemented.

7.1.12 Hash Tables

`(require 'hash-table)`

predicate->hash *pred*   [Function]

    Returns a hash function (like `hashq`, `hashv`, or `hash`) corresponding to the equality
predicate *pred*. *pred* should be `eq?`, `eqv?`, `equal?`, `=`, `char=?`, `char-ci=?`, `string=?`,
or `string-ci=?`.

A hash table is a vector of association lists.

make-hash-table *k*   [Function]

    Returns a vector of *k* empty (association) lists.

Hash table functions provide utilities for an associative database. These functions take an
equality predicate, *pred*, as an argument. *pred* should be `eq?`, `eqv?`, `equal?`, `=`, `char=?`,
`char-ci=?`, `string=?`, or `string-ci=?`.

predicate->hash-asso *pred*   [Function]

    Returns a hash association function of 2 arguments, *key* and *hashtab*, corresponding
to *pred*. The returned function returns a key-value pair whose key is *pred*-equal to
its first argument or `#f` if no key in *hashtab* is *pred*-equal to the first argument.

hash-inquirer *pred*   [Function]

    Returns a procedure of 2 arguments, *hashtab* and *key*, which returns the value asso-
ciated with *key* in *hashtab* or `#f` if *key* does not appear in *hashtab*.

hash-associator *pred*   [Function]

    Returns a procedure of 3 arguments, *hashtab*, *key*, and *value*, which modifies *hashtab*
so that *key* and *value* associated. Any previous value associated with *key* will be lost.

hash-remover *pred*   [Function]

    Returns a procedure of 2 arguments, *hashtab* and *key*, which modifies *hashtab* so that
the association whose key is *key* is removed.

`hash-map` *proc hash-table*   [Function]
> Returns a new hash table formed by mapping *proc* over the keys and values of *hash-table*. *proc* must be a function of 2 arguments which returns the new value part.

`hash-for-each` *proc hash-table*   [Function]
> Applies *proc* to each pair of keys and values of *hash-table*. *proc* must be a function of 2 arguments. The returned value is unspecified.

`hash-rehasher` *pred*   [Function]
> `hash-rehasher` accepts a hash table predicate and returns a function of two arguments *hashtab* and *new-k* which is specialized for that predicate.
>
> This function is used for nondestrutively resizing a hash table. *hashtab* should be an existing hash-table using *pred*, *new-k* is the size of a new hash table to be returned. The new hash table will have all of the associations of the old hash table.

7.1.13 Macroless Object System

`(require 'object)`

   This is the Macroless Object System written by Wade Humeniuk (whumeniu@datap.ca). Conceptual Tributes: Section 3.14 [Yasos], page 35, MacScheme's %object, CLOS, Lack of R4RS macros.

7.1.14 Concepts

OBJECT   An object is an ordered association-list (by `eq?`) of methods (procedures). Methods can be added (`make-method!`), deleted (`unmake-method!`) and retrieved (`get-method`). Objects may inherit methods from other objects. The object binds to the environment it was created in, allowing closures to be used to hide private procedures and data.

GENERIC-METHOD
>   A generic-method associates (in terms of `eq?`) object's method. This allows scheme function style to be used for objects. The calling scheme for using a generic method is (`generic-method object param1 param2 ...`).

METHOD   A method is a procedure that exists in the object. To use a method get-method must be called to look-up the method. Generic methods implement the get-method functionality. Methods may be added to an object associated with any scheme obj in terms of eq?

GENERIC-PREDICATE
>   A generic method that returns a boolean value for any scheme obj.

PREDICATE
>   A object's method asscociated with a generic-predicate. Returns `#t`.

7.1.15 Procedures

`make-object` *ancestor ...*   [Function]
> Returns an object. Current object implementation is a tagged vector. *ancestor*s are optional and must be objects in terms of object?. *ancestor*s methods are included

in the object. Multiple *ancestor*s might associate the same generic-method with a
method. In this case the method of the *ancestor* first appearing in the list is the one
returned by `get-method`.

`object?` *obj*   [Function]

Returns boolean value whether *obj* was created by make-object.

`make-generic-method` *exception-procedure*   [Function]

Returns a procedure which be associated with an object's methods. If *exception-
procedure* is specified then it is used to process non-objects.

`make-generic-predicate`   [Function]

Returns a boolean procedure for any scheme object.

`make-method!` *object generic-method method*   [Function]

Associates *method* to the *generic-method* in the object. The *method* over-
rides any previous association with the *generic-method* within the object.
Using `unmake-method!` will restore the object's previous association with the
generic-method. method must be a procedure.

`make-predicate!` *object generic-preciate*   [Function]

Makes a predicate method associated with the *generic-predicate*.

`unmake-method!` *object generic-method*   [Function]

Removes an object's association with a *generic-method* .

`get-method` *object generic-method*   [Function]

Returns the object's method associated (if any) with the *generic-method*. If no asso-
ciated method exists an error is flagged.

7.1.16 Examples

```
(require 'object)

(define instantiate (make-generic-method))

(define (make-instance-object . ancestors)
  (define self (apply make-object
                      (map (lambda (obj) (instantiate obj)) ancestors)))
  (make-method! self instantiate (lambda (self) self))
  self)

(define who (make-generic-method))
(define imigrate! (make-generic-method))
(define emigrate! (make-generic-method))
(define describe (make-generic-method))
(define name (make-generic-method))
(define address (make-generic-method))
(define members (make-generic-method))
```

```
(define society
  (let ()
    (define self (make-instance-object))
    (define population '())
    (make-method! self imigrate!
                  (lambda (new-person)
                    (if (not (eq? new-person self))
                        (set! population (cons new-person population)))))
    (make-method! self emigrate!
                  (lambda (person)
                    (if (not (eq? person self))
                        (set! population
                              (comlist:remove-if (lambda (member)
                                                   (eq? member person))
                                                 population)))))
    (make-method! self describe
                  (lambda (self)
                    (map (lambda (person) (describe person)) population)))
    (make-method! self who
                  (lambda (self) (map (lambda (person) (name person))
                                      population)))
    (make-method! self members (lambda (self) population))
    self))

(define (make-person %name %address)
  (define self (make-instance-object society))
  (make-method! self name (lambda (self) %name))
  (make-method! self address (lambda (self) %address))
  (make-method! self who (lambda (self) (name self)))
  (make-method! self instantiate
                (lambda (self)
                  (make-person (string-append (name self) "-son-of")
                               %address)))
  (make-method! self describe
                (lambda (self) (list (name self) (address self))))
  (imigrate! self)
  self)
```

7.1.16.1 Inverter Documentation

Inheritance:

```
        <inverter>::(<number> <description>)
```

Generic-methods

```
        <inverter>::value        ⇒  <number>::value
        <inverter>::set-value!  ⇒  <number>::set-value!
        <inverter>::describe     ⇒  <description>::describe
```

```
            <inverter>::help
            <inverter>::invert
            <inverter>::inverter?
```

7.1.16.2 Number Documention

Inheritance

```
            <number>::()
```

   Slots

```
            <number>::<x>
```

Generic Methods

```
            <number>::value
            <number>::set-value!
```

7.1.16.3 Inverter code

```
  (require 'object)

  (define value (make-generic-method (lambda (val) val)))
  (define set-value! (make-generic-method))
  (define invert (make-generic-method
                   (lambda (val)
                     (if (number? val)
                         (/ 1 val)
                         (error "Method not supported:" val)))))
  (define noop (make-generic-method))
  (define inverter? (make-generic-predicate))
  (define describe (make-generic-method))
  (define help (make-generic-method))

  (define (make-number x)
    (define self (make-object))
    (make-method! self value (lambda (this) x))
    (make-method! self set-value!
                  (lambda (this new-value) (set! x new-value)))
    self)

  (define (make-description str)
    (define self (make-object))
    (make-method! self describe (lambda (this) str))
    (make-method! self help (lambda (this) "Help not available"))
    self)

  (define (make-inverter)
    (let* ((self (make-object
                   (make-number 1)
                   (make-description "A number which can be inverted")))
```

```
        (<value> (get-method self value)))
    (make-method! self invert (lambda (self) (/ 1 (<value> self))))
    (make-predicate! self inverter?)
    (unmake-method! self help)
    (make-method! self help
                  (lambda (self)
                    (display "Inverter Methods:") (newline)
                    (display "  (value inverter) ==> n") (newline)))
    self))

;;;;; Try it out

(define invert! (make-generic-method))

(define x (make-inverter))

(make-method! x invert! (lambda (x) (set-value! x (/ 1 (value x)))))

(value x)                         ⇒ 1
(set-value! x 33)                 ⇒ undefined
(invert! x)                       ⇒ undefined
(value x)                         ⇒ 1/33

(unmake-method! x invert!)        ⇒ undefined

(invert! x)                       [error]   ERROR: Method not supported: x
```

7.1.17 Priority Queues

`(require 'priority-queue)`

This algorithm for priority queues is due to *Introduction to Algorithms* by T. Cormen, C. Leiserson, R. Rivest. 1989 MIT Press.

make-heap *pred<?*   [Function]
    Returns a binary heap suitable which can be used for priority queue operations.

heap-length *heap*   [Function]
    Returns the number of elements in *heap*.

heap-insert! *heap item*   [Procedure]
    Inserts *item* into *heap*. *item* can be inserted multiple times. The value returned is unspecified.

heap-extract-max! *heap*   [Procedure]
    Returns the item which is larger than all others according to the *pred<?* argument to `make-heap`. If there are no items in *heap*, an error is signaled.

7.1.18 Queues

`(require 'queue)`

A *queue* is a list where elements can be added to both the front and rear, and removed from the front (i.e., they are what are often called *dequeues*). A queue may also be used like a stack.

`make-queue`   [Function]
>   Returns a new, empty queue.

`queue?` *obj*   [Function]
>   Returns #t if *obj* is a queue.

`queue-empty?` *q*   [Function]
>   Returns #t if the queue *q* is empty.

`queue-push!` *q datum*   [Procedure]
>   Adds *datum* to the front of queue *q*.

`enqueue!` *q datum*   [Procedure]
>   Adds *datum* to the rear of queue *q*.

`dequeue!` *q*   [Procedure]
`queue-pop!` *q*   [Procedure]
>   Both of these procedures remove and return the datum at the front of the queue. `queue-pop!` is used to suggest that the queue is being used like a stack.

All of the following functions raise an error if the queue *q* is empty.

`dequeue-all!` *q*   [Procedure]
>   Removes and returns (the list) of all contents of queue *q*.

`queue-front` *q*   [Function]
>   Returns the datum at the front of the queue *q*.

`queue-rear` *q*   [Function]
>   Returns the datum at the rear of the queue *q*.

7.1.19 Records

`(require 'record)`

The Record package provides a facility for user to define their own record data types.

`make-record-type` *type-name field-names*   [Function]
>   Returns a *record-type descriptor*, a value representing a new data type disjoint from all others. The *type-name* argument must be a string, but is only used for debugging purposes (such as the printed representation of a record of the new type). The *field-names* argument is a list of symbols naming the *fields* of a record of the new type. It is an error if the list contains any duplicates. It is unspecified how record-type descriptors are represented.

`record-constructor` *rtd* [*field-names*]   [Function]
>   Returns a procedure for constructing new members of the type represented by *rtd*. The returned procedure accepts exactly as many arguments as there are symbols in the

given list, *field-names*; these are used, in order, as the initial values of those fields in a new record, which is returned by the constructor procedure. The values of any fields not named in that list are unspecified. The *field-names* argument defaults to the list of field names in the call to `make-record-type` that created the type represented by *rtd*; if the *field-names* argument is provided, it is an error if it contains any duplicates or any symbols not in the default list.

`record-predicate` *rtd*   [Function]

Returns a procedure for testing membership in the type represented by *rtd*. The returned procedure accepts exactly one argument and returns a true value if the argument is a member of the indicated record type; it returns a false value otherwise.

`record-accessor` *rtd field-name*   [Function]

Returns a procedure for reading the value of a particular field of a member of the type represented by *rtd*. The returned procedure accepts exactly one argument which must be a record of the appropriate type; it returns the current value of the field named by the symbol *field-name* in that record. The symbol *field-name* must be a member of the list of field-names in the call to `make-record-type` that created the type represented by *rtd*.

`record-modifier` *rtd field-name*   [Function]

Returns a procedure for writing the value of a particular field of a member of the type represented by *rtd*. The returned procedure accepts exactly two arguments: first, a record of the appropriate type, and second, an arbitrary Scheme value; it modifies the field named by the symbol *field-name* in that record to contain the given value. The returned value of the modifier procedure is unspecified. The symbol *field-name* must be a member of the list of field-names in the call to `make-record-type` that created the type represented by *rtd*.

In May of 1996, as a product of discussion on the `rrrs-authors` mailing list, I rewrote `record.scm` to portably implement type disjointness for record data types.

As long as an implementation's procedures are opaque and the `record` code is loaded before other programs, this will give disjoint record types which are unforgeable and incorruptible by R4RS procedures.

As a consequence, the procedures `record?`, `record-type-descriptor`, `record-type-name`.and `record-type-field-names` are no longer supported.

7.2 Sorting and Searching

7.2.1 Common List Functions

`(require 'common-list-functions)`

The procedures below follow the Common LISP equivalents apart from optional arguments in some cases.

7.2.1.1 List construction

`make-list` *k*   [Function]

make-list *k init*   [Function]
>   make-list creates and returns a list of *k* elements. If *init* is included, all elements in
>   the list are initialized to *init*.
>
>   Example:
>
> ```
> (make-list 3)
>    ⇒ (#<unspecified> #<unspecified> #<unspecified>)
> (make-list 5 'foo)
>    ⇒ (foo foo foo foo foo)
> ```

list* *obj1 obj2 ...*   [Function]
>   Works like `list` except that the cdr of the last pair is the last argument unless there
>   is only one argument, when the result is just that argument. Sometimes called `cons*`.
>   E.g.:
>
> ```
> (list* 1)
>    ⇒ 1
> (list* 1 2 3)
>    ⇒ (1 2 . 3)
> (list* 1 2 '(3 4))
>    ⇒ (1 2 3 4)
> (list* args '())
>    ≡ (list args)
> ```

copy-list *lst*   [Function]
>   copy-list makes a copy of *lst* using new pairs and returns it. Only the top level
>   of the list is copied, i.e., pairs forming elements of the copied list remain `eq?` to the
>   corresponding elements of the original; the copy is, however, not `eq?` to the original,
>   but is `equal?` to it.
>
>   Example:
>
> ```
> (copy-list '(foo foo foo))
>    ⇒ (foo foo foo)
> (define q '(foo bar baz bang))
> (define p q)
> (eq? p q)
>    ⇒ #t
> (define r (copy-list q))
> (eq? q r)
>    ⇒ #f
> (equal? q r)
>    ⇒ #t
> (define bar '(bar))
> (eq? bar (car (copy-list (list bar 'foo))))
>    ⇒ #t
> ```

7.2.1.2 Lists as sets

`eqv?` is used to test for membership by procedures which treat lists as sets.

adjoin *e l*   [Function]

> adjoin returns the adjoint of the element *e* and the list *l*. That is, if *e* is in *l*, `adjoin`
> returns *l*, otherwise, it returns (`cons e l`).
>
> Example:
>
>   (adjoin 'baz '(bar baz bang))
>   ⇒ (bar baz bang)
>   (adjoin 'foo '(bar baz bang))
>   ⇒ (foo bar baz bang)

union *l1 l2*   [Function]

> union returns a list of all elements that are in *l1* or *l2*. Duplicates between *l1* and *l2*
> are culled. Duplicates within *l1* or within *l2* may or may not be removed.
>
> Example:
>
>   (union '(1 2 3 4) '(5 6 7 8))
>   ⇒ (1 2 3 4 5 6 7 8)
>   (union '(0 1 2 3 4) '(3 4 5 6))
>   ⇒ (5 6 0 1 2 3 4)

intersection *l1 l2*   [Function]

> intersection returns a list of all elements that are in both *l1* and *l2*.
>
> Example:
>
>   (intersection '(1 2 3 4) '(3 4 5 6))
>   ⇒ (3 4)
>   (intersection '(1 2 3 4) '(5 6 7 8))
>   ⇒ ()

set-difference *l1 l2*   [Function]

> set-difference returns a list of all elements that are in *l1* but not in *l2*.
>
> Example:
>
>   (set-difference '(1 2 3 4) '(3 4 5 6))
>   ⇒ (1 2)
>   (set-difference '(1 2 3 4) '(1 2 3 4 5 6))
>   ⇒ ()

subset? *list1 list2*   [Function]

> Returns #t if every element of *list1* is `eqv?` an element of *list2*; otherwise returns #f.
>
> Example:
>
>   (subset? '(1 2 3 4) '(3 4 5 6))
>   ⇒ #f
>   (subset? '(1 2 3 4) '(6 5 4 3 2 1 0))
>   ⇒ #t

member-if *pred lst*   [Function]

> member-if returns the list headed by the first element of *lst* to satisfy (`pred`
> `element`). Member-if returns #f if *pred* returns #f for every *element* in *lst*.

Example:

```
(member-if vector? '(a 2 b 4))
    ⇒ #f
(member-if number? '(a 2 b 4))
    ⇒ (2 b 4)
```

some *pred lst1 lst2 . . .*   [Function]

pred is a boolean function of as many arguments as there are list arguments to `some` i.e., *lst* plus any optional arguments. *pred* is applied to successive elements of the list arguments in order. `some` returns `#t` as soon as one of these applications returns `#t`, and is `#f` if none returns `#t`. All the lists should have the same length.

Example:

```
(some odd? '(1 2 3 4))
    ⇒ #t

(some odd? '(2 4 6 8))
    ⇒ #f

(some > '(1 3) '(2 4))
    ⇒ #f
```

every *pred lst1 lst2 . . .*   [Function]

`every` is analogous to `some` except it returns `#t` if every application of *pred* is `#t` and `#f` otherwise.

Example:

```
(every even? '(1 2 3 4))
    ⇒ #f

(every even? '(2 4 6 8))
    ⇒ #t

(every > '(2 3) '(1 4))
    ⇒ #f
```

notany *pred lst1 . . .*   [Function]

`notany` is analogous to `some` but returns `#t` if no application of *pred* returns `#t` or `#f` as soon as any one does.

notevery *pred lst1 . . .*   [Function]

`notevery` is analogous to `some` but returns `#t` as soon as an application of *pred* returns `#f`, and `#f` otherwise.

Example:

```
(notevery even? '(1 2 3 4))
    ⇒ #t

(notevery even? '(2 4 6 8))
```

> ⇒ #f

`list-of??` *predicate*   [Function]
>   Returns a predicate which returns true if its argument is a list every element of which
>   satisfies *predicate*.

`list-of??` *predicate low-bound high-bound*   [Function]
>   *low-bound* and *high-bound* are non-negative integers. `list-of??` returns a predicate
>   which returns true if its argument is a list of length between *low-bound* and *high-bound* (inclusive); every element of which satisfies *predicate*.

`list-of??` *predicate bound*   [Function]
>   *bound* is an integer. If *bound* is negative, `list-of??` returns a predicate which
>   returns true if its argument is a list of length greater than `(- bound)`; every element
>   of which satisfies *predicate*. Otherwise, `list-of??` returns a predicate which returns
>   true if its argument is a list of length less than or equal to *bound*; every element of
>   which satisfies *predicate*.

`find-if` *pred lst*   [Function]
>   `find-if` searches for the first *element* in *lst* such that (`pred element`) returns `#t`. If
>   it finds any such *element* in *lst*, *element* is returned. Otherwise, `#f` is returned.
>
>   Example:
>
>    ```
>    (find-if number? '(foo 1 bar 2))
>        ⇒ 1
>
>    (find-if number? '(foo bar baz bang))
>        ⇒ #f
>
>    (find-if symbol? '(1 2 foo bar))
>        ⇒ foo
>    ```

`remove` *elt lst*   [Function]
>   `remove` removes all occurrences of *elt* from *lst* using `eqv?` to test for equality and
>   returns everything that's left. N.B.: other implementations (Chez, Scheme->C and
>   T, at least) use `equal?` as the equality test.
>
>   Example:
>
>    ```
>    (remove 1 '(1 2 1 3 1 4 1 5))
>        ⇒ (2 3 4 5)
>
>    (remove 'foo '(bar baz bang))
>        ⇒ (bar baz bang)
>    ```

`remove-if` *pred lst*   [Function]
>   `remove-if` removes all *element*s from *lst* where (`pred element`) is `#t` and returns
>   everything that's left.
>
>   Example:
>
>    ```
>    (remove-if number? '(1 2 3 4))
>    ```

$\Rightarrow$ ()

```
(remove-if even? '(1 2 3 4 5 6 7 8))
   ⇒ (1 3 5 7)
```

remove-if-not *pred lst*   [Function]

> **remove-if-not** removes all *element*s from *lst* for which (`pred element`) is `#f` and returns everything that's left.
>
> Example:
>
> ```
> (remove-if-not number? '(foo bar baz))
>    ⇒ ()
> (remove-if-not odd? '(1 2 3 4 5 6 7 8))
>    ⇒ (1 3 5 7)
> ```

has-duplicates? *lst*   [Function]

> returns `#t` if 2 members of *lst* are `equal?`, `#f` otherwise.
>
> Example:
>
> ```
> (has-duplicates? '(1 2 3 4))
>    ⇒ #f
>
> (has-duplicates? '(2 4 3 4))
>    ⇒ #t
> ```

The procedure `remove-duplicates` uses `member` (rather than `memv`).

remove-duplicates *lst*   [Function]

> returns a copy of *lst* with its duplicate members removed. Elements are considered duplicate if they are `equal?`.
>
> Example:
>
> ```
> (remove-duplicates '(1 2 3 4))
>    ⇒ (1 2 3 4)
>
> (remove-duplicates '(2 4 3 4))
>    ⇒ (2 4 3)
> ```

7.2.1.3 Lists as sequences

position *obj lst*   [Function]

> **position** returns the 0-based position of *obj* in *lst*, or `#f` if *obj* does not occur in *lst*.
>
> Example:
>
> ```
> (position 'foo '(foo bar baz bang))
>    ⇒ 0
> (position 'baz '(foo bar baz bang))
>    ⇒ 2
> (position 'oops '(foo bar baz bang))
>    ⇒ #f
> ```

`reduce` *p* *lst*   [Function]

    `reduce` combines all the elements of a sequence using a binary operation (the combination is left-associative). For example, using `+`, one can add up all the elements. `reduce` allows you to apply a function which accepts only two arguments to more than 2 objects. Functional programmers usually refer to this as *foldl*. `collect:reduce` (see Section 7.1.10 [Collections], page 210) provides a version of `collect` generalized to collections.

    Example:

```
(reduce + '(1 2 3 4))
    ⇒ 10
(define (bad-sum . l) (reduce + l))
(bad-sum 1 2 3 4)
    ≡ (reduce + (1 2 3 4))
    ≡ (+ (+ (+ 1 2) 3) 4)
⇒ 10
(bad-sum)
    ≡ (reduce + ())
    ⇒ ()
(reduce string-append '("hello" "cruel" "world"))
    ≡ (string-append (string-append "hello" "cruel") "world")
    ⇒ "hellocruelworld"
(reduce anything '())
    ⇒ ()
(reduce anything '(x))
    ⇒ x
```

    What follows is a rather non-standard implementation of `reverse` in terms of `reduce` and a combinator elsewhere called *C*.

```
;;; Contributed by Jussi Piitulainen (jpiitula @ ling.helsinki.fi)

(define commute
  (lambda (f)
    (lambda (x y)
      (f y x))))

(define reverse
  (lambda (args)
    (reduce-init (commute cons) '() args)))
```

`reduce-init` *p* *init* *lst*   [Function]

    `reduce-init` is the same as reduce, except that it implicitly inserts *init* at the start of the list. `reduce-init` is preferred if you want to handle the null list, the one-element, and lists with two or more elements consistently. It is common to use the operator's idempotent as the initializer. Functional programmers usually call this *foldl*.

    Example:

```
(define (sum . l) (reduce-init + 0 l))
```

```
(sum 1 2 3 4)
   ≡ (reduce-init + 0 (1 2 3 4))
   ≡ (+ (+ (+ (+ 0 1) 2) 3) 4)
   ⇒ 10
(sum)
   ≡ (reduce-init + 0 '())
   ⇒ 0

(reduce-init string-append "@" '("hello" "cruel" "world"))
≡
(string-append (string-append (string-append "@" "hello")
                                  "cruel")
               "world")
⇒ "@hellocruelworld"
```

Given a differentiation of 2 arguments, diff, the following will differentiate by any number of variables.

```
(define (diff* exp . vars)
  (reduce-init diff exp vars))
```

Example:

```
;;; Real-world example:  Insertion sort using reduce-init.

(define (insert l item)
  (if (null? l)
      (list item)
      (if (< (car l) item)
          (cons (car l) (insert (cdr l) item))
          (cons item l))))
(define (insertion-sort l) (reduce-init insert '() l))

(insertion-sort '(3 1 4 1 5)
   ≡ (reduce-init insert () (3 1 4 1 5))
   ≡ (insert (insert (insert (insert (insert () 3) 1) 4) 1) 5)
   ≡ (insert (insert (insert (insert (3)) 1) 4) 1) 5)
   ≡ (insert (insert (insert (1 3) 4) 1) 5)
   ≡ (insert (insert (1 3 4) 1) 5)
   ≡ (insert (1 1 3 4) 5)
   ⇒ (1 1 3 4 5)
```

last *lst n*   [Function]

last returns the last *n* elements of *lst*. *n* must be a non-negative integer.

Example:

```
(last '(foo bar baz bang) 2)
   ⇒ (baz bang)
(last '(1 2 3) 0)
```

⇒ ()

butlast *lst n*   [Function]

    `butlast` returns all but the last *n* elements of *lst*.

    Example:

```
(butlast '(a b c d) 3)
    ⇒ (a)
(butlast '(a b c d) 4)
    ⇒ ()
```

`last` and `butlast` split a list into two parts when given identical arguments.

```
(last '(a b c d e) 2)
   ⇒ (d e)
(butlast '(a b c d e) 2)
   ⇒ (a b c)
```

nthcdr *n lst*   [Function]

    `nthcdr` takes *n* `cdrs` of *lst* and returns the result. Thus (`nthcdr 3 lst`) ≡ (`cdddr lst`)

    Example:

```
(nthcdr 2 '(a b c d))
    ⇒ (c d)
(nthcdr 0 '(a b c d))
    ⇒ (a b c d)
```

butnthcdr *n lst*   [Function]

    `butnthcdr` returns all but the nthcdr *n* elements of *lst*.

    Example:

```
(butnthcdr 3 '(a b c d))
    ⇒ (a b c)
(butnthcdr 4 '(a b c d))
    ⇒ (a b c d)
```

`nthcdr` and `butnthcdr` split a list into two parts when given identical arguments.

```
(nthcdr 2 '(a b c d e))
   ⇒ (c d e)
(butnthcdr 2 '(a b c d e))
   ⇒ (a b)
```

butnth *n lst*   [Function]

    `butnth` returns a list of all but the nth element of *lst*.

    Example:

```
(butnth 2 '(a b c d))
    ⇒ (a b d)
(butnth 4 '(a b c d))
    ⇒ (a b c d)
```

7.2.1.4 Destructive list operations

These procedures may mutate the list they operate on, but any such mutation is undefined.

nconc *args*   [Procedure]

> **nconc** destructively concatenates its arguments. (Compare this with **append**, which copies arguments rather than destroying them.) Sometimes called **append!** (see Section 7.4.4 [Rev2 Procedures], page 249).
>
> Example: You want to find the subsets of a set. Here's the obvious way:
>
> ```
> (define (subsets set)
>   (if (null? set)
>       '(())
>       (append (map (lambda (sub) (cons (car set) sub))
>                    (subsets (cdr set)))
>               (subsets (cdr set)))))
> ```
>
> But that does way more consing than you need. Instead, you could replace the **append** with **nconc**, since you don't have any need for all the intermediate results.
>
> Example:
>
> ```
> (define x '(a b c))
> (define y '(d e f))
> (nconc x y)
>     ⇒ (a b c d e f)
> x
>     ⇒ (a b c d e f)
> ```
>
> **nconc** is the same as **append!** in **sc2.scm**.

nreverse *lst*   [Procedure]

> **nreverse** reverses the order of elements in *lst* by mutating **cdr**s of the list. Sometimes called **reverse!**.
>
> Example:
>
> ```
> (define foo '(a b c))
> (nreverse foo)
>     ⇒ (c b a)
> foo
>     ⇒ (a)
> ```
>
> Some people have been confused about how to use **nreverse**, thinking that it doesn't return a value. It needs to be pointed out that
>
> ```
> (set! lst (nreverse lst))
> ```
>
> is the proper usage, not
>
> ```
> (nreverse lst)
> ```
>
> The example should suffice to show why this is the case.

delete *elt lst*   [Procedure]
delete-if *pred lst*   [Procedure]

`delete-if-not` *pred lst*   [Procedure]

> Destructive versions of `remove` `remove-if`, and `remove-if-not`.
>
> Example:
>
> ```
> (define lst (list 'foo 'bar 'baz 'bang))
> (delete 'foo lst)
>     ⇒ (bar baz bang)
> lst
>     ⇒ (foo bar baz bang)
>
> (define lst (list 1 2 3 4 5 6 7 8 9))
> (delete-if odd? lst)
>     ⇒ (2 4 6 8)
> lst
>     ⇒ (1 2 4 6 8)
> ```
>
> Some people have been confused about how to use `delete`, `delete-if`, and
> `delete-if`, thinking that they don't return a value. It needs to be pointed out that
>
> ```
> (set! lst (delete el lst))
> ```
>
> is the proper usage, not
>
> ```
> (delete el lst)
> ```
>
> The examples should suffice to show why this is the case.

7.2.1.5 Non-List functions

`and?` *arg1 ...*   [Function]

> `and?` checks to see if all its arguments are true. If they are, `and?` returns `#t`, otherwise,
> `#f`. (In contrast to `and`, this is a function, so all arguments are always evaluated and
> in an unspecified order.)
>
> Example:
>
> ```
> (and? 1 2 3)
>     ⇒ #t
> (and #f 1 2)
>     ⇒ #f
> ```

`or?` *arg1 ...*   [Function]

> `or?` checks to see if any of its arguments are true. If any is true, `or?` returns `#t`, and
> `#f` otherwise. (To `or` as `and?` is to `and`.)
>
> Example:
>
> ```
> (or? 1 2 #f)
>     ⇒ #t
> (or? #f #f #f)
>     ⇒ #f
> ```

`atom?` *object*   [Function]

> Returns `#t` if *object* is not a pair and `#f` if it is pair. (Called `atom` in Common LISP.)

```
(atom? 1)
    ⇒ #t
(atom? '(1 2))
    ⇒ #f
(atom? #(1 2))   ; dubious!
    ⇒ #t
```

7.2.2 Tree operations

`(require 'tree)`

These are operations that treat lists a representations of trees.

`subst` *new old tree*	[Function]
`substq` *new old tree*	[Function]
`substv` *new old tree*	[Function]
`subst` *new old tree equ?*	[Function]

`subst` makes a copy of *tree*, substituting *new* for every subtree or leaf of *tree* which is `equal?` to *old* and returns a modified tree. The original *tree* is unchanged, but may share parts with the result.

`substq` and `substv` are similar, but test against *old* using `eq?` and `eqv?` respectively. If `subst` is called with a fourth argument, *equ?* is the equality predicate.

Examples:

```
(substq 'tempest 'hurricane '(shakespeare wrote (the hurricane)))
    ⇒ (shakespeare wrote (the tempest))
(substq 'foo '() '(shakespeare wrote (twelfth night)))
    ⇒ (shakespeare wrote (twelfth night . foo) . foo)
(subst '(a . cons) '(old . pair)
        '((old . spice) ((old . shoes) old . pair) (old . pair)))
    ⇒ ((old . spice) ((old . shoes) a . cons) (a . cons))
```

`copy-tree` *tree*	[Function]

Makes a copy of the nested list structure *tree* using new pairs and returns it. All levels are copied, so that none of the pairs in the tree are `eq?` to the original ones – only the leaves are.

Example:

```
(define bar '(bar))
(copy-tree (list bar 'foo))
    ⇒ ((bar) foo)
(eq? bar (car (copy-tree (list bar 'foo))))
    ⇒ #f
```

7.2.3 Chapter Ordering

`(require 'chapter-order)`

The 'chap:' functions deal with strings which are ordered like chapter numbers (or letters) in a book. Each section of the string consists of consecutive numeric or consecutive aphabetic characters of like case.

`chap:`**`string<?`** *string1 string2*   [Function]

    Returns #t if the first non-matching run of alphabetic upper-case or the first non-matching run of alphabetic lower-case or the first non-matching run of numeric characters of *string1* is `string<?` than the corresponding non-matching run of characters of *string2*.

```
(chap:string<? "a.9" "a.10")                    ⇒ #t
(chap:string<? "4c" "4aa")                      ⇒ #t
(chap:string<? "Revised^{3.99}" "Revised^{4}")  ⇒ #t
```

`chap:`**`string>?`** *string1 string2*   [Function]
`chap:`**`string<=?`** *string1 string2*   [Function]
`chap:`**`string>=?`** *string1 string2*   [Function]

    Implement the corresponding chapter-order predicates.

`chap:`**`next-string`** *string*   [Function]

    Returns the next string in the *chapter order*. If *string* has no alphabetic or numeric characters, (`string-append` *string* `"0"`) is returnd. The argument to chap:next-string will always be `chap:string<?` than the result.

```
(chap:next-string "a.9")            ⇒ "a.10"
(chap:next-string "4c")             ⇒ "4d"
(chap:next-string "4z")             ⇒ "4aa"
(chap:next-string "Revised^{4}")    ⇒ "Revised^{5}"
```

7.2.4 Sorting

(`require` `'sort`) or (`require` `'srfi-95`)

    [by Richard A. O'Keefe, 1991]

    I am providing this source code with no restrictions at all on its use (but please retain D.H.D.Warren's credit for the original idea).

    The code of `merge` and `merge!` could have been quite a bit simpler, but they have been coded to reduce the amount of work done per iteration. (For example, we only have one `null?` test per iteration.)

    I gave serious consideration to producing Common-LISP-compatible functions. However, Common LISP's `sort` is our `sort!` (well, in fact Common LISP's `stable-sort` is our `sort!`; merge sort is *fast* as well as stable!) so adapting CL code to Scheme takes a bit of work anyway. I did, however, appeal to CL to determine the *order* of the arguments.

    The standard functions `<`, `>`, `char<?`, `char>?`, `char-ci<?`, `char-ci>?`, `string<?`, `string>?`, `string-ci<?`, and `string-ci>?` are suitable for use as comparison functions. Think of (`less?` `x y`) as saying when x must *not* precede y.

    [Addendum by Aubrey Jaffer, 2006]

    These procedures are stable when called with predicates which return #f when applied to identical arguments.

    The `sorted?`, `merge`, and `merge!` procedures consume asymptotic time and space no larger than $O(N)$, where N is the sum of the lengths of the sequence arguments. The `sort`

and `sort!` procedures consume asymptotic time and space no larger than $O(N*log(N))$, where N is the length of the sequence argument.

All five functions take an optional *key* argument corresponding to a CL-style '`&key`' argument. A *less?* predicate with a *key* argument behaves like:

    (lambda (x y) (less? (key x) (key y)))

All five functions will call the *key* argument at most once per element.

The '`!`' variants sort in place; `sort!` returns its *sequence* argument.

`sorted?` *sequence less?*   [Function]
`sorted?` *sequence less? key*   [Function]
> Returns `#t` when the sequence argument is in non-decreasing order according to *less?* (that is, there is no adjacent pair `... x y ...` for which (`less?` `y` `x`)).
>
> Returns `#f` when the sequence contains at least one out-of-order pair. It is an error if the sequence is not a list or array (including vectors and strings).

`merge` *list1 list2 less?*   [Function]
`merge` *list1 list2 less? key*   [Function]
> Merges two sorted lists, returning a freshly allocated list as its result.

`merge!` *list1 list2 less?*   [Function]
`merge!` *list1 list2 less? key*   [Function]
> Merges two sorted lists, re-using the pairs of *list1* and *list2* to build the result. The result will be either *list1* or *list2*.

`sort` *sequence less?*   [Function]
`sort` *sequence less? key*   [Function]
> Accepts a list or array (including vectors and strings) for *sequence*; and returns a completely new sequence which is sorted according to *less?*. The returned sequence is the same type as the argument *sequence*. Given valid arguments, it is always the case that:
>
>   (sorted? (sort sequence less?) less?) ⇒ #t

`sort!` *sequence less?*   [Function]
`sort!` *sequence less? key*   [Function]
> Returns list, array, vector, or string *sequence* which has been mutated to order its elements according to *less?*. Given valid arguments, it is always the case that:
>
>   (sorted? (sort! sequence less?) less?) ⇒ #t

7.2.5 Topological Sort

(require 'topological-sort) or (require 'tsort)

The algorithm is inspired by Cormen, Leiserson and Rivest (1990) *Introduction to Algorithms*, chapter 23.

`tsort` *dag pred*   [Function]
`topological-sort` *dag pred*   [Function]
> where

dag   is a list of sublists. The car of each sublist is a vertex. The cdr is the adjacency list of that vertex, i.e. a list of all vertices to which there exists an edge from the car vertex.

pred   is one of `eq?`, `eqv?`, `equal?`, `=`, `char=?`, `char-ci=?`, `string=?`, or `string-ci=?`.

Sort the directed acyclic graph *dag* so that for every edge from vertex *u* to *v*, *u* will come before *v* in the resulting list of vertices.

Time complexity: O (|V| + |E|)

Example (from Cormen):

> Prof. Bumstead topologically sorts his clothing when getting dressed. The first argument to `tsort` describes which garments he needs to put on before others. (For example, Prof Bumstead needs to put on his shirt before he puts on his tie or his belt.) `tsort` gives the correct order of dressing:

```
(require 'tsort)
(tsort '((shirt tie belt)
         (tie jacket)
         (belt jacket)
         (watch)
         (pants shoes belt)
         (undershorts pants shoes)
         (socks shoes))
       eq?)
⇒
(socks undershorts pants shoes watch shirt belt tie jacket)
```

7.2.6 Hashing

`(require 'hash)`

  These hashing functions are for use in quickly classifying objects. Hash tables use these functions.

hashq *obj k*   [Function]
hashv *obj k*   [Function]
hash *obj k*   [Function]

> Returns an exact non-negative integer less than *k*. For each non-negative integer less than *k* there are arguments *obj* for which the hashing functions applied to *obj* and *k* returns that integer.

> For `hashq`, (eq? obj1 obj2) implies (= (hashq obj1 k) (hashq obj2)).

> For `hashv`, (eqv? obj1 obj2) implies (= (hashv obj1 k) (hashv obj2)).

> For `hash`, (equal? obj1 obj2) implies (= (hash obj1 k) (hash obj2)).

> `hash`, `hashv`, and `hashq` return in time bounded by a constant. Notice that items having the same `hash` implies the items have the same `hashv` implies the items have the same `hashq`.

7.2.7 Space-Filling Curves

7.2.7.1 Multidimensional Space-Filling Curves

`(require 'space-filling)`

The algorithms and cell properties are described in `http://people.csail.mit.edu/jaffer/Geometry/RMDSFF.pdf`

`make-cell` *type rank side precession*   [Function]
`make-cell` *type rank side*   [Function]
`make-cell` *type rank*   [Function]

> *type* must be the symbol `diagonal`, `adjacent`, or `centered`. *rank* must be an integer larger than 1. *side*, if present, must be an even integer larger than 1 if *type* is `adjacent` or an odd integer larger than 2 otherwise; *side* defaults to the smallest value. *precession*, if present, must be an integer between 0 and *side^rank*-1; it is relevant only when *type* is `diagonal` or `centered`.

`make-cell` *Hamiltonian-path-vector precession*   [Function]
`make-cell` *Hamiltonian-path-vector*   [Function]

> *type* must be a vector of *side^rank* lists of *rank* of integers encoding the coordinate positions of a Hamiltonian path on the *rank*-dimensional grid of points starting and ending on corners of the grid. The starting corner must be the origin (all-zero coordinates). If the side-length is even, then the ending corner must be non-zero in only one coordinate; otherwise, the ending corner must be the furthest diagonally opposite corner from the origin.

> `make-cell` returns a data object suitable for passing as the first argument to `integer->coordinates` or `coordinates->integer`.

> Hilbert, Peano, and centered Peano cells are generated respectively by:

```
(make-cell 'adjacent rank 2)    ; Hilbert
(make-cell 'diagonal rank 3)    ; Peano
(make-cell 'centered rank 3)    ; centered Peano
```

In the conversion procedures, if the cell is `diagonal` or `adjacent`, then the coordinates and scalar must be nonnegative integers. If `centered`, then the integers can be negative.

`integer->coordinates` *cell u*   [Function]

> `integer->coordinates` converts the integer *u* to a list of coordinates according to *cell*.

`coordinates->integer` *cell v*   [Function]

> `coordinates->integer` converts the list of coordinates *v* to an integer according to *cell*.

coordinates->integer and *integer->coordinates* are inverse functions when passed the same *cell* argument.

7.2.7.2 Hilbert Space-Filling Curve

`(require 'hilbert-fill)`

The *Hilbert Space-Filling Curve* is a one-to-one mapping between a unit line segment and an *n*-dimensional unit cube. This implementation treats the nonnegative integers either as fractional bits of a given width or as nonnegative integers.

The integer procedures map the non-negative integers to an arbitrarily large *n*-dimensional cube with its corner at the origin and all coordinates are non-negative.

For any exact nonnegative integer *scalar* and exact integer *rank* > 2,

```
(= scalar (hilbert-coordinates->integer
          (integer->hilbert-coordinates scalar rank)))
                                    ⇒ #t
```

When treating integers as *k* fractional bits,

```
(= scalar (hilbert-coordinates->integer
          (integer->hilbert-coordinates scalar rank k)) k)
                                    ⇒ #t
```

integer->hilbert-coordinates *scalar rank*   [Function]
> Returns a list of *rank* integer coordinates corresponding to exact non-negative integer *scalar*. The lists returned by `integer->hilbert-coordinates` for *scalar* arguments 0 and 1 will differ in the first element.

integer->hilbert-coordinates *scalar rank k*   [Function]
> *scalar* must be a nonnegative integer of no more than `rank*k` bits.
>
> `integer->hilbert-coordinates` Returns a list of *rank* *k*-bit nonnegative integer coordinates corresponding to exact non-negative integer *scalar*. The curves generated by `integer->hilbert-coordinates` have the same alignment independent of *k*.

hilbert-coordinates->integer *coords*   [Function]
hilbert-coordinates->integer *coords k*   [Function]
> Returns an exact non-negative integer corresponding to *coords*, a list of non-negative integer coordinates.

7.2.7.3 Gray code

A *Gray code* is an ordering of non-negative integers in which exactly one bit differs between each pair of successive elements. There are multiple Gray codings. An n-bit Gray code corresponds to a Hamiltonian cycle on an n-dimensional hypercube.

Gray codes find use communicating incrementally changing values between asynchronous agents. De-laminated Gray codes comprise the coordinates of Hilbert space-filling curves.

integer->gray-code *k*   [Function]
> Converts *k* to a Gray code of the same `integer-length` as *k*.

gray-code->integer *k*   [Function]
> Converts the Gray code *k* to an integer of the same `integer-length` as *k*.
>
> For any non-negative integer *k*,

```
(eqv? k (gray-code->integer (integer->gray-code k)))
```

`=` *k1 k2*	[Function]
`gray-code<?` *k1 k2*	[Function]
`gray-code>?` *k1 k2*	[Function]
`gray-code<=?` *k1 k2*	[Function]
`gray-code>=?` *k1 k2*	[Function]

> These procedures return #t if their Gray code arguments are (respectively): equal, monotonically increasing, monotonically decreasing, monotonically nondecreasing, or monotonically nonincreasing.
>
> For any non-negative integers *k1* and *k2*, the Gray code predicate of (`integer->gray-code` k1) and (`integer->gray-code` k2) will return the same value as the corresponding predicate of *k1* and *k2*.

7.2.7.4 Bitwise Lamination

`delaminate-list` *count ks*   [Function]

> Returns a list of *count* integers comprised of the *j*th bit of the integers *ks* where *j* ranges from *count*-1 to 0.

```
(map (lambda (k) (number->string k 2))
    (delaminate-list 4 '(7 6 5 4 0 0 0 0)))
    ⇒ ("0" "11110000" "11000000" "10100000")
```

> `delaminate-list` is its own inverse:

```
(delaminate-list 8 (delaminate-list 4 '(7 6 5 4 0 0 0 0)))
    ⇒ (7 6 5 4 0 0 0 0)
```

7.2.7.5 Peano Space-Filling Curve

`(require 'peano-fill)`

`natural->peano-coordinates` *scalar rank*   [Function]

> Returns a list of *rank* nonnegative integer coordinates corresponding to exact nonnegative integer *scalar*. The lists returned by `natural->peano-coordinates` for *scalar* arguments 0 and 1 will differ in the first element.

`peano-coordinates->natural` *coords*   [Function]

> Returns an exact nonnegative integer corresponding to *coords*, a list of nonnegative integer coordinates.

`integer->peano-coordinates` *scalar rank*   [Function]

> Returns a list of *rank* integer coordinates corresponding to exact integer *scalar*. The lists returned by `integer->peano-coordinates` for *scalar* arguments 0 and 1 will differ in the first element.

`peano-coordinates->integer` *coords*   [Function]

> Returns an exact integer corresponding to *coords*, a list of integer coordinates.

7.2.7.6 Sierpinski Curve

`(require 'sierpinski)`

`make-sierpinski-indexer` *max-coordinate*   [Function]

Returns a procedure (eg hash-function) of 2 numeric arguments which preserves *nearness* in its mapping from NxN to N.

max-coordinate is the maximum coordinate (a positive integer) of a population of points. The returned procedures is a function that takes the x and y coordinates of a point, (non-negative integers) and returns an integer corresponding to the relative position of that point along a Sierpinski curve. (You can think of this as computing a (pseudo-) inverse of the Sierpinski spacefilling curve.)

Example use: Make an indexer (hash-function) for integer points lying in square of integer grid points [0,99]x[0,99]:

```
(define space-key (make-sierpinski-indexer 100))
```

Now let's compute the index of some points:

```
(space-key 24 78)                    ⇒ 9206
(space-key 23 80)                    ⇒ 9172
```

Note that locations (24, 78) and (23, 80) are near in index and therefore, because the Sierpinski spacefilling curve is continuous, we know they must also be near in the plane. Nearness in the plane does not, however, necessarily correspond to nearness in index, although it *tends* to be so.

Example applications:

- Sort points by Sierpinski index to get heuristic solution to *travelling salesman problem.* For details of performance, see L. Platzman and J. Bartholdi, "Spacefilling curves and the Euclidean travelling salesman problem", JACM 36(4):719–737 (October 1989) and references therein.

- Use Sierpinski index as key by which to store 2-dimensional data in a 1-dimensional data structure (such as a table). Then locations that are near each other in 2-d space will tend to be near each other in 1-d data structure; and locations that are near in 1-d data structure will be near in 2-d space. This can significantly speed retrieval from secondary storage because contiguous regions in the plane will tend to correspond to contiguous regions in secondary storage. (This is a standard technique for managing CAD/CAM or geographic data.)

7.2.8 Soundex

```
(require 'soundex)
```

`soundex` *name*   [Function]

Computes the *soundex* hash of *name*. Returns a string of an initial letter and up to three digits between 0 and 6. Soundex supposedly has the property that names that sound similar in normal English pronunciation tend to map to the same key.

Soundex was a classic algorithm used for manual filing of personal records before the advent of computers. It performs adequately for English names but has trouble with other languages.

See Knuth, Vol. 3 *Sorting and searching*, pp 391–2

To manage unusual inputs, `soundex` omits all non-alphabetic characters. Consequently, in this implementation:

```
(soundex <string of blanks>)     ⇒ ""
(soundex "")                     ⇒ ""
```

Examples from Knuth:

```
(map soundex '("Euler" "Gauss" "Hilbert" "Knuth"
                       "Lloyd" "Lukasiewicz"))
       ⇒ ("E460" "G200" "H416" "K530" "L300" "L222")


(map soundex '("Ellery" "Ghosh" "Heilbronn" "Kant"
                        "Ladd" "Lissajous"))
       ⇒ ("E460" "G200" "H416" "K530" "L300" "L222")
```

Some cases in which the algorithm fails (Knuth):

```
(map soundex '("Rogers" "Rodgers"))      ⇒ ("R262" "R326")


(map soundex '("Sinclair" "St. Clair")) ⇒ ("S524" "S324")


(map soundex '("Tchebysheff" "Chebyshev")) ⇒ ("T212" "C121")
```

7.2.9 String Search

`(require 'string-search)`

string-index *string char*   [Procedure]
string-index-ci *string char*   [Procedure]
> Returns the index of the first occurence of *char* within *string*, or `#f` if the *string* does not contain a character *char*.

string-reverse-index *string char*   [Procedure]
string-reverse-index-ci *string char*   [Procedure]
> Returns the index of the last occurence of *char* within *string*, or `#f` if the *string* does not contain a character *char*.

substring? *pattern string*   [Procedure]
substring-ci? *pattern string*   [Procedure]
> Searches *string* to see if some substring of *string* is equal to *pattern*. `substring?` returns the index of the first character of the first substring of *string* that is equal to *pattern*; or `#f` if *string* does not contain *pattern*.

```
(substring? "rat" "pirate") ⇒  2
(substring? "rat" "outrage") ⇒  #f
(substring? "" any-string) ⇒  0
```

find-string-from-port? *str in-port max-no-chars*   [Procedure]
> Looks for a string *str* within the first *max-no-chars* chars of the input port *in-port*.

find-string-from-port? *str in-port*   [Procedure]
> When called with two arguments, the search span is limited by the end of the input stream.

`find-string-from-port?` *str in-port char*   [Procedure]

> Searches up to the first occurrence of character *char* in *str*.

`find-string-from-port?` *str in-port proc*   [Procedure]

> Searches up to the first occurrence of the procedure *proc* returning non-false when
> called with a character (from *in-port*) argument.
>
> When the *str* is found, `find-string-from-port?` returns the number of characters
> it has read from the port, and the port is set to read the first char after that (that is,
> after the *str*) The function returns `#f` when the *str* isn't found.
>
> `find-string-from-port?` reads the port *strictly* sequentially, and does not perform
> any buffering. So `find-string-from-port?` can be used even if the *in-port* is open
> to a pipe or other communication channel.

`string-subst` *txt old1 new1 ...*   [Function]

> Returns a copy of string *txt* with all occurrences of string *old1* in *txt* replaced with
> *new1*; then *old2* replaced with *new2* .... Matches are found from the left. Matches
> do not overlap.

`count-newlines` *str*   [Function]

> Returns the number of '`#\newline`' characters in string *str*.

7.2.10 Sequence Comparison

`(require 'diff)`

`diff:edit-length` implements the algorithm:

The values returned by `diff:edit-length` can be used to gauge the degree of match between two sequences.

`diff:edits` and `diff:longest-common-subsequence` combine the algorithm with the divide-and-conquer method outlined in:

If the items being sequenced are text lines, then the computed edit-list is equivalent to the output of the *diff* utility program. If the items being sequenced are words, then it is like the lesser known *spiff* program.

`diff:longest-common-subsequence` *array1 array2 p-lim*   [Function]
`diff:longest-common-subsequence` *array1 array2*   [Function]

> *array1* and *array2* are one-dimensional arrays.
>
> The non-negative integer *p-lim*, if provided, is maximum number of deletions of the
> shorter sequence to allow. `diff:longest-common-subsequence` will return `#f` if more
> deletions would be necessary.
>
> `diff:longest-common-subsequence` returns a one-dimensional array of length
> `(quotient (- (+ len1 len2) (diff:edit-length array1 array2)) 2)` holding the
> longest sequence common to both *array*s.

`diff:edits` *array1 array2 p-lim*   [Function]
`diff:edits` *array1 array2*   [Function]

> *array1* and *array2* are one-dimensional arrays.

The non-negative integer *p-lim*, if provided, is maximum number of deletions of the shorter sequence to allow. `diff:edits` will return `#f` if more deletions would be necessary.

`diff:edits` returns a vector of length (`diff:edit-length` *array1* *array2*) composed of a shortest sequence of edits transformaing *array1* to *array2*.

Each edit is an integer:

k > 0   Inserts (`array-ref` *array1* (`+ -1` *j*)) into the sequence.

k < 0   Deletes (`array-ref` *array2* (`- -1` *k*)) from the sequence.

`diff:edit-length` *array1* *array2* *p-lim*   [Function]
`diff:edit-length` *array1* *array2*   [Function]
    array1 and *array2* are one-dimensional arrays.

    The non-negative integer *p-lim*, if provided, is maximum number of deletions of the shorter sequence to allow. `diff:edit-length` will return `#f` if more deletions would be necessary.

    `diff:edit-length` returns the length of the shortest sequence of edits transformaing *array1* to *array2*.

```
(diff:longest-common-subsequence "fghiejcklm" "fgehijkpqrlm")
⇒ "fghijklm"

(diff:edit-length "fghiejcklm" "fgehijkpqrlm")
⇒ 6

(diff:edits "fghiejcklm" "fgehijkpqrlm")
⇒ #A:fixZ32b(3 -5 -7 8 9 10)
       ; e   c   h p q   r
```

7.3 Procedures

Anything that doesn't fall neatly into any of the other categories winds up here.

7.3.1 Type Coercion

(require 'coerce)

`type-of` *obj*   [Function]
    Returns a symbol name for the type of *obj*.

`coerce` *obj* *result-type*   [Function]
    Converts and returns *obj* of type `char`, `number`, `string`, `symbol`, `list`, or `vector` to *result-type* (which must be one of these symbols).

7.3.2 String-Case

(require 'string-case)

`string-upcase` *str*   [Procedure]
`string-downcase` *str*   [Procedure]

string-capitalize *str*   [Procedure]
>   The obvious string conversion routines. These are non-destructive.

string-upcase! *str*   [Function]
string-downcase! *str*   [Function]
string-capitalize! *str*   [Function]
>   The destructive versions of the functions above.

string-ci->symbol *str*   [Function]
>   Converts string *str* to a symbol having the same case as if the symbol had been `read`.

symbol-append *obj1* ...   [Function]
>   Converts *obj1* ... to strings, appends them, and converts to a symbol which is
>   returned. Strings and numbers are converted to read's symbol case; the case of
>   symbol characters is not changed. #f is converted to the empty string (symbol).

StudlyCapsExpand *str delimiter*   [Function]
StudlyCapsExpand *str*   [Function]
>   *delimiter* must be a string or character. If absent, *delimiter* defaults to '-'.
>   `StudlyCapsExpand` returns a copy of *str* where *delimiter* is inserted between each
>   lower-case character immediately followed by an upper-case character; and between
>   two upper-case characters immediately followed by a lower-case character.

```
(StudlyCapsExpand "aX" " ")   ⇒ "a X"
(StudlyCapsExpand "aX" "..")  ⇒ "a..X"
(StudlyCapsExpand "AX")       ⇒ "AX"
(StudlyCapsExpand "Ax")       ⇒ "Ax"
(StudlyCapsExpand "AXLE")     ⇒ "AXLE"
(StudlyCapsExpand "aAXACz")   ⇒ "a-AXA-Cz"
(StudlyCapsExpand "AaXACz")   ⇒ "Aa-XA-Cz"
(StudlyCapsExpand "AAaXACz")  ⇒ "A-Aa-XA-Cz"
(StudlyCapsExpand "AAaXAC")   ⇒ "A-Aa-XAC"
```

7.3.3 String Ports

`(require 'string-port)`

call-with-output-string *proc*   [Procedure]
>   *proc* must be a procedure of one argument. This procedure calls *proc* with one
>   argument: a (newly created) output port. When the function returns, the string
>   composed of the characters written into the port is returned.

call-with-input-string *string proc*   [Procedure]
>   *proc* must be a procedure of one argument. This procedure calls *proc* with one
>   argument: an (newly created) input port from which *string*'s contents may be read.
>   When *proc* returns, the port is closed and the value yielded by the procedure *proc* is
>   returned.

7.3.4 Line I/O

`(require 'line-i/o)`

`read-line`   [Function]
`read-line` *port*   [Function]
> Returns a string of the characters up to, but not including a newline or end of file, updating *port* to point to the character following the newline. If no characters are available, an end of file object is returned. The *port* argument may be omitted, in which case it defaults to the value returned by `current-input-port`.

`read-line!` *string*   [Procedure]
`read-line!` *string port*   [Procedure]
> Fills *string* with characters up to, but not including a newline or end of file, updating the *port* to point to the last character read or following the newline if it was read. If no characters are available, an end of file object is returned. If a newline or end of file was found, the number of characters read is returned. Otherwise, `#f` is returned. The *port* argument may be omitted, in which case it defaults to the value returned by `current-input-port`.

`write-line` *string*   [Function]
`write-line` *string port*   [Function]
> Writes *string* followed by a newline to the given *port* and returns an unspecified value. The *Port* argument may be omitted, in which case it defaults to the value returned by `current-input-port`.

`system->line` *command tmp*   [Function]
`system->line` *command*   [Function]
> *command* must be a string. The string *tmp*, if supplied, is a path to use as a temporary file. `system->line` calls `system` with *command* as argument, redirecting stdout to file *tmp*. `system->line` returns a string containing the first line of output from *tmp*.
>
> `system->line` is intended to be a portable method for getting one-line results from programs like `pwd`, `whoami`, `hostname`, `which`, `identify`, and `cksum`. Its behavior when called with programs which generate lots of output is unspecified.

7.3.5 Multi-Processing

`(require 'process)`

> This module implements asynchronous (non-polled) time-sliced multi-processing in the SCM Scheme implementation using procedures `alarm` and `alarm-interrupt`. Until this is ported to another implementation, consider it an example of writing schedulers in Scheme.

`add-process!` *proc*   [Procedure]
> Adds proc, which must be a procedure (or continuation) capable of accepting accepting one argument, to the `process:queue`. The value returned is unspecified. The argument to *proc* should be ignored. If *proc* returns, the process is killed.

`process:schedule!`   [Procedure]
> Saves the current process on `process:queue` and runs the next process from `process:queue`. The value returned is unspecified.

`kill-process!`   [Procedure]

> Kills the current process and runs the next process from `process:queue`. If there are no more processes on `process:queue`, `(slib:exit)` is called (see Section 2.4 [System], page 15).

7.3.6 Metric Units

`(require 'metric-units)`

> `http://people.csail.mit.edu/jaffer/MIXF`

Metric Interchange Format is a character string encoding for numerical values and units which:

- is unambiguous in all locales;
- uses only [TOG] "Portable Character Set" characters matching "Basic Latin" characters in Plane 0 of the Universal Character Set [UCS];
- is transparent to [UTF-7] and [UTF-8] UCS transformation formats;
- is human readable and writable;
- is machine readable and writable;
- incorporates SI prefixes and units;
- incorporates [ISO 6093] numbers; and
- incorporates [IEC 60027-2] binary prefixes.

In the expression for the value of a quantity, the unit symbol is placed after the numerical value. A dot (PERIOD, '.') is placed between the numerical value and the unit symbol.

Within a compound unit, each of the base and derived symbols can optionally have an attached SI prefix.

Unit symbols formed from other unit symbols by multiplication are indicated by means of a dot (PERIOD, '.') placed between them.

Unit symbols formed from other unit symbols by division are indicated by means of a SOLIDUS ('/') or negative exponents. The SOLIDUS must not be repeated in the same compound unit unless contained within a parenthesized subexpression.

The grouping formed by a prefix symbol attached to a unit symbol constitutes a new inseparable symbol (forming a multiple or submultiple of the unit concerned) which can be raised to a positive or negative power and which can be combined with other unit symbols to form compound unit symbols.

The grouping formed by surrounding compound unit symbols with parentheses ('(' and ')') constitutes a new inseparable symbol which can be raised to a positive or negative power and which can be combined with other unit symbols to form compound unit symbols.

Compound prefix symbols, that is, prefix symbols formed by the juxtaposition of two or more prefix symbols, are not permitted.

Prefix symbols are not used with the time-related unit symbols min (minute), h (hour), d (day). No prefix symbol may be used with dB (decibel). Only submultiple prefix symbols may be used with the unit symbols L (liter), Np (neper), o (degree), oC (degree Celsius),

rad (radian), and sr (steradian). Submultiple prefix symbols may not be used with the unit symbols t (metric ton), r (revolution), or Bd (baud).

A unit exponent follows the unit, separated by a CIRCUMFLEX ('^'). Exponents may be positive or negative. Fractional exponents must be parenthesized.

7.3.6.1 SI Prefixes

```
Factor      Name    Symbol  |  Factor     Name    Symbol
======      ====    ======  |  ======     ====    ======
  1e24      yotta     Y     |   1e-1      deci       d
  1e21      zetta     Z     |   1e-2      centi      c
  1e18      exa       E     |   1e-3      milli      m
  1e15      peta      P     |   1e-6      micro      u
  1e12      tera      T     |   1e-9      nano       n
  1e9       giga      G     |   1e-12     pico       p
  1e6       mega      M     |   1e-15     femto      f
  1e3       kilo      k     |   1e-18     atto       a
  1e2       hecto     h     |   1e-21     zepto      z
  1e1       deka      da    |   1e-24     yocto      y
```

7.3.6.2 Binary Prefixes

These binary prefixes are valid only with the units B (byte) and bit. However, decimal prefixes can also be used with bit; and decimal multiple (not submultiple) prefixes can also be used with B (byte).

```
         Factor          (power-of-2)   Name   Symbol
         ======          ============   ====   ======
1.152921504606846976e18   (2^60)        exbi     Ei
  1.125899906842624e15    (2^50)        pebi     Pi
    1.099511627776e12     (2^40)        tebi     Ti
       1.073741824e9      (2^30)        gibi     Gi
          1.048576e6      (2^20)        mebi     Mi
             1.024e3      (2^10)        kibi     Ki
```

7.3.6.3 Unit Symbols

```
     Type of Quantity        Name        Symbol   Equivalent
     ================        ====        ======   ==========
  time                    second           s
  time                    minute          min  = 60.s
  time                    hour             h   = 60.min
  time                    day              d   = 24.h
  frequency               hertz           Hz     s^-1
  signaling rate          baud            Bd     s^-1
  length                  meter            m
  volume                  liter            L      dm^3
  plane angle             radian          rad
  solid angle             steradian        sr     rad^2
  plane angle             revolution     * r    = 6.283185307179586.rad
```

```
    plane angle                degree         * o   = 2.777777777777778e-3.r
    information capacity       bit              bit
    information capacity       byte, octet      B   = 8.bit
    mass                       gram             g
    mass                       ton              t     Mg
    mass              unified atomic mass unit  u   = 1.66053873e-27.kg
    amount of substance        mole             mol
    catalytic activity         katal            kat   mol/s
    thermodynamic temperature kelvin            K
    centigrade temperature     degree Celsius   oC
    luminous intensity         candela          cd
    luminous flux              lumen            lm    cd.sr
    illuminance                lux              lx    lm/m^2
    force                      newton           N     m.kg.s^-2
    pressure, stress           pascal           Pa    N/m^2
    energy, work, heat         joule            J     N.m
    energy                     electronvolt     eV  = 1.602176462e-19.J
    power, radiant flux        watt             W     J/s
    logarithm of power ratio   neper            Np
    logarithm of power ratio   decibel        * dB  = 0.1151293.Np
    electric current           ampere           A
    electric charge            coulomb          C     s.A
    electric potential, EMF    volt             V     W/A
    capacitance                farad            F     C/V
    electric resistance        ohm              Ohm   V/A
    electric conductance       siemens          S     A/V
    magnetic flux              weber            Wb    V.s
    magnetic flux density      tesla            T     Wb/m^2
    inductance                 henry            H     Wb/A
    radionuclide activity      becquerel        Bq    s^-1
    absorbed dose energy       gray             Gy    m^2.s^-2
    dose equivalent            sievert          Sv    m^2.s^-2
```

* The formulas are:

- $r/rad = 8 * atan(1)$
- $o/r = 1 / 360$
- $db/Np = ln(10) / 20$

si:conversion-factor *to-unit from-unit*   [Function]

> If the strings *from-unit* and *to-unit* express valid unit expressions for quantities of the same unit-dimensions, then the value returned by `si:conversion-factor` will be such that multiplying a numerical value expressed in *from-unit*s by the returned conversion factor yields the numerical value expressed in *to-unit*s.
>
> Otherwise, `si:conversion-factor` returns:
>
> -3   if neither *from-unit* nor *to-unit* is a syntactically valid unit.
>
> -2   if *from-unit* is not a syntactically valid unit.

-1   if *to-unit* is not a syntactically valid unit.

0   if linear conversion (by a factor) is not possible.

```
(si:conversion-factor "km/s" "m/s" ) ⇒ 0.001
(si:conversion-factor "N"    "m/s" ) ⇒ 0
(si:conversion-factor "moC"  "oC"  ) ⇒ 1000
(si:conversion-factor "mK"   "oC"  ) ⇒ 0
(si:conversion-factor "rad"  "o"   ) ⇒ 0.0174533
(si:conversion-factor "K"    "o"   ) ⇒ 0
(si:conversion-factor "K"    "K"   ) ⇒ 1
(si:conversion-factor "oK"   "oK"  ) ⇒ -3
(si:conversion-factor ""     "s/s" ) ⇒ 1
(si:conversion-factor "km/h" "mph" ) ⇒ -2
```

7.4 Standards Support

7.4.1 RnRS

The `r2rs`, `r3rs`, `r4rs`, and `r5rs` features attempt to provide procedures and macros to bring a Scheme implementation to the desired version of Scheme.

`r2rs`   [Feature]
   Requires features implementing procedures and optional procedures specified by *Revised^2 Report on the Algorithmic Language Scheme*; namely `rev3-procedures` and `rev2-procedures`.

`r3rs`   [Feature]
   Requires features implementing procedures and optional procedures specified by *Revised^3 Report on the Algorithmic Language Scheme*; namely `rev3-procedures`.

   Note: SLIB already mandates the `r3rs` procedures which can be portably implemented in `r4rs` implementations.

`r4rs`   [Feature]
   Requires features implementing procedures and optional procedures specified by *Revised^4 Report on the Algorithmic Language Scheme*; namely `rev4-optional-procedures`.

`r5rs`   [Feature]
   Requires features implementing procedures and optional procedures specified by *Revised^5 Report on the Algorithmic Language Scheme*; namely `values`, `macro`, and `eval`.

7.4.2 With-File

```
(require 'with-file)
```

`with-input-from-file` *file thunk*   [Function]
`with-output-to-file` *file thunk*   [Function]
   Description found in R4RS.

7.4.3 Transcripts

```
(require 'transcript)
```

transcript-on *filename*   [Function]
transcript-off *filename*   [Function]
    Redefines `read-char`, `read`, `write-char`, `write`, `display`, and `newline`.

7.4.4 Rev2 Procedures

```
(require 'rev2-procedures)
```

    The procedures below were specified in the *Revised^2 Report on Scheme*. **N.B.**: The symbols `1+` and `-1+` are not *R4RS* syntax. Scheme->C, for instance, chokes on this module.

substring-move-left! *string1 start1 end1 string2 start2*   [Procedure]
substring-move-right! *string1 start1 end1 string2 start2*   [Procedure]
    string1 and *string2* must be a strings, and *start1*, *start2* and *end1* must be exact integers satisfying

        0 <= *start1* <= *end1* <= (string-length *string1*)
        0 <= *start2* <= *end1* - *start1* + *start2* <= (string-length *string2*)

    `substring-move-left!` and `substring-move-right!` store characters of *string1* beginning with index *start1* (inclusive) and ending with index *end1* (exclusive) into *string2* beginning with index *start2* (inclusive).

    `substring-move-left!` stores characters in time order of increasing indices. `substring-move-right!` stores characters in time order of increasing indeces.

substring-fill! *string start end char*   [Procedure]
    Fills the elements *start–end* of *string* with the character *char*.

string-null? *str*   [Function]
    ≡ (= 0 (string-length *str*))

append! *pair1 ...*   [Procedure]
    Destructively appends its arguments. Equivalent to `nconc`.

1+ *n*   [Function]
    Adds 1 to *n*.

-1+ *n*   [Function]
    Subtracts 1 from *n*.

<?   [Function]
<=?   [Function]
=?   [Function]
>?   [Function]
>=?   [Function]
    These are equivalent to the procedures of the same name but without the trailing '?'.

7.4.5 Rev4 Optional Procedures

`(require 'rev4-optional-procedures)`

For the specification of these optional procedures, See Section "Standard procedures" in *Revised(4) Scheme*.

`list-tail` *l p*   [Function]

`string-copy`   [Function]

`string-fill!` *s obj*   [Procedure]

`vector-fill!` *s obj*   [Procedure]

7.4.6 Multi-argument / and -

`(require 'multiarg/and-)`

For the specification of these optional forms, See Section "Numerical operations" in *Revised(4) Scheme*.

`/` *dividend divisor1 . . .*   [Function]

`-` *minuend subtrahend1 . . .*   [Function]

7.4.7 Multi-argument Apply

`(require 'multiarg-apply)`

For the specification of this optional form, See Section "Control features" in *Revised(4) Scheme*.

`apply` *proc arg1 . . .*   [Function]

7.4.8 Rationalize

`(require 'rationalize)`

`rationalize` *x e*   [Function]
> Computes the correct result for exact arguments (provided the implementation supports exact rational numbers of unlimited precision); and produces a reasonable answer for inexact arguments when inexact arithmetic is implemented using floating-point.

`Rationalize` has limited use in implementations lacking exact (non-integer) rational numbers. The following procedures return a list of the numerator and denominator.

`find-ratio` *x e*   [Function]
> `find-ratio` returns the list of the *simplest* numerator and denominator whose quotient differs from *x* by no more than *e*.
>
> ```
> (find-ratio 3/97 .0001)              ⇒ (3 97)
> (find-ratio 3/97 .001)               ⇒ (1 32)
> ```

```
find-ratio-between x y                                                    [Function]
```
> `find-ratio-between` returns the list of the *simplest* numerator and denominator between *x* and *y*.
>
> ```
> (find-ratio-between 2/7 3/5)        ⇒ (1 2)
> (find-ratio-between -3/5 -2/7)      ⇒ (-1 2)
> ```

7.4.9 Promises

```
(require 'promise)
```

```
make-promise proc                                                        [Function]
```

```
force promise                                                            [Function]
```
> `(require 'delay)` provides `force` and `delay`:

```
delay obj                                                                [Macro]
```
> Change occurrences of `(delay expression)` to
>
> ```
> (make-promise (lambda () expression))
> ```
>
> (see Section "Control features" in *Revised(4) Scheme*).

7.4.10 Dynamic-Wind

```
(require 'dynamic-wind)
```

This facility is a generalization of Common LISP `unwind-protect`, designed to take into account the fact that continuations produced by `call-with-current-continuation` may be reentered.

```
dynamic-wind thunk1 thunk2 thunk3                                         [Procedure]
```
> The arguments *thunk1*, *thunk2*, and *thunk3* must all be procedures of no arguments (thunks).
>
> `dynamic-wind` calls *thunk1*, *thunk2*, and then *thunk3*. The value returned by *thunk2* is returned as the result of `dynamic-wind`. *thunk3* is also called just before control leaves the dynamic context of *thunk2* by calling a continuation created outside that context. Furthermore, *thunk1* is called before reentering the dynamic context of *thunk2* by calling a continuation created inside that context. (Control is inside the context of *thunk2* if *thunk2* is on the current return stack).
>
> **Warning:** There is no provision for dealing with errors or interrupts. If an error or interrupt occurs while using `dynamic-wind`, the dynamic environment will be that in effect at the time of the error or interrupt.

7.4.11 Eval

```
(require 'eval)
```

```
eval expression environment-specifier                                    [Function]
```
> Evaluates *expression* in the specified environment and returns its value. *Expression* must be a valid Scheme expression represented as data, and *environment-specifier*

must be a value returned by one of the three procedures described below. Implementations may extend `eval` to allow non-expression programs (definitions) as the first argument and to allow other values as environments, with the restriction that `eval` is not allowed to create new bindings in the environments associated with `null-environment` or `scheme-report-environment`.

```
(eval '(* 7 3) (scheme-report-environment 5))
```
⇒ 21

```
(let ((f (eval '(lambda (f x) (f x x))
               (null-environment))))
  (f + 10))
```
⇒ 20

scheme-report-environment *version*   [Function]
null-environment *version*   [Function]
null-environment   [Function]

Version must be an exact non-negative integer n corresponding to a version of one of the Revised^n Reports on Scheme. `Scheme-report-environment` returns a specifier for an environment that contains the set of bindings specified in the corresponding report that the implementation supports. `Null-environment` returns a specifier for an environment that contains only the (syntactic) bindings for all the syntactic keywords defined in the given version of the report.

Not all versions may be available in all implementations at all times. However, an implementation that conforms to version n of the Revised^n Reports on Scheme must accept version n. An error is signalled if the specified version is not available.

The effect of assigning (through the use of `eval`) a variable bound in a `scheme-report-environment` (for example `car`) is unspecified. Thus the environments specified by `scheme-report-environment` may be immutable.

interaction-environment   [Function]

This optional procedure returns a specifier for the environment that contains implementation-defined bindings, typically a superset of those listed in the report. The intent is that this procedure will return the environment in which the implementation would evaluate expressions dynamically typed by the user.

Here are some more `eval` examples:

```
(require 'eval)
⇒ #<unspecified>
(define car 'volvo)
⇒ #<unspecified>
car
⇒ volvo
(eval 'car (interaction-environment))
⇒ volvo
(eval 'car (scheme-report-environment 5))
⇒ #<primitive-procedure car>
```

```
(eval '(eval 'car (interaction-environment))
      (scheme-report-environment 5))
⇒ volvo
(eval '(eval '(set! car 'buick) (interaction-environment))
      (scheme-report-environment 5))
⇒ #<unspecified>
car
⇒ buick
(eval 'car (scheme-report-environment 5))
⇒ #<primitive-procedure car>
(eval '(eval 'car (interaction-environment))
      (scheme-report-environment 5))
⇒ buick
```

7.4.12 Values

`(require 'values)`

values *obj* ...   [Function]
>   `values` takes any number of arguments, and passes (returns) them to its continuation.

call-with-values *thunk proc*   [Function]
>   *thunk* must be a procedure of no arguments, and *proc* must be a procedure.
>   `call-with-values` calls *thunk* with a continuation that, when passed some values,
>   calls *proc* with those values as arguments.
>
>   Except for continuations created by the `call-with-values` procedure, all continua-
>   tions take exactly one value, as now; the effect of passing no value or more than one
>   value to continuations that were not created by the `call-with-values` procedure is
>   unspecified.

7.4.13 SRFI

`(require 'srfi)`

Implements *Scheme Request For Implementation* (SRFI) as described at `http://srfi.schemers.org/`

cond-expand <*clause1*> <*clause2*> ...   [Macro]
>   *Syntax:* Each <clause> should be of the form
>
>   (<feature> <expression1> ...)
>
>   where <feature> is a boolean expression composed of symbols and 'and', 'or', and
>   'not' of boolean expressions. The last <clause> may be an "else clause," which has
>   the form
>
>   (else <expression1> <expression2> ...).
>
>   The first clause whose feature expression is satisfied is expanded. If no feature ex-
>   pression is satisfied and there is no else clause, an error is signaled.
>
>   SLIB `cond-expand` is an extension of SRFI-0, `http://srfi.schemers.org/srfi-0/srfi-0.html`.

- SRFI-2 Section 3.12 [Guarded LET* special form], page 34,
- SRFI-8 Section 3.11 [Binding to multiple values], page 34,
- SRFI-9 Section 3.8 [Define-Record-Type], page 33,
- SRFI-11 Section 3.11 [Binding to multiple values], page 34,
- SRFI-23 (`define error slib:error`)
- SRFI-28 Section 4.2 [Format], page 46,
- SRFI-39 Section 3.10 [Parameter Objects], page 34,
- SRFI-47 Section 7.1.1 [Arrays], page 196,
- SRFI-59 Section 2.1 [Vicinity], page 11,
- SRFI-60 Section 5.1 [Bit-Twiddling], page 103,
- SRFI-61 Section 3.13 [Guarded COND Clause], page 34,
- SRFI-63 Section 7.1.1 [Arrays], page 196,
- SRFI-94 Section 5.3 [Irrational Integer Functions], page 108, and Section 5.4 [Irrational Real Functions], page 109,
- SRFI-95 Section 7.2.4 [Sorting], page 233,
- SRFI-96 Chapter 2 [Universal SLIB Procedures], page 11,

7.4.13.1 SRFI-1

(require 'srfi-1)

Implements the *SRFI-1 list-processing library* as described at `http://srfi.schemers.org/srfi-1/srfi-1.html`

Constructors

xcons *d a*   [Function]
> (define (xcons d a) (cons a d)).

list-tabulate *len proc*   [Function]
> Returns a list of length *len*. Element *i* is (`proc i`) for 0 <= *i* < *len*.

cons* *obj1 obj2*   [Function]

list-copy *flist*   [Function]

iota *count start step*   [Function]
iota *count start*   [Function]
iota *count*   [Function]
> Returns a list of *count* numbers: (*start*, *start+step*, ..., *start+(count-1)*step*).

circular-list *obj1 obj2 ...*   [Function]
> Returns a circular list of *obj1*, *obj2*, ...

Predicates

proper-list? *obj*   [Function]

circular-list? *x*   [Function]

dotted-list? *obj*   [Function]

null-list? *obj*   [Function]

not-pair? *obj*   [Function]

list= *=pred list* . . .   [Function]

Selectors

first *pair*   [Function]

second *pair*   [Function]

third *pair*   [Function]

fourth *pair*   [Function]

fifth *pair*   [Function]
sixth *pair*   [Function]
seventh *pair*   [Function]
eighth *pair*   [Function]
ninth *pair*   [Function]
tenth *pair*   [Function]

car+cdr *pair*   [Function]

drop *lst k*   [Function]
take *lst k*   [Function]
take! *lst k*   [Function]

take-right *lst k*   [Function]

drop-right *lst k*   [Function]

drop-right! *lst k*   [Procedure]

split-at *lst k*   [Function]
split-at! *lst k*   [Function]

last *lst k* . . .   [Function]

Miscellaneous

length+ *clist*   [Function]

concatenate *lists*   [Function]
concatenate! *lists*   [Function]

reverse! *lst*   [Procedure]

append-reverse *rev-head tail*   [Function]
append-reverse! *rev-head tail*   [Function]

zip *list1 list2* . . .   [Function]

unzip1 *lst*   [Function]
unzip2 *lst*   [Function]
unzip3 *lst*   [Function]
unzip4 *lst*   [Function]
unzip5 *lst*   [Function]

count *pred list1 list2* . . .   [Function]

Fold and Unfold

`fold` *kons knil clist1 clist2 . . .*   [Function]

`fold-right` *kons knil clist1 clist2 . . .*   [Function]

`pair-fold` *kons knil clist1 clist2 . . .*   [Function]

`pair-fold-right` *kons knil clist1 clist2 . . .*   [Function]

`reduce` *arg . . .*   [Function]

`map!` *f clist1 clist2 . . .*   [Procedure]

`pair-for-each` *f clist1 clist2 . . .*   [Function]

Filtering and Partitioning

`filter` *pred list*   [Function]

`filter!` *pred list*   [Procedure]

`partition` *pred list*   [Function]

`remove` *pred list*   [Function]

`partition!` *pred list*   [Procedure]

`remove!` *pred list*   [Procedure]

Searching

`find` *pred clist*   [Function]

`find-tail` *pred clist*   [Function]

`span` *pred list*   [Function]

`span!` *pred list*   [Procedure]

`break` *pred list*   [Function]

`break!` *pred list*   [Procedure]

`any` *pred clist1 clist2 . . .*   [Function]

`list-index` *pred clist1 clist2 . . .*   [Function]

`member` *obj list =*   [Function]
`member` *obj list*   [Function]

Deleting

`delete-duplicates` *x list =*   [Function]
`delete-duplicates` *x list*   [Function]

`delete-duplicates!` *x list =*   [Procedure]
`delete-duplicates!` *x list*   [Procedure]

Association lists

`assoc` *obj alist pred*	[Function]
`assoc` *obj alist*	[Function]

`alist-cons` *key datum alist*	[Function]

`alist-copy` *alist*	[Function]

`alist-delete` *key alist* =	[Function]
`alist-delete` *key alist*	[Function]

`alist-delete!` *key alist* =	[Procedure]
`alist-delete!` *key alist*	[Procedure]

Set operations

`lset<=` = *list1* . . .   [Function]

> Determine if a transitive subset relation exists between the lists *list1* . . ., using = to determine equality of list members.

`lset=` = *list1 list2* . . .   [Function]

`lset-adjoin` *list elt1* . . .   [Function]

`lset-union` = *list1* . . .   [Function]

`lset-intersection` = *list1 list2* . . .   [Function]

`lset-difference` = *list1 list2* . . .   [Function]

`lset-xor` = *list1* . . .   [Function]

`lset-diff+intersection` = *list1 list2* . . .   [Function]

These are linear-update variants. They are allowed, but not required, to use the cons cells in their first list parameter to construct their answer. `lset-union!` is permitted to recycle cons cells from any of its list arguments.

`lset-intersection!` = *list1 list2* . . .   [Procedure]

`lset-difference!` = *list1 list2* . . .   [Procedure]

`lset-union!` = *list1* . . .   [Procedure]

`lset-xor!` = *list1* . . .   [Procedure]

`lset-diff+intersection!` = *list1 list2* . . .   [Procedure]

7.5 Session Support

If (`provided?` `'abort`):

`abort`   [Function]

> Resumes the top level Read-Eval-Print loop. If provided, `abort` is used by the `break` and `debug` packages.

7.5.1 Repl

`(require 'repl)`

Here is a read-eval-print-loop which, given an eval, evaluates forms.

`repl:top-level` *repl:eval*   [Procedure]

reads, `repl:eval`s and `write`s expressions from `(current-input-port)` to `(current-output-port)` until an end-of-file is encountered. `load`, `slib:eval`, `slib:error`, and `repl:quit` dynamically bound during `repl:top-level`.

`repl:quit`   [Procedure]

Exits from the invocation of `repl:top-level`.

The `repl:` procedures establish, as much as is possible to do portably, a top level environment supporting macros. `repl:top-level` uses `dynamic-wind` to catch error conditions and interrupts. If your implementation supports this you are all set.

Otherwise, if there is some way your implementation can catch error conditions and interrupts, then have them call `slib:error`. It will display its arguments and reenter `repl:top-level`. `slib:error` dynamically bound by `repl:top-level`.

To have your top level loop always use macros, add any interrupt catching lines and the following lines to your Scheme init file:

```
(require 'macro)
(require 'repl)
(repl:top-level macro:eval)
```

7.5.2 Quick Print

`(require 'qp)`

When displaying error messages and warnings, it is paramount that the output generated for circular lists and large data structures be limited. This section supplies a procedure to do this. It could be much improved.

Notice that the neccessity for truncating output eliminates Common-Lisp's Section 4.2 [Format], page 46, from consideration; even when variables `*print-level*` and `*print-level*` are set, huge strings and bit-vectors are *not* limited.

`qp` *arg1 ...*   [Procedure]
`qpn` *arg1 ...*   [Procedure]
`qpr` *arg1 ...*   [Procedure]

`qp` writes its arguments, separated by spaces, to `(current-output-port)`. `qp` compresses printing by substituting '`...`' for substructure it does not have sufficient room to print. `qpn` is like `qp` but outputs a newline before returning. `qpr` is like `qpn` except that it returns its last argument.

`*qp-width*`   [Variable]

qp-width is the largest number of characters that `qp` should use. If *qp-width* is #f, then all items will be `writen`. If *qp-width* is 0, then all items except procedures will be `writen`; procedures will be indicated by '`#[proc]`'.

7.5.3 Debug

`(require 'debug)`

Requiring `debug` automatically requires `trace` and `break`.

An application with its own datatypes may want to substitute its own printer for `qp`. This example shows how to do this:

```
(define qpn (lambda args) ...)
(provide 'qp)
(require 'debug)
```

trace-all *file* . . .   [Procedure]
> Traces (see Section 7.5.5 [Trace], page 260) all procedures `define`d at top-level in `file` . . . .

track-all *file* . . .   [Procedure]
> Tracks (see Section 7.5.5 [Trace], page 260) all procedures `define`d at top-level in `file` . . . .

stack-all *file* . . .   [Procedure]
> Stacks (see Section 7.5.5 [Trace], page 260) all procedures `define`d at top-level in `file` . . . .

break-all *file* . . .   [Procedure]
> Breakpoints (see Section 7.5.4 [Breakpoints], page 259) all procedures `define`d at top-level in `file` . . . .

7.5.4 Breakpoints

`(require 'break)`

init-debug   [Function]
> If your Scheme implementation does not support `break` or `abort`, a message will appear when you `(require 'break)` or `(require 'debug)` telling you to type `(init-debug)`. This is in order to establish a top-level continuation. Typing `(init-debug)` at top level sets up a continuation for `break`.

breakpoint *arg1* . . .   [Function]
> Returns from the top level continuation and pushes the continuation from which it was called on a continuation stack.

continue   [Function]
> Pops the topmost continuation off of the continuation stack and returns an unspecified value to it.

continue *arg1* . . .   [Function]
> Pops the topmost continuation off of the continuation stack and returns *arg1* . . . to it.

break *proc1* . . .   [Macro]
> Redefines the top-level named procedures given as arguments so that `breakpoint` is called before calling *proc1* . . . .

break   [Macro]
>   With no arguments, makes sure that all the currently broken identifiers are broken
>   (even if those identifiers have been redefined) and returns a list of the broken identi-
>   fiers.

unbreak *proc1 ...*   [Macro]
>   Turns breakpoints off for its arguments.

unbreak   [Macro]
>   With no arguments, unbreaks all currently broken identifiers and returns a list of
>   these formerly broken identifiers.

These are *procedures* for breaking. If defmacros are not natively supported by your
implementation, these might be more convenient to use.

breakf *proc*   [Function]
breakf *proc name*   [Function]
>   To break, type
>
>   (set! symbol (breakf symbol))
>
>   or
>
>   (set! symbol (breakf symbol 'symbol))
>
>   or
>
>   (define symbol (breakf function))
>
>   or
>
>   (define symbol (breakf function 'symbol))

unbreakf *proc*   [Function]
>   To unbreak, type
>
>   (set! symbol (unbreakf symbol))

7.5.5 Tracing

`(require 'trace)`

This feature provides three ways to monitor procedure invocations:

stack   Pushes the procedure-name when the procedure is called; pops when it returns.

track   Pushes the procedure-name and arguments when the procedure is called; pops
            when it returns.

trace   Pushes the procedure-name and prints 'CALL *procedure-name arg1* ...' when
            the procdure is called; pops and prints 'RETN *procedure-name value*' when the
            procedure returns.

debug:max-count   [Variable]
>   If a traced procedure calls itself or untraced procedures which call it, stack, track,
>   and trace will limit the number of stack pushes to *debug:max-count*.

`print-call-stack`   [Function]
`print-call-stack` *port*   [Function]
>   Prints the call-stack to *port* or the current-error-port.

`trace` *proc1* . . .   [Macro]
>   Traces the top-level named procedures given as arguments.

`trace`   [Macro]
>   With no arguments, makes sure that all the currently traced identifiers are traced
>   (even if those identifiers have been redefined) and returns a list of the traced identifiers.

`track` *proc1* . . .   [Macro]
>   Traces the top-level named procedures given as arguments.

`track`   [Macro]
>   With no arguments, makes sure that all the currently tracked identifiers are tracked
>   (even if those identifiers have been redefined) and returns a list of the tracked iden-
>   tifiers.

`stack` *proc1* . . .   [Macro]
>   Traces the top-level named procedures given as arguments.

`stack`   [Macro]
>   With no arguments, makes sure that all the currently stacked identifiers are stacked
>   (even if those identifiers have been redefined) and returns a list of the stacked iden-
>   tifiers.

`untrace` *proc1* . . .   [Macro]
>   Turns tracing, tracking, and off for its arguments.

`untrace`   [Macro]
>   With no arguments, untraces all currently traced identifiers and returns a list of these
>   formerly traced identifiers.

`untrack` *proc1* . . .   [Macro]
>   Turns tracing, tracking, and off for its arguments.

`untrack`   [Macro]
>   With no arguments, untracks all currently tracked identifiers and returns a list of
>   these formerly tracked identifiers.

`unstack` *proc1* . . .   [Macro]
>   Turns tracing, stacking, and off for its arguments.

`unstack`   [Macro]
>   With no arguments, unstacks all currently stacked identifiers and returns a list of
>   these formerly stacked identifiers.

These are *procedures* for tracing. If defmacros are not natively supported by your
implementation, these might be more convenient to use.

`tracef` *proc*	[Function]
`tracef` *proc name*	[Function]
`trackf` *proc*	[Function]
`trackf` *proc name*	[Function]
`stackf` *proc*	[Function]
`stackf` *proc name*	[Function]

> To trace, type
>
> ```
> (set! symbol (tracef symbol))
> ```
>
> or
>
> ```
> (set! symbol (tracef symbol 'symbol))
> ```
>
> or
>
> ```
> (define symbol (tracef function))
> ```
>
> or
>
> ```
> (define symbol (tracef function 'symbol))
> ```

`untracef` *proc*   [Function]

> Removes tracing, tracking, or stacking for *proc*. To untrace, type
>
> ```
> (set! symbol (untracef symbol))
> ```

7.6 System Interface

If (`provided?` '`getenv`):

`getenv` *name*   [Function]

> Looks up *name*, a string, in the program environment. If *name* is found a string of its value is returned. Otherwise, `#f` is returned.

If (`provided?` '`system`):

`system` *command-string*   [Function]

> Executes the *command-string* on the computer and returns the integer status code. This behaves the same as the POSIX `system` call.

If (`provided?` '`program-arguments`):

`program-arguments`   [Function]

> Returns a list of strings, the first of which is the program name followed by the command-line arguments.

7.6.1 Directories

(`require` '`directory`)

`current-directory`   [Function]

> `current-directory` returns a string containing the absolute file name representing the current working directory. If this string cannot be obtained, #f is returned.
>
> If `current-directory` cannot be supported by the platform, then #f is returned.

make-directory *name*   [Function]
>   Creates a sub-directory *name* of the current-directory. If successful, `make-directory`
>   returns #t; otherwise #f.

directory-for-each *proc directory*   [Function]
>   *proc* must be a procedure taking one argument. 'Directory-For-Each' applies *proc*
>   to the (string) name of each file in *directory*. The dynamic order in which *proc* is
>   applied to the filenames is unspecified. The value returned by 'directory-for-each'
>   is unspecified.

directory-for-each *proc directory pred*   [Function]
>   Applies *proc* only to those filenames for which the procedure *pred* returns a non-false
>   value.

directory-for-each *proc directory match*   [Function]
>   Applies *proc* only to those filenames for which (`filename:match?? match`) would
>   return a non-false value (see Section "Filenames" in *SLIB*).

```
(require 'directory)
(directory-for-each print "." "[A-Z]*.scm")
⊣
"Bev2slib.scm"
"Template.scm"
```

directory*-for-each *proc path-glob*   [Function]
>   *path-glob* is a pathname whose last component is a (wildcard) pattern (see
>   Section "Filenames" in *SLIB*). *proc* must be a procedure taking one argument.
>   'directory*-for-each' applies *proc* to the (string) name of each file in the current
>   directory. The dynamic order in which *proc* is applied to the filenames is unspecified.
>   The value returned by 'directory*-for-each' is unspecified.

7.6.2 Transactions

If `system` is provided by the Scheme implementation, the *transact* package provides functions for file-locking and file-replacement transactions.

```
(require 'transact)
```

File Locking

Unix file-locking is focussed on write permissions for segments of a existing file. While this might be employed for (binary) database access, it is not used for everyday contention (between users) for text files.

Microsoft has several file-locking protocols. Their model denies write access to a file if any reader has it open. This is too restrictive. Write access is denied even when the reader has reached end-of-file. And tracking read access (which is much more common than write access) causes havoc when remote hosts crash or disconnect.

It is bizarre that the concept of multi-user contention for modifying files has not been adequately addressed by either of the large operating system development efforts. There is further irony that both camps support contention detection and resolution only through weak conventions of some their document editing programs.

The *file-lock* procedures implement a transaction method for file replacement compatible with the methods used by the GNU *emacs* text editor on Unix systems and the Microsoft *Word* editor.

Both protocols employ what I term a *certificate* containing the user, hostname, time, and (on Unix) process-id. Intent to replace *file* is indicated by adding to *file*'s directory a certificate object whose name is derived from *file*.

The Microsoft Word certificate is contained in a 162 byte file named for the visited *file* with a '~$' prefix. Emacs/Unix creates a symbolic link to a certificate named for the visited *file* prefixed with '.#'. Because Unix systems can import Microsoft file systems, these routines maintain and check both Emacs and Word certificates.

file-lock-owner *path*   [Function]
>   Returns the string '`user@hostname`' associated with the lock owner of file *path* if locked; and #f otherwise.

file-lock! *path email*   [Procedure]
file-lock! *path*   [Procedure]
>   *path* must be a string naming the file to be locked. If supplied, *email* must be a string formatted as '`user@hostname`'. If absent, *email* defaults to the value returned by `user-email-address`.
>
>   If *path* is already locked, then `file-lock!` returns '`#f`'. If *path* is unlocked, then `file-lock!` returns the certificate string associated with the new lock for file *path*.

file-unlock! *path certificate*   [Procedure]
>   *path* must be a string naming the file to be unlocked. *certificate* must be the string returned by `file-lock!` for *path*.
>
>   If *path* is locked with *certificate*, then `file-unlock!` removes the locks and returns '`#t`'. Otherwise, `file-unlock!` leaves the file system unaltered and returns '`#f`'.

describe-file-lock *path prefix*   [Function]
describe-file-lock *path*   [Function]
>   *path* must be a string naming a file. Optional argument *prefix* is a string printed before each line of the message. `describe-file-lock` prints to (`current-error-port`) that *path* is locked for writing and lists its lock-files.

```
(describe-file-lock "my.txt" ">> ")
⊣
>> "my.txt" is locked for writing by 'luser@no.com.4829:1200536423'
>> (lock files are "~$my.txt" and ".#my.txt")
```

File Transactions

emacs:backup-name *path backup-style*   [Function]
>   *path* must be a string. *backup-style* must be a symbol. Depending on *backup-style*, `emacs:backup-name` returns:
>
>   none   #f
>
>   simple   the string "*path*~"

numbered   the string "*path.˜n˜*", where *n* is one greater than the highest number appearing in a filename matching "*path.˜*˜*". *n* defaults to 1 when no filename matches.

existing   the string "*path.˜n˜*" if a numbered backup already exists in this directory; otherwise. "*path˜*"

orig   the string "*path*.orig"

bak   the string "*path*.bak"

transact-file-replacement *proc path backup-style certificate*   [Function]
transact-file-replacement *proc path backup-style*   [Function]
transact-file-replacement *proc path*   [Function]

path must be a string naming an existing file. *backup-style* is one of the symbols none, simple, numbered, existing, orig, bak or #f; with meanings described above; or a string naming the location of a backup file. *backup-style* defaults to #f. If supplied, *certificate* is the certificate with which *path* is locked.

proc must be a procedure taking two string arguments:

- *path*, the original filename (to be read); and
- a temporary file-name.

If *path* is locked by other than *certificate*, or if *certificate* is supplied and *path* is not locked, then **transact-file-replacement** returns #f. If *certificate* is not supplied, then, **transact-file-replacement** creates temporary (Emacs and Word) locks for *path* during the transaction. The lock status of *path* will be restored before **transact-file-replacement** returns.

transact-file-replacement calls *proc* with *path* (which should not be modified) and a temporary file path to be written. If *proc* returns any value other than #t, then the file named by *path* is not altered and **transact-file-replacement** returns #f. Otherwise, **emacs:backup-name** is called with *path* and *backup-style*. If it returns a string, then *path* is renamed to it.

Finally, the temporary file is renamed *path*. **transact-file-replacement** returns #t if *path* was successfully replaced; and #f otherwise.

Identification

user-email-address   [Function]

user-email-address returns a string of the form '**username@hostname**'. If this e-mail address cannot be obtained, #f is returned.

7.6.3 CVS

(require 'cvs)

cvs-files *directory/*   [Function]

Returns a list of the local pathnames (with prefix *directory/*) of all CVS controlled files in *directory/* and in *directory/*'s subdirectories.

cvs-directories *directory/*   [Function]
>   Returns a list of all of *directory/* and all *directory/*'s CVS controlled subdirectories.

cvs-root *path/*   [Function]
>   Returns the (string) contents of *path*/CVS/Root; or (`getenv "CVSROOT"`) if Root
>   doesn't exist.

cvs-repository *directory/*   [Function]
>   Returns the (string) contents of *directory*/CVS/Root appended with *direc-
>   tory*/CVS/Repository; or #f if *directory*/CVS/Repository doesn't exist.

cvs-set-root! *new-root directory/*   [Procedure]
>   Writes *new-root* to file CVS/Root of *directory/*.

cvs-set-roots! *new-root directory/*   [Procedure]
>   Writes *new-root* to file CVS/Root of *directory/* and all its CVS subdirectories.

cvs-vet *directory/*   [Function]
>   Signals an error if CVS/Repository or CVS/Root files in *directory/* or any subdirec-
>   tory do not match.

7.7 Extra-SLIB Packages

Several Scheme packages have been written using SLIB. There are several reasons why a
package might not be included in the SLIB distribution:

- Because it requires special hardware or software which is not universal.
- Because it is large and of limited interest to most Scheme users.
- Because it has copying terms different enough from the other SLIB packages that its
  inclusion would cause confusion.
- Because it is an application program, rather than a library module.
- Because I have been too busy to integrate it.

Once an optional package is installed (and an entry added to `*catalog*`), the `require`
mechanism allows it to be called up and used as easily as any other SLIB package. Some
optional packages (for which `*catalog*` already has entries) available from SLIB sites are:

SLIB-PSD is a portable debugger for Scheme (requires emacs editor).
>   `http://groups.csail.mit.edu/mac/ftpdir/scm/slib-psd1-3.tar.gz`
>   `ftp://ftp.cs.indiana.edu/pub/scheme-repository/utl/slib-psd1-3.`
>   `tar.gz`
>
>   With PSD, you can run a Scheme program in an Emacs buffer, set breakpoints,
>   single step evaluation and access and modify the program's variables. It works
>   by instrumenting the original source code, so it should run with any R4RS
>   compliant Scheme. It has been tested with SCM, Elk 1.5, and the sci interpreter
>   in the Scheme->C system, but should work with other Schemes with a minimal
>   amount of porting, if at all. Includes documentation and user's manual. Written
>   by Pertti Kellomäki, the Lisp Pointers article describing PSD (Lisp Pointers

VI(1):15-23, January-March 1993) is available at
`http://www.cs.tut.fi/staff/pk/scheme/psd/article/article.html`

SCHELOG
  is an embedding of Prolog in Scheme.
  `http://www.ccs.neu.edu/~dorai/schelog/schelog.html`

JFILTER is a Scheme program which converts text among the JIS, EUC, and Shift-JIS
  Japanese character sets.
  `http://www.math.u-toyama.ac.jp/~iwao/Scheme/Jfilter`

8 About SLIB

More people than I can name have contributed to SLIB. Thanks to all of you!

> SLIB 3b6, released February 2020.
> Aubrey Jaffer <agj@alum.mit.edu>

> Current information about SLIB can be found on SLIB's *WWW* home page:
> http://people.csail.mit.edu/jaffer/SLIB

> SLIB is part of the GNU project.

8.1 Installation

There are five parts to installation:

- Unpack the SLIB distribution.

- Install documentation and `slib` script.

- Configure the Scheme implementation(s) to locate the SLIB directory and implementation directories.

- Arrange for each Scheme implementation to load its SLIB initialization file.

- Build the SLIB catalog for each Scheme implementation.

8.1.1 Unpacking the SLIB Distribution

If the SLIB distribution is a GNU/Linux RPM, it will create the SLIB directory `/usr/share/slib`.

If the SLIB distribution is a ZIP file, unzip the distribution to create the SLIB directory. Locate this `slib` directory either in your home directory (if only you will use this SLIB installation); or put it in a location where libraries reside on your system. On unix systems this might be `/usr/share/slib`, `/usr/local/lib/slib`, or `/usr/lib/slib`. If you know where SLIB should go on other platforms, please inform agj@alum.mit.edu.

8.1.2 Install documentation and slib script

```
make infoz
make install
```

8.1.3 Configure Scheme Implementation to Locate SLIB

If the Scheme implementation supports `getenv`, then the value of the shell environment variable *SCHEME_LIBRARY_PATH* will be used for (`library-vicinity`) if it is defined. Currently, Bigloo, Chez, Elk, Gambit, Gauche, Guile, Jscheme, Larceny, MITScheme, MzScheme, RScheme, S7, STk, VSCM, and SCM support `getenv`. Scheme48 supports `getenv` but does not use it for determining `library-vicinity`. (That is done from the Makefile.)

The (`library-vicinity`) can also be set from the SLIB initialization file or by implementation-specific means.

8.1.4 Configure Scheme Implementation to Locate and Implementation Directory

Support for locating an implementation's auxiliary directory is uneven among implementations. Also, the person installing SLIB may not have write permission to some of these directories (necessary for writing slibcat). Therefore, those implementations supporting `getenv` (except SCM and Scheme48) provide a means for specifying the `implementation-vicinity` through environment variables. Define the indicated environment variable to the pathname (with trailing slash or backslash) of the desired directory. Do not use `slib/` as an implementation-vicinity!

Bigloo	BIGLOO_IMPLEMENTATION_PATH
Chez	CHEZ_IMPLEMENTATION_PATH
ELK	ELK_IMPLEMENTATION_PATH
Gambit	GAMBIT_IMPLEMENTATION_PATH
Guile	GUILE_IMPLEMENTATION_PATH
Jscheme	JSCHEME_IMPLEMENTATION_PATH
MIT-Scheme	MITSCHEME_IMPLEMENTATION_PATH
MzScheme	MZSCHEME_IMPLEMENTATION_PATH
RScheme	RSCHEME_IMPLEMENTATION_PATH
S7	S7_IMPLEMENTATION_PATH
STk	STK_IMPLEMENTATION_PATH
Vscm	VSCM_IMPLEMENTATION_PATH

8.1.5 Loading SLIB Initialization File

If you use the `slib` script to start your SLIB session, then this step is unnecessary.

Check the manifest in `README` to find a configuration file for your Scheme implementation. Initialization files for most IEEE P1178 compliant Scheme Implementations are included with this distribution.

You should check the definitions of `software-type`, `scheme-implementation-version`, `implementation-vicinity`, and `library-vicinity` in the initialization file. There are comments in the file for how to configure it.

Once this is done, modify the startup file for your Scheme implementation to `load` this initialization file.

8.1.6 Build New SLIB Catalog for the Implementation

When SLIB is first used from an implementation, a file named `slibcat` is written to the `implementation-vicinity` for that implementation. Because users may lack permission to write in `implementation-vicinity`, it is good practice to build the new catalog when installing SLIB.

To build (or rebuild) the catalog, start the Scheme implementation (with SLIB), then:

```
(require 'new-catalog)
```

The catalog also supports color-name dictionaries. With an SLIB-installed scheme implementation, type:

```
(require 'color-names)
(make-slib-color-name-db)
```

```
(require 'new-catalog)
(slib:exit)
```

8.1.7 Implementation-specific Instructions

Multiple implementations of Scheme can all use the same SLIB directory. Simply configure each implementation's initialization file as outlined above.

SCM   [Implementation]
>   The SCM implementation does not require any initialization file as SLIB support is already built into SCM. See the documentation with SCM for installation instructions.

Larceny   [Implementation]
>   Starting with version 0.96, Larceny contains its own SLIB initialization file, loaded by (require 'srfi-96). If SCHEME_LIBRARY_PATH is not set, then Larceny looks for an slib subdirectory of a directory in the list returned by (current-require-path)
>
>       ```
>       larceny -- -e "(require 'srfi-96)"
>       ```

Gauche-0.9   [Implementation]
>   Gauche also supports SLIB. It finds SLIB at installation time; (use slib) to enable.
>
>       ```
>       gosh -u slib
>       ```

ELK   [Implementation]
>       ```
>       elk -i -l ${SCHEME_LIBRARY_PATH}elk.init
>       ```

PLT Scheme   [Implementation]
DrScheme   [Implementation]
MzScheme   [Implementation]
>   The init.ss file in the _slibinit_ collection is an SLIB initialization file. To run SLIB in MzScheme:
>
>       ```
>       mzscheme -f ${SCHEME_LIBRARY_PATH}mzscheme.init
>       ```

MIT Scheme   [Implementation]
>       ```
>       scheme -load ${SCHEME_LIBRARY_PATH}mitscheme.init
>       ```

Gambit-C 3.0   [Implementation]
>       ```
>       gsi -:s ${SCHEME_LIBRARY_PATH}gambit.init -
>       ```

SISC   [Implementation]
>       ```
>       sisc -e "(load \"${SCHEME_LIBRARY_PATH}sisc.init\")" --
>       ```

Kawa   [Implementation]
>       ```
>       kawa -f ${SCHEME_LIBRARY_PATH}kawa.init --
>       ```

Guile   [Implementation]
>   Guile versions 1.6 and earlier link to an archaic SLIB version. In RedHat or Fedora installations:
>
>       ```
>       rm /usr/share/guile/slib
>       ```

```
ln -s ${SCHEME_LIBRARY_PATH} /usr/share/guile/slib
```

In Debian installations:

```
rm /usr/share/guile/1.6/slib
ln -s ${SCHEME_LIBRARY_PATH} /usr/share/guile/1.6/slib
```

${SCHEME_LIBRARY_PATH} is where SLIB gets installed.

Guile before version 1.8 with SLIB can then be started thus:

```
guile -l ${SCHEME_LIBRARY_PATH}guile.init
```

Guile version 1.8 and after with SLIB can then be started thus:

```
guile -l ${SCHEME_LIBRARY_PATH}guile.init \
      -l ${SCHEME_LIBRARY_PATH}guile.use
```

The Guile manual has a different way of installing SLIB:

`http://www.gnu.org/software/guile/manual/html_node/SLIB-installation.html`

Scheme48   [Implementation]

To make a Scheme48 image for an installation under `<prefix>`,

1. `cd` to the SLIB directory
2. type `make prefix=<prefix> slib48`.
3. To install the image, type `make prefix=<prefix> install48`. This will also create a shell script with the name `slib48` which will invoke the saved image.

VSCM   [Implementation]

From: Matthias Blume <blume @ cs.Princeton.EDU>
Date: Tue, 1 Mar 1994 11:42:31 -0500

Disclaimer: The code below is only a quick hack. If I find some time to spare I might get around to make some more things work.

You have to provide `vscm.init` as an explicit command line argument. Since this is not very nice I would recommend the following installation procedure:

1. run scheme
2. `(load "vscm.init")`
3. `(slib:dump "dumpfile")`
4. mv dumpfile place-where-vscm-standard-bootfile-resides. For example:

   `mv dumpfile /usr/local/vscm/lib/scheme-boot`

   In this case vscm should have been compiled with flag:

   -DDEFAULT_BOOTFILE='"/usr/local/vscm/lib/scheme-boot"'

   See Makefile (definition of DDP) for details.

S7   [Implementation]

S7 is not a standalone implementation, but runs as the extension language for the Snd sound editor. `${SCHEME_LIBRARY_PATH}s7.init` can be loaded from the Snd init file or on the Snd command line thus:

```
snd -l ${SCHEME_LIBRARY_PATH}s7.init
```

8.2 The SLIB script

SLIB comes with shell script for Unix platforms.

slib [scheme | scm | gsi | mzscheme | guile
      | scheme48 | scmlit | elk | sisc | kawa]

Starts an interactive Scheme-with-SLIB session.

The optional argument to the `slib` script is the Scheme implementation to run. Absent the argument, it searches for implementations in the above order.

8.3 Porting

If there is no initialization file for your Scheme implementation, you will have to create one. Your Scheme implementation must be largely compliant with

IEEE Std 1178-1990,
Revised^4 Report on the Algorithmic Language Scheme, or
Revised^5 Report on the Algorithmic Language Scheme

in order to support SLIB.[1]

`http://cvs.savannah.gnu.org/viewcvs/*checkout*/scm/scm/r4rstest.scm` is a file which checks compliance with much of R4RS.

`Template.scm` is an example configuration file. The comments inside will direct you on how to customize it to reflect your system. Give your new initialization file the implementation's name with `.init` appended. For instance, if you were porting `foo-scheme` then the initialization file might be called `foo.init`.

Your customized version should then be loaded as part of your scheme implementation's initialization. It will load `require.scm` from the library; this will allow the use of `provide`, `provided?`, and `require` along with the *vicinity* functions (these functions are documented in the sections Section 1.1 [Feature], page 1, and Section 1.2 [Require], page 2). The rest of the library will then be accessible in a system independent fashion.

Please mail new working configuration files to `agj@alum.mit.edu` so that they can be included in the SLIB distribution.

8.4 Compiled and Implementation-Specific Features

Often an implementation can implement an SLIB feature more efficiently than the R4RS-compliant source code in SLIB. Alternatively, implementations with compilers can compile SLIB source code into binary files which run faster than loading source code.

Additionally, the SLIB catalog can be augmented with extra-SLIB features which can be loaded by the implementation. The catalog format is described in See Section 1.3 [Library Catalogs], page 3.

---

[1] If you are porting a *Revised^3 Report on the Algorithmic Language Scheme* implementation, then you will need to finish writing `sc4sc3.scm` and `load` it from your initialization file.

These implementation-specific modifications are made when a new catalog is created (see Section 1.4 [Catalog Creation], page 3). If `mkimpcat.scm` in `implementation-invicinity` exists, it is loaded. That should produce the file `implcat` in `implementation-invicinity`, whose associations will override those of SLIB. `implcat` is copied into `slibcat` in `implementation-vicinity` as part of the catalog creation process; modifications to `implcat` after that will have no effect.

8.5 Coding Guidelines

All library packages are written in IEEE P1178 Scheme and assume that a configuration file and `require.scm` package have already been loaded. Other versions of Scheme can be supported in library packages as well by using, for example, `(provided? 'r3rs)` or `(require 'r3rs)` (see Section 1.2 [Require], page 2).

If a procedure defined in a module is called by other procedures in that module, then those procedures should instead call an alias defined in that module:

```
(define module-name:foo foo)
```

The module name and ':' should prefix that symbol for the internal name. Do not export internal aliases.

A procedure is exported from a module by putting Schmooz-style comments (see Section 4.15 [Schmooz], page 101) or ';@' at the beginning of the line immediately preceding the definition (`define`, `define-syntax`, or `defmacro`). Modules, exports and other relevant issues are discussed in Section 1.6 [Compiling Scheme], page 6.

Code submitted for inclusion in SLIB should not duplicate (more than one) routines already in SLIB files. Use `require` to force those library routines to be used by your package.

Documentation should be provided in Emacs Texinfo format if possible, but documentation must be provided.

Your package will be released sooner with SLIB if you send me a file which tests your code. Please run this test *before* you send me the code!

8.5.1 Modifications

Please document your changes. A line or two for `ChangeLog` is sufficient for simple fixes or extensions. Look at the format of `ChangeLog` to see what information is desired. Please send me `diff` files from the latest SLIB distribution (remember to send `diff`s of `slib.texi` and `ChangeLog`). This makes for less email traffic and makes it easier for me to integrate when more than one person is changing a file (this happens a lot with `slib.texi` and '`*.init`' files).

If someone else wrote a package you want to significantly modify, please try to contact the author, who may be working on a new version. This will insure against wasting effort on obsolete versions.

Please *do not* reformat the source code with your favorite beautifier, make 10 fixes, and send me the resulting source code. I do not have the time to fish through 10000 diffs to find your 10 real fixes.

8.6 Copyrights

This section has instructions for SLIB authors regarding copyrights.

Each package in SLIB must either be in the public domain, or come with a statement of terms permitting users to copy, redistribute and modify it. The comments at the beginning of `require.scm` and `macwork.scm` illustrate copyright and appropriate terms.

If your code or changes amount to less than about 10 lines, you do not need to add your copyright or send a disclaimer.

8.6.1 Putting code into the Public Domain

In order to put code in the public domain you should sign a copyright disclaimer and send it to the SLIB maintainer. Contact agj@alum.mit.edu for the address to mail the disclaimer to.

> I, <*my-name*>, hereby affirm that I have placed the software package <*name*> in the public domain.

> I affirm that I am the sole author and sole copyright holder for the software package, that I have the right to place this software package in the public domain, and that I will do nothing to undermine this status in the future.

> *signature and date*

This wording assumes that you are the sole author. If you are not the sole author, the wording needs to be different. If you don't want to be bothered with sending a letter every time you release or modify a module, make your letter say that it also applies to your future revisions of that module.

Make sure no employer has any claim to the copyright on the work you are submitting. If there is any doubt, create a copyright disclaimer and have your employer sign it. Mail the signed disclaimer to the SLIB maintainer. Contact agj@alum.mit.edu for the address to mail the disclaimer to. An example disclaimer follows.

8.6.2 Explicit copying terms

If you submit more than about 10 lines of code which you are not placing into the Public Domain (by sending me a disclaimer) you need to:

- Arrange that your name appears in a copyright line for the appropriate year. Multiple copyright lines are acceptable.
- With your copyright line, specify any terms you require to be different from those already in the file.
- Make sure no employer has any claim to the copyright on the work you are submitting. If there is any doubt, create a copyright disclaimer and have your employer sign it. Mail the signed disclaim to the SLIB maintainer. Contact agj@alum.mit.edu for the address to mail the disclaimer to.

8.6.3 Example: Company Copyright Disclaimer

This disclaimer should be signed by a vice president or general manager of the company. If you can't get at them, anyone else authorized to license out software produced there will do. Here is a sample wording:

<*employer*> Corporation hereby disclaims all copyright interest in the program <*program*> written by <*name*>.

<*employer*> Corporation affirms that it has no other intellectual property interest that would undermine this release, and will do nothing to undermine it in the future.

<*signature and date*>,
<*name*>, <*title*>, <*employer*> Corporation

8.7 About this manual

- Entries that are labeled as Functions are called for their return values. Entries that are labeled as Procedures are called primarily for their side effects.

- Examples in this text were produced using the `scm` Scheme implementation.

- At the beginning of each section, there is a line that looks like:

      (require 'feature)

  Include this line in your code prior to using the package.

8.7.1 GNU Free Documentation License

Version 1.3, 3 November 2008

Copyright © 2000, 2001, 2002, 2007, 2008 Free Software Foundation, Inc.
`http://fsf.org/`

Everyone is permitted to copy and distribute verbatim copies
of this license document, but changing it is not allowed.

0. PREAMBLE

   The purpose of this License is to make a manual, textbook, or other functional and useful document *free* in the sense of freedom: to assure everyone the effective freedom to copy and redistribute it, with or without modifying it, either commercially or non-commercially. Secondarily, this License preserves for the author and publisher a way to get credit for their work, while not being considered responsible for modifications made by others.

   This License is a kind of "copyleft", which means that derivative works of the document must themselves be free in the same sense. It complements the GNU General Public License, which is a copyleft license designed for free software.

   We have designed this License in order to use it for manuals for free software, because free software needs free documentation: a free program should come with manuals providing the same freedoms that the software does. But this License is not limited to software manuals; it can be used for any textual work, regardless of subject matter or whether it is published as a printed book. We recommend this License principally for works whose purpose is instruction or reference.

1. APPLICABILITY AND DEFINITIONS

   This License applies to any manual or other work, in any medium, that contains a notice placed by the copyright holder saying it can be distributed under the terms of this License. Such a notice grants a world-wide, royalty-free license, unlimited in

duration, to use that work under the conditions stated herein. The "Document", below, refers to any such manual or work. Any member of the public is a licensee, and is addressed as "you". You accept the license if you copy, modify or distribute the work in a way requiring permission under copyright law.

A "Modified Version" of the Document means any work containing the Document or a portion of it, either copied verbatim, or with modifications and/or translated into another language.

A "Secondary Section" is a named appendix or a front-matter section of the Document that deals exclusively with the relationship of the publishers or authors of the Document to the Document's overall subject (or to related matters) and contains nothing that could fall directly within that overall subject. (Thus, if the Document is in part a textbook of mathematics, a Secondary Section may not explain any mathematics.) The relationship could be a matter of historical connection with the subject or with related matters, or of legal, commercial, philosophical, ethical or political position regarding them.

The "Invariant Sections" are certain Secondary Sections whose titles are designated, as being those of Invariant Sections, in the notice that says that the Document is released under this License. If a section does not fit the above definition of Secondary then it is not allowed to be designated as Invariant. The Document may contain zero Invariant Sections. If the Document does not identify any Invariant Sections then there are none.

The "Cover Texts" are certain short passages of text that are listed, as Front-Cover Texts or Back-Cover Texts, in the notice that says that the Document is released under this License. A Front-Cover Text may be at most 5 words, and a Back-Cover Text may be at most 25 words.

A "Transparent" copy of the Document means a machine-readable copy, represented in a format whose specification is available to the general public, that is suitable for revising the document straightforwardly with generic text editors or (for images composed of pixels) generic paint programs or (for drawings) some widely available drawing editor, and that is suitable for input to text formatters or for automatic translation to a variety of formats suitable for input to text formatters. A copy made in an otherwise Transparent file format whose markup, or absence of markup, has been arranged to thwart or discourage subsequent modification by readers is not Transparent. An image format is not Transparent if used for any substantial amount of text. A copy that is not "Transparent" is called "Opaque".

Examples of suitable formats for Transparent copies include plain ASCII without markup, Texinfo input format, LaTeX input format, SGML or XML using a publicly available DTD, and standard-conforming simple HTML, PostScript or PDF designed for human modification. Examples of transparent image formats include PNG, XCF and JPG. Opaque formats include proprietary formats that can be read and edited only by proprietary word processors, SGML or XML for which the DTD and/or processing tools are not generally available, and the machine-generated HTML, PostScript or PDF produced by some word processors for output purposes only.

The "Title Page" means, for a printed book, the title page itself, plus such following pages as are needed to hold, legibly, the material this License requires to appear in the title page. For works in formats which do not have any title page as such, "Title Page"

means the text near the most prominent appearance of the work's title, preceding the beginning of the body of the text.

The "publisher" means any person or entity that distributes copies of the Document to the public.

A section "Entitled XYZ" means a named subunit of the Document whose title either is precisely XYZ or contains XYZ in parentheses following text that translates XYZ in another language. (Here XYZ stands for a specific section name mentioned below, such as "Acknowledgements", "Dedications", "Endorsements", or "History".) To "Preserve the Title" of such a section when you modify the Document means that it remains a section "Entitled XYZ" according to this definition.

The Document may include Warranty Disclaimers next to the notice which states that this License applies to the Document. These Warranty Disclaimers are considered to be included by reference in this License, but only as regards disclaiming warranties: any other implication that these Warranty Disclaimers may have is void and has no effect on the meaning of this License.

2. VERBATIM COPYING

You may copy and distribute the Document in any medium, either commercially or noncommercially, provided that this License, the copyright notices, and the license notice saying this License applies to the Document are reproduced in all copies, and that you add no other conditions whatsoever to those of this License. You may not use technical measures to obstruct or control the reading or further copying of the copies you make or distribute. However, you may accept compensation in exchange for copies. If you distribute a large enough number of copies you must also follow the conditions in section 3.

You may also lend copies, under the same conditions stated above, and you may publicly display copies.

3. COPYING IN QUANTITY

If you publish printed copies (or copies in media that commonly have printed covers) of the Document, numbering more than 100, and the Document's license notice requires Cover Texts, you must enclose the copies in covers that carry, clearly and legibly, all these Cover Texts: Front-Cover Texts on the front cover, and Back-Cover Texts on the back cover. Both covers must also clearly and legibly identify you as the publisher of these copies. The front cover must present the full title with all words of the title equally prominent and visible. You may add other material on the covers in addition. Copying with changes limited to the covers, as long as they preserve the title of the Document and satisfy these conditions, can be treated as verbatim copying in other respects.

If the required texts for either cover are too voluminous to fit legibly, you should put the first ones listed (as many as fit reasonably) on the actual cover, and continue the rest onto adjacent pages.

If you publish or distribute Opaque copies of the Document numbering more than 100, you must either include a machine-readable Transparent copy along with each Opaque copy, or state in or with each Opaque copy a computer-network location from which the general network-using public has access to download using public-standard network protocols a complete Transparent copy of the Document, free of added material. If

you use the latter option, you must take reasonably prudent steps, when you begin distribution of Opaque copies in quantity, to ensure that this Transparent copy will remain thus accessible at the stated location until at least one year after the last time you distribute an Opaque copy (directly or through your agents or retailers) of that edition to the public.

It is requested, but not required, that you contact the authors of the Document well before redistributing any large number of copies, to give them a chance to provide you with an updated version of the Document.

4. MODIFICATIONS

   You may copy and distribute a Modified Version of the Document under the conditions of sections 2 and 3 above, provided that you release the Modified Version under precisely this License, with the Modified Version filling the role of the Document, thus licensing distribution and modification of the Modified Version to whoever possesses a copy of it. In addition, you must do these things in the Modified Version:

   A. Use in the Title Page (and on the covers, if any) a title distinct from that of the Document, and from those of previous versions (which should, if there were any, be listed in the History section of the Document). You may use the same title as a previous version if the original publisher of that version gives permission.

   B. List on the Title Page, as authors, one or more persons or entities responsible for authorship of the modifications in the Modified Version, together with at least five of the principal authors of the Document (all of its principal authors, if it has fewer than five), unless they release you from this requirement.

   C. State on the Title page the name of the publisher of the Modified Version, as the publisher.

   D. Preserve all the copyright notices of the Document.

   E. Add an appropriate copyright notice for your modifications adjacent to the other copyright notices.

   F. Include, immediately after the copyright notices, a license notice giving the public permission to use the Modified Version under the terms of this License, in the form shown in the Addendum below.

   G. Preserve in that license notice the full lists of Invariant Sections and required Cover Texts given in the Document's license notice.

   H. Include an unaltered copy of this License.

   I. Preserve the section Entitled "History", Preserve its Title, and add to it an item stating at least the title, year, new authors, and publisher of the Modified Version as given on the Title Page. If there is no section Entitled "History" in the Document, create one stating the title, year, authors, and publisher of the Document as given on its Title Page, then add an item describing the Modified Version as stated in the previous sentence.

   J. Preserve the network location, if any, given in the Document for public access to a Transparent copy of the Document, and likewise the network locations given in the Document for previous versions it was based on. These may be placed in the "History" section. You may omit a network location for a work that was published at least four years before the Document itself, or if the original publisher of the version it refers to gives permission.

K. For any section Entitled "Acknowledgements" or "Dedications", Preserve the Title of the section, and preserve in the section all the substance and tone of each of the contributor acknowledgements and/or dedications given therein.

L. Preserve all the Invariant Sections of the Document, unaltered in their text and in their titles. Section numbers or the equivalent are not considered part of the section titles.

M. Delete any section Entitled "Endorsements". Such a section may not be included in the Modified Version.

N. Do not retitle any existing section to be Entitled "Endorsements" or to conflict in title with any Invariant Section.

O. Preserve any Warranty Disclaimers.

If the Modified Version includes new front-matter sections or appendices that qualify as Secondary Sections and contain no material copied from the Document, you may at your option designate some or all of these sections as invariant. To do this, add their titles to the list of Invariant Sections in the Modified Version's license notice. These titles must be distinct from any other section titles.

You may add a section Entitled "Endorsements", provided it contains nothing but endorsements of your Modified Version by various parties—for example, statements of peer review or that the text has been approved by an organization as the authoritative definition of a standard.

You may add a passage of up to five words as a Front-Cover Text, and a passage of up to 25 words as a Back-Cover Text, to the end of the list of Cover Texts in the Modified Version. Only one passage of Front-Cover Text and one of Back-Cover Text may be added by (or through arrangements made by) any one entity. If the Document already includes a cover text for the same cover, previously added by you or by arrangement made by the same entity you are acting on behalf of, you may not add another; but you may replace the old one, on explicit permission from the previous publisher that added the old one.

The author(s) and publisher(s) of the Document do not by this License give permission to use their names for publicity for or to assert or imply endorsement of any Modified Version.

5. COMBINING DOCUMENTS

You may combine the Document with other documents released under this License, under the terms defined in section 4 above for modified versions, provided that you include in the combination all of the Invariant Sections of all of the original documents, unmodified, and list them all as Invariant Sections of your combined work in its license notice, and that you preserve all their Warranty Disclaimers.

The combined work need only contain one copy of this License, and multiple identical Invariant Sections may be replaced with a single copy. If there are multiple Invariant Sections with the same name but different contents, make the title of each such section unique by adding at the end of it, in parentheses, the name of the original author or publisher of that section if known, or else a unique number. Make the same adjustment to the section titles in the list of Invariant Sections in the license notice of the combined work.

In the combination, you must combine any sections Entitled "History" in the various original documents, forming one section Entitled "History"; likewise combine any sections Entitled "Acknowledgements", and any sections Entitled "Dedications". You must delete all sections Entitled "Endorsements."

6. COLLECTIONS OF DOCUMENTS

You may make a collection consisting of the Document and other documents released under this License, and replace the individual copies of this License in the various documents with a single copy that is included in the collection, provided that you follow the rules of this License for verbatim copying of each of the documents in all other respects.

You may extract a single document from such a collection, and distribute it individually under this License, provided you insert a copy of this License into the extracted document, and follow this License in all other respects regarding verbatim copying of that document.

7. AGGREGATION WITH INDEPENDENT WORKS

A compilation of the Document or its derivatives with other separate and independent documents or works, in or on a volume of a storage or distribution medium, is called an "aggregate" if the copyright resulting from the compilation is not used to limit the legal rights of the compilation's users beyond what the individual works permit. When the Document is included in an aggregate, this License does not apply to the other works in the aggregate which are not themselves derivative works of the Document.

If the Cover Text requirement of section 3 is applicable to these copies of the Document, then if the Document is less than one half of the entire aggregate, the Document's Cover Texts may be placed on covers that bracket the Document within the aggregate, or the electronic equivalent of covers if the Document is in electronic form. Otherwise they must appear on printed covers that bracket the whole aggregate.

8. TRANSLATION

Translation is considered a kind of modification, so you may distribute translations of the Document under the terms of section 4. Replacing Invariant Sections with translations requires special permission from their copyright holders, but you may include translations of some or all Invariant Sections in addition to the original versions of these Invariant Sections. You may include a translation of this License, and all the license notices in the Document, and any Warranty Disclaimers, provided that you also include the original English version of this License and the original versions of those notices and disclaimers. In case of a disagreement between the translation and the original version of this License or a notice or disclaimer, the original version will prevail.

If a section in the Document is Entitled "Acknowledgements", "Dedications", or "History", the requirement (section 4) to Preserve its Title (section 1) will typically require changing the actual title.

9. TERMINATION

You may not copy, modify, sublicense, or distribute the Document except as expressly provided under this License. Any attempt otherwise to copy, modify, sublicense, or distribute it is void, and will automatically terminate your rights under this License.

However, if you cease all violation of this License, then your license from a particular copyright holder is reinstated (a) provisionally, unless and until the copyright holder explicitly and finally terminates your license, and (b) permanently, if the copyright holder fails to notify you of the violation by some reasonable means prior to 60 days after the cessation.

Moreover, your license from a particular copyright holder is reinstated permanently if the copyright holder notifies you of the violation by some reasonable means, this is the first time you have received notice of violation of this License (for any work) from that copyright holder, and you cure the violation prior to 30 days after your receipt of the notice.

Termination of your rights under this section does not terminate the licenses of parties who have received copies or rights from you under this License. If your rights have been terminated and not permanently reinstated, receipt of a copy of some or all of the same material does not give you any rights to use it.

10. FUTURE REVISIONS OF THIS LICENSE

The Free Software Foundation may publish new, revised versions of the GNU Free Documentation License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns. See `http://www.gnu.org/copyleft/`.

Each version of the License is given a distinguishing version number. If the Document specifies that a particular numbered version of this License "or any later version" applies to it, you have the option of following the terms and conditions either of that specified version or of any later version that has been published (not as a draft) by the Free Software Foundation. If the Document does not specify a version number of this License, you may choose any version ever published (not as a draft) by the Free Software Foundation. If the Document specifies that a proxy can decide which future versions of this License can be used, that proxy's public statement of acceptance of a version permanently authorizes you to choose that version for the Document.

11. RELICENSING

"Massive Multiauthor Collaboration Site" (or "MMC Site") means any World Wide Web server that publishes copyrightable works and also provides prominent facilities for anybody to edit those works. A public wiki that anybody can edit is an example of such a server. A "Massive Multiauthor Collaboration" (or "MMC") contained in the site means any set of copyrightable works thus published on the MMC site.

"CC-BY-SA" means the Creative Commons Attribution-Share Alike 3.0 license published by Creative Commons Corporation, a not-for-profit corporation with a principal place of business in San Francisco, California, as well as future copyleft versions of that license published by that same organization.

"Incorporate" means to publish or republish a Document, in whole or in part, as part of another Document.

An MMC is "eligible for relicensing" if it is licensed under this License, and if all works that were first published under this License somewhere other than this MMC, and subsequently incorporated in whole or in part into the MMC, (1) had no cover texts or invariant sections, and (2) were thus incorporated prior to November 1, 2008.

The operator of an MMC Site may republish an MMC contained in the site under CC-BY-SA on the same site at any time before August 1, 2009, provided the MMC is eligible for relicensing.

ADDENDUM: How to use this License for your documents

To use this License in a document you have written, include a copy of the License in the document and put the following copyright and license notices just after the title page:

```
Copyright (C)  year  your name.
Permission is granted to copy, distribute and/or modify this document
under the terms of the GNU Free Documentation License, Version 1.3
or any later version published by the Free Software Foundation;
with no Invariant Sections, no Front-Cover Texts, and no Back-Cover
Texts.  A copy of the license is included in the section entitled ``GNU
Free Documentation License''.
```

If you have Invariant Sections, Front-Cover Texts and Back-Cover Texts, replace the "with. . . Texts." line with this:

```
with the Invariant Sections being list their titles, with
the Front-Cover Texts being list, and with the Back-Cover Texts
being list.
```

If you have Invariant Sections without Cover Texts, or some other combination of the three, merge those two alternatives to suit the situation.

If your document contains nontrivial examples of program code, we recommend releasing these examples in parallel under your choice of free software license, such as the GNU General Public License, to permit their use in free software.

Procedure and Macro Index

E

F

G

H

I

J

K

L

M

N

O

P

S

T

U

V

W

X

Y

Z

Variable Index

Concept and Feature Index

F

G

H

I

J

K

L

M

N